



Ruby on Rails Tutorial

Learn Rails by Example

Michael Hartl

Ruby on Rails Tutorial

Learn Rails by Example

Michael Hartl

Contents

1	From zero to deploy	11
1.1	Introduction	13
1.1.1	Comments for various readers	14
1.1.2	“Scaling” Rails	16
1.1.3	Conventions in this book	16
1.2	Up and running	17
1.2.1	Development environments	17
1.2.2	Ruby, RubyGems, and Rails	19
1.2.3	The first application	21
1.2.4	Model-view-controller (MVC)	22
1.2.5	script/server	24
1.3	Version control with Git	31
1.3.1	Installation and setup	31
1.3.2	Adding and committing	33
1.3.3	What good does Git do you?	34
1.3.4	GitHub	35
1.3.5	Branch, edit, commit, merge	37
1.4	Deploying	41
1.4.1	Heroku setup	42
1.4.2	Heroku deployment, step one	43
1.4.3	Heroku deployment, step two	43
1.4.4	Heroku commands	43
1.5	Conclusion	46
2	A demo app	47
2.1	Planning the application	47
2.1.1	Modeling users	48
2.1.2	Modeling microposts	48
2.2	The Users resource	49
2.2.1	A user tour	50
2.2.2	MVC in action	57
2.2.3	Weaknesses of this Users resource	62
2.3	The Microposts resource	63
2.3.1	A micropost microtour	63
2.3.2	Putting the <i>micro</i> in microposts	65
2.3.3	A user has many microposts	68

2.3.4	Inheritance hierarchies	69
2.3.5	Deploying the demo app	71
2.4	Conclusion	73
3	Mostly static pages	75
3.1	Static pages	76
3.1.1	Truly static pages	77
3.1.2	Static pages with Rails	80
3.2	Our first tests	83
3.2.1	Testing tools	84
3.2.2	TDD: Red, Green, Refactor	86
3.3	Slightly dynamic pages	97
3.3.1	Testing a title change	97
3.3.2	Passing title tests	99
3.3.3	Instance variables and Embedded Ruby	101
3.3.4	Eliminating duplication with layouts	105
3.4	Conclusion	107
3.5	Exercises	107
4	Rails-flavored Ruby	111
4.1	Motivation	111
4.1.1	A <code>title</code> helper	111
4.1.2	Cascading Style Sheets	113
4.2	Strings and methods	115
4.2.1	Comments	116
4.2.2	Strings	116
4.2.3	Objects and message passing	119
4.2.4	Method definitions	121
4.2.5	Back to the <code>title</code> helper	122
4.3	Other data structures	123
4.3.1	Arrays and ranges	123
4.3.2	Blocks	126
4.3.3	Hashes and symbols	127
4.3.4	CSS revisited	130
4.4	Ruby classes	131
4.4.1	Constructors	131
4.4.2	Class inheritance	132
4.4.3	Modifying built-in classes	134
4.4.4	A controller class	135
4.4.5	A user class	136
4.5	Exercises	138
5	Filling in the layout	139
5.1	Adding some structure	139
5.1.1	Site navigation	141
5.1.2	Custom CSS	145
5.1.3	Partials	152
5.2	Layout links	155

5.2.1	Integration tests	157
5.2.2	Rails routes	159
5.2.3	Named routes	162
5.3	User signup: A first step	163
5.3.1	Users controller	165
5.3.2	Signup URL	168
5.4	Conclusion	170
5.5	Exercises	170
6	Modeling and viewing users, part I	173
6.1	User model	173
6.1.1	Database migrations	175
6.1.2	The model file	179
6.1.3	Creating user objects	181
6.1.4	Finding user objects	184
6.1.5	Updating user objects	185
6.2	User validations	187
6.2.1	Validating presence	187
6.2.2	Length validation	193
6.2.3	Format validation	194
6.2.4	Uniqueness validation	196
6.3	Viewing users	201
6.3.1	Debug and Rails environments	201
6.3.2	User model, view, controller	204
6.3.3	A Users resource	208
6.4	Conclusion	211
6.5	Exercises	211
7	Modeling and viewing users, part II	213
7.1	Insecure passwords	213
7.1.1	Password validations	213
7.1.2	A password migration	217
7.1.3	An Active Record callback	220
7.2	Secure passwords	222
7.2.1	A secure password test	223
7.2.2	Some secure password theory	224
7.2.3	Implementing <code>has_password?</code>	225
7.2.4	An authenticate method	228
7.3	Better user views	230
7.3.1	Testing the user show page (with factories)	231
7.3.2	A name and a Gravatar	234
7.3.3	A user sidebar	240
7.4	Conclusion	242
7.4.1	Git commit	242
7.4.2	Heroku deploy	244
7.5	Exercises	244

8 Sign up	247
8.1 Signup form	247
8.1.1 Using <code>form_for</code>	248
8.1.2 The form HTML	252
8.2 Signup failure	256
8.2.1 Testing failure	256
8.2.2 A working form	260
8.2.3 Signup error messages	263
8.2.4 Filtering parameter logging	264
8.3 Signup success	266
8.3.1 Testing success	268
8.3.2 The finished signup form	269
8.3.3 The flash	269
8.3.4 The first signup	272
8.4 RSpec integration tests	272
8.4.1 Webrat	275
8.4.2 Users signup failure should not make a new user	276
8.4.3 Users signup success should make a new user	279
8.5 Conclusion	280
8.6 Exercises	281
9 Sign in, sign out	285
9.1 Sessions	285
9.1.1 Sessions controller	286
9.1.2 Signin form	288
9.2 Signin failure	291
9.2.1 Reviewing form submission	291
9.2.2 Failed signin (test and code)	294
9.3 Signin success	297
9.3.1 The completed <code>create</code> action	299
9.3.2 Remember me	300
9.3.3 Cookies	304
9.3.4 Current user	305
9.4 Signing out	311
9.4.1 Destroying sessions	311
9.4.2 Signin upon signup	313
9.4.3 Changing the layout links	314
9.4.4 Signin/out integration tests	317
9.5 Conclusion	319
9.6 Exercises	319
10 Updating, showing, and deleting users	321
10.1 Updating users	321
10.1.1 Edit form	321
10.1.2 Enabling edits	328
10.2 Protecting pages	331
10.2.1 Requiring signed-in users	331
10.2.2 Requiring the right user	334

10.2.3 An expectation bonus	336
10.2.4 Friendly forwarding	337
10.3 Showing users	339
10.3.1 User index	341
10.3.2 Sample users	343
10.3.3 Pagination	347
10.3.4 Partial refactoring	352
10.4 Destroying users	353
10.4.1 Administrative users	353
10.4.2 The <code>destroy</code> action	357
10.5 Conclusion	361
10.6 Exercises	361
11 User microposts	363
11.1 A Micropost model	363
11.1.1 The basic model	363
11.1.2 User/Micropost associations	366
11.1.3 Micropost refinements	370
11.1.4 Micropost validations	373
11.2 Showing microposts	375
11.2.1 Augmenting the user show page	375
11.2.2 Sample microposts	381
11.3 Manipulating microposts	383
11.3.1 Access control	387
11.3.2 Creating microposts	389
11.3.3 A proto-feed	396
11.3.4 Destroying microposts	403
11.3.5 Testing the new home page	406
11.4 Conclusion	408
11.5 Exercises	409
12 Following users	413

About the author

[Michael Hartl](#) is a programmer, educator, and entrepreneur. Michael was coauthor of *RailsSpace*, a best-selling Rails tutorial book published in 2007, and was cofounder and lead developer of [Insoshi](#), a popular social networking platform in Ruby on Rails. Previously, he taught theoretical and computational physics at the [California Institute of Technology](#) (Caltech) for six years, where he received the Lifetime Achievement Award for Excellence in Teaching in 2000. Michael is a graduate of [Harvard College](#), has a [Ph.D. in Physics](#) from [Caltech](#), and is an alumnus of the [Y Combinator](#) program.

A book-in-progress

Ruby on Rails Tutorial is nearly done. The most recently completed book chapter is [Chapter 11](#); I estimate that the final chapter, [Chapter 12](#), will be done by May 24, 2010. If you have comments or suggestions, I welcome your [feedback](#). You should also subscribe via [RSS](#) or [by email](#) to get the latest project updates.

Ruby on Rails Tutorial Screencasts

To complement the written book, I plan to produce an associated screencast series. If you'd like to be notified when the screencasts are ready, you should [subscribe to the news feed](#) or [sign up for email notifications](#).

Acknowledgments

This book owes a lot to my previous Rails book, *RailsSpace*, and hence to my coauthor [Aurelius Prochazka](#). I'd like to thank Aure both for the work he did on that book and for his support of this one. I'd also like to thank the [Rails Activists](#) for their support and enthusiasm. As for the rest: "your name here", as they say. Help me out and you might just get a nod. Rails newbies: [tell me what's confusing](#); Rails experts: [tell me what sucks](#). (But please be nice. :-)

Copyright and license

Ruby on Rails Tutorial: Learn Rails by Example. Copyright © 2010 by Michael Hartl. All source code in *Ruby on Rails Tutorial* is available under the [MIT License](#) and the [Beerware License](#).

Copyright (c) 2010 Michael Hartl

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

```
/*
 * -----
 * "THE BEERWARE LICENSE" (Revision 42):
 * Michael Hartl wrote this code. As long as you retain this notice, you can
 * do whatever you want with this stuff. If we meet someday, and you think
 * this stuff is worth it, you can buy me a beer in return.
 * -----
 */
```


Chapter 1

From zero to deploy

Welcome to *Ruby on Rails Tutorial: Learn Rails by Example*. The goal of this book is to be the best answer to the question, “If I want to learn web development with [Ruby on Rails](#), where should I start?” Although the Ruby on Rails web framework benefits from a wealth of learning resources—including books, blogs, and screencasts—many of them assume a substantial background in Rails. *Ruby on Rails Tutorial* is designed to give you this background. By the time you finish this book, you will have a solid foundation in Rails programming, ready to benefit from more advanced Rails resources.

Ruby on Rails Tutorial follows essentially the same approach as my previous Rails book,¹ teaching web development with Rails by building a sample application from scratch. This allows you to see the many different pieces of Ruby on Rails in context, motivated by real problems, as well as how those pieces fit together. Moreover, although Rails is a large and fast-moving framework, in *Ruby on Rails Tutorial* we’ll focus on the smaller, more stable set of core Rails techniques that have crystallized in the last couple years, so the knowledge you gain here will not soon be obsolete. We’ll also take care not to rely on the generated code called *scaffolding* (Box 1.1), so that the Rails knowledge you gain will be adaptable to new problems.

In this chapter, we’ll get started with Ruby on Rails by installing all the necessary software and setting up our development environment (Section 1.2). We’ll then create our first Rails application, called (appropriately enough) `first_app`. *Rails Tutorial* emphasizes good software development practices, so immediately after creating our fresh new Rails project we’ll put it under version control with Git (Section 1.3). And, believe it or not, in this chapter we’ll even put our first app on the wider web by *deploying* it to production (Section 1.4).

In Chapter 2, we’ll make a second project, whose purpose will be to demonstrate the basic workings of a Rails application. To get up and running quickly, we’ll build this *demo app* (called `demo_app`) using scaffolding to generate code; since this code is both ugly and complex, Chapter 2 will focus on interacting with the demo app through its *URLs*² using a web browser.

In Chapter 3, we’ll create a *sample application* (called `sample_app`), this time writing all the code from scratch. We’ll develop the sample app using *test-driven development* (TDD),³ getting started in Chapter 3 by creating static pages and then adding a little dynamic content. We’ll take a quick detour in Chapter 4 to learn a little about the Ruby language underlying Rails. Then, in Chapter 5 through Chapter 10, we’ll complete the foundation for the sample application by making a site layout, a user data model, and a full registration and authentication system. Finally, in Chapter 11 and Chapter 12 we’ll add microblogging

¹*RailsSpace*, by Michael Hartl and Aurelius Prochazka (Addison-Wesley, 2007).

²URL stands for Uniform Resource Locator. In practice it is usually equivalent to “the thing you see in the address bar of your browser”. By the way, the current preferred term is *URI*, for Uniform Resource Identifier, but popular usage still tilts toward URL.

³We’ll be using [RSpec](#) for our tests, and in this context TDD is also known as Behavior Driven Development, or BDD. (Frankly, I’m not convinced there’s much of a difference.)

and social features to make a working example site.

The final sample application will bear more than a passing resemblance to a certain popular [social microblogging site](#)—a site which, coincidentally, is also written in Rails. Though of necessity our efforts will focus on this specific sample application, the emphasis throughout *Rails Tutorial* will be on general principles, so that you will have a solid foundation no matter what kinds of web applications you want to build.

Box 1.1. Scaffolding: Quicker, easier, more seductive

From the beginning, Rails has benefited from a palpable sense of excitement, starting with the famous [15-minute weblog video](#) by Rails creator David Heinemeier Hansson, now updated as the [15-minute weblog using Rails 2](#) by Ryan Bates. These videos are a great way to get a taste of Rails' power, and I recommend watching them. But be warned: they accomplish their amazing fifteen-minute feat using a feature called *scaffolding*, which relies heavily on *generated code*, magically created by the Rails `generate` command.

When writing a Ruby on Rails tutorial, it is tempting to rely on the scaffolding approach—it's [quicker, easier, more seductive](#). But the complexity and sheer amount of code in the scaffolding can be utterly overwhelming to a beginning Rails developer; you may be able to use it, but you probably won't understand it. Following the scaffolding approach risks turning you into a virtuoso script generator with little (and brittle) actual knowledge of Rails.

In *Rails Tutorial*, we'll take the (nearly) polar opposite approach: although [Chapter 2](#) will develop a small demo app using scaffolding, the core of *Rails Tutorial* is the sample app, which we'll start writing in [Chapter 3](#). At each stage of developing the sample application, we will generate *small, bite-sized* pieces of code—simple enough to understand, yet novel enough to be challenging. The cumulative effect will be a deeper, more flexible knowledge of Rails, giving you a good background for writing nearly any type of web application.

1.1 Introduction

Since its debut in 2004, Ruby on Rails has rapidly become one of the most powerful and popular frameworks for building dynamic web applications. Rails users run the gamut from scrappy startups to huge companies: [Posterous](#), [UserVoice](#), [37signals](#), [Shopify](#), [Scribd](#), [Twitter](#), [Hulu](#), the [Yellow Pages](#)—the [list of sites using Rails](#) goes on and on. There are also many web development shops that specialize in Rails, such as [ENTP](#), [thoughtbot](#), [Pivotal Labs](#), and [Hashrocket](#), plus innumerable independent consultants, trainers, and contractors.

What makes Rails so great? Rails owes much of its success to its elegant and compact design. By exploiting the malleability of the underlying [Ruby](#) language, Rails effectively creates a [domain-specific language](#) for writing web applications. As a result, many common web programming tasks—such as generating HTML, making data models, and routing URLs—are easy with Rails, and the resulting application code is concise and readable.

Rails also adapts rapidly to new developments in web technology and framework design. For example, Rails was one of the first frameworks to fully digest and implement the REST architectural style for structuring web applications (which we'll be learning about throughout this tutorial). And when other frameworks develop successful new techniques, Rails creator [David Heinemeier Hansson](#) and the [Rails core team](#) don't hesitate to incorporate their ideas. Perhaps the most dramatic example is the merger of Rails and Merb, a rival Ruby web framework, so that Rails now benefits from Merb's modular design, stable [API](#), and improved performance.⁴ (Anyone who has attended a talk by Merb developer and Rails core team member [Yehuda Katz](#) can't help but notice what an *extremely* good idea it was to bring the Merb team on board.)

Finally, Rails benefits from an unusually enthusiastic and diverse community. The results include well-attended [conferences](#), a huge number of [plugins](#) and [gems](#) (self-contained solutions to specific problems such as pagination or image upload), a rich variety of informative blogs, and a cornucopia of discussion forums and IRC channels. The large number of Rails programmers also makes it easier to handle the inevitable

⁴These changes are slated for release as part of Rails 3, which is currently in beta.

application errors: the “Google the error message” algorithm nearly always produces a relevant blog post or discussion-forum thread.

1.1.1 Comments for various readers

Rails Tutorial contains integrated tutorials not only for Rails, but also for the underlying Ruby language, as well as for [HTML](#), [CSS](#), some [JavaScript](#), and even a little [SQL](#). This means that, no matter where you currently are in your knowledge of web development, by the time you finish this tutorial you will be ready for more advanced Rails resources, as well as for the more systematic treatments of the other subjects mentioned.

Although this book has no formal prerequisites, you should of course have at least *some* computer experience. If you’ve never even used a text editor before, it will be tough going, but with enough determination you can probably soldier through. If, on the other hand, your `.emacs` file is so complex it could make a grown man cry, there is still plenty of material to keep you challenged. *Rails Tutorial* is designed to teach Rails development no matter what your background is, but your path and reading experience will depend on your particular circumstances.

All readers: One common question when learning Rails is whether to learn Ruby first. The answer depends on your personal learning style. If you prefer to learn everything systematically from the ground up, then learning Ruby first might work well for you, and there are several book recommendations in this section to get you started. On the other hand, many beginning Rails developers are excited about making *web* applications, and would rather not slog through a 500-page book on pure Ruby before ever writing a single web page. Moreover, the subset of Ruby needed by Rails developers is different from what you’ll find in a pure-Ruby introduction, whereas *Rails Tutorial* focuses on exactly that subset. If your primary interest is making web applications, I recommend starting with *Rails Tutorial* and then reading a book on pure Ruby next. It’s not an all-or-nothing proposition, though: if you start reading *Rails Tutorial* and feel your (lack of) Ruby knowledge holding you back, feel free to switch to a Ruby book and come back when you feel ready. You might also consider getting a taste of Ruby by following a short online tutorial, such as can be found at <http://www.ruby-lang.org/> or <http://rubylearning.com/>.

Another common question is whether to use tests from the start. As noted in the introduction, *Rails Tutorial* uses test-driven development (also called test-first development), which in my view is the best way to develop Rails applications, but it does introduce a substantial amount of overhead and complexity. If you find yourself getting bogged down by the tests, feel free to skip them on first reading.⁵ Indeed, some readers may find the inclusion of so many moving parts—such as tests, version control, and deployment—a bit overwhelming at first, and if you find yourself expending excessive energy on any of these steps, *don’t hesitate to skip them*. Although I have included only material I consider essential to developing professional-grade Rails applications, only the core application code is strictly necessary the first time through.

Inexperienced programmers (non-designers): *Rails Tutorial* doesn’t assume any background other than general computer knowledge, so if you have limited programming experience this book is a good place to start. Please bear in mind that it is only the first step on a long journey; web development has many moving parts, including HTML/CSS, JavaScript, databases (including SQL), version control, and deployment. This book contains short introductions to these subjects, but there is much more to learn.

Inexperienced programmers (designers): Your design skills give you a big leg up, since you probably already know HTML and CSS. After finishing this book you will be in an excellent position to work with existing Rails projects and possibly start some of your own. You may find the programming material challenging, but the Ruby language is unusually friendly to beginners, especially those with an artistic bent.

⁵In practice, this will involve omitting all files with `spec` in their name, as we will start to see in Section 3.2.2.

After finishing *Ruby on Rails Tutorial*, I recommend that newer programmers read [Beginning Ruby](#) by Peter Cooper, which shares the same basic instructional philosophy as *Rails Tutorial*. I also recommend [The Ruby Way](#) by Hal Fulton. Finally, to gain a deeper understanding of Rails I recommend [The Rails Way](#)⁶ by Obie Fernandez.

Web applications, even relatively simple ones, are by their nature fairly complex. If you are completely new to web programming and find *Rails Tutorial* overwhelming, it could be that you're not quite ready to make web applications yet. In that case, I'd suggest learning the basics of HTML and CSS and then giving *Rails Tutorial* another go. (Unfortunately, I don't have a personal recommendation here, but [Head First HTML](#) looks promising.) You might also consider reading the first few chapters of [Beginning Ruby](#), which starts with sample applications much smaller than a full-blown web app.

Experienced programmers new to web development: Your previous experience means you probably already understand ideas like classes, methods, data structures, etc., which is a big advantage. Be warned that if your background is in C/C++ or Java, you may find Ruby a bit of an odd duck, and it might take time to get used to it; just stick with it and eventually you'll be fine. (Ruby even lets you put semicolons at the ends of lines if you miss them too much. :-) *Rails Tutorial* covers all the web-specific ideas you'll need, so don't worry if you don't currently know a PUT from a POST.

Experienced web developers new to Rails: You have a great head start, especially if you have used a dynamic language such as PHP or (even better) Python. The basics of what we cover will likely be familiar, but test-driven development may be new to you, as may be the structured REST style favored by Rails. Ruby has its own idiosyncrasies, so those will likely be new, too.

Experienced Ruby programmers: The set of Ruby programmers who don't know Rails is a small one nowadays, but if you are a member of this elite group you can fly through this book and then move on to [The Rails Way](#) by Obie Fernandez.

Inexperienced Rails programmers: You've perhaps read some other tutorials and made a few small Rails apps yourself. Based on reader feedback, I'm confident that you can still get a lot out of this book. Among other things, the techniques here may be more up-to-date than the ones you picked up when you originally learned Rails.

Experienced Rails programmers: This book is unnecessary for you, but many experienced Rails developers have expressed surprise at how much they learned from this book, and you might enjoy seeing Rails from a different perspective.

After finishing *Ruby on Rails Tutorial*, I recommend that experienced (non-Ruby) programmers read [The Well-Grounded Rubyist](#) by David A. Black, which is an excellent in-depth discussion of Ruby from the ground up, or [The Ruby Way](#) by Hal Fulton, which is also fairly advanced but takes a more topical approach. Then move on to [The Rails Way](#) to deepen your Rails expertise.

At the end of this process, no matter where you started, you will be ready for the more intermediate-to-advanced Rails resources. Here are some I particularly recommend:

- [Railscasts](#): Excellent free Rails screencasts
- [PeepCode](#), [Pragmatic.tv](#), [EnvyCasts](#): Excellent commercial screencasters

⁶A new version of *The Rails Way*, covering Rails 3, is due out soon. You might want to wait for that to come out before buying it.

- [Rails Guides](#): Good topical and up-to-date Rails references. *Rails Tutorial* refers frequently to the *Rails Guides* for more in-depth treatment of specific topics.
- [Rails blogs](#): Too many to list, but there are tons of good ones.

1.1.2 “Scaling” Rails

Before moving on with the rest of the introduction, I’d like to take a moment to address the one issue that dogged the Rails framework the most in its early days: the supposed inability of Rails to “scale”—i.e., to handle large amounts of traffic. Part of this issue relied on a misconception; [you scale a site, not a framework](#), and Rails, as awesome as it is, is only a framework. So the real question should have been, “Can a site built with Rails scale?” In any case, the question has now been definitively answered in the affirmative: some of the most heavily trafficked sites in the world use Rails. Actually *doing* the scaling is beyond the scope of just Rails, but rest assured that if *your* application ever needs to handle the load of Hulu or the Yellow Pages, Rails won’t stop you from taking over the world.

1.1.3 Conventions in this book

The conventions in this book are mostly self-explanatory; in this section I’ll mention some that may not be. First, both the [HTML](#) and [PDF](#) editions of this book are full of links, both to internal sections (such as [Section 1.2](#)) and to external sites (such as the main [Ruby on Rails download](#) page).⁷ Sometimes you may encounter an undefined internal link, such as [Section ??](#), but don’t be alarmed; this is a book-in-progress, and these undefined cross-references are usually intentional reminders of gaps to fill.

Second, your humble author is a Linux/OS X kind of guy, and hasn’t used Windows as his primary OS for more than a decade; as a result, *Rails Tutorial* has an unmistakable Unix flavor.⁸ For example, in this book all command line examples use a Unix-style command line prompt (a dollar sign):

```
$ echo "hello, world"
hello, world
```

Rails comes with lots of commands that can be run directly under Unix but often require an explicit call to the Ruby interpreter under Windows;⁹ for example, in [Section 1.2.5](#) we’ll run a local development web server as follows:

```
# Unix (Linux, OS X, etc.)
$ script/server
# Windows (explicitly uses the Ruby interpreter)
> ruby script/server
```

⁷When reading *Rails Tutorial*, you may find it convenient to follow an internal section link to look at the reference and then immediately go back to where you were before. This is easy when reading the book as a web page, since you can just use the Back button of your browser, but both Adobe Reader and OS X’s Preview allow you to do this with the PDF as well. In Reader, you can right-click on the document and select “Previous View” to go back. In Preview, use the Go menu: [Go > Back](#).

⁸Indeed, the entire Rails community has this flavor. In a full room at [RailsConf](#) you’ll see a handful of PCs in a sea of MacBooks—with probably half the PCs running Linux. You can certainly develop Rails apps on Microsoft Windows, but you’ll definitely be in the minority. Since I have limited Windows Rails dev experience myself, I’d especially appreciate it if Windows readers could [give me feedback](#) about what works and what doesn’t, so that we can get this tutorial to work on all common platforms.

⁹The Ruby interpreter, called [ruby](#), parses and executes Ruby source code files.

Rails Tutorial will also use Unix-style forward slashes as directory separators; my Rails Tutorial sample app, for instance, lives in

```
/Users/mhartl/rails_projects/first_app
```

The root directory for any given app is known as the Rails root, and henceforth all directories will be relative to this directory. For example, the `config` directory of my sample application is in

```
/Users/mhartl/rails_projects/first_app/config
```

This means that when referring to the file

```
/Users/mhartl/rails_projects/first_app/config/routes.rb
```

I'll omit the Rails root and write `config/routes.rb` for brevity.

Finally, *Rails Tutorial* often shows output from various programs (shell commands, version control status, Ruby programs, etc.); because of the innumerable small differences between different computer systems, the output you see may not always agree exactly with what is shown in the text. This is normal, and is not any cause for concern. As long as the application pages load and the test suite passes, you're good to go.

1.2 Up and running

It's time now to get going with a Ruby on Rails development environment and our first application. There is quite a bit of overhead here, especially if you don't have extensive programming experience, so don't get discouraged if it takes a while to get started. It's not just you; every developer goes through it (often more than once), but rest assured that the effort will be richly rewarded.

1.2.1 Development environments

Considering various idiosyncratic customizations, there are probably as many development environments as there are Rails programmers, but there are at least two broad themes: text editor/command line environments, and integrated development environments (IDEs). Let's consider the latter first.

IDEs

There is no shortage of Rails IDEs; indeed, the main Ruby on Rails site names four: [RadRails](#), [RubyMine](#), [3rd Rail](#), and [NetBeans](#). All are cross-platform, and I've heard good things about several of them. I encourage you to try them and see if they work for you, but I have a confession to make: I have never found an IDE that met all my Rails development needs—and for some projects I haven't even been able to get them to work at all.

Text editors and command lines

What are we to use to develop Rails apps, if not some awesome all-in-one IDE? I'd guess the majority of Rails developers opt for the same solution I've chosen: use a *text editor* to edit text, and a *command line* to issue commands ([Figure 1.1](#)). Which combination you use depends on your tastes and your platform:

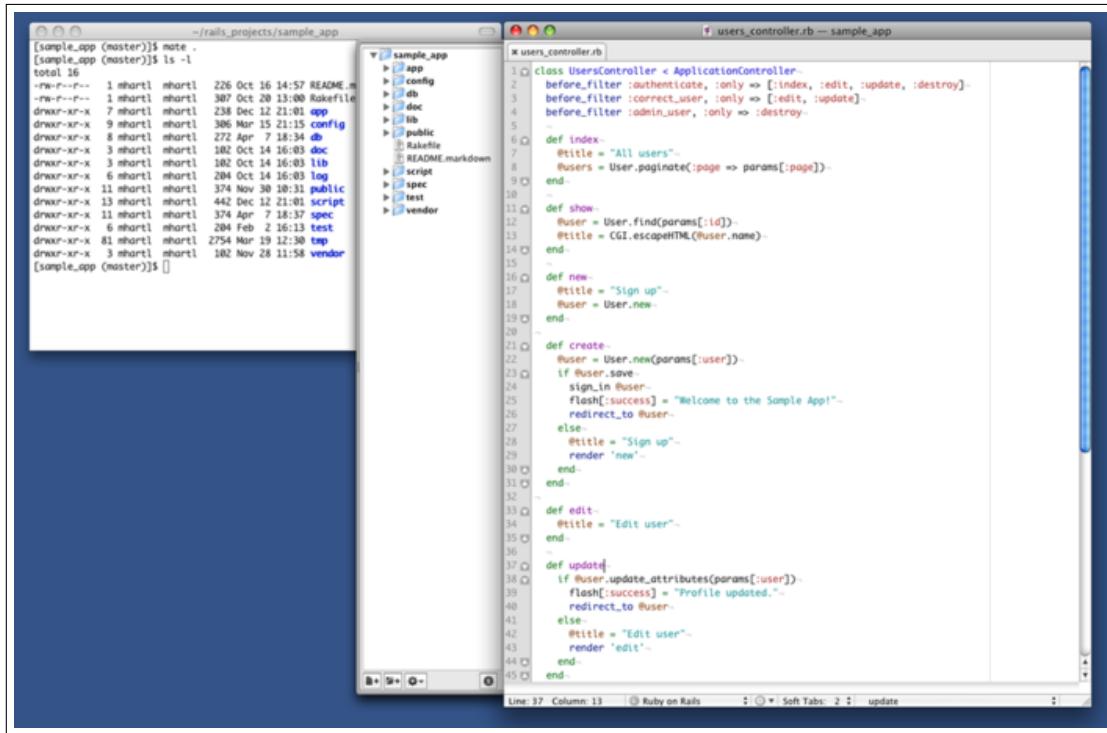


Figure 1.1: A text editor/command line development environment (TextMate/iTerm). ([full size](#))

- **Macintosh OS X:** Like many Rails developers, I prefer [TextMate](#). Other options include Emacs and [MacVim](#) (launched with the command `macvim`), the excellent Macintosh version of Vim.¹⁰ I use iTerm for my command line terminal; others prefer the native Terminal app.
- **Linux:** Your editor options are basically the same as OS X, minus TextMate. I'd recommend graphical Vim (gVim), [gedit](#) (with the [GMate](#) plugins), or [Kate](#). As far as command lines go, you're totally set: every Linux distribution comes with at least one command line terminal application (and often several).
- **Windows:** Unfortunately, I can't make any personal recommendations here, but you can do what I did: drop “[rails windows](#)” into Google to see what the latest thinking is on setting up a Rails development environment on Windows. Two combinations look especially promising: [Vim for Windows](#) with [Console](#) (recommended by [Akita On Rails](#)) or the [E Text Editor](#) with Console and [Cygwin](#) (recommended by [Ben Kittrell](#)). Rails Tutorial readers have suggested looking at [Komodo Edit](#) (cross-platform) and the [Sublime Text](#) editor (Windows only) as well. No matter which editor you choose, I recommend trying Cygwin, which provides the equivalent of a Unix terminal under Windows; most of the command-line examples in the book should work with minimum modification.

If you go with some flavor of Vim, be sure to tap into the thriving community of [Vim-using Rails hackers](#). See especially the [rails.vim](#) enhancements and the [NERD tree](#) project drawer.

¹⁰The vi editor is one of the most ancient yet powerful weapons in the Unix arsenal, and Vim is “vi improved”.

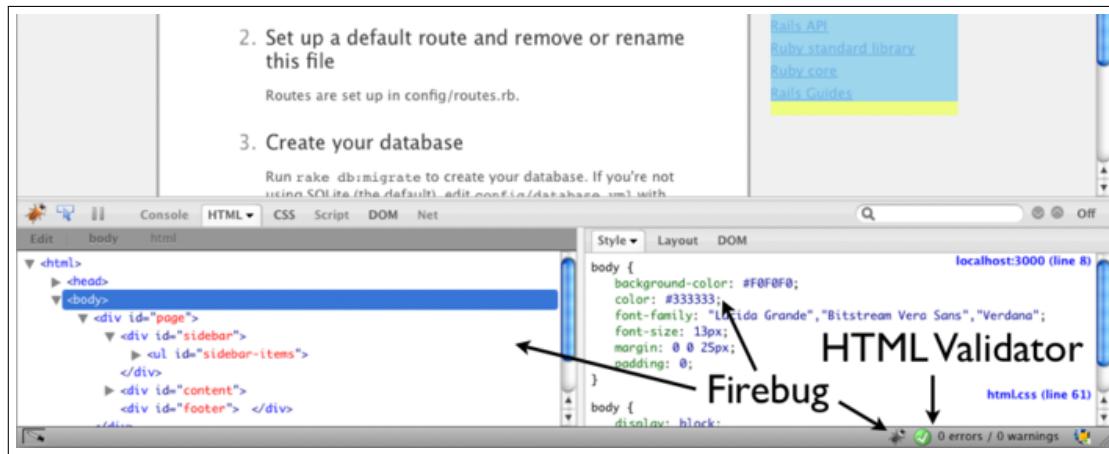


Figure 1.2: Firefox running with the [Firebug](#) and [HTML Validator](#) add-ons. ([full size](#))

Browsers

Although there are many web browsers to choose from, the vast majority of Rails programmers use either Firefox or Safari when developing. I prefer Firefox, and the screenshots in Rails Tutorial will generally be of a Firefox browser. One of the big reasons I like Firefox is the amazing [Firebug](#) add-on, which lets you perform all sorts of magic, such as dynamically inspecting (and even editing) the HTML structure and CSS rules on any page. (Apparently Safari has some sort of equivalent, but I already know and love Firebug.)

I also recommend the [Firefox HTML Validator](#), which automatically checks to make sure that your pages are valid HTML markup. (If you install Validator, you should go to the add-on Preferences to disable it for all pages, and then visit [localhost](#) and explicitly enable it for local pages. Otherwise it will try to validate every page you visit.) I have never been able to consistently keep a large project's HTML validated without HTML Validator. Figure 1.2 shows both add-ons in action.

A note about tools

In the process of getting your development environment up and running, you may find that you spend a *lot* of time getting everything just right. The learning process for editors and IDEs is particularly long; you can spend weeks on TextMate or Vim tutorials alone. If you're new to this game, I want to assure you that *spending time learning tools is normal*. Everyone goes through it. Sometimes it is frustrating, and it's easy to get impatient when you have an awesome web app in your head and you *just want to learn Rails already*, but have to spend a week learning some weird ancient Unix editor just to get started. But a craftsman has to know his tools; in the end the reward is worth the effort.

1.2.2 Ruby, RubyGems, and Rails

Now it's time to install Ruby and Rails. The canonical up-to-date source for this step is the [Ruby on Rails download page](#). I'll assume you can go there now; parts of this book can be read profitably offline, but not this part. I'll just inject some of my own comments on the steps.

Install Ruby

It's possible that your system already has Ruby; try running

```
$ ruby -v  
ruby 1.8.7
```

to see the version number. The Rails team recommends Ruby 1.8.7, but 1.8.6 works for the time being. (Rails 3 will require Ruby 1.8.7 and will work best with Ruby 1.9.2.)

If you have to install Ruby, I *strongly* recommend doing a search for “install ruby OS X 10.5” (or whatever is appropriate for your system), since there are tons of options depending on your exact setup. Be warned that installing Ruby might take some trickery. If you run into any errors, just Google them.

More advanced users may want to consider installing Ruby using [Ruby Version Manager](#), also known as [RVM](#), which allows you to install and manage multiple versions of Ruby on the same machine. This is particularly important if you want to experiment with Rails 3.0 beta, which can't run in parallel with Rails 2.3.8 (installed in [Section 1.2.2](#)). If you want to go this route, I suggest using RVM to install both Ruby 1.8.7/Rails 2.3.8 and Ruby 1.9.2/Rails 3.0.pre; see the [RVM documentation](#) for further details. If you run into any problems with RVM, you can often find its creator, Wayne Seguin, on the RVM IRC channel (#rvm on freenode.net).¹¹

Note: If you're running Ubuntu Linux, at this point you might try [these instructions for installing Rails on Karmic Koala](#) rather than following the steps below.

Install RubyGems

RubyGems is a package manager for Ruby projects, and there are tons of great libraries (including Rails) available as Ruby packages, or *gems*. Installing RubyGems should be easy once you install Ruby. You should [download RubyGems](#), extract/unzip it, and then go to the `rubygems` directory and run the setup program:

```
$ [sudo] ruby setup.rb
```

Here `sudo` executes the command `ruby setup.rb` as an administrative user, which has access to files and directories that normal users can't touch; I have put it in brackets to indicate that using `sudo` may or may not be necessary for your particular system. Most Unix/Linux/OS X systems require `sudo` by default, unless you are using RVM as suggested in [Section 1.2.2](#). Note that you should *not* actually type any brackets; you should run either

```
$ sudo ruby setup.rb
```

or

¹¹If you haven't used IRC before, I suggest you start by searching the web for “irc client ;your platform;”. I use [Colloquy](#), an IRC client for OS X.

```
$ ruby setup.rb
```

depending on your system.

If you already have RubyGems installed, you might want to update your system to the latest version:

```
$ [sudo] gem update --system
```

Install Rails

This part should be *really* easy once you install RubyGems:

```
$ [sudo] gem install rails -v 2.3.8
```

Here the optional version flag `-v` just ensures that you get the exact version of Rails used in this tutorial.¹² (Rails 3 is already in beta, and I will update the tutorial as soon as possible after its official release.) To verify that this worked, run the following command:¹³

```
$ rails -v
Rails 2.3.8
```

1.2.3 The first application

Virtually all Rails applications start the same way, with the `rails` command. This handy program creates a skeleton Rails application in a directory of your choice. To get started, make a directory for your Rails projects and then run the `rails` command to make the initial sample application:

Listing 1.1. Running the `rails` script to generate a sample application.

```
$ mkdir rails_projects
$ cd rails_projects
$ rails first_app
  create
  create  app/controllers
  create  app/helpers
  create  app/models
  create  app/views/layouts
```

¹²Because gem updates often cause minor but potentially confusing breakage, in this tutorial I'll usually include an explicit version number known to work. Feel free to experiment, though; if you want to live on the edge, omit `-v` and everything after it—just promise not to come crying to me if it breaks. :-)

¹³If you're running Linux and get an `unrecognized option '-v'` error, this is a known issue running Rails on some versions of Debian.

File/Directory	Purpose
app/	Core application (app) code, including models, views, controllers, and helpers
config/	Application configuration
db/	Files to manipulate the database
doc/	Documentation for the application
lib/	Library modules
log/	Application log files
public/	Data accessible to the public (e.g., web browsers), including images and cascading style sheets (CSS)
script/	Scripts provided by Rails for generating code, opening console sessions, or starting a local web server
test/	Application tests (made obsolete by the spec/ directory in Section 3.2.1)
tmp/	Temporary files
vendor/	Third-party code such as plugins and gems
README	A brief description of the application
Rakefile	Utility tasks available via the <code>rake</code> command

Table 1.1: A summary of the default Rails directory structure.

```

create config/environments
create config/initializers
create config/locales
create db
create doc
create lib
create lib/tasks
create log
create public/images
create public/javascripts
create public/stylesheets
.
.
.

```

Notice how many files and directories the `rails` command creates. This standard directory and file structure ([Figure 1.3](#)) is one of the many advantages of Rails; it immediately gets you from zero to a functional (if minimal) application. Moreover, since the structure is common to all Rails apps, you can immediately get your bearings when looking at someone else’s code. A summary of the default Rails files appears in [Table 1.1](#); we’ll learn about most of these files and directories throughout the rest of this book.

1.2.4 Model-view-controller (MVC)

Even at this early stage, it’s helpful to get a high-level overview of how Rails applications work ([Figure 1.4](#)). You might have noticed that the standard Rails application structure ([Figure 1.3](#)) has an application directory called `app/` with three subdirectories: `models`, `views`, and `controllers`. This is a hint that Rails follows the **model-view-controller** (MVC) architectural pattern, which enforces a separation between “domain logic” (also called “business logic”) from the input and presentation logic associated with a graphical user interface (GUI). In the case of web applications, the “domain logic” typically consists of data models for things like users, articles, and products, and the GUI is just a web page in a web browser.

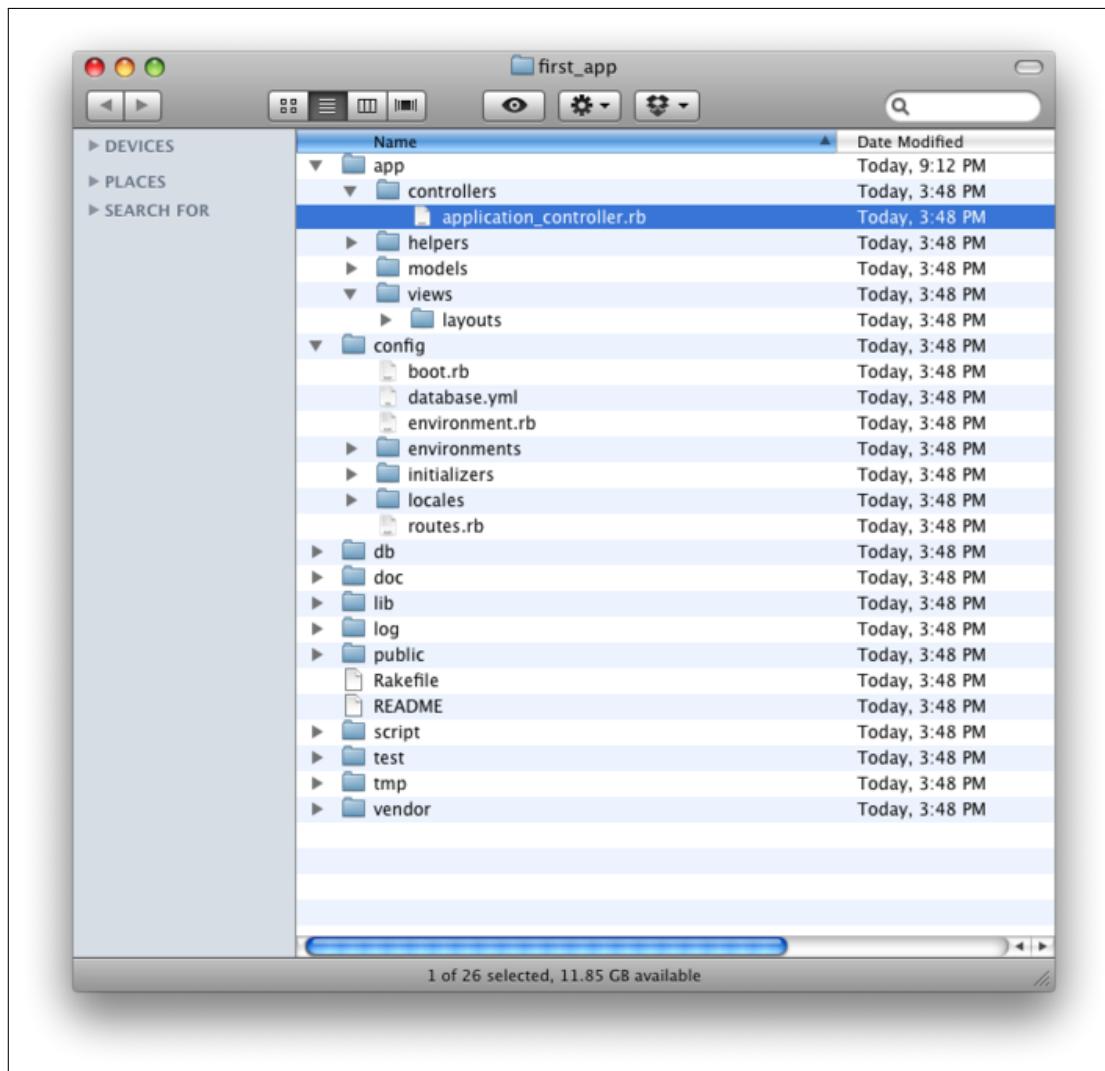


Figure 1.3: The directory structure for a newly hatched Rails app. ([full size](#))

When interacting with a Rails application, a browser sends a *request*, which is received by a web server and passed on to a Rails *controller*, which is in charge of what to do next. In some cases, the controller will immediately render a *view*, which is a template that gets converted to HTML and sent back to the browser. More commonly for dynamic sites, the controller interacts with a *model*, which is a Ruby object that represents an element of the site (such as a user) and is in charge of communicating with the database. After invoking the model, the controller then renders the view and returns the complete web page to the browser as HTML.

If this discussion seems a bit abstract right now, worry not; we'll refer back to this section frequently. In addition, Section 2.2.2 has a more detailed discussion of MVC in the context of the demo app. Finally, the sample app will use all aspects of MVC; we'll cover controllers and views starting in Section 3.1.2, models starting in Section 6.1, and we'll see all three working together in Section 6.3.2.

1.2.5 `script/server`

Thanks to running `rails` in Listing 1.1, we already have an application we can run—but how? Happily, Rails comes with a command-line program, or *script*, that runs a *local* web server,¹⁴ visible only from your development machine:¹⁵

```
$ cd first_app/
$ script/server
=> Booting WEBrick
=> Rails 2.3.8 application starting on http://0.0.0.0:3000
=> Call with -d to detach
=> Ctrl-C to shutdown server
```

This tells us that the application is running on `port number` 3000¹⁶ at the address `0.0.0.0`. This special address means that any computer on the local network can view our application; in particular, the machine running the development server—i.e., the local development machine—can view the application using the address `localhost:3000`.¹⁷ We can see the result of visiting `http://localhost:3000/` in Figure 1.5.

Unfortunately, we're not quite done; click on the link `About your application's environment` and you'll see your very first Rails error message¹⁸ (Figure 1.6). Looking in the web server log (Figure 1.7), we eventually see the incriminating lines:

```
Status: 500 Internal Server Error
no such file to load -- sqlite3
```

¹⁴The default Rails web server is *WEBrick*, a pure-Ruby server that isn't suitable for production use but is fine in development. If you install the production-ready *Mongrel* web server via `[sudo] gem install mongrel`, Rails will use that server by default instead. Either way works.

¹⁵Recall from Section 1.1.3 that Windows users might have to type `ruby script/server` instead.

¹⁶Normally, web sites run on port 80, but this usually requires special privileges, so Rails picks a less restricted higher-numbered port for the development server.

¹⁷You can also access the application by visiting `0.0.0.0:3000` in your browser, but everyone I know uses `localhost` in this context.

¹⁸It won't be your last. :-)

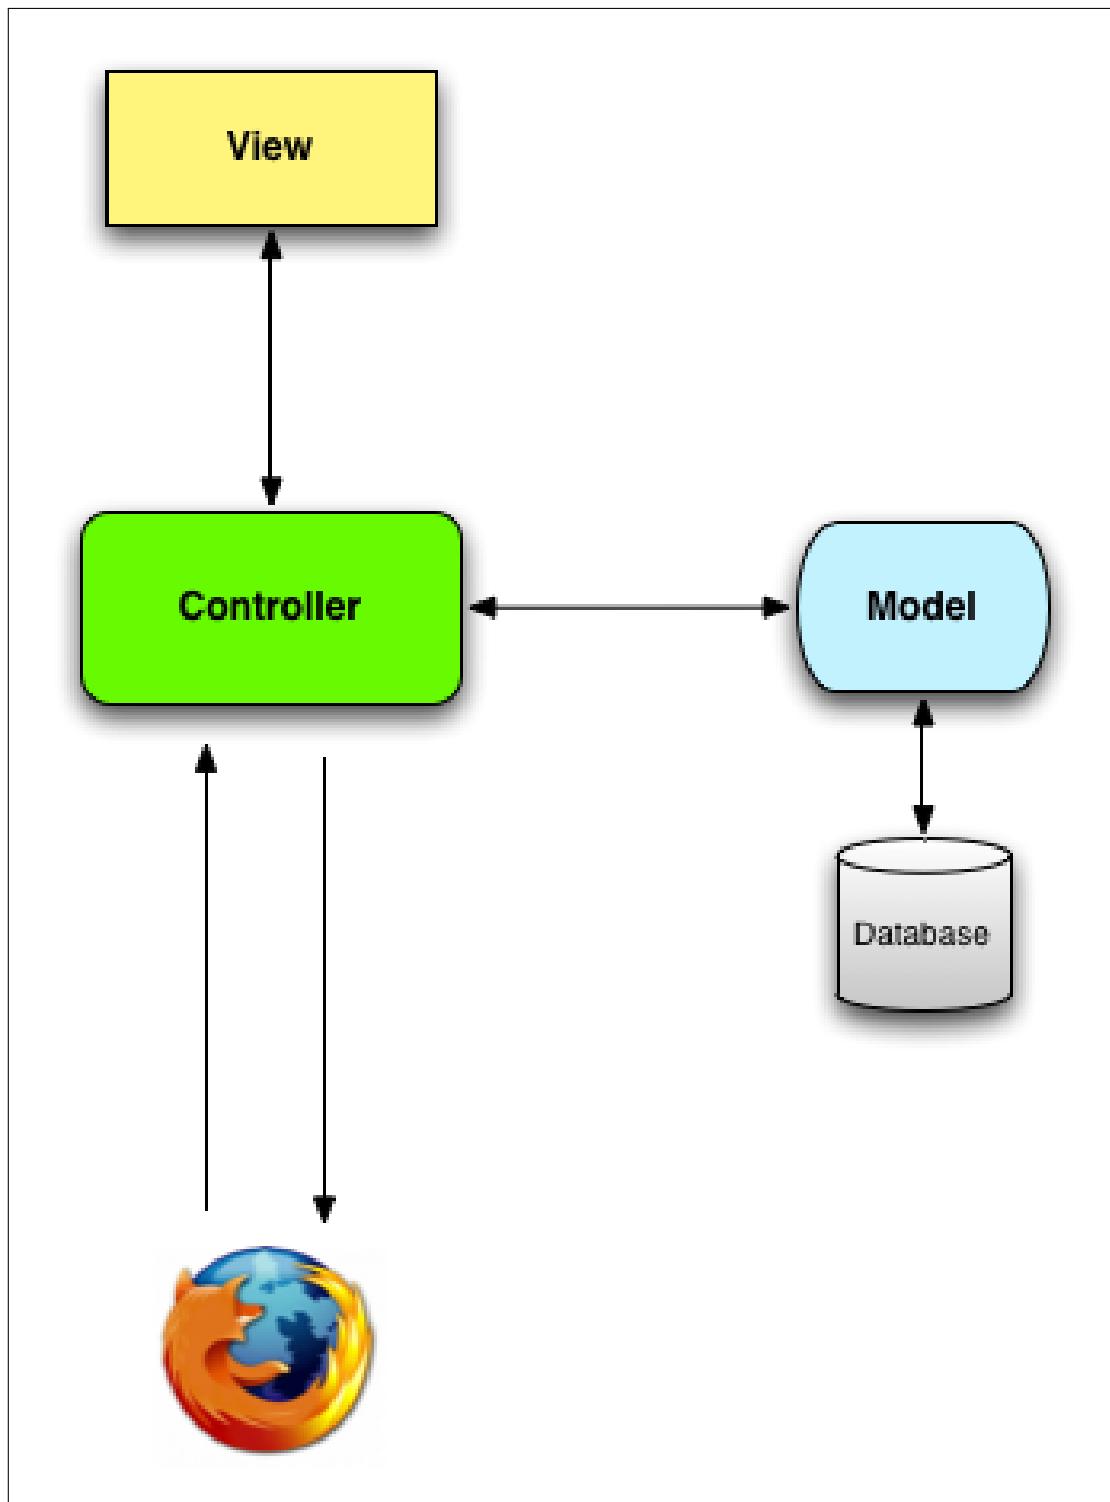


Figure 1.4: A schematic representation of the model-view-controller (MVC) architecture. ([full size](#))
[Ruby on Rails Tutorial](#) [Send Feedback](#)

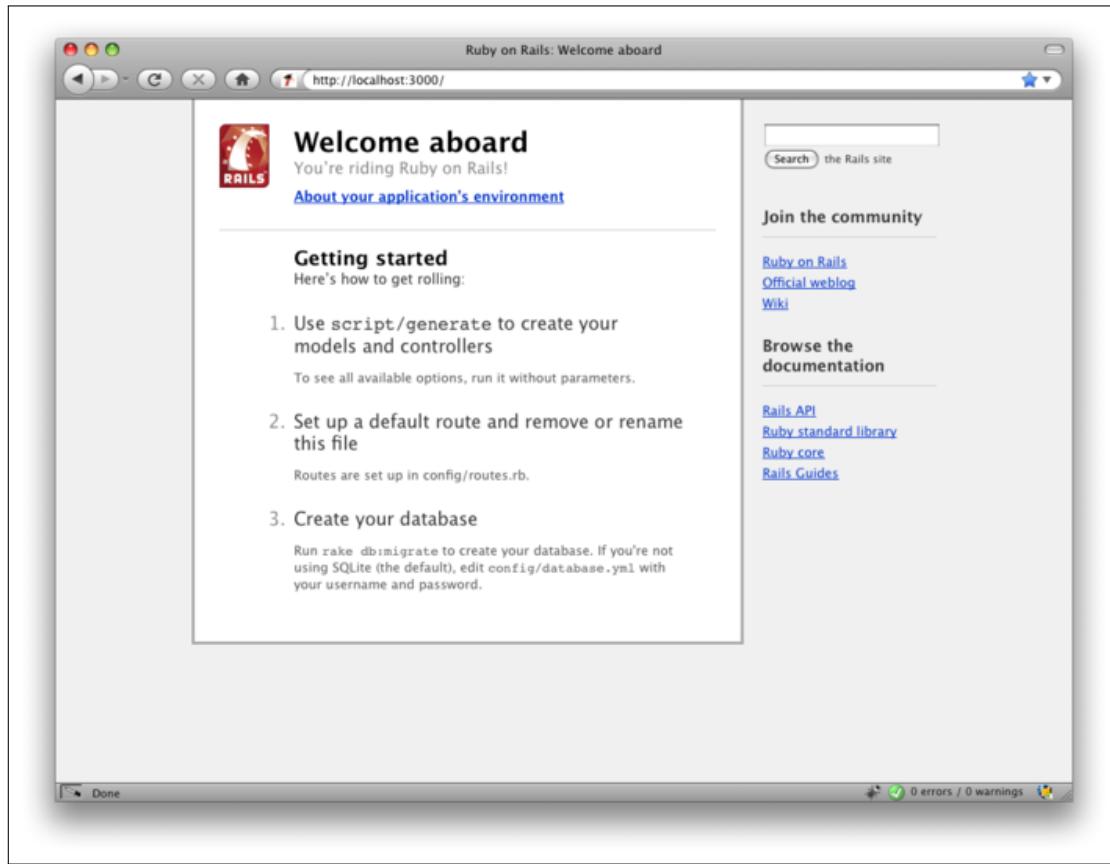


Figure 1.5: The default Rails page (<http://localhost:3000/>). (full size)

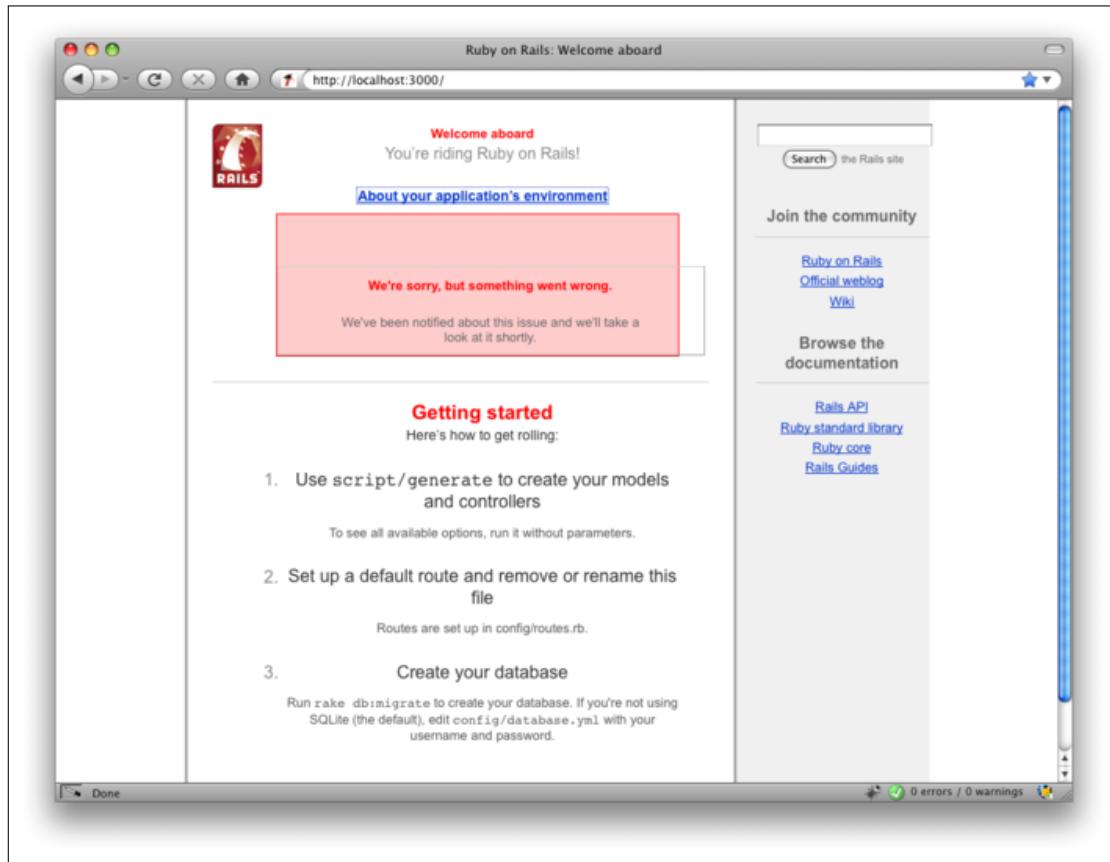
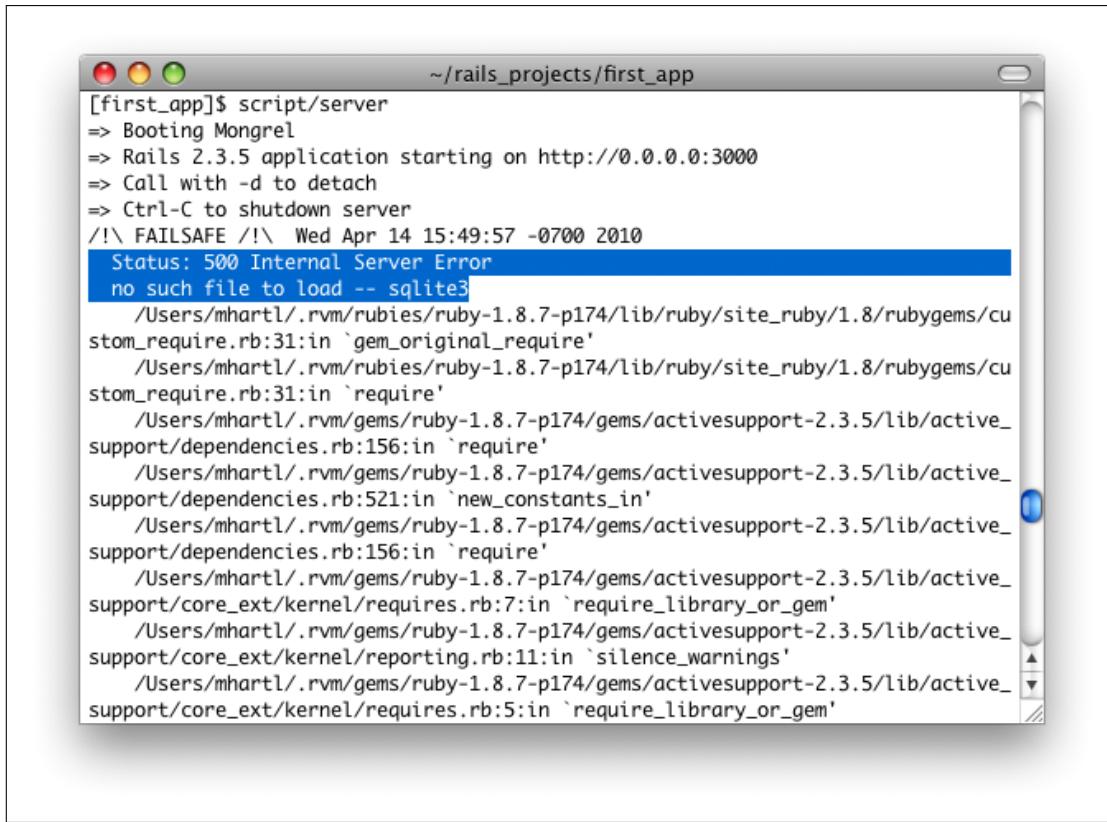


Figure 1.6: The default Rails page (<http://localhost:3000/>) with an error. ([full size](#))



```
[first_app]$ script/server
=> Booting Mongrel
=> Rails 2.3.5 application starting on http://0.0.0.0:3000
=> Call with -d to detach
=> Ctrl-C to shutdown server
/!\ FAILSAFE /!\  Wed Apr 14 15:49:57 -0700 2010
  Status: 500 Internal Server Error
    no such file or load -- sqlite3
      /Users/mhartl/.rvm/rubies/ruby-1.8.7-p174/lib/ruby/site_ruby/1.8/rubygems/cu
      stom_require.rb:31:in `gem_original_require'
        /Users/mhartl/.rvm/rubies/ruby-1.8.7-p174/lib/ruby/site_ruby/1.8/rubygems/cu
      stom_require.rb:31:in `require'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/dependencies.rb:156:in `require'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/dependencies.rb:521:in `new_constants_in'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/dependencies.rb:156:in `require'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/core_ext/kernel/requirements.rb:7:in `require_library_or_gem'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/core_ext/kernel/reporting.rb:11:in `silence_warnings'
        /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/activesupport-2.3.5/lib/active_
      support/core_ext/kernel/requirements.rb:5:in `require_library_or_gem'
```

Figure 1.7: The incriminating lines in the server log. ([full size](#))

This is a hint that we're missing SQLite, the default database used by Rails. Even though we won't need a database until [Chapter 6](#), since Rails is *opinionated software* it still expects a database to be installed, so let's take a few seconds to set it up. First, we need the SQLite gem for Ruby:

```
$ [sudo] gem install sqlite3-ruby -v 1.2.5
Building native extensions. This could take a while...
.
.
```

If you're running Linux you might get an error like this:

```
Building native extensions. This could take a while...
ERROR: Error installing sqlite3-ruby:
ERROR: Failed to build gem native extension.
```

In this case, install `sqlite3` and `libsqllite3-dev` first:

```
$ [sudo] apt-get install sqlite3
$ [sudo] apt-get install libsqlite3-dev
$ [sudo] gem install sqlite3-ruby
```

Once you've installed SQLite, all that's left is to set up the (empty) database using *Rake* ([Box 1.2](#)):

```
$ rake db:migrate
```

This does nothing but create an empty SQLite database for now,¹⁹ but we'll see in [Section 2.2.1](#) (and in more detail starting in [Chapter 6](#)) how to use `rake db:migrate` to create a database for an application's data model.

Now hit Ctrl-C to quit the server, and then restart it with `script/server`. You should see the result in [Figure 1.8](#). Of course, we don't need the default Rails page in the long run, but it's nice to see it working for now. We'll remove the default Rails page (and replace it with a custom home page) in [Section 5.2.2](#).

How would you know to install SQLite if I weren't here to tell you? The best bet is usually to drop the error message into Google (maybe with the word *rails* thrown in) and keep searching until you find an answer. If that doesn't work, emailing the [Ruby on Rails mailing list](#) is probably worth a try.

At this point, in principle you're ready to start developing your Rails application, so I suppose you could skip ahead to [Chapter 3](#). On the other hand, it would be irresponsible of me not to *very strongly recommend* that you follow the rest of the chapter—especially the next section on version control. A couple years ago maybe you could get away with writing a Rails tutorial book without integrating version control from the start, but not any more.

¹⁹Technically, this first invocation of `rake db:migrate` does add one table, `schema_migrations`, which Rails uses internally to track changes to the data model. But don't worry about that right now.

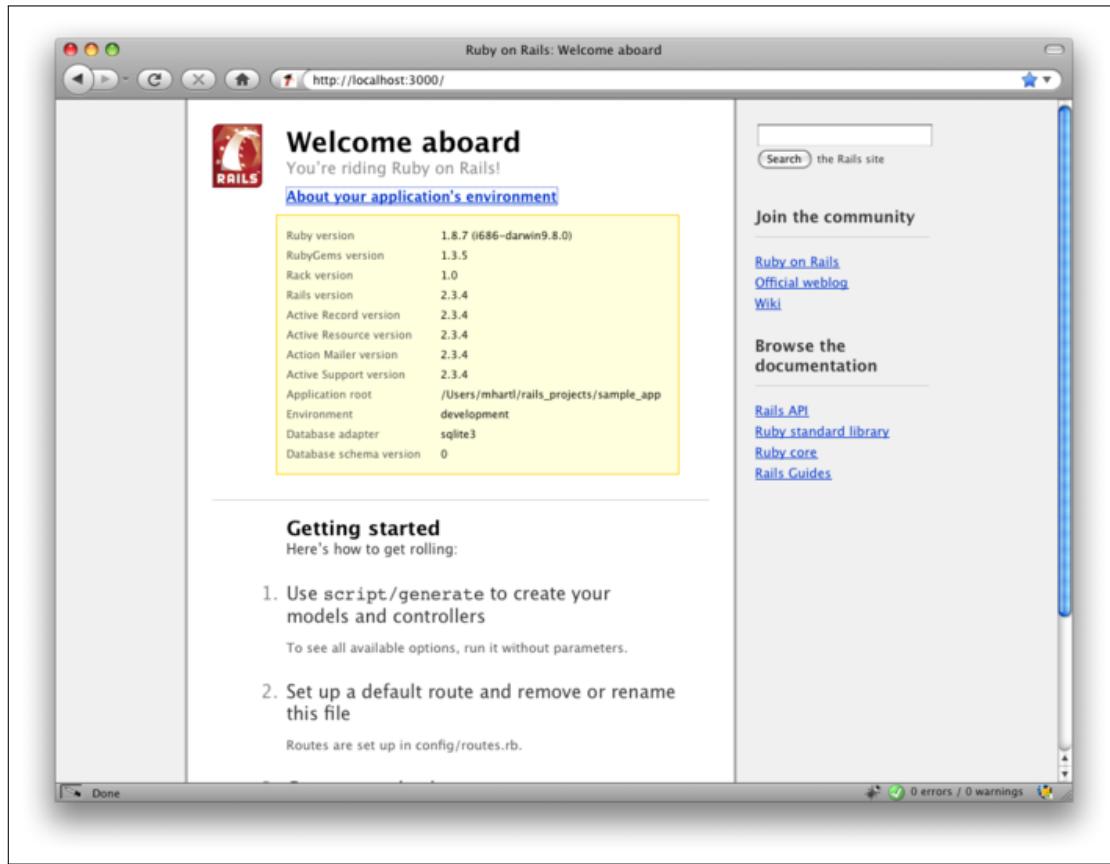


Figure 1.8: The default page <http://localhost:3000/> with the SQLite gem installed. ([full size](#))

Box 1.2. Rake

In the Unix tradition, the `make` utility has played an important role in building executable programs from source code; many a computer hacker has committed to muscle memory the line

```
$ ./configure && make && sudo make install
```

commonly used to compile code on Unix systems (including Linux and Mac OS X).

Rake is *Ruby make*, a make-like language written in Ruby. Rails uses Rake extensively, especially for the innumerable little administrative tasks necessary when developing database-backed web applications. The `rake db:migrate` command is probably the most common, but there are many others; you can see a list of database tasks using `-T db`:

```
$ rake -T db
```

To see all the Rake tasks available, run

```
$ rake -T
```

The list is likely to be overwhelming, but don't worry, you don't have to know all (or even most) of these commands. By the end of *Rails Tutorial*, you'll know all the most important ones.

1.3 Version control with Git

Now that we have a fresh and working Rails application, we'll take a moment for a step that, while technically optional, would be viewed by many Rails developers as practically essential, namely, placing our application source code under *version control*.

Version control systems allow us to track changes to our project's code, collaborate more easily, and roll back any inadvertent errors (such as accidentally deleting files). There are many options for version control, but the Rails community has largely standardized on Git, a version control system originally developed by Linus Torvalds to host the Linux kernel. Git is a large subject, and we'll only be scratching the surface in this book, but there are many good free resources online; I especially recommend *Pro Git* by Scott Chacon (Apress, 2009).

Putting your source code under version control with Git is *strongly* recommended, not only because it's nearly a universal practice in the Rails world, but also because it will allow you to share your code more easily ([Section 1.3.4](#)) and deploy your application right here in the first chapter ([Section 1.4](#)).

1.3.1 Installation and setup

The first step is to install Git if you don't have it already. You should follow the instructions for your platform at the [Installing Git section of Pro Git](#).

First-time system setup

After installing Git, you should perform a set of one-time setup steps. These are *system* setups, meaning you only have to do them once per computer:

```
$ git config --global user.name "Your Name"  
$ git config --global user.email youremail@example.com
```

You can optionally set the editor Git will use for commit messages, and there's a subtlety if you use a graphical editor such as TextMate, gVim, or MacVim; you need to use a flag to make sure that the editor stays attached to the shell instead of detaching immediately.²⁰

```
$ git config --global core.editor "mate -w"
```

Replace "**mate -w**" with "**gvim -f**" for gVim or "**mvim -f**" for MacVim.

First-time repository setup

Now we come to some steps that are necessary each time you create a new *repository* (which only happens once in this book, but is likely to happen again some day). First navigate to the root directory of the first app and initialize a new repository:

```
$ git init  
Initialized empty Git repository in /Users/mhartl/rails_projects/first_app/.git/
```

The next step is to add the project files to the repository. There's a minor complication, though: by default Git tracks the changes of *all* the files, but there are some files we don't want to track. For example, Rails creates log files to record the behavior of the application; these files change frequently, and we don't want our version control system to have to update them constantly. Git has a simple mechanism to ignore such files: simply include a file called **.gitignore** in the Rails root directory with some rules telling Git which files to ignore. To ignore all log files, which live in the **log/** directory, we use the asterisk wildcard and ignore all files that end in **.log** by including **log/*.log** in **.gitignore** (and so on for the other kinds of files we want Git to ignore). So, create a new **.gitignore** file using your favorite editor:²¹

```
$ mate .gitignore
```

Then fill the file with the contents of Listing 1.2.

Listing 1.2. A good starting **.gitignore** file for a Rails project.

²⁰Normally this is a feature, since it lets you continue to use the command line after launching your editor, but Git interprets the detachment as closing the file with an empty commit message, which prevents the commit from going through. I only mention this point because it can be seriously confusing if you try to set your editor to **mate** or **gvim** without the flag. If you find this note confusing, feel free to ignore it.

²¹Here I'm using **mate** to launch TextMate; of course, if you're not using TextMate you should substitute the command for your editor (e.g., **emacs**, **vim**, **gvim**, or **mvim**).

```
log/*.log
tmp/*
tmp/**/*
doc/api
doc/app
db/*.sqlite3
*.swp
*~
.DS_Store
```

The lines in Listing 1.2 tell Git to ignore log files, Rails temp files, documentation files, SQLite databases, Vim and Emacs swap files, and (for you OS X users) the weird `.DS_Store` directories created by the Mac Finder application. Don't worry if you wouldn't have guessed these lines; Understanding just why each of these file types should be ignored comes from having practical experience with both Rails and Git. By using this `.gitignore` from the start, you can benefit immediately from the collective Git experience of the Rails community.

1.3.2 Adding and committing

Finally, we'll add the files in your new Rails project to Git and then commit the results. You can add all the files (apart from those that match the ignore patterns in `.gitignore`) as follows:²²

```
$ git add .
```

Here the dot ‘`.`’ represents the current directory, and Git is smart enough to add the files *recursively*, so it automatically includes all the subdirectories. This command adds the project files to a *staging area*, which contains pending changes to your project; you can see which files are in the staging area using the `status` command:²³

```
$ git status
# On branch master
#
# Initial commit
#
# Changes to be committed:
#   (use "git rm --cached <file>..." to unstage)
#
#       new file:   README
#       new file:   Rakefile
#
```

²²Windows users may get the message `warning: CRLF will be replaced by LF in .gitignore`. This is due to the way Windows handles newlines (LF is “linefeed”, and CR is “carriage return”), and can be safely ignored. If the message bothers you, try running `git config core.autocrlf false` at the command line to turn it off.

²³If in the future any unwanted files start showing up when you type `git status`, just add them to your `.gitignore` file from Listing 1.2.

```
.
```

(The results are long, so I've used vertical dots to indicate omitted output.)

To tell Git you want to keep the changes, use the `commit` command:

```
$ git commit -m "Initial commit"
[master (root-commit) df0a62f] Initial commit
42 files changed, 8461 insertions(+), 0 deletions(-)
create mode 100644 README
create mode 100644 Rakefile
.
.
```

The `-m` flag lets you add a message for the commit; if you omit `-m`, Git will open the editor you set in [Section 1.3.1](#) and have you enter the message there.

It is important to note that Git commits are *local*, recorded only on the machine on which the commits occur. This is in contrast to the popular open-source version control system called Subversion, in which a commit necessarily makes changes on a remote repository. Git divides a Subversion-style commit into its two logical pieces: a local recording of the changes (`git commit`) and a push of the changes up to a remote repository (`git push`). We'll see an example of the push step in [Section 1.3.5](#).

By the way, you can see a list of your commit messages using the `log` command:

```
$ git log
commit df0a62f3f091e53ffa799309b3e32c27b0b38eb4
Author: Michael Hartl <michael@michaelhartl.com>
Date:   Thu Oct 15 11:36:21 2009 -0700

    Initial commit
```

To exit `git log`, you may have to type `q` to quit.

1.3.3 What good does Git do you?

It's probably not entirely clear at this point why putting your source under version control does you any good, so let me give just one example. (We'll see many others in the chapters ahead.) Suppose you've made some accidental changes, such as (D'oh!) deleting the critical `app/controllers/` directory:

```
$ ls app/controllers/
application_controller.rb
$ rm -rf app/controllers/
$ ls app/controllers/
ls: app/controllers/: No such file or directory
```

Here we're using the Unix `ls` command to list the contents of the `app/controllers/` directory and the `rm` command to remove it. The `-rf` flag means “recursive force”, which recursively removes all files, directories, subdirectories, and so on, without asking for explicit confirmation of each deletion.

Let's check the status to see what's up:

```
$ git status
# On branch master
# Changed but not updated:
#   (use "git add/rm <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
#       deleted:    app/controllers/application_controller.rb
#
no changes added to commit (use "git add" and/or "git commit -a")
```

We see here that a couple files have been deleted, but the changes have only been *staged*; they haven't been committed yet. This means we can still undo the changes easily by having Git check out the previous commit with the `checkout` command (and a `-f` flag to force overwriting the current changes):

```
$ git checkout -f
$ git status
# On branch master
nothing to commit (working directory clean)
$ ls app/controllers/
application_controller.rb
```

The missing directory and file are back. That's a relief!

1.3.4 GitHub

Now that you've put your project under version control with Git, it's time to push your code up to [GitHub](#), a social code site optimized for hosting and sharing Git repositories. Putting a copy of your Git repository at GitHub serves two purposes: it's a full backup of your code (including the full history of commits), and it makes any future collaboration much easier. This step is optional, but being a GitHub member will open the door to participating in a wide variety of Ruby and Rails projects (GitHub has high adoption rates in the Ruby and Rails communities, and in fact is itself written in Rails).

GitHub has a variety of paid plans, but for open source code their services are free, so sign up for a [free GitHub account](#) if you don't have one already. (You might have to read about [SSH keys](#) first.) After signing up, you'll see a page like the one in [Figure 1.9](#). Click on [create a repository](#) and fill in the information as in [Figure 1.10](#). After submitting the form, push up your first application as follows:

```
$ git remote add origin git@github.com:<username>/first_app.git
$ git push origin master
```

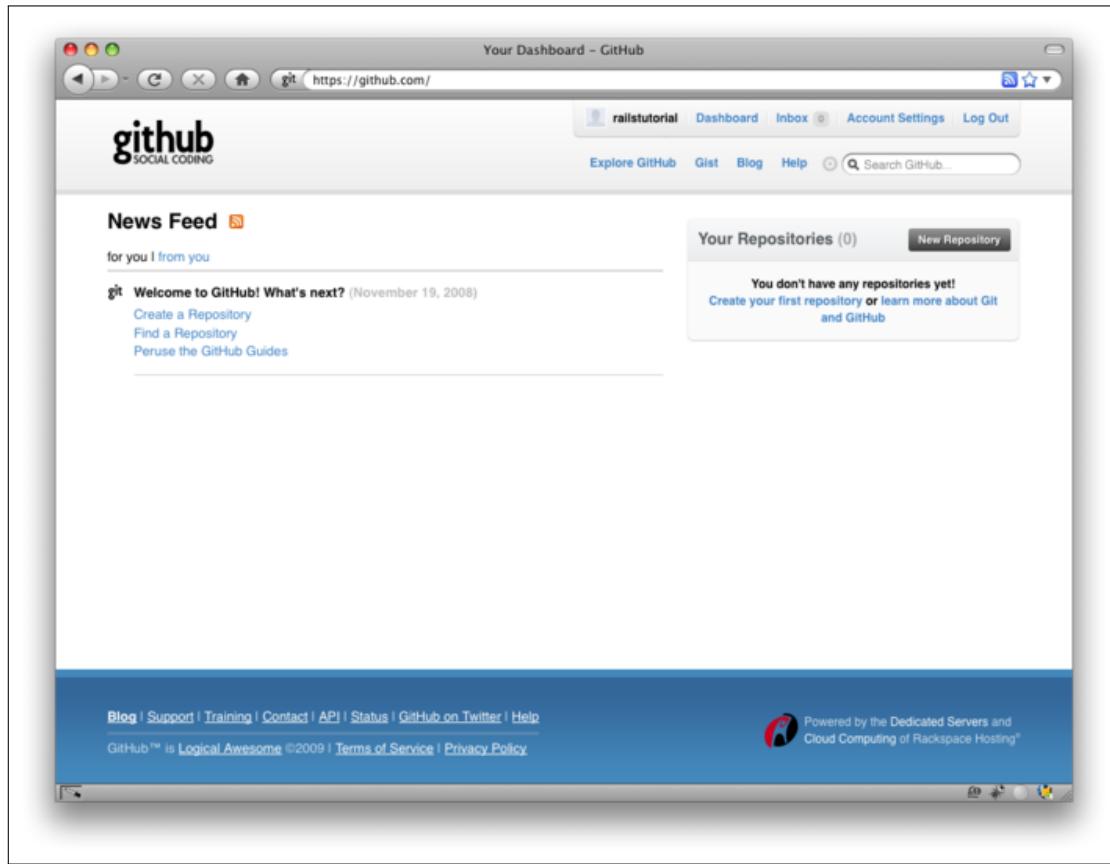


Figure 1.9: The first GitHub page after account creation. ([full size](#))

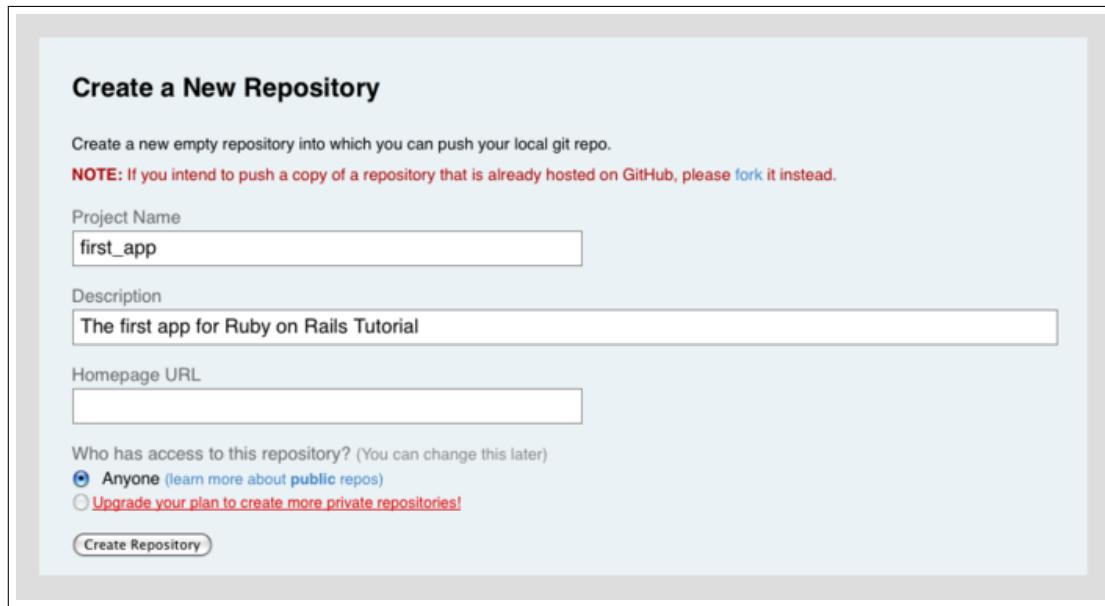


Figure 1.10: Creating the first app repository at GitHub. ([full size](#))

These commands tell Git that you want to add GitHub as the origin for your main (*master*) branch and then push your repository up to GitHub. Of course, you should replace <username> with your actual username. For example, the command I ran for the `railstutorial` user was

```
$ git remote add origin git@github.com:railstutorial/first_app.git
```

The result is a page at GitHub for the first application repository, with file browsing, full commit history, and lots of other goodies (Figure 1.11).

1.3.5 Branch, edit, commit, merge

If you've followed the steps in Section 1.3.4, you might notice that GitHub automatically shows the contents of the `README` file on the main repository page. In our case, since the project is a Rails application generated using the `rails` command, the `README` file is the one that comes with Rails (Figure 1.12). This isn't very helpful, so in this section we'll make our first edit by changing the `README` to describe our project rather than the Rails framework itself. In the process, we'll see a first example of the branch, edit, commit, merge workflow that I recommend using with Git.

Branch

Git is incredibly good at making *branches*, which are essentially copies of a repository where we can make (possibly experimental) changes without modifying the parent files. In most cases, the parent repository is the *master* branch, and we can create a new topic branch by using `checkout` with the `-b` flag:

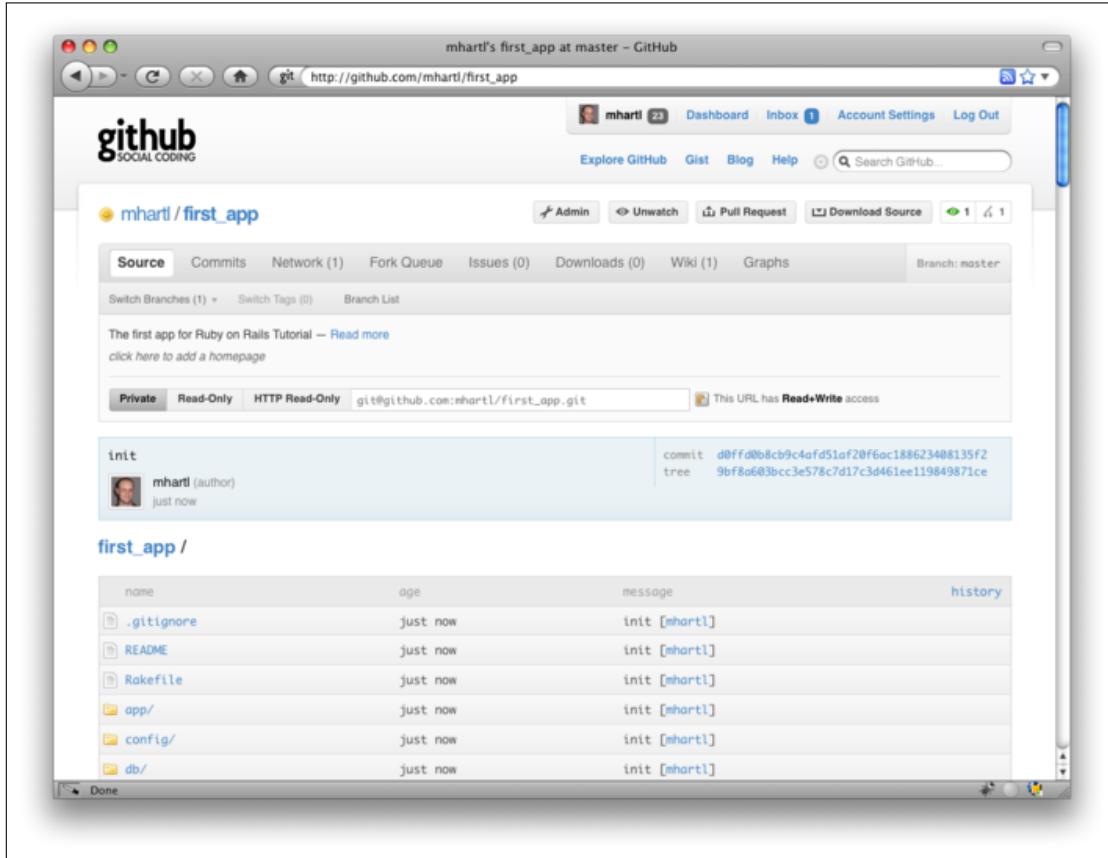


Figure 1.11: A GitHub repository page. ([full size](#))

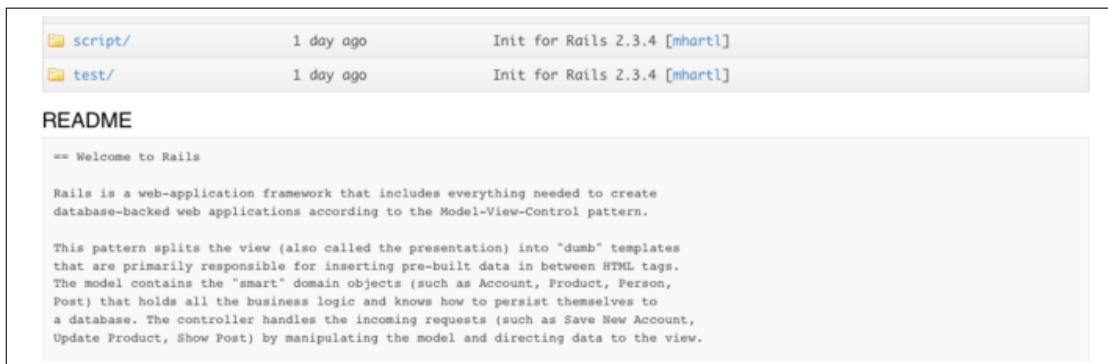


Figure 1.12: The initial (rather useless) **README** file for our project at GitHub. ([full size](#))

```
$ git checkout -b modify-README
Switched to a new branch 'modify-README'
$ git branch
master
* modify-README
```

Here the second command, `git branch`, just lists all the local branches, and the asterisk `*` identifies which branch we're currently on. Note that `git checkout -b modify-README` both creates a new branch and switches to it, as indicated by the asterisk in front of the `modify-README` branch.

The full value of branching only becomes clear when working on a project with multiple developers,²⁴ but branches are helpful even for a single-developer tutorial such as this one. In particular, the master branch is insulated from any changes we make to the topic branch, so even if we *really* screw things up we can always abandon the changes by checking out the master branch and deleting the topic branch. We'll see how to do this at the end of the section.

By the way, for a change as small as this one I wouldn't normally bother with a new branch, but it's never too early to start practicing good habits.

Edit

After creating the topic branch, we'll edit it to make it a little more descriptive. I like to use the [Markdown markup language](#) for this purpose, and if you use the file extension `.markdown` then GitHub will automatically format it nicely for you. So, first we'll use Git's version of the Unix `mv` ("move") command to change the name, and then fill it in with the contents of Listing 1.3:

```
$ git mv README README.markdown
$ mate README.markdown
```

Listing 1.3. The new `README` file, `README.markdown`.

```
# Ruby on Rails Tutorial: first application

This is the first application for
[*Ruby on Rails Tutorial: Learn Rails by Example*] (http://www.railstutorial.org/)
by [Michael Hartl] (http://www.michaelhartl.com/).
```

Commit

With the changes made, we can take a look at the status of our branch:

```
$ git status
# On branch modify-README
```

²⁴See the chapter [Git Branching in Pro Git](#) for details.

```
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
#       renamed:    README -> README.markdown
#
# Changed but not updated:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
#       modified:   README.markdown
#
```

At this point, we could use `git add .` as in [Section 1.3.2](#), but Git provides the `-a` flag as a shortcut for the (very common) case of committing all modifications to existing files (or files created using `git mv`, which don't count as new files to Git):

```
$ git commit -a -m "Improved the README file"
2 files changed, 5 insertions(+), 243 deletions(-)
delete mode 100644 README
create mode 100644 README.markdown
```

Be careful about using the `-a` flag improperly; if you have added any new files to the project since the last commit, you still have tell Git about them using `git add` first.

Merge

Now that we've finished making our changes, we're ready to *merge* the results back into our master branch:²⁵

```
$ git checkout master
Switched to branch 'master'
$ git merge modify-README
Updating 34f06b7..2c92bef
Fast forward
 README      | 243 -----
 README.markdown | 5 +
 2 files changed, 5 insertions(+), 243 deletions(-)
 delete mode 100644 README
 create mode 100644 README.markdown
```

Note that the Git output frequently includes things like `34f06b7`, which are related to Git's internal representation of repositories. Your exact results will differ in these details, but otherwise should essentially match the output shown above.

After you've merged in the changes, you can tidy up your branches by deleting the topic branch using `git branch -d` if you're done with it:

²⁵Experienced Git users will recognize the wisdom of running `git rebase master` before switching to the master branch, but this step will not be necessary in this book.



Figure 1.13: The improved **README** file formatted with Markdown. ([full size](#))

```
$ git branch -d modify-README  
Deleted branch modify-README (was 2c92bef) .
```

This step is optional, and in fact it's quite common leave the topic branch intact. This way you can switch back and forth between the topic and master branches, merging in changes every time you reach a natural stopping point.

As mentioned above, it's also possible to abandon your topic branch changes, in this case with **git branch -D**:

```
# For illustration only; don't do this unless you mess up a branch  
$ git checkout -b topic-branch  
$ <really screw up the branch>  
$ git add .  
$ git commit -a -m "Screwed up"  
$ git checkout master  
$ git branch -D topic-branch
```

Unlike the **-d** flag, the **-D** flag will delete the branch even though we haven't merged in the changes.

Push

Now that we've updated the **README**, we can push the changes up to GitHub to see the result:²⁶

```
$ git push
```

As promised, GitHub nicely formats the new file using Markdown (Figure 1.13).

1.4 Deploying

Even at this early stage, we're already going to deploy our (still-empty) Rails application to production. This step is optional, but deploying early and often allows us to catch any deployment problems early in

²⁶When collaborating on a project with other developers, you should run **git pull** before this step to pull in any remote changes.

our development cycle. The alternative—deploying only after laborious effort sealed away in a development environment—often leads to terrible integration headaches when launch time comes.²⁷

Deploying Rails applications used to be a pain, but the Rails deployment ecosystem has matured rapidly in the past few years, and now there are several great options. These include shared hosts or virtual private servers running [Phusion Passenger](#) (a module for the Apache and Nginx²⁸ web servers), full-service deployment companies such as [Engine Yard](#) and [Rails Machine](#), and cloud deployment services such as [Engine Yard Cloud](#) and [Heroku](#).

My favorite Rails deployment option is Heroku, which is a hosted platform built specifically for deploying Rails and other Ruby web applications.²⁹ Heroku makes deploying Rails applications ridiculously easy—as long as your source code is under version control with Git. (This is yet another reason to follow the Git setup steps in [Section 1.3](#) if you haven’t already.) The rest of this section is dedicated to deploying our first application to Heroku.

1.4.1 Heroku setup

After [signing up for a Heroku account](#), install the Heroku gem:

```
$ [sudo] gem install heroku
```

As with GitHub ([Section 1.3.4](#)), when using Heroku you will need to [create SSH keys](#) if you haven’t already, and then tell Heroku your [public key](#) so that you can use Git to push the sample application repository up to their servers:

```
$ heroku keys:add
```

Finally, use the **heroku** gem to create a place on the Heroku servers for the sample app to live ([Listing 1.4](#)).

Listing 1.4. Creating a new application at Heroku.

```
$ heroku create
Created http://severe-fire-61.herokuapp.com/ | git@heroku.com:severe-fire-61.git
Git remote heroku added
```

Yes, that’s it. The **heroku** gem creates a new subdomain just for our application, available for immediate viewing. There’s nothing there yet, though, so let’s get busy deploying.

²⁷Though it shouldn’t matter for the example applications in *Rails Tutorial*, if you’re worried about accidentally making your app public too soon there are several options; see [Section 1.4.4](#) for one.

²⁸Pronounced “Engine X”.

²⁹Heroku works with any Ruby web platform that uses [Rack middleware](#), which provides a standard interface between web frameworks and web servers. Adoption of the Rack interface has been extraordinarily strong in the Ruby community, including frameworks as varied as [Sinatra](#), [Ramaze](#), [Camping](#), and Rails, which means that Heroku basically supports any Ruby web app.

1.4.2 Heroku deployment, step one

To deploy to Heroku, the first step is to use Git to push the application to Heroku:

```
$ git push heroku master
```

Note: I've had some reports of error pages when using Rails 2.3.8 on Heroku. I'm hoping it gets sorted out soon, but for now you can fix the issue as follows. First, create a new file called `.gems` in the `RAILS_ROOT` directory with the following contents:

```
rails --version 2.3.8
```

Then commit and push:

```
$ git add .gems
$ git commit -a -m "Added a .gems file to use Rails 2.3.8"
$ git push heroku
```

Now everything should work (see step two below).

1.4.3 Heroku deployment, step two

There is no step two! We're already done ([Figure 1.14](#)). To see your newly deployed application, you can visit the address that you saw when you ran `heroku create` (i.e., [Listing 1.4](#), but with the address for your app, not the address for mine).³⁰ You can also use a command provided by the `heroku` gem that automatically opens your browser with the right address:

```
$ heroku open
```

Once you've deployed successfully, Heroku provides a beautiful interface for administering and configuring your application ([Figure 1.15](#)).

1.4.4 Heroku commands

There are tons of [Heroku commands](#), and we'll barely scratch the surface in this book. Let's take a minute to show just one of them by renaming the application as follows:

```
$ heroku rename railstutorial
```

³⁰Because of the details of their setup, the “About your application’s environment” link doesn’t work on Heroku; instead, you get an error much like the one in [Figure 1.6](#). Don’t worry; this is normal. The error will go away when we remove the default Rails page in [Section 5.2.2](#).

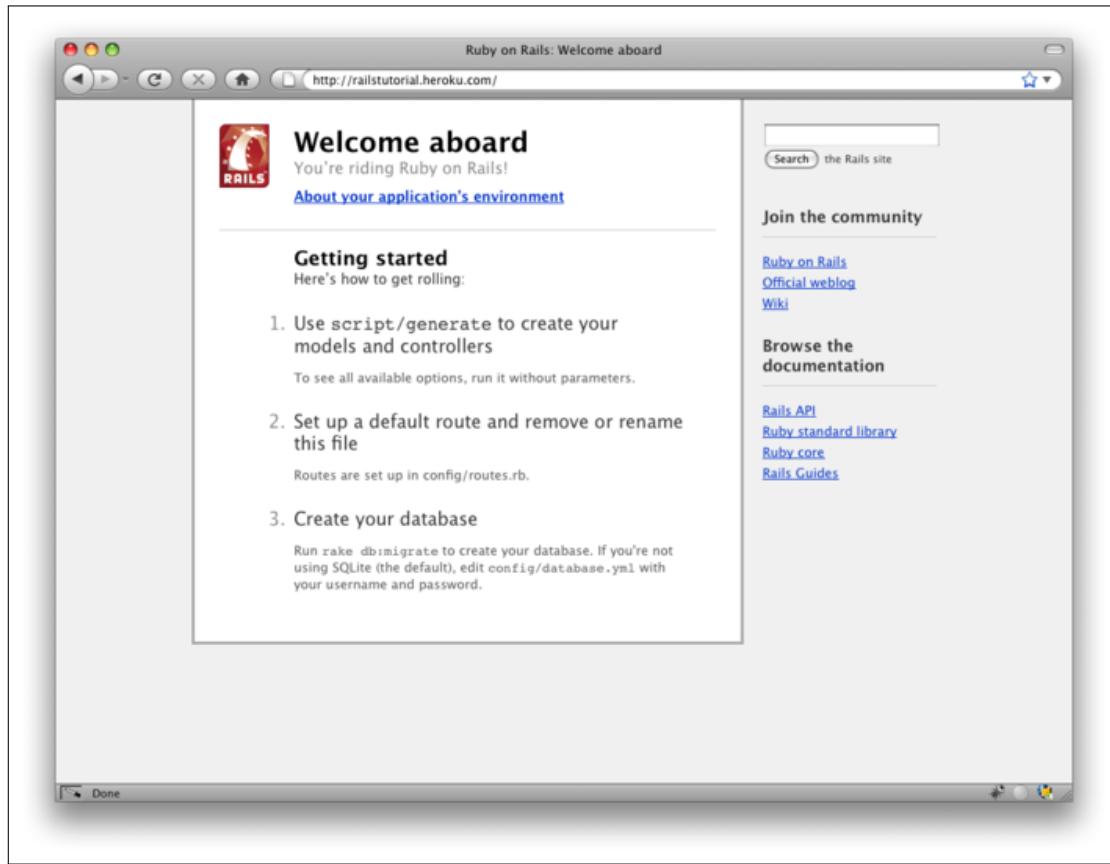


Figure 1.14: The first Rails Tutorial application running on Heroku. (full size)

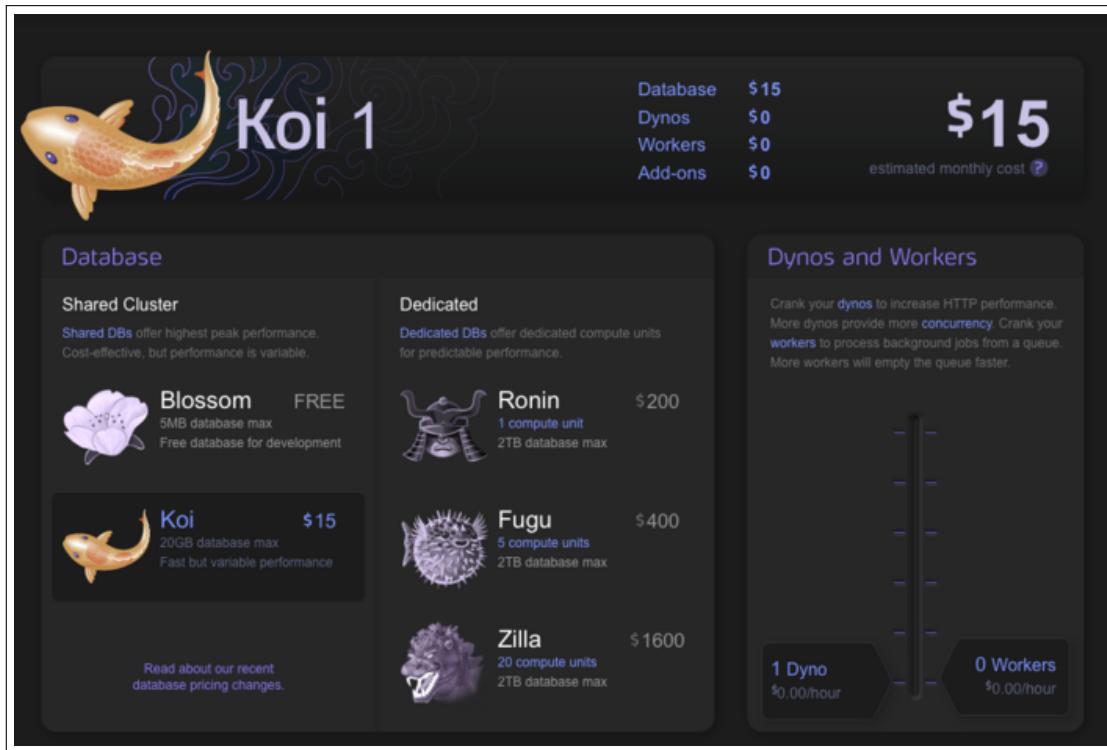


Figure 1.15: The beautiful interface at Heroku. (full size)

Don't use this name yourself; it's already taken by me! In fact, you probably shouldn't bother with this step right now; using the default address supplied by Heroku is fine. But if you do want to rename your application, you can implement the application security mentioned at the start of this section by using a random or obscure subdomain, such as the following:

```
hwpcbmze.herokuapp.com  
seyjhflo.herokuapp.com  
jhyicevg.herokuapp.com
```

With a random subdomain like this, someone could visit your site only if you gave them the address. (By the way, as a preview of Ruby's compact awesomeness, here's the code I used to generate the random subdomains:

```
('a'..'z').to_a.shuffle[0..7].join
```

Pretty sweet.)

In addition to supporting subdomains, Heroku also supports custom domains. (In fact, the [Ruby on Rails Tutorial site](#) lives at Heroku; if you're reading this book online, you're looking at a Heroku-hosted site right now!) See the [Heroku documentation](#) for more information about custom domains and other Heroku topics.

1.5 Conclusion

We've come a long way in this chapter: installation, development environment setup, version control, and deployment. If you want to share your progress at this point, feel free to send a tweet or Facebook status update with something like this:

I'm learning Ruby on Rails with @railstutorial! <http://bit.ly/railstutorial>

All that's left is to, you know, actually start learning Rails. Let's get to it!

Chapter 2

A demo app

In this chapter, we'll develop a simple demonstration application to show off some of the power of Rails. The purpose is to get a high-level overview of Ruby on Rails programming (and web development in general) by rapidly generating an application using *scaffold generators*.¹ As discussed in [Box 1.1](#), the rest of the book will take the opposite approach, developing a full application incrementally and explaining each new concepts as it arises, but for a quick overview (and some instant gratification) there is no substitute for scaffolding. The resulting demo app will allow us to interact with it through its URLs, giving us insight into the structure of a Rails application, including a first example of the *REST architecture* favored by Rails.

As with the forthcoming sample application, the demo app will consist of *users* and their associated *microposts* (thus constituting a minimalist Twitter-style app). The functionality will be utterly under-developed, and many of the steps will seem like magic, but worry not: the full sample app will develop a similar application from the ground up starting in [Chapter 3](#), and I will provide plentiful forward-references to later material. In the mean time, have patience and a little faith—the whole point of this tutorial is to take you *beyond* this superficial, scaffold-driven approach to achieve a deeper understanding of Rails.

2.1 Planning the application

In this section we'll outline our plans for the demo application. As in [Section 1.2.3](#), we'll start by generating the application skeleton using the `rails` command:

```
$ cd ~/rails_projects  
$ rails demo_app  
$ cd demo_app
```

Next, we'll copy the `.gitignore` file from [Listing 1.2](#) to the demo app's Rails root directory:

```
$ git init  
$ cp ~/rails_projects/first_app/.gitignore .  
$ git add .  
$ git commit -m "Initial commit"
```

¹I urge you not to look too closely at the generated code; at this stage, it will only serve to confuse you.

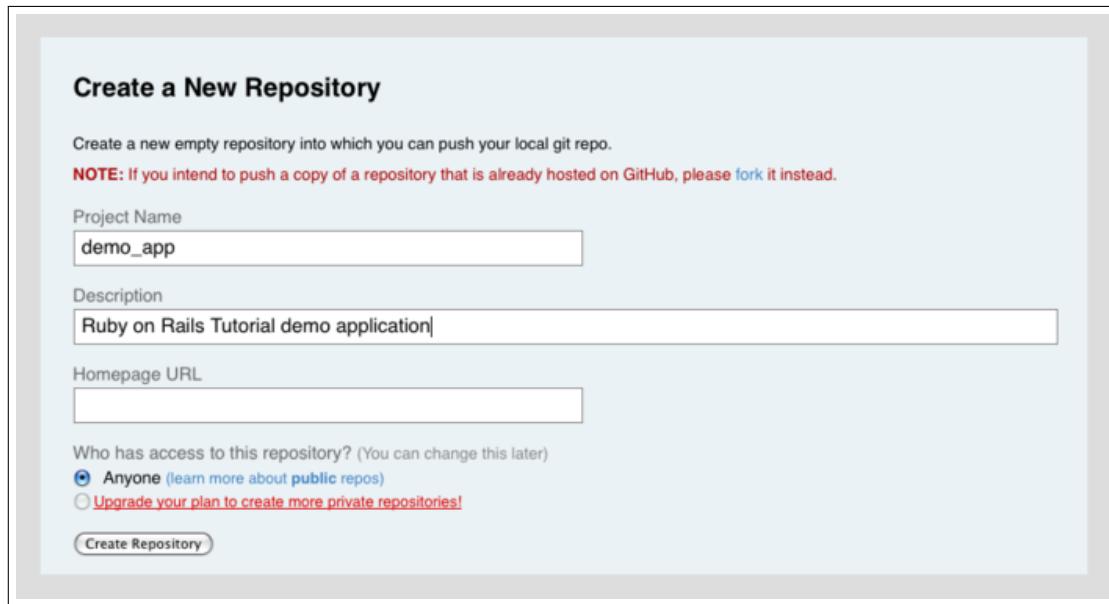


Figure 2.1: Creating a demo app repository at GitHub. ([full size](#))

You can also optionally create a new repository (Figure 2.1) and push it up to GitHub:

```
$ git remote add origin git@github.com:<username>/demo_app.git
$ git push origin master
```

Now we're ready to start making the app itself. The typical first step when making a web application is to create a *data model*, which is a representation of the structures needed by our application. In our case, the demo app will be a stripped-down microblog, with only users and short (micro) posts. Thus, we'll begin with a model for *users* of the app (Section 2.1.1), and then we'll add a model for *microposts* (Section 2.1.2).

2.1.1 Modeling users

There are as many choices for a user data model as there are different registration forms on the web; we'll go with a distinctly minimalist approach. Users of our demo app will have a unique `integer` identifier called `id`, a publicly viewable `name` (of type `string`), and an `email` address (also a `string`) that will double as a username. A summary of the data model for users appears in Figure 2.2.

As we'll see starting in Section 6.1.1, the label `users` in Figure 2.2 corresponds to a *table* in a database, and the `id`, `name`, and `email` attributes are *columns* in that table.

2.1.2 Modeling microposts

The core of the micropost data model is even simpler than the one for users: a micropost has only an `id` and a `content` field for the micropost's text (of type `string`).² There's an additional complication, though:

²When modeling longer posts, such as those for a normal (non-micro) blog, you should use the `text` type in place of `string`.

users	
id	integer
name	string
email	string

Figure 2.2: The data model for users.

microposts	
id	integer
content	string
user_id	integer

Figure 2.3: The data model for microposts.

we want to *associate* each micropost with a particular user; we'll accomplish this by recording the `user_id` of the owner of the post. The results are shown in Figure 2.3.

We'll see in Section 2.3.3 (and more fully in Chapter 11) how this `user_id` attribute allows us to succinctly express the notion that a user potentially has many associated microposts.

2.2 The Users resource

In this section, we'll implement the users data model in Section 2.1.1, along with a web interface to that model. The combination will constitute a *Users resource*, which will allow us to think of users as objects that can be created, read, updated, and deleted through the web via the [HTTP protocol](#).

As promised in the introduction, our Users resource will be created by a scaffold generator program, which comes standard with each Rails project. The argument of the `scaffold` command is the name of the resource (in this case, `Users`), together with optional parameters for the data model's attributes:³

```
$ script/generate scaffold User name:string email:string
```

By including `name:string` and `email:string`, we have arranged for the User model to have the form shown in Figure 2.2. (Note that there is no need to include a parameter for `id`; it is created automatically by Rails.⁴)

To proceed with the demo application, we first need to *migrate* the database (a process mentioned briefly in Section 1.2.5):

³The name of the scaffold follows the convention of *models*, which are singular, rather than resources and controllers, which are plural. Thus, we have `User` instead `Users`.

⁴The user id is needed as the *primary key* in the database.

URL	Action	Purpose
/users	index	page to list all users
/users/1	show	page to show user with id 1
/users/new	new	page to make a new user
/users/1/edit	edit	page to edit user with id 1

Table 2.1: The correspondence between pages and URLs for the Users resource.

```
$ rake db:migrate
```

This simply updates the database with our new `users` data model. We'll learn more about database migrations starting in [Section 6.1.1](#).

With that, we can run the local web server using

```
$ script/server
```

and the demo application should be ready to go at <http://localhost:3000/>.

2.2.1 A user tour

Visiting the root url <http://localhost:3000/> shows the same default Rails page shown in [Figure 1.5](#), but in generating the Users resource scaffolding we have also created a large number of pages for manipulating users. For example, the page for listing all users is at `/users`, and the page for making a new user is at `/users/new`.⁵ The rest of this section is dedicated to taking a whirlwind tour through these user pages. As we proceed, it may help to refer to [Table 2.1](#), which shows the correspondence between pages and URLs.

We start with the page to show all the users in our application, called `index`; as you might expect, initially there are no users at all ([Figure 2.4](#)).

To make a new user, we visit the `new` page, as shown in [Figure 2.5](#). (In [Chapter 8](#), this will become the user signup page.)

We can create a user by entering name and email values in the text fields and then clicking the Create button. The result is the user `show` page, as seen in [Figure 2.6](#). (The green welcome message is accomplished using the `flash`, which we'll learn about in [Section 8.3.3](#).) Note that the URL is `/users/1`; as you might suspect, the number `1` is simply the user's `id` attribute from [Figure 2.2](#). In [Section 7.3](#), this page will become the user's profile.

To change a user's information, we visit the `edit` page ([Figure 2.7](#)). By modifying the user information and clicking the Update button, we arrange to change the information for the user in the demo application ([Figure 2.8](#)). (As we'll see in detail starting in [Chapter 6](#), this user data is stored in a database back-end.) We'll add user edit/update functionality to the sample application in [Section 10.1](#).

Now we'll create a second user by revisiting the `new` page and submitting a second set of user information; the resulting user `index` is shown in [Figure 2.9](#). [Section 10.3](#) will develop the user index into a more polished page for showing all users.

⁵Since the `http://localhost:3000` part of the address is implicit whenever we are developing locally, I'll usually omit it from now on.

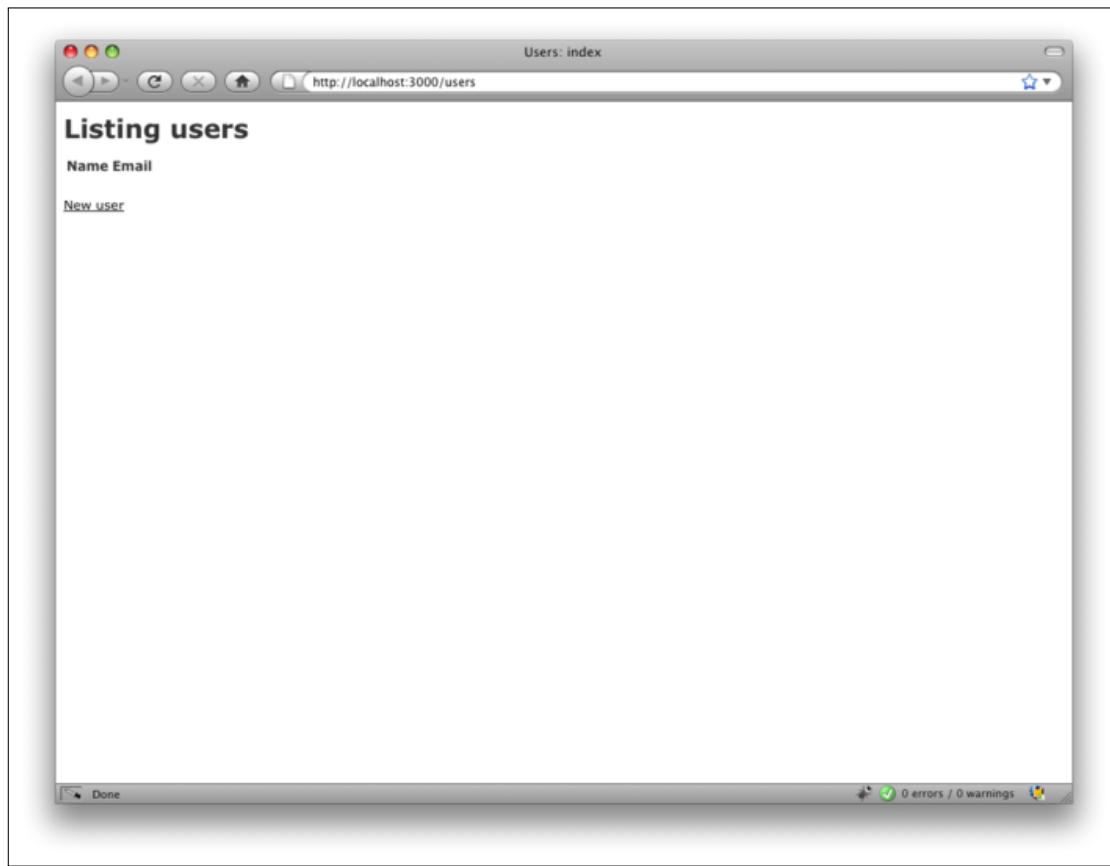


Figure 2.4: The initial index page for the Users resource ([/users](#)). (full size)

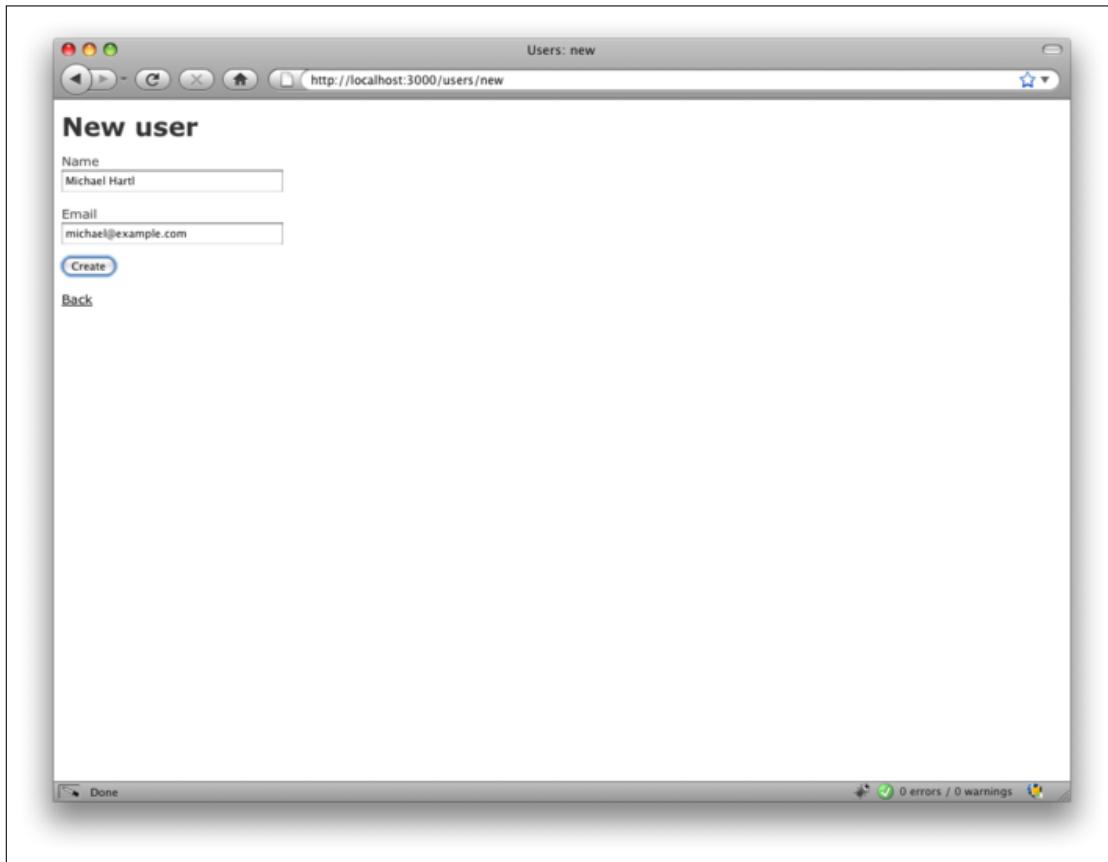


Figure 2.5: The new user page (</users/new>). (full size)

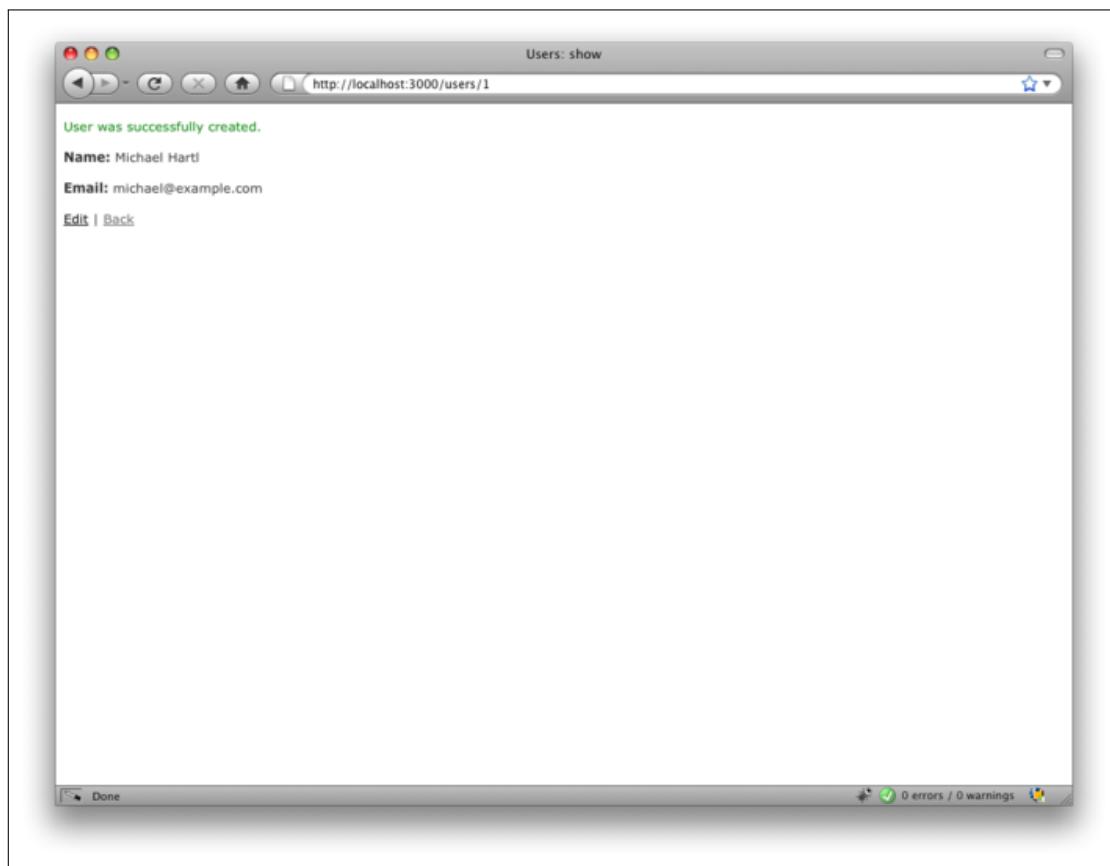


Figure 2.6: The page to show a user (</users/1>). [\(full size\)](#)

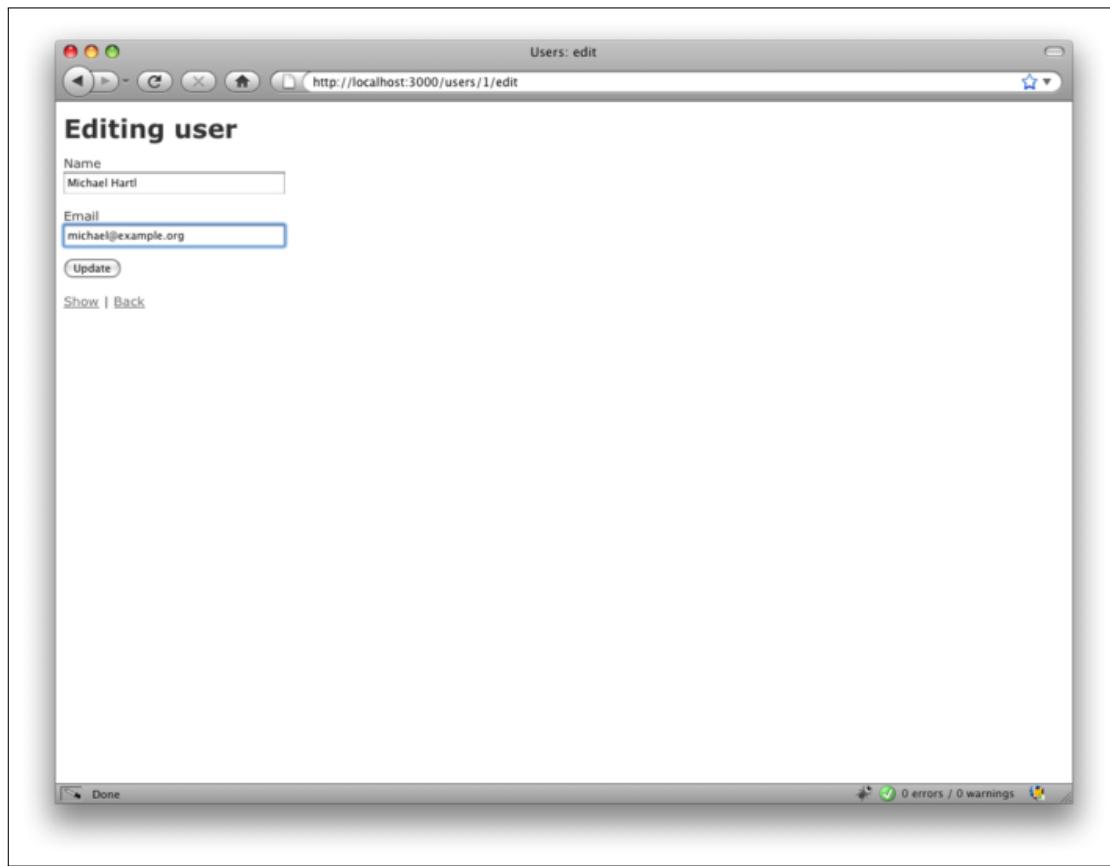


Figure 2.7: The user edit page (</users/1/edit>). [\(full size\)](#)

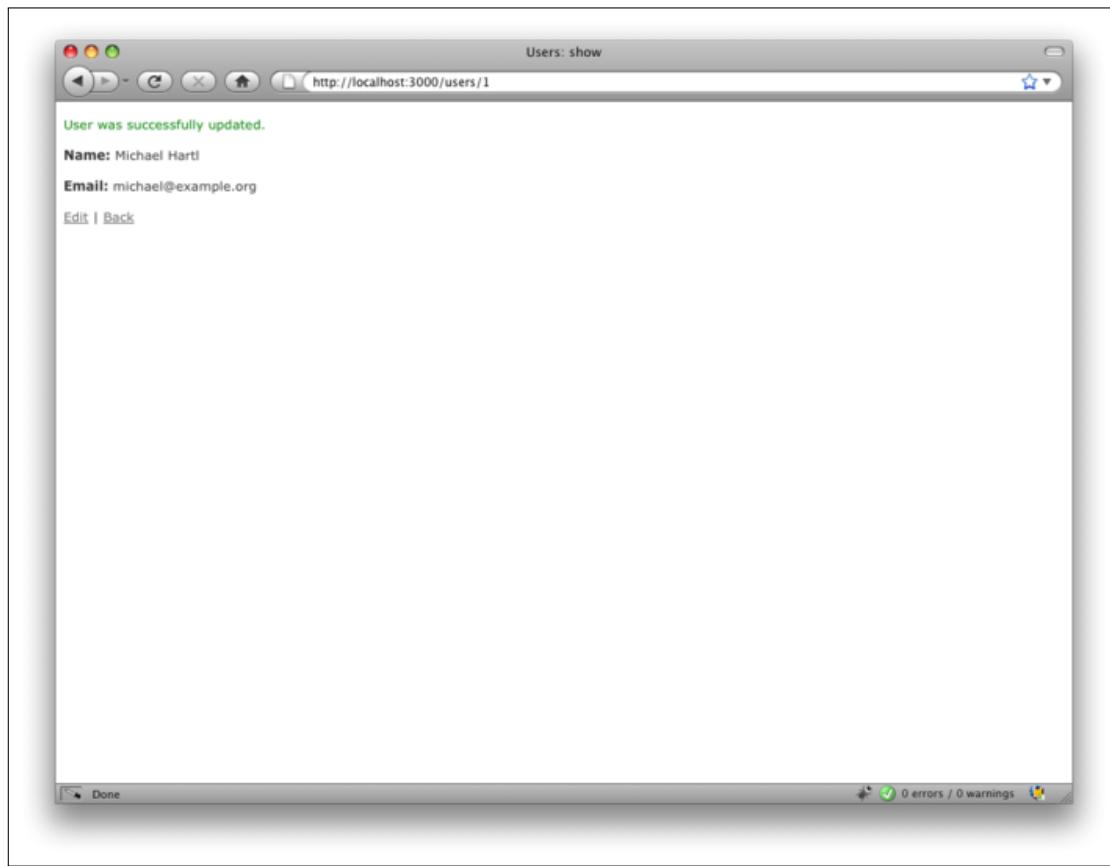


Figure 2.8: A user with updated information. ([full size](#))

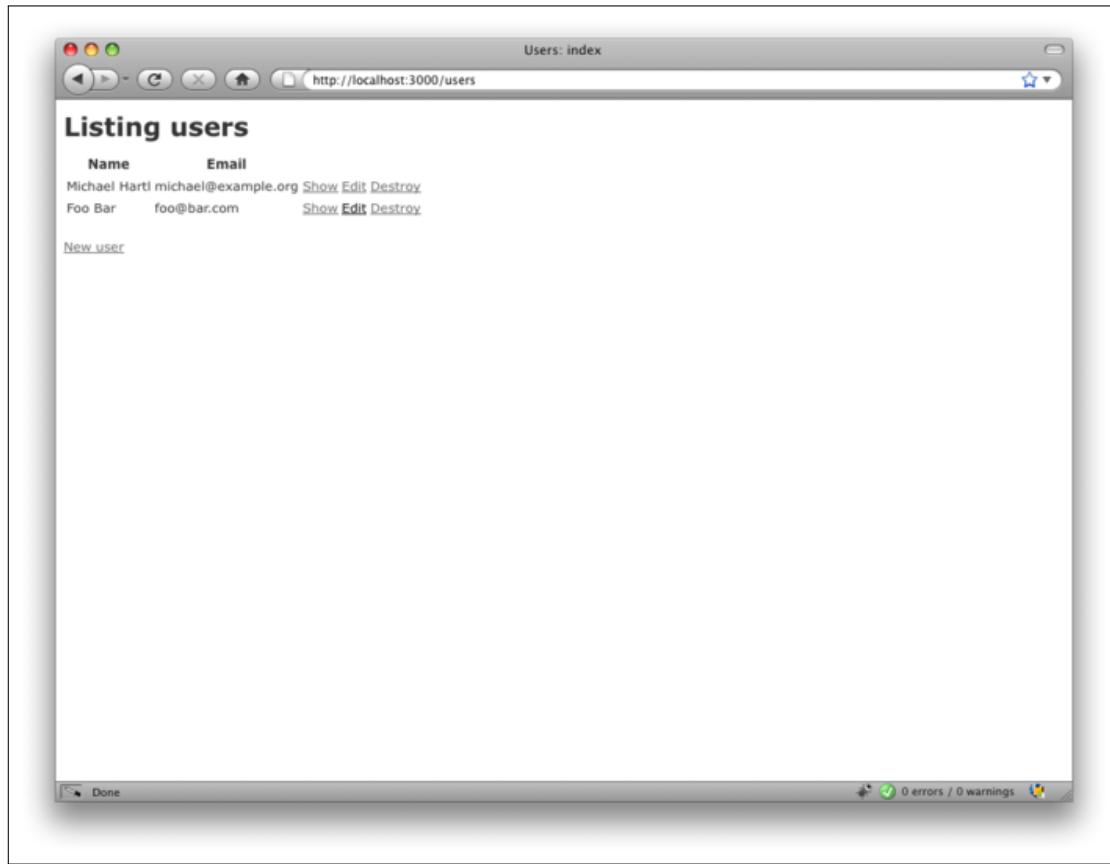


Figure 2.9: The user index page ([/users](#)) with a second user. ([full size](#))

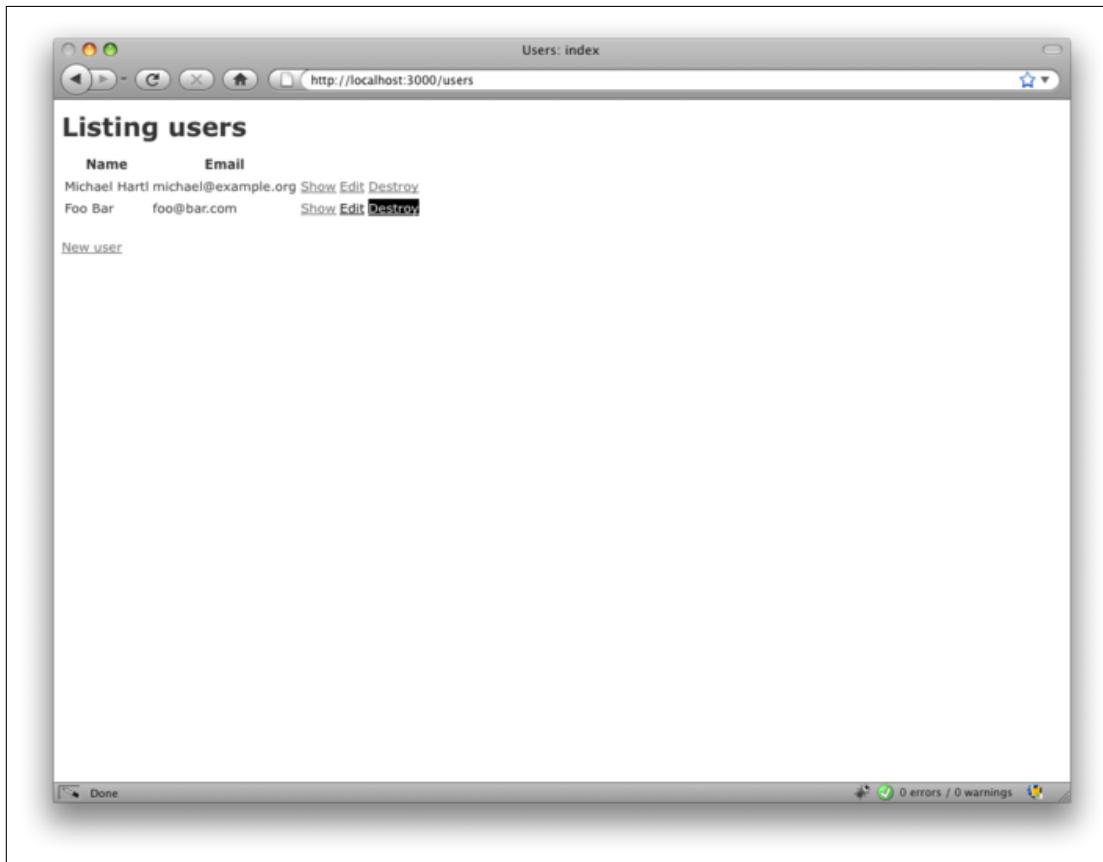


Figure 2.10: Destroying a user. ([full size](#))

Having shown how to create, show, and edit users, we come finally to destroying them (Figure 2.10). You should verify that clicking on the link in Figure 2.10 destroys the second user, yielding an index page with only one user. Section 10.4 adds user deletion to the sample app, taking care to restrict its use to a special class of administrative users.

2.2.2 MVC in action

Now that we've completed a quick overview of the Users resource, let's examine one particular part of it in the context of the Model-View-Controller (MVC) pattern introduced in Section 1.2.4. Our strategy will be to describe the results of a typical browser hit—a visit to the user index page at `/users`—in terms of MVC (Figure 2.11).

The steps in Figure 2.11

1. The browser issues a request for the `/users` URL.
2. Rails routes `/users` to the `index` action in the Users controller.
3. The `index` action asks the User model to retrieve all users (`User.all`).

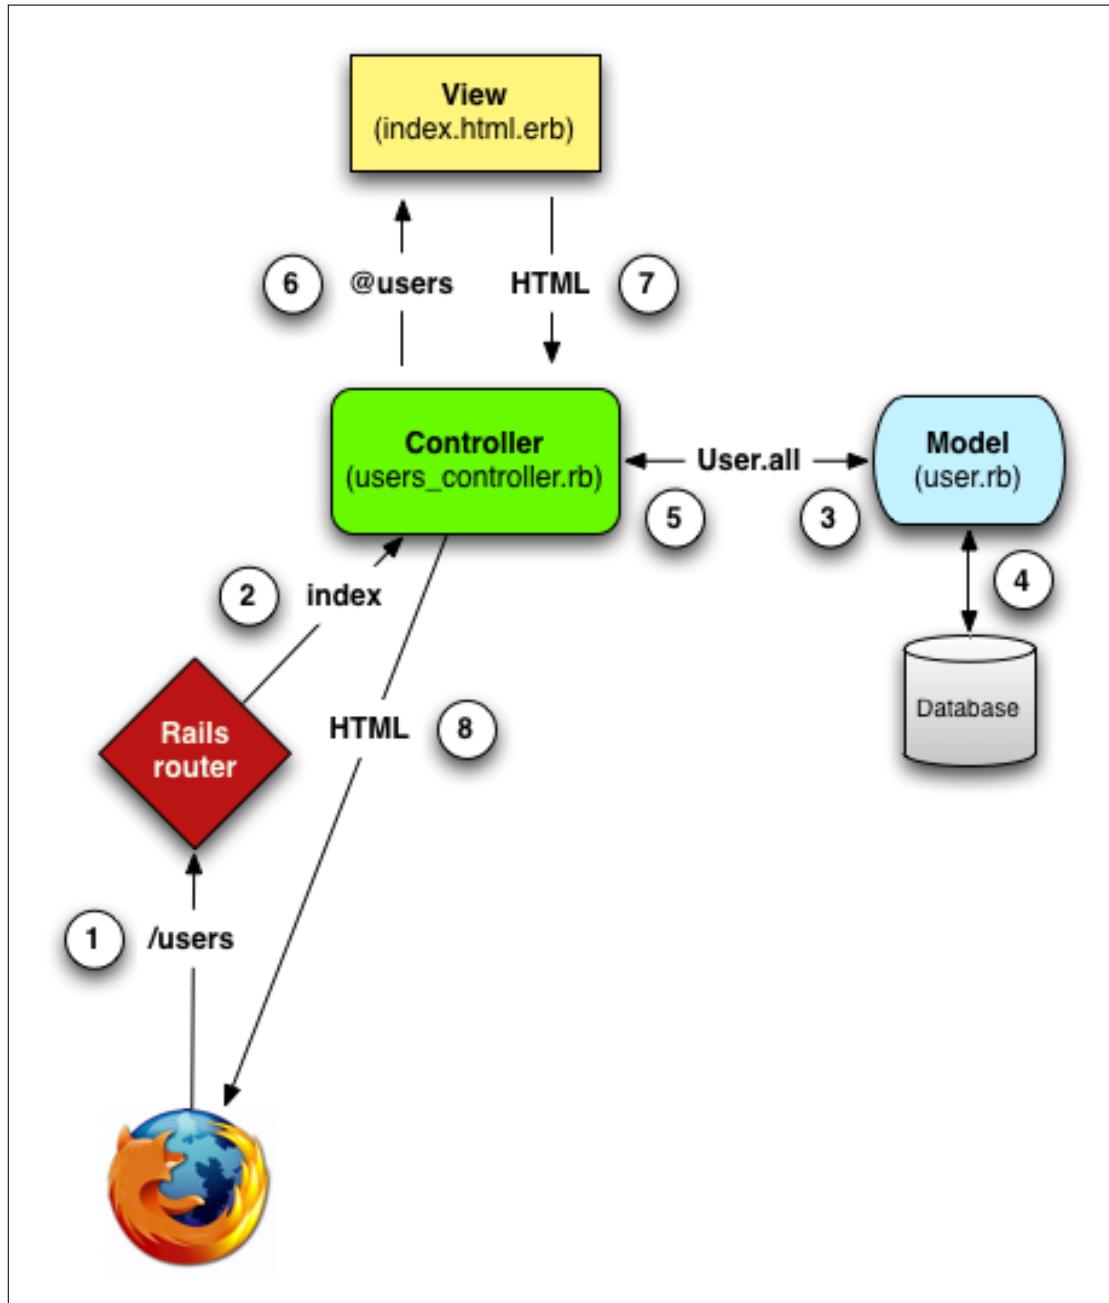


Figure 2.11: A detailed diagram of MVC in Rails. ([full size](#))

4. The User model pulls all the users from the database.
5. The User model returns the list of users to the controller.
6. The controller captures the users in the `@users` variable, which is passed to the `index` view.
7. The view uses Embedded Ruby to render the page as HTML.
8. The controller passes the HTML back to the browser.⁶

We start with a request issued from the browser—i.e., the result of typing a URL in the address bar or clicking on a link (step 1 in [Figure 2.11](#)). This request hits the *Rails router* (step 2), which dispatches to the proper *controller action* based on the URL (and, as we’ll see in [Box 3.1](#), the type of request). The code to create the mapping of user URLs to controller actions for the Users resource appears in [Listing 2.1](#);⁷ this code effectively sets up the table of URL/action pairs seen in [Table 2.1](#).

Listing 2.1. The Rails routes, with a rule for the Users resource.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  map.resources :users
  .
  .
  .
end
```

The pages from the tour in [Section 2.2.1](#) correspond to *actions* in the Users *controller*, which is a collection of related actions; the controller generated by the scaffolding is shown schematically in [Listing 2.2](#). Note the notation `class UsersController < ApplicationController`; this is an example of a Ruby *class* with *inheritance*. (We’ll discuss inheritance briefly in [Section 2.3.4](#) and cover both subjects in more detail in [Section 4.4](#).)

Listing 2.2. The Users controller in schematic form.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def index
    .
    .
    .
  end

  def show
    .
  end
```

⁶Some references indicate that the view returns the HTML directly to the browser (via a web server such as Apache or Nginx). Regardless of the implementation details, I prefer to think of the controller as a central hub through which all the application’s information flows.

⁷The strange notation `:users` is a *symbol*, which we’ll learn about in [Section 4.3.3](#).

```
  .
  end

  def new
  .
  .
  .

  end

  def create
  .
  .
  .

  end

  def edit
  .
  .
  .

  end

  def update
  .
  .
  .

  end

  def destroy
  .
  .
  .

  end
end
```

You may notice that there are more actions than there are pages; the `index`, `show`, `new`, and `edit` actions all correspond to pages from Section 2.2.1, but there are additional `create`, `update`, and `destroy` actions as well. These actions don't typically render pages (although they sometimes do); instead, their main purpose is to modify information about users in the database. This full suite of controller actions, summarized in Table 2.2, represents the implementation of the REST architecture in Rails (Box 2.1). Note from Table 2.2 that there is some overlap in the URLs; for example, both the user `show` action and the `update` action correspond to the URL `/users/1`. The difference between them is the HTTP request method they respond to. We'll learn more about HTTP request methods in starting in Section 3.2.2.

Box 2.1. REpresentational State Transfer (REST)

If you read much about Ruby on Rails web development, you'll see a lot of references to "REST", which is an acronym for REpresentational State Transfer. REST is an architectural style for developing distributed, networked systems and software applications such as the World Wide Web and web applications. Although

HTTP request	URL	Action	Purpose
GET	/users	index	page to list all users
GET	/users/1	show	page to show user with id 1
GET	/users/new	new	page to make a new user
POST	/users	create	create a new user
GET	/users/1/edit	edit	page to edit user with id 1
PUT	/users/1	update	update user with id 1
DELETE	/users/1	destroy	delete user with id 1

Table 2.2: RESTful routes provided by the Users resource in Listing 2.1.

REST theory is rather abstract, in the context of Rails applications REST means that most application components (such as users and microposts) are modeled as *resources* that can be created, read, updated, and deleted—operations that correspond both to the [CRUD operations of relational databases](#) and the four fundamental [HTTP request methods](#): POST, GET, PUT, and DELETE. (We’ll learn more about HTTP requests in Section 3.2.2 and especially Box 3.1.)

As a Rails application developer, the RESTful style of development helps you make choices about which controllers and actions to write: you simply structure the application using resources that get created, read, updated, and deleted. In the case of users and microposts, this process is straightforward, since they are naturally resources in their own right. In Chapter 12, we’ll see an example where REST principles allow us to model a subtler problem, “following users”, in a natural and convenient way.

To examine the relationship between the Users controller and the User model, let’s focus on a simplified version of the **index** action, shown in Listing 2.3.⁸

Listing 2.3. The simplified user **index** action for the demo application.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def index
    @users = User.all
  end
  .
  .
  .
end
```

This **index** action has the line `@users = User.all` (step 3), which asks the User model to retrieve a list of all the users from the database (step 4), and then places them in the variable `@users` (pronounced “at-users”) (step 5). The User model itself appears in Listing 2.4; although it is rather plain, it comes equipped with a large amount of functionality because of inheritance (Section 2.3.4 and Section 4.4). In particular, by using the Rails library called *Active Record*, the code in Listing 2.4 arranges for `User.all` to return all the users.

⁸The scaffold code is ugly and confusing, so I’ve suppressed it.

Listing 2.4. The User model for the demo application.

`app/models/user.rb`

```
class User < ActiveRecord::Base
end
```

Once the `@users` variable is defined, the controller calls the `view` (step 6), shown in [Listing 2.5](#). Variables that start with the `@` sign, called *instance variables*, are automatically available in the view; in this case, the `index.html.erb` view in [Listing 2.5](#) iterates through the `@users` list and outputs a line of HTML for each one.⁹

Listing 2.5. The view for the user index. (You are not expected to understand it now.)

`app/views/users/index.html.erb`

```
<h1>Listing users</h1>



| Name             | Email             |                                                   |                                                                   |                                                                                                             |
|------------------|-------------------|---------------------------------------------------|-------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <%= user.name %> | <%= user.email %> | <a href="#">&lt;%= link_to 'Show', user %&gt;</a> | <a href="#">&lt;%= link_to 'Edit', edit_user_path(user) %&gt;</a> | <a href="#">&lt;%= link_to 'Destroy', user, :confirm =&gt; 'Are you sure?', :method =&gt; :delete %&gt;</a> |

  
>

<%= link_to 'New user', new_user_path %>
```

The view converts its contents to HTML (step 7), which is then returned by the controller to the browser for display (step 8).

2.2.3 Weaknesses of this Users resource

Though good for getting a general overview of Rails, the scaffold Users resource suffers from a number of severe weaknesses.

⁹Remember, you aren't supposed to understand this code right now. It is shown only for purposes of illustration.

- **No data validations.** Our User model accepts data such as blank names and invalid email addresses without complaint.
- **No authentication.** We have no notion signing in or out, and no way to prevent any user from performing any operation.
- **No tests.** This isn't technically true—the scaffolding includes rudimentary tests—but the generated tests are ugly and inflexible, and they don't test for data validation, authentication, or any other custom requirements.
- **No layout.** There is no consistent site styling or navigation.
- **No real understanding.** If you understand the scaffold code, you probably shouldn't be reading this book. :-)

2.3 The Microposts resource

Having generated and explored the Users resource, we turn now to the associated Microposts resource. Throughout this section, I recommend comparing the elements of the Microposts resource with the analogous user elements from [Section 2.2](#); you should see that the two resources parallel each other in many ways. The RESTful structure of Rails applications is best absorbed by this sort of repetition of form; indeed, seeing the parallel structure of Users and Microposts even at this early stage is one of the prime motivations for this chapter. (As we'll see, writing applications more robust than the toy example in this chapter takes considerable effort—we won't see the Microposts resource again until [Chapter 11](#)—and I didn't want to defer its first appearance quite that far.)

2.3.1 A micropost microtour

As with the Users resource, we'll generate scaffold code for the Microposts resource using `script/generate scaffold`, in this case implementing the data model from [Figure 2.3](#):¹⁰

```
$ script/generate scaffold Micropost content:string user_id:integer
```

To update our database with the new data model, we need to run a migration as in [Section 2.2](#):

```
$ rake db:migrate
```

Now we are in a position to create microposts in the same way we created users in [Section 2.2.1](#). As you might guess, the scaffold generator has updated the Rails routes file with a rule for Microposts resource, as seen in [Listing 11.21](#).¹¹ As with users, the `map.resources :microposts` routing rule maps a micropost URLs to actions in the Microposts controller, as seen in [Table 2.3](#).

¹⁰As with the User scaffold, the scaffold generator for microposts follows the singular convention of Rails models; thus, we have `generate Micropost`.

¹¹The scaffold code may have extra newlines compared to Listing 11.21; this is not a cause for concern, as Ruby ignores extra newlines.

HTTP request	URL	Action	Purpose
GET	/microposts	index	page to list all microposts
GET	/microposts/1	show	page to show micropost with id 1
GET	/microposts/new	new	page to make a new micropost
POST	/microposts	create	create a new micropost
GET	/microposts/1/edit	edit	page to edit micropost with id 1
PUT	/microposts/1	update	update micropost with id 1
DELETE	/microposts/1	destroy	delete micropost with id 1

Table 2.3: RESTful routes provided by the Microposts resource in Listing 2.6.

Listing 2.6. The Rails routes, with a new rule for Microposts resources.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  map.resources :users
  map.resources :microposts
  .
  .
  .
end
```

The Microposts controller itself appears in schematic form Listing 2.7. Note that, apart from having **MicropostsController** in place of **UsersController**, Listing 2.7 is *identical* to the code in Listing 2.2. This is a reflection of the REST architecture common to both resources.

Listing 2.7. The Microposts controller in schematic form.

`app/controllers/microposts_controller.rb`

```
class MicropostsController < ApplicationController

  def index
    .
    .
    .
  end

  def show
    .
    .
    .
  end

  def new
    .
    .
    .
  end
```

```
end

def create
.
.
.

end

def edit
.
.
.

end

def update
.
.

end

def destroy
.
.

end
end
```

To make some actual microposts, we enter information at the new microposts page, [/microposts/new](#), as seen in [Figure 2.12](#).

At this point, go ahead and create a micropost or two, taking care to make sure that at least one has a `user_id` of 1 to match the id of the first user created in [Section 2.2.1](#). The result should look something like [Figure 2.13](#).

2.3.2 Putting the *micro* in microposts

Any *micropost* worthy of the name should have some means of enforcing the length of the post. Implementing this constraint in Rails is easy with *validations*; to accept microposts with at most 140 characters (à la Twitter), we use a length validation, as shown in [Listing 2.8](#).

Listing 2.8. Constraining microposts to at most 140 characters with a length validation.

`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  validates_length_of :content, :maximum => 140
end
```

At this point, the code in [Listing 2.8](#) may look rather mysterious—we'll cover validations more thoroughly starting in [Section 6.2](#)—but its effects are readily apparent if we go to the new micropost page and enter more

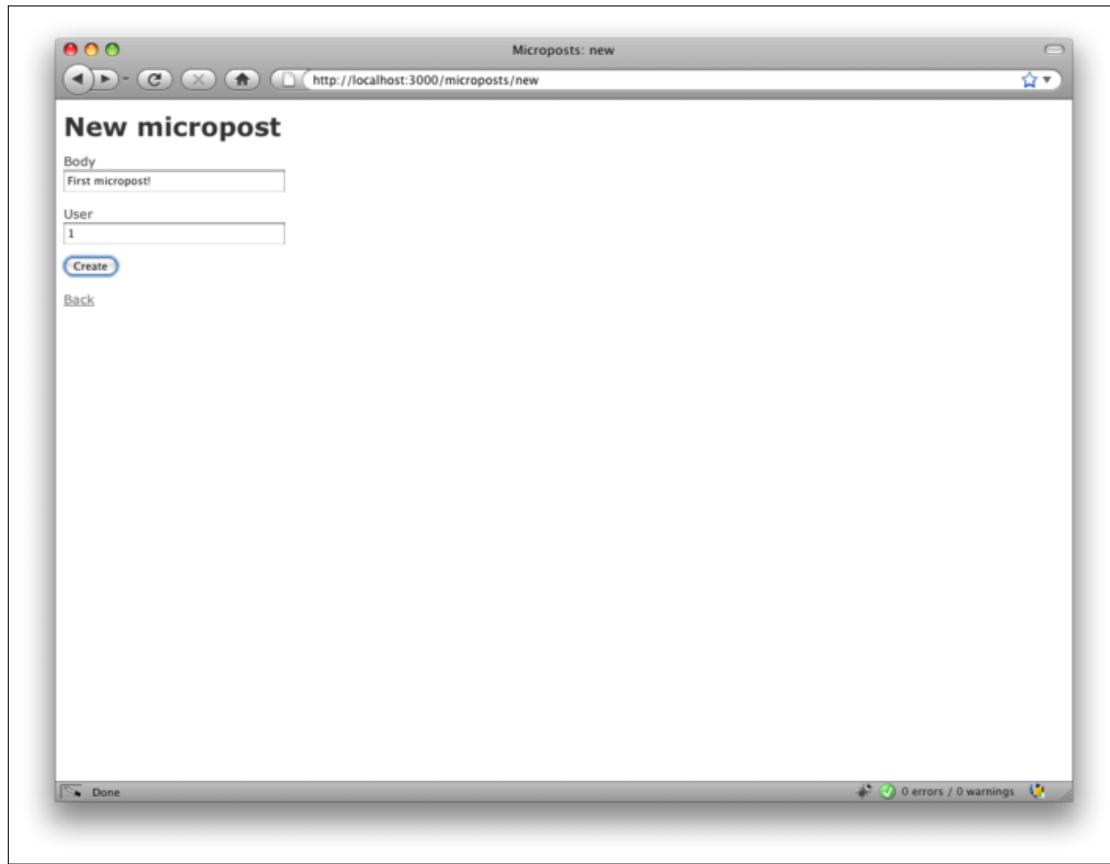


Figure 2.12: The new micropost page ([/microposts/new](#)). [\(full size\)](#)

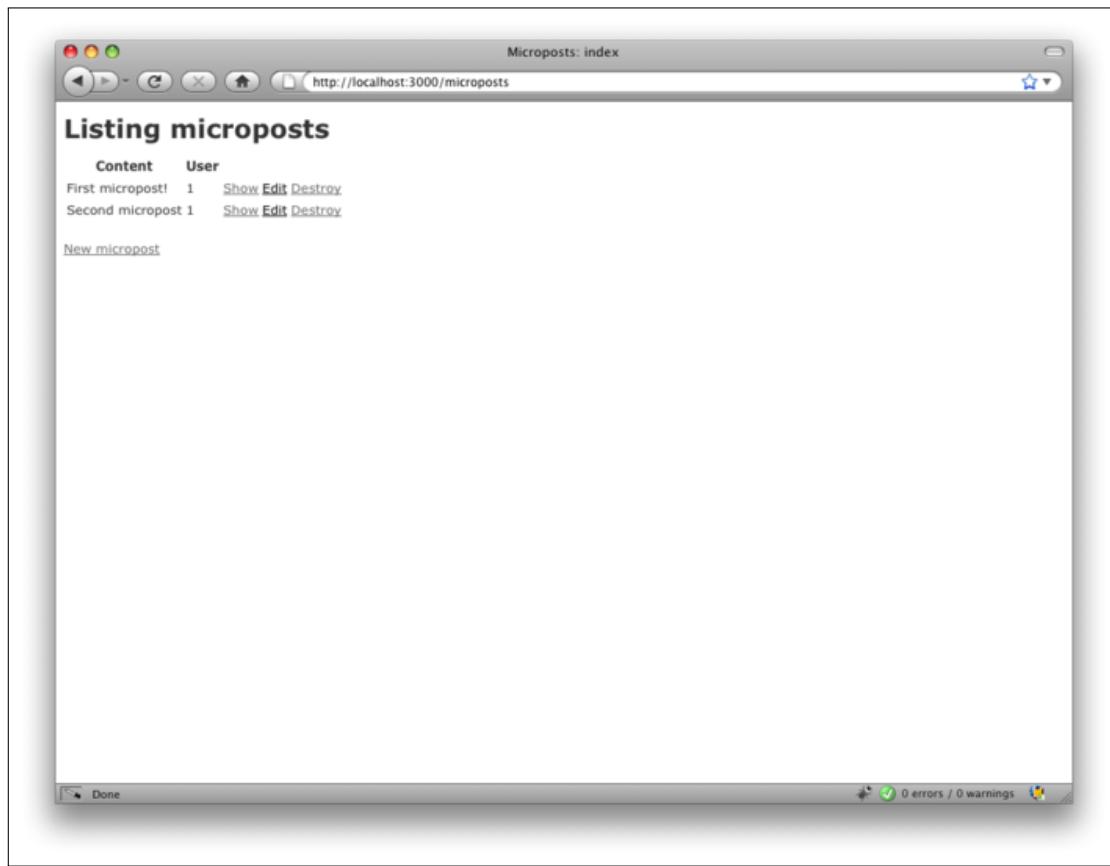


Figure 2.13: The micropost index page ([/microposts](#)). ([full size](#))

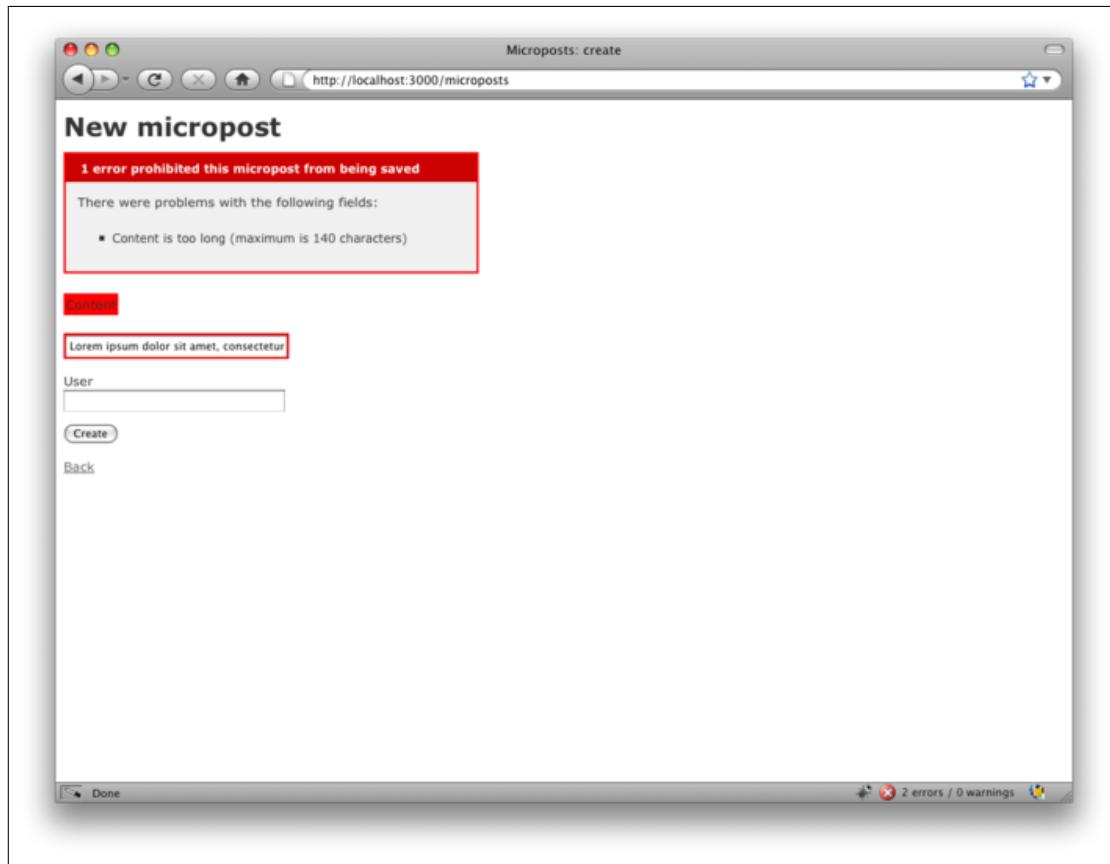


Figure 2.14: Error messages for a failed micropost creation. ([full size](#))

than 140 characters for the content of the post. As seen in Figure 2.14, Rails renders *error messages* indicating that the micropost's content is too long.¹² (We'll learn more about error messages in Section 8.2.3.)

2.3.3 A user `has_many` microposts

One of the most powerful features of Rails is the ability to form *associations* between different data models. In the case of our User model, each user potentially has many microposts. We can express this in code by updating the User and Micropost models as in Listing 2.9 and Listing 2.10.

Listing 2.9. A user has many microposts.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  has_many :microposts
end
```

¹²You might notice that the HTML Validator indicates an error; annoyingly, the default Rails error messages are not valid HTML.

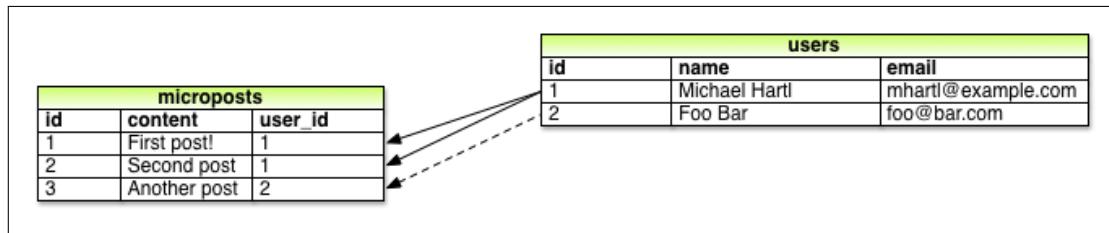


Figure 2.15: The association between microposts and users.

Listing 2.10. A micropost belongs to a user.

`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  belongs_to :user

  validates_length_of :content, :maximum => 140
end
```

We can visualize the result of this association in Figure 2.15. Because of the `user_id` column in the `microposts` table, Rails (using Active Record) can infer the microposts associated with each user.

In Chapter 11 and Chapter 12, we will use the association of users and microposts both to display all a user's microposts and to construct a Twitter-like micropost feed. For now, we can examine the implications of the user-micropost association by using the `console`, which is a useful tool for interacting with Rails applications. We first invoke the console with `script/console` at the command line, and then retrieve the first user from the database using `User.first` (putting the results in the variable `first_user`):

```
$ script/console
>> first_user = User.first
=> #<User id: 1, name: "Michael Hartl", email: "michael@example.org",
created_at: "2010-04-03 02:01:31", updated_at: "2010-04-03 02:01:31">
>> first_user.microposts
=> [#<Micropost id: 1, content: "First micropost!", user_id: 1, created_at:
"2010-04-03 02:37:37", updated_at: "2010-04-03 02:37:37">, #<Micropost id: 2,
content: "Second micropost", user_id: 1, created_at: "2010-04-03 02:38:54",
updated_at: "2010-04-03 02:38:54">]
```

Here we have accessed the user's microposts using the code `first_user.microposts`: with this code, Active Record automatically returns all the microposts with `user_id` equal to the id of `first_user` (in this case, 1). We'll learn much more about the association facilities in Active Record in Chapter 11 and Chapter 12.

2.3.4 Inheritance hierarchies

We end our discussion of the demo application with a brief description of the controller and model class hierarchies in Rails. This discussion will only make much sense if you have some experience with object-

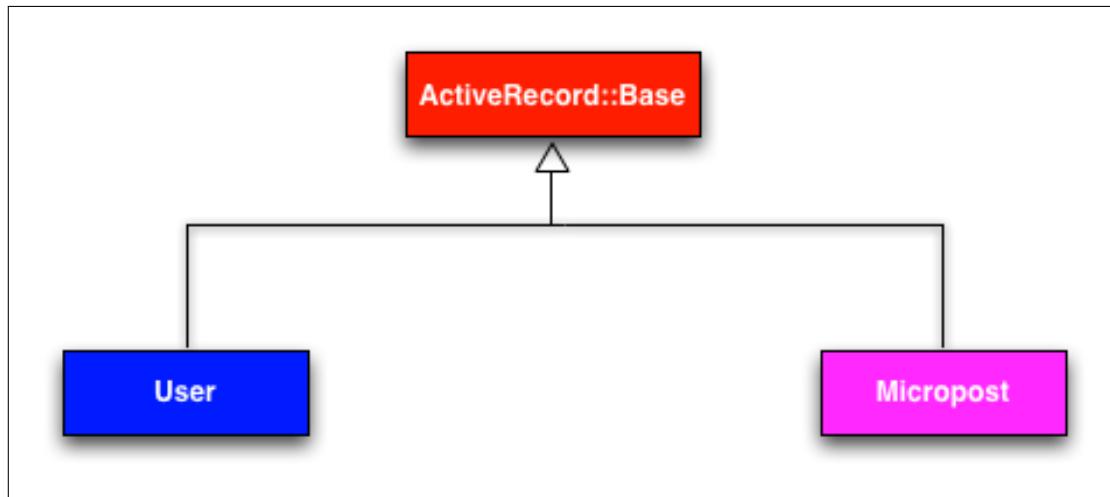


Figure 2.16: The inheritance hierarchy for the User and Micropost models.

oriented programming (OOP); if you haven't studied OOP, feel free to skip this section. In particular, if you are unfamiliar with *classes* (discussed in Section 4.4), I suggest looping back to this section at a later time.

We start with the inheritance structure for models. Comparing Listing 2.11 and Listing 2.12, we see that both the User model and the Micropost model inherit (via the left angle bracket `<`) from `ActiveRecord::Base`, which is the base class for models provided by ActiveRecord; a diagram summarizing this relationship appears in Figure 2.16. It is by inheriting from `ActiveRecord::Base` that our model objects gain the ability to communicate with the database, treat the database columns as Ruby attributes, and so on.

Listing 2.11. The `User` class, with inheritance.

`app/models/user.rb`

```

class User < ActiveRecord::Base
  .
  .
end
  
```

Listing 2.12. The `Micropost` class, with inheritance.

`app/models/micropost.rb`

```

class Micropost < ActiveRecord::Base
  .
  .
end
  
```

The inheritance structure for controllers is only slightly more complicated. Comparing Listing 2.13 and Listing 2.14, we see that both the Users controller and the Microposts controller inherit from the Application

controller. Examining Listing 2.15, we see that `ApplicationController` itself inherits from `ActionController::Base`; this is the base class for controllers provided by the Rails library Action Pack. The relationships between these classes is illustrated in Figure 2.17.

Listing 2.13. The `UsersController` class, with inheritance.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
end
```

Listing 2.14. The `MicropostsController` class, with inheritance.

`app/controllers/microposts_controller.rb`

```
class MicropostsController < ApplicationController
  .
  .
  .
end
```

Listing 2.15. The `ApplicationController` class, with inheritance.

`app/controllers/application_controller.rb`

```
class ApplicationController < ActionController::Base
  .
  .
  .
end
```

As with model inheritance, by inheriting ultimately from `ActionController::Base` both the `Users` and `Microposts` controllers gain a large amount of functionality, such as the ability to manipulate model objects, filter inbound HTTP requests, and render views as HTML. Since all Rails controllers inherit from `ApplicationController`, rules defined in the Application controller automatically apply to every action in the application. For example, in Section 8.2.4 we'll define a rule in the Application controller to filter passwords from all the Rails log files, thereby avoiding a serious potential security breach.

2.3.5 Deploying the demo app

With the completion of the Microposts resource, now is a good time to push the repository up to GitHub:¹³

¹³Ordinarily, you should make smaller, more frequent commits, but for the purposes of this chapter a single big commit at the end is just fine.

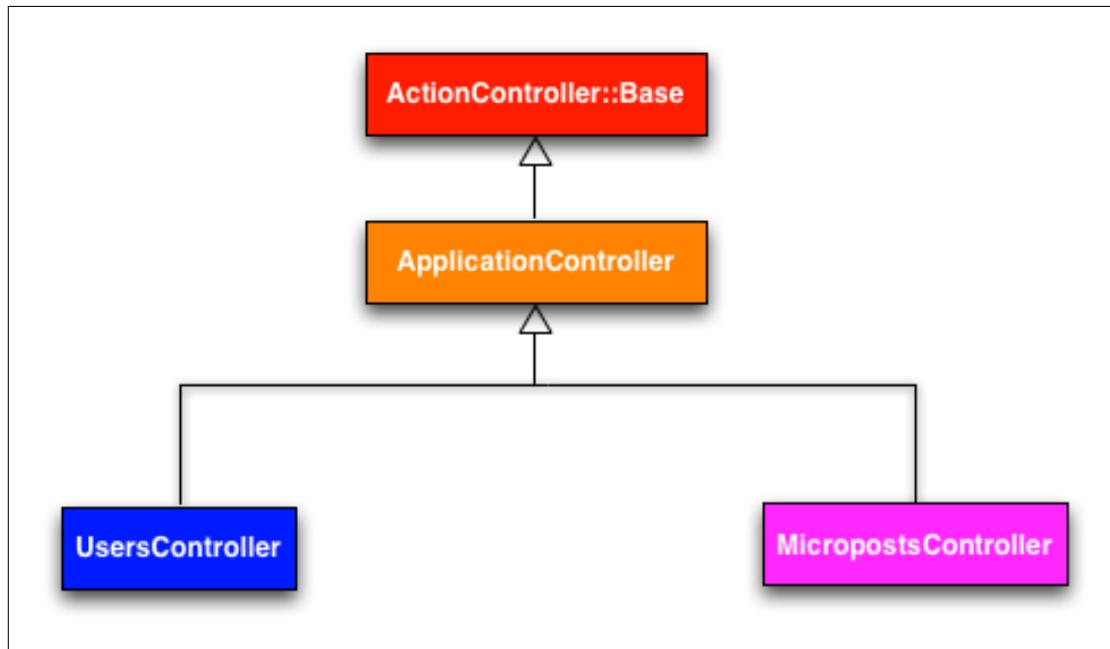


Figure 2.17: The inheritance hierarchy for the Users and Microposts controllers.

```

$ git add .
$ git commit -a -m "Done with the demo app"
$ git push
  
```

You can also deploy the demo app to Heroku:

```

$ heroku create
$ git push heroku master
$ heroku rake db:migrate
  
```

Note the final line here, which runs the database migrations on the Heroku server. This updates the database at Heroku with the necessary user/micropost data model. If you want to push the *data* up, too, you can do so using the `taps` gem and `db:push`:

```

$ [sudo] gem install taps
$ heroku db:push
  
```

2.4 Conclusion

We've come now to the end of the 30,000-foot view of a Rails application. The demo app developed in this chapter has several strengths and a host of weaknesses.

Strengths

- High-level overview of Rails
- Introduction to MVC
- First taste of the REST architecture
- Beginning data modeling
- A live, database-backed web application in production¹⁴

Weaknesses

- No custom layout or styling
- No static pages (like “Home” or “About”)
- No user passwords
- No user images
- No signing in
- No security
- No automatic user/micropost association
- No notion of “following” or “followed”
- No micropost feed
- No test-driven development
- **No real understanding**

The rest of this tutorial is dedicated to building on the strengths and eliminating the weaknesses.

¹⁴If you deployed to Heroku in [Section 2.3.5](#).

Chapter 3

Mostly static pages

In this chapter, we will begin developing the sample application that we will serve as our example throughout the rest of this tutorial. Although the sample app will eventually have users, microposts, and a full login and authentication framework, we will begin with a seemingly limited topic: the creation of static pages. Despite its seeming simplicity, making static pages is a highly instructive exercise, rich in implications—a perfect start for our nascent application.

Although Rails is designed for making database-backed dynamic websites, it also excels at making the kind of static pages we might make with raw HTML files. In fact, using Rails even for static pages yields a distinct advantage: we can easily add just a *small* amount of dynamic content. In this chapter we'll learn how. Along the way, we'll get our first taste of *automated testing*, which will help us be more confident that our code is correct. Moreover, having a good test suite will allow us to *refactor* our code with confidence, changing its form without changing its function.

As in [Chapter 2](#), before getting started we need to create a new Rails project, this time called `sample_app`:

```
$ cd ~/rails_projects  
$ rails sample_app  
$ cd sample_app
```

As before, we'll copy the `.gitignore` file from [Listing 1.2](#) to the sample app's Rails root directory:

```
$ git init  
$ cp ~/rails_projects/first_app/.gitignore .  
$ git add .  
$ git commit -m "Initial commit"
```

Since we'll be using this sample app throughout the rest of the book, it's a good idea to make a repository at GitHub ([Figure 3.1](#)) and push it up:

```
$ git remote add origin git@github.com:<username>/sample_app.git  
$ git push origin master
```

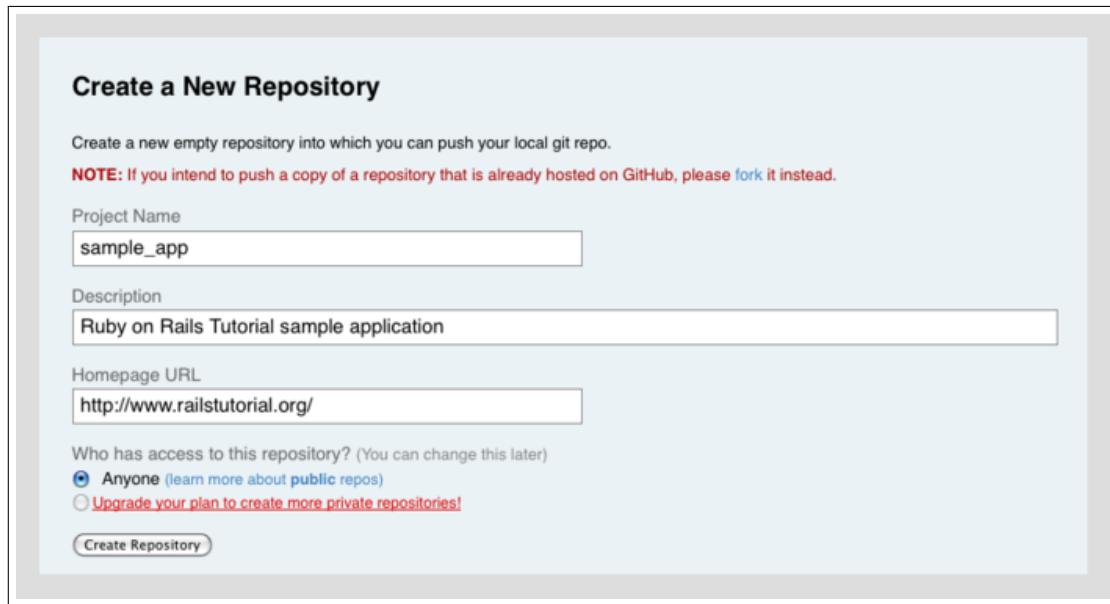


Figure 3.1: Creating the sample app repository at GitHub. ([full size](#))

Of course, we can optionally deploy the app to Heroku even at this early stage:

```
$ heroku create
$ git push heroku master
```

As you proceed through the rest of the book, I recommend pushing and deploying the application regularly:

```
$ git push
$ git push heroku
```

With that, we're ready to get started developing the sample application.

3.1 Static pages

Rails has two main ways of making static web pages. First, Rails can handle *truly* static pages consisting of raw HTML files. Second, Rails allows us to define *views* containing raw HTML, which Rails can *render* so that the web server can send it to the browser.

In order to get our bearings, it's helpful to recall the Rails directory structure from [Section 1.2.3](#) (Figure 1.3). In this section, we'll be working mainly in the `app/controllers` and `app/views` directories. (In Section 3.2, we'll even add a new directory of our own.)

3.1.1 Truly static pages

We start with truly static pages. Recall from [Section 1.2.5](#) that every Rails application comes with a minimal working application thanks to the `rails` script, with a default welcome page at the address `http://localhost:3000/` ([Figure 1.5](#)).

To learn where this page comes from, take a look at the file `public/index.html` ([Figure 3.2](#)). Because the file contains its own stylesheet information, it's a little messy, but it gets the job done: by default, Rails serves any files in the `public` directory directly to the browser.¹ In the case of the special `index.html` file, you don't even have to indicate the file in the URL, as `index.html` is the default. You can include it if you want, though; the addresses

`http://localhost:3000/`

and

`http://localhost:3000/index.html`

are equivalent.

As you might expect, if we want we can make our own static HTML files and put them in the same `public` directory as `index.html`. For example, let's create a file with a friendly greeting ([Listing 3.1](#)):²

```
$ mate public/hello.html
```

Listing 3.1. A typical HTML file, with a friendly greeting.

`public/hello.html`

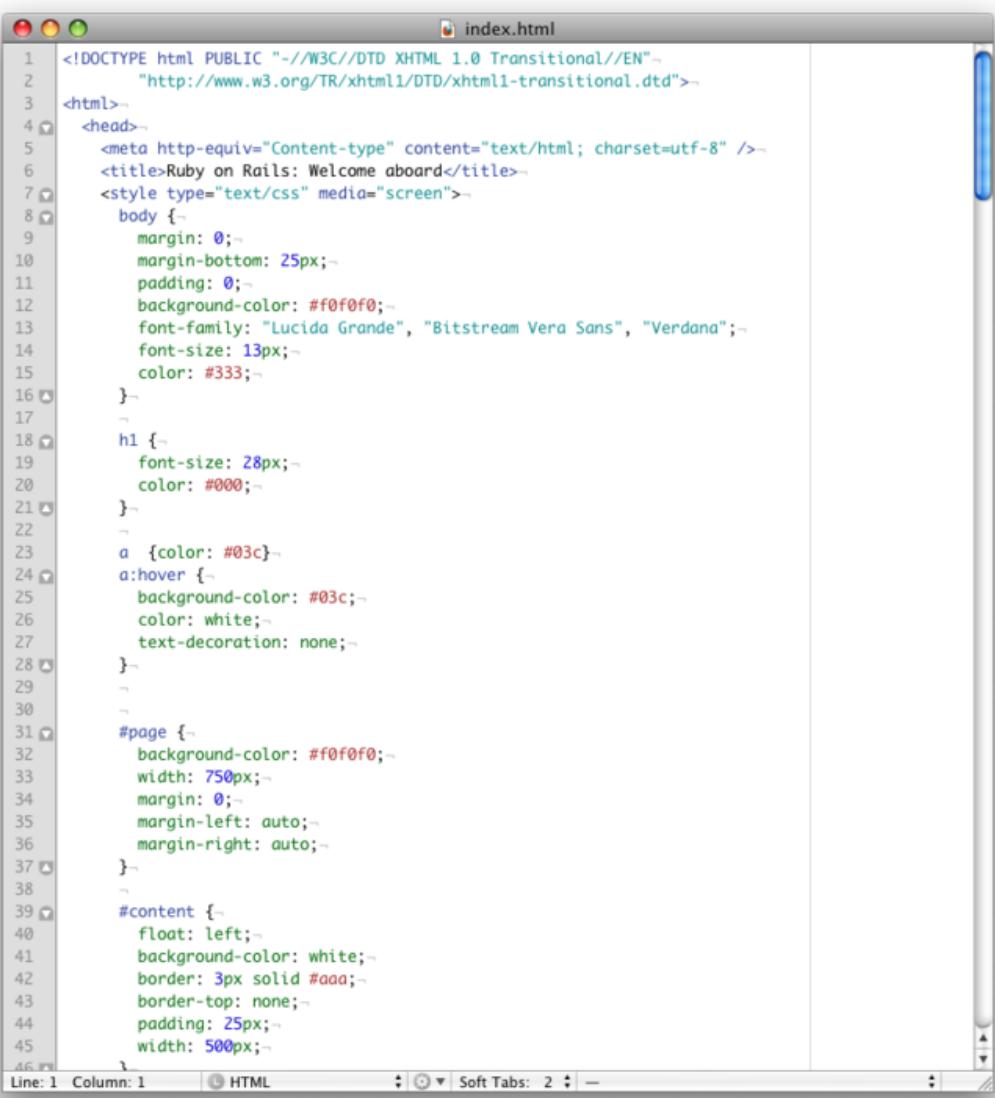
```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html>
  <head>
    <title>Greeting</title>
  </head>
  <body>
    <p>Hello, world!</p>
  </body>
</html>
```

We see in [Listing 3.1](#) the typical structure of an HTML document: a *document type*, or doctype, declaration at the top to tell browsers which version of HTML we're using;³ a `head` section, in this case with "Greeting" inside a `title` tag; and a `body` section, in this case with "Hello, world!" inside a `p` (paragraph) tag. (The indentation is optional—HTML is not sensitive to whitespace, and ignores both tabs and spaces—but it makes the document's structure easier to see.) As promised, when visiting the address `http://localhost:3000/hello.html`, Rails renders it straightforward ([Figure 3.3](#)). Note that the title displayed at the top of the browser window in [Figure 3.3](#) is just the contents inside the `title` tag, namely, "Greeting".

¹In fact, Rails ensures that requests for such files never hit the main Rails stack; they are delivered directly from the filesystem.

²As usual, replace `mate` with the command for your text editor.

³HTML changes with time; by explicitly making a doctype declaration we make it likelier that browsers will render our pages properly in the future.



The screenshot shows a Mac OS X TextEdit window titled "index.html". The window contains the source code for the "index.html" file. The code is a combination of HTML and CSS. The CSS is written in a syntax highlighting style, with colors used for different elements like tags, attributes, and styles. The code includes doctype declarations, meta tags, and styles for body, h1, a, and page elements.

```
1 <!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"-
2     "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">-
3 <html>-
4     <head>-
5         <meta http-equiv="Content-type" content="text/html; charset=utf-8" />-
6         <title>Ruby on Rails: Welcome aboard</title>-
7         <style type="text/css" media="screen">-
8             body {-
9                 margin: 0;-
10                margin-bottom: 25px;-
11                padding: 0;-
12                background-color: #f0f0f0;-
13                font-family: "Lucida Grande", "Bitstream Vera Sans", "Verdana";-
14                font-size: 13px;-
15                color: #333;-
16            }-
17            h1 {-
18                font-size: 28px;-
19                color: #000;-
20            }-
21            a {color: #03c;-
22            a:hover {-
23                background-color: #03c;-
24                color: white;-
25                text-decoration: none;-
26            }-
27            }-
28            }-
29            }-
30            }-
31            #page {-
32                background-color: #f0f0f0;-
33                width: 750px;-
34                margin: 0;-
35                margin-left: auto;-
36                margin-right: auto;-
37            }-
38            }-
39            #content {-
40                float: left;-
41                background-color: white;-
42                border: 3px solid #aaa;-
43                border-top: none;-
44                padding: 25px;-
45                width: 500px;-
46            }-

```

Figure 3.2: The `public/index.html` file. (full size)

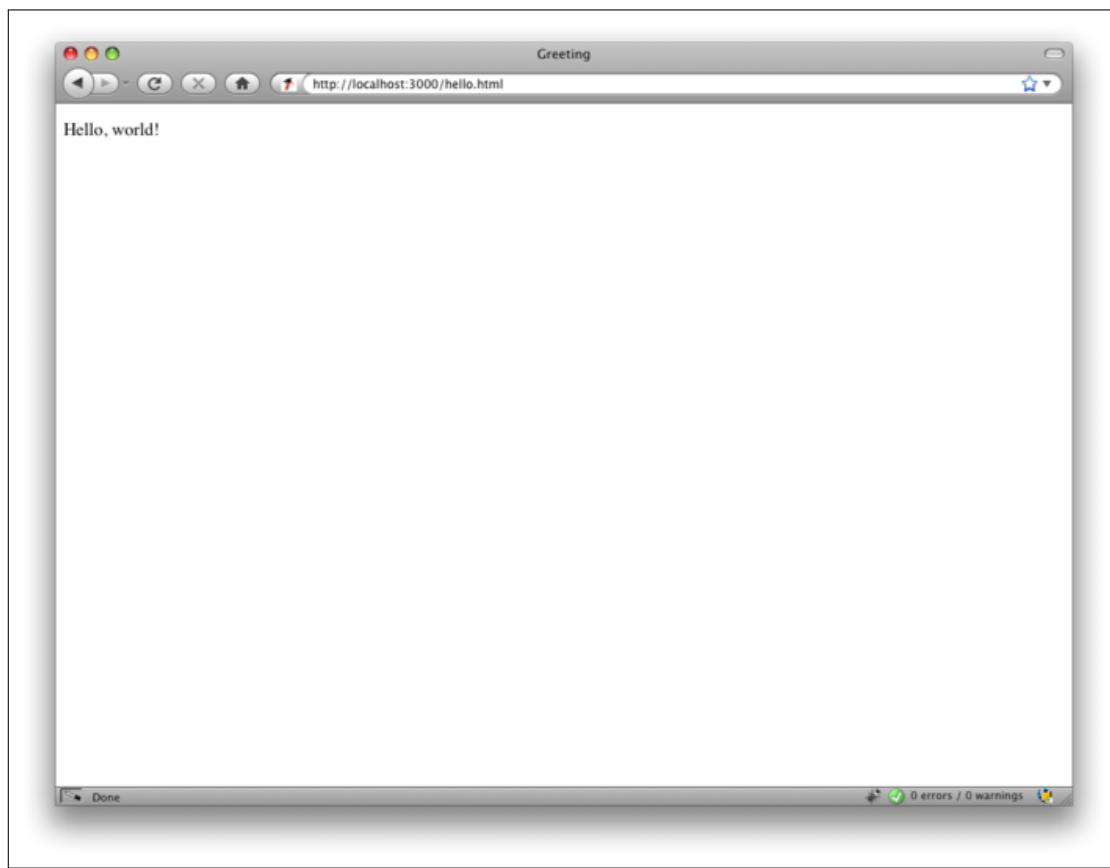


Figure 3.3: Our very own static HTML file (<http://localhost:3000/hello.html>). [\(full size\)](#)

Since this file is just for demonstration purposes, we don't really want it to be part of our sample application, so it's probably best to remove it once the thrill of creating it has worn off:

```
$ rm public/hello.html
```

We'll leave the `index.html` file alone for now, but of course eventually we should remove it: we don't want the root of our application to be the Rails default page shown in [Figure 1.5](#). We'll see in [Section 5.2](#) how to change the address `http://localhost:3000/` to point to something other than `public/index.html`.

3.1.2 Static pages with Rails

The ability to return static HTML files is nice, but it's not particularly useful for making dynamic web applications. In this section, we'll take a first step toward making dynamic pages by creating a set of Rails *actions*, which are a more powerful way to define URLs than static files.⁴ Rails actions come bundled together inside *controllers* (the C in MVC from [Section 1.2.4](#)), which contain sets of actions related by a common purpose. We got a glimpse of controllers in [Chapter 2](#), and will come to a deeper understanding once we explore the [REST architecture](#) more fully (starting in [Chapter 6](#)); in essence, a controller is a container for a group of (possibly dynamic) web pages.

To get started, recall from [Section 1.3.5](#) that, when using Git, it's a good practice to do our work on a separate branch topic branch rather than the master branch. If you're using Git for version control, you should run the following command:

```
$ git checkout -b static-pages
```

Rails comes with a script for making controllers called `generate`; all it needs to work its magic is the controller's name. Since we're making this controller to handle (mostly) static pages, we'll just call it the Pages controller, and plan to make actions for a Home page, a Contact page, and an About page. The `generate` script takes an optional list of actions, so we'll include some of our initial actions directly on the command line:⁵

Listing 3.2. Generating a Pages controller.

```
$ script/generate controller Pages home contact
exists  app/controllers/
exists  app/helpers/
create  app/views/pages
exists  test/functional/
```

⁴Our method for making static pages is probably the simplest, but it's not the only way. The optimal method really depends on your needs; if you expect a *large* number of static pages, using a Pages controller can get quite cumbersome, but in our sample app we'll only need a few. See this [blog post on simple pages at has_many :through](#) for a survey of techniques for making static pages with Rails. *Warning:* the discussion is fairly advanced, so you might want to wait a while before trying to understand it.

⁵Windows users will probably have to run `generate` explicitly with Ruby using `ruby script/generate`

Throughout *Rails Tutorial*, I'll omit this explicit `ruby` for brevity.

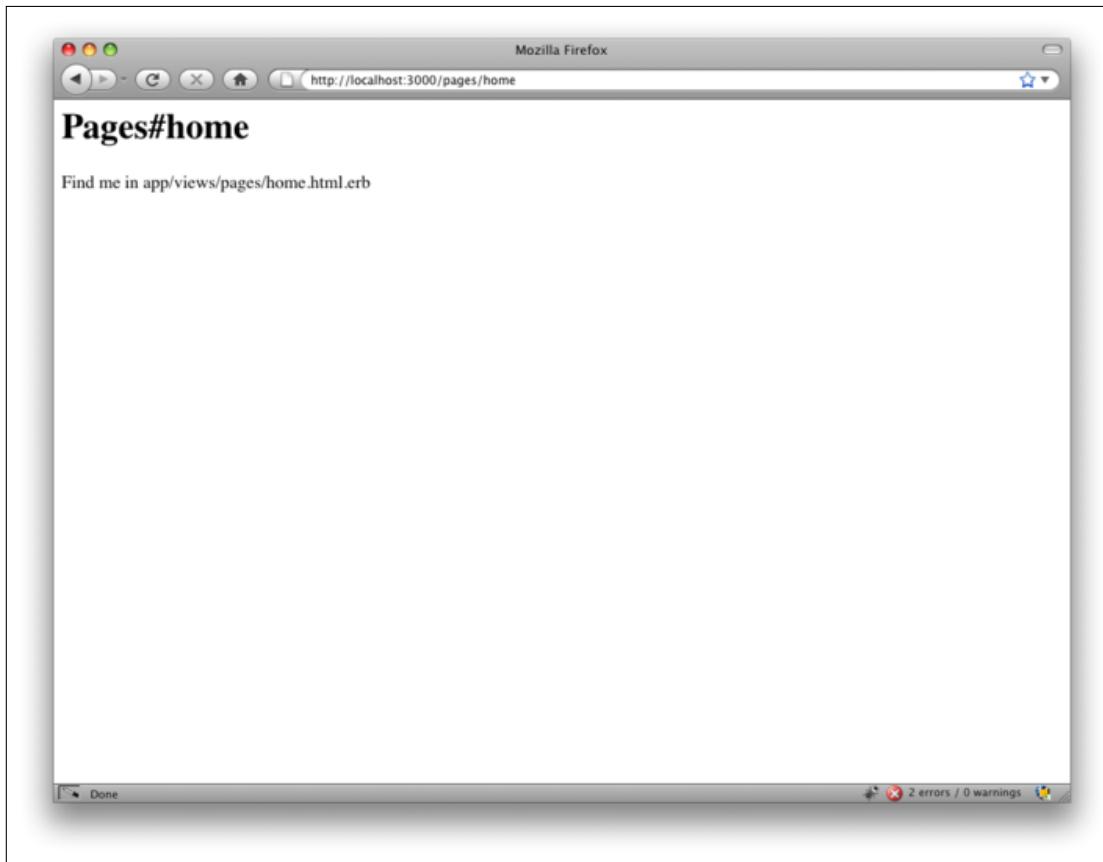


Figure 3.4: The raw home view ([/pages/home](#)) generated by Rails. ([full size](#))

```
create  test/unit/helpers/
create  app/controllers/pages_controller.rb
create  test/functional/pages_controller_test.rb
create  app/helpers/pages_helper.rb
create  test/unit/helpers/pages_helper_test.rb
create  app/views/pages/home.html.erb
create  app/views/pages/contact.html.erb
```

Here, I've intentionally “forgotten” the **about** page so that we can see how to add it in by hand (Section 3.2). By default, Rails ensures that the actions inside controllers can be found at a URL of the following form:

`http://localhost:3000/controller/action`

In our case, this means that when we generate a **home** action inside the Pages controller we automatically get a page at the address `/pages/home`. To see the results, kill the server by hitting Ctrl-C, run **script/server**, and then navigate to `/pages/home` (Figure 3.4).

To understand where this page comes from, let's start by taking a look at the Pages controller in a text editor; you should see something like Listing 3.3. (You may note that, unlike the demo Users and Microposts controllers from Chapter 2, the Pages controller does not follow the REST conventions.)

Listing 3.3. The Pages controller made by Listing 3.2.

`app/controllers/pages_controller.rb`

```
class PagesController < ApplicationController
  def home
  end

  def contact
  end
end
```

We see here that `pages_controller.rb` defines a *class* called `PagesController`. Classes are simply a convenient way to organize *functions* (also called *methods*) like the `home` and `contact` actions, which are defined using the `def` keyword. The angle bracket `<` indicates that `PagesController` *inherits* from the Rails class `ApplicationController`; as we'll see momentarily, this means that our pages come equipped with a large amount of Rails-specific functionality. (We'll learn more about both classes and inheritance in Section 4.4.)

In the case of the Pages controller, both its methods are initially empty:

```
def home
end

def contact
end
```

In plain Ruby, these methods would simply do nothing. In Rails, the situation is different; `PagesController` is a Ruby class, but because it inherits from `ApplicationController` the behavior of its methods is specific to Rails: when visiting the URL `/pages/home`, Rails looks in the Pages controller and executes the code in the `home` action, and then renders the *view* (the V in MVC from Section 1.2.4) corresponding to the action. In the present case, the `home` action is empty, so all hitting `/pages/home` does is render the view. So, what does a view look like, and how do we find it?

If you take another look at the output in Listing 3.2, you might be able to guess the correspondence between actions and views: an action like `home` has a corresponding view called `home.html.erb`. We'll learn in Section 3.3 what the `.erb` part means; from the `.html` part you probably won't be surprised that it basically looks like HTML (Listing 3.4).

Listing 3.4. The generated view for the Home page.

`app/views/pages/home.html.erb`

```
<h1>Pages#home</h1>
<p>Find me in app/views/pages/home.html.erb</p>
```

The view for the `contact` action is analogous (Listing 3.5).

Listing 3.5. The generated view for the Contact page.

`app/views/pages/contact.html.erb`

```
<h1>Pages#contact</h1>
<p>Find me in app/views/pages/contact.html.erb</p>
```

Both of these views are just placeholders: they have a top-level heading (inside the `h1` tag) and a paragraph (`p` tag) with the full path to the relevant file. We'll add some (very slightly) dynamic content starting in Section 3.3, but as they stand these views underscore an important point: Rails views can simply contain static HTML. As far as the browser is concerned, the raw HTML files from Section 3.1.1 and the controller/action method of delivering pages are indistinguishable: all the browser ever sees is HTML.

In the remainder of this chapter, we'll first add the `about` action we "forgot" in Section 3.1.2, add a very small amount of dynamic content, and then take the first steps toward styling the pages with CSS. Before moving on, if you're using Git it's a good idea to add the files for the Pages controller to the repository at this time:

```
$ git add .
$ git commit -am "Added a Pages controller"
```

You may recall from Section 1.3.5 that we used the Git command `git commit -a -m "Message"`, with flags for "all changes" (`-a`) and a message (`-m`); Git also lets us roll the two flags into one as `-am`, and I'll stick with this more compact formulation throughout the rest of this book.

3.2 Our first tests

If you ask five Rails developers how to test any given piece of code, you'll get about fifteen different answers—but they'll all agree that you should definitely be writing tests. It's in this spirit that we'll write tests for our sample application, writing good tests without worrying too much about making them perfect. We'll also make the increasingly common choice about *when* to write tests by writing them *before* the application code—an approach known as *test-driven development*, or TDD. Our specific example will be to add an About page to our sample site. Fortunately, adding the extra page is not hard—you might even be able to guess the answer based on the examples in the previous section—which means that we can focus on testing, which contains quite a few new ideas.

At first, testing for the existence of a page might seem like overkill, but experience shows that it is not. So many things can go wrong when writing software that having a good test suite is invaluable to assure quality. Moreover, it is common for computer programs—and especially web applications—to be constantly extended, and any time you make a change you risk introducing errors. Writing tests doesn't guarantee that these bugs won't happen, but it makes them much more likely to be caught (and fixed) when they occur. Furthermore, by writing tests for bugs that *do* happen, we can make them much less likely to recur.

(As noted in Section 1.1.1, if you find the tests overwhelming, go ahead and skip them on first reading. Once you have a stronger grasp of Rails and Ruby, you can loop back and learn testing on a second pass.)

3.2.1 Testing tools

To write tests for our sample application, we'll be using a framework called **RSpec**, which is a *domain-specific language* for describing the behavior of code, together with a program (called **spec**) to verify the desired behavior. Designed for testing any Ruby program, RSpec has gained significant traction in the Rails community. Obie Fernandez, author of *The Rails Way*, has called RSpec "the Rails Way", and I agree.⁶

Installing RSpec

Since RSpec comes as a gem, installing it is easy:

```
$ [sudo] gem install rspec -v 1.3.0
$ [sudo] gem install rspec-rails -v 1.3.2
```

If you're using Ruby 1.8.7 or lower, RSpec should now be operational, but it won't work immediately if you're using Ruby 1.9. RSpec is built on top of the `Test::Unit` library, the default testing framework for Rails; up through Ruby 1.8.7, `Test::Unit` ships as part of the Ruby Standard Library, but (for reasons unknown to me) it is no longer included in Ruby 1.9. Moreover, RSpec doesn't yet work with the most recent standalone version of `Test::Unit`. This means that, if you are using Ruby 1.9, you should uninstall any existing versions of `Test::Unit` and then install the version RSpec needs:⁷

```
$ ruby -v
1.9.1
$ [sudo] gem uninstall test-unit
$ [sudo] gem install test-unit -v 1.2.3
```

To get RSpec started in a Rails application, we just need to run a script in the Rails root directory that generates some initial files:

```
$ script/generate rspec
```

Then add the files to the repository:

```
$ git add .
$ git status
# On branch static-pages
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
#       new file:   lib/tasks/rspec.rake
#       new file:   script/autospec
```

⁶The `Shoulda` testing framework is a good alternate choice—the Other Rails Way, so to speak.

⁷Much thanks to reader Ron Green for going well beyond the call of duty to identify and solve this thorny problem.

```
#      new file: script/spec
#      new file: spec/rcov.opts
#      new file: spec/spec.opts
#      new file: spec/spec_helper.rb
#
$ git commit -am "Added RSpec files"
```

Installing Autotest

Although not strictly necessary for testing, I've found that the Autotest tool is extraordinarily valuable for test-driven development. Autotest (part of the [ZenTest](#) suite) continuously runs your test suite in the background based on the specific file changes you make. The result is instant feedback on the status of your tests. We'll learn more about Autotest when we see it in action ([Section 3.2.2](#)), but for now you should install it if you want to try it:

```
$ [sudo] gem install autotest-rails -v 4.1.0
```

The next steps depend on your platform. I'll go through the steps for OS X, since that's what I use, and then give references to blog posts that discuss Autotest on Linux and Windows. On OS X, you should [install Growl](#) (if you don't have it already) and then install the **autotest-fsevent** and **autotest-growl** gems.⁸

```
# On OS X
$ [sudo] gem install autotest-fsevent -v 0.1.1
$ [sudo] gem install autotest-growl -v 0.2.0
```

Then make an Autotest configuration file in your Rails root directory and fill it with the contents of Listing 3.6:

```
$ mate .autotest
```

(Note: this will create an Autotest configuration for the sample application only; if you want to share this Autotest configuration with other Rails or Ruby projects, you should create the **.autotest** file in your *home* directory instead:

```
$ mate ~/.autotest
```

where ~ (tilde) is the Unix symbol for “home directory”.)

⁸The most recent version of **autotest-fsevent**, 0.1.3, outputs verbose error messages on many versions of OS X, including mine. I recommend using version 0.1.1 until the problem is fixed.

Listing 3.6. The `.autotest` configuration file for Autotest on OS X.

```
require "autotest/growl"
require "autotest/fsevent"
```

If you’re running Linux with the Gnome desktop, you should try the steps at [Automate Everything](#), which sets up on Linux a system similar to Growl notifications on OS X. Windows users should try installing [Growl for Windows](#) and then follow the instructions at the [GitHub page for autotest-growl](#).

With that, we’re ready to get testing!

3.2.2 TDD: Red, Green, Refactor

In test-driven development, we first write a *failing* test: in our case, a piece of code that expresses the idea that there “should be an `about`” page. Then we get the test to pass, in our case by adding the `about` action and corresponding view. The reason we don’t typically do the reverse—implement first, then test—is to make sure that we actually test for the feature we’re adding. Before I started using TDD, I was amazed to discover how often my “tests” actually tested the wrong thing, or even tested nothing at all. By making sure that the test fails first and *then* passes, we can be more confident that the test is doing the right thing.

It’s important to understand that TDD is not always the right tool for the job. In particular, when you aren’t at all sure how to solve a given programming problem, it’s often useful to skip the tests and write only application code, just to get a sense of what the solution will look like. (In the language of [Extreme Programming \(XP\)](#), this exploratory step is called a *spike*.) Once you see the general shape of the solution, you can then use TDD to implement a more polished version.

One way to proceed in test-driven development is a cycle known as “Red, Green, Refactor”. The first step, Red, refers to writing a failing test, which many test tools indicate with the color red. The next step, Green, refers to a passing test, indicated with the color (wait for it) green. Once we have a passing test (or set of tests), we are free to *refactor* our code, changing the form (eliminating duplication, for example) without changing the function.

We don’t have any colors yet, so let’s get started toward Red. RSpec (and testing in general) can be a little intimidating at first, so we’ll make use of the `generate rspec_controller` script to get ourselves started. It works just like the `generate controller` script we saw in Section 3.1.2:

```
$ script/generate rspec_controller Pages home contact
exists app/controllers/
exists app/helpers/
exists app/views/pages
create spec/controllers/
create spec/helpers/
create spec/views/pages
create spec/controllers/pages_controller_spec.rb
create spec/helpers/pages_helper_spec.rb
identical app/controllers/pages_controller.rb
identical app/helpers/pages_helper.rb
create spec/views/pages/home.html.erb_spec.rb
identical app/views/pages/home.html.erb
create spec/views/pages/contact.html.erb_spec.rb
identical app/views/pages/contact.html.erb
```

Note that this step has created a directory called `spec` in our Rails root. This is to differentiate the RSpec tests from the ones that ship with Rails by default (using a framework called `Test::Unit`), which live in the `test` directory. Now that we are set up with RSpec, we won't need that directory any longer, so it's best to remove it:

```
$ rm -rf test/
```

In the output from `generate rspec_controller` above, you can see from lines like

```
identical app/controllers/pages_controller.rb
```

that this script does nothing to the files we created in [Section 3.1.2](#), but it would have created them if they didn't already exist.

By the way, I'm not partial to separate tests for views, which I've found to be quite brittle, so let's remove them, together with the rarely used helper spec:

```
$ rm -rf spec/views
$ rm -f spec/helpers/pages_helper_spec.rb
```

We'll handle tests for views and helpers directly in the controller tests starting in [Section 3.3](#).

To get started with RSpec, take a look at the Pages controller spec⁹ we just generated ([Listing 3.7](#)).

Listing 3.7. The generated Pages controller spec.

`spec/controllers/pages_controller_spec.rb`

```
require 'spec_helper'

describe PagesController do

  describe "GET 'home'" do
    it "should be successful" do
      get 'home'
      response.should be_success
    end
  end

  describe "GET 'contact'" do
    it "should be successful" do
      get 'contact'
      response.should be_success
    end
  end
end
```

⁹In the context of RSpec, tests are often called *specs*, but for simplicity I'll usually stick to the term "test"—*except* when referring to a file such as `pages_controller_spec`, in which case I'll write "Pages controller spec".

(Here I've deleted one of the generated examples and just left the ones we'll be using.)

This code is pure Ruby, but even if you've studied Ruby before it probably won't look very familiar. This is because RSpec uses the general malleability of Ruby to define a *domain-specific language* (DSL) built just for testing. The important point is that *you do not need to understand RSpec's syntax to be able to use RSpec*. It may seem like magic at first, but RSpec is designed to read more or less like English, and if you follow the examples from the `generate` script and the other examples in this tutorial you'll pick it up fairly quickly.

Let's focus on the first test to get a sense of what it does:

```
describe "GET 'home'" do
  it "should be successful" do
    get 'home'
    response.should be_success
  end
end
```

The first line indicates that we are describing a GET operation for the `home` action. This is just a description, and it can be anything you want; RSpec doesn't care, but you and other human readers probably do. In this case, GET refers to one of the *HTTP verbs*, which are things that you can do with the hypertext transfer protocol (Box 3.1). Then the spec says that when you visit the home page, it should be successful. As with the first line, what goes inside the quote marks is irrelevant to RSpec, and is intended to be descriptive to human readers.

The third line, `get 'home'`, is the first line that really does something. Inside of RSpec, this line *actually submits a GET request*; in other words, it acts like a browser and hits a page, in this case `/pages/home`. (It knows to hit the Pages controller automatically because this is a Pages controller test; it knows to hit the home page because we tell it to explicitly.) Finally, the fourth line says that the *response* of our application should indicate success (i.e., it should return a *status code* of 200; see Box 3.1).

Box 3.1. GET, et cetera.

The hypertext transfer protocol ([HTTP](#)) defines four basic operations, corresponding to the four verbs *get*, *post*, *put*, and *delete*. These refer to operations between a *client* computer (typically running a web browser such as Firefox or Safari) and a *server* (typically running a web server such as Apache or Nginx). (It's important to understand that, when developing Rails applications on a local computer, the client and server are the same physical machine, but in general they are different.) An emphasis on HTTP verbs is typical of web frameworks (including Rails) influenced by the *REST architecture*, which we'll start learning about in Chapter 8.

GET is the most common HTTP operation, used for *reading* data on the web; it just means “get a page”, and every time you visit a site like [google.com](#) or [craigslist.org](#) your browser is submitting a GET request. POST is the next most common operation; it is the request sent by your browser when you submit a form. In Rails applications, POST requests are typically used for *creating* things (although HTTP also allows POST to perform updates); for example, the POST request sent when you submit a registration form creates a new user on the remote site. The other two verbs, PUT and DELETE, are designed for *updating* and *destroying* things on the remote server. These requests are less common than GET and POST since browsers are incapable of sending them natively, but some web frameworks (including Ruby on Rails) have clever ways of making it *seem* like browsers are issuing such requests.

After the client sends a request, the server *responds* with a numerical code indicating the [HTTP status](#) of the response. For example, a status code of 200 means “success”, and a status code of 301 means “permanent

redirect". If you [install curl](#), a command-line client that can issue HTTP requests, you can see this directly at, e.g., www.railstutorial.org (where the --head flag prevents `curl` from returning the whole page):

```
$ curl --head www.railstutorial.org
HTTP/1.1 200 OK
.
.
```

Here the Rails Tutorial site indicates that the request was successful by returning the status 200 OK. In contrast, railstutorial.org is permanently redirected (to www.railstutorial.org, naturally), indicated by status code 301 (a "301 redirect"):

```
$ curl --head railstutorial.org
HTTP/1.1 301 Moved Permanently
Location: http://www.railstutorial.org/
.
.
```

When we write `response.should be_success` in an RSpec test, RSpec verifies that our application's response to the request is status code 200.

Now it's time to run our tests. There are several different and mostly equivalent ways to do this.¹⁰ One way to run all the tests is to use the `spec` script at the command line as follows:¹¹

```
$ spec spec/
....
Finished in 0.183973 seconds
2 examples, 0 failures
```

(If any test fails, be sure that you've migrated the database with `rake db:migrate` as described in [Section 1.2.5](#).) In this command the first `spec` is a program provided by RSpec, while `spec/` is the *directory* whose specs you want to run. You can, if you like, run only the specs in a subdirectory, for example, just the controller specs:

```
$ spec spec/controllers/
....
```

¹⁰Most IDEs also have an interface to testing, but as noted in [Section 1.2.1](#) I have limited experience with those tools.

¹¹You can also run `rake spec`, which is basically equivalent. When something goes wrong, `spec spec/` will show the [stack trace](#), whereas `rake spec` doesn't show the stack trace by default. Since the stack trace is often useful in debugging, `spec spec/` is a prudent default. I confess, though, that out of habit I usually type `rake spec`, and switch to `spec spec/` if something goes wrong.

```
Finished in 0.18124 seconds  
2 examples, 0 failures
```

You can also run a single file:

```
$ spec spec/controllers/pages_controller_spec.rb  
....  
Finished in 0.18328 seconds  
2 examples, 0 failures
```

The results of all three commands are the same since the Pages controller spec is currently our only test file. Throughout the rest of this book, I won't usually show the output of running the tests, but you should run **spec spec/** (or one of its variants) regularly as you follow along—or, better yet, use Autotest to run the test suite automatically. Speaking of which...

If you've installed Autotest, you can run it on your RSpec tests as follows:

```
$ autotest
```

If you're using a Mac with Growl notifications enabled, you should be able to replicate my setup, shown in [Figure 3.5](#). With Autotest running in the background and Growl notifications telling you the status of your tests, TDD can be positively addictive.

Red

Now let's get to Red part of the Red-Green cycle by writing a failing test for the **about** page. Following the models from [Listing 3.7](#), you can probably guess the right test ([Listing 3.8](#)).

Listing 3.8. The Pages controller spec with a failing test for the About page.

```
spec/controllers/pages_controller_spec.rb
```

```
require 'spec_helper'  
  
describe PagesController do  
  
  describe "GET 'home'" do  
    it "should be successful" do  
      get 'home'  
      response.should be_success  
    end  
  end  
  
  describe "GET 'contact'" do  
    it "should be successful" do
```

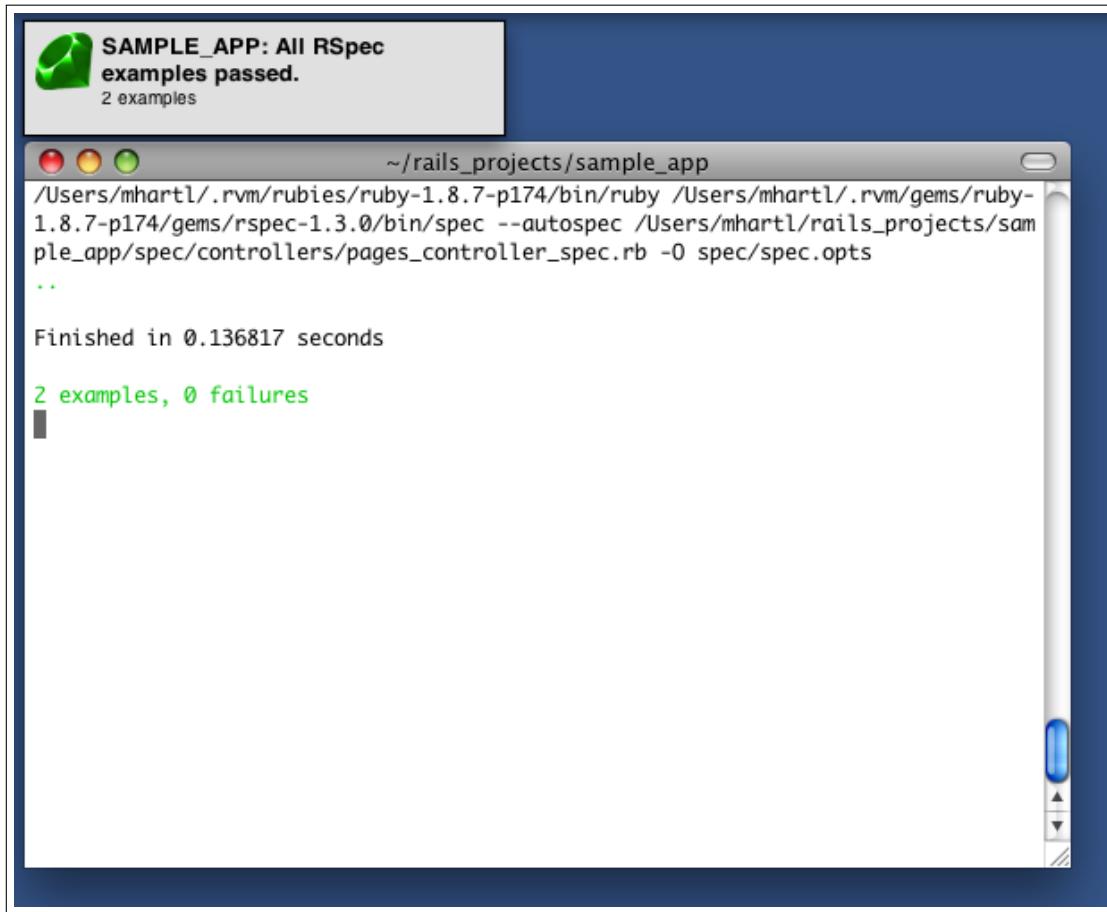
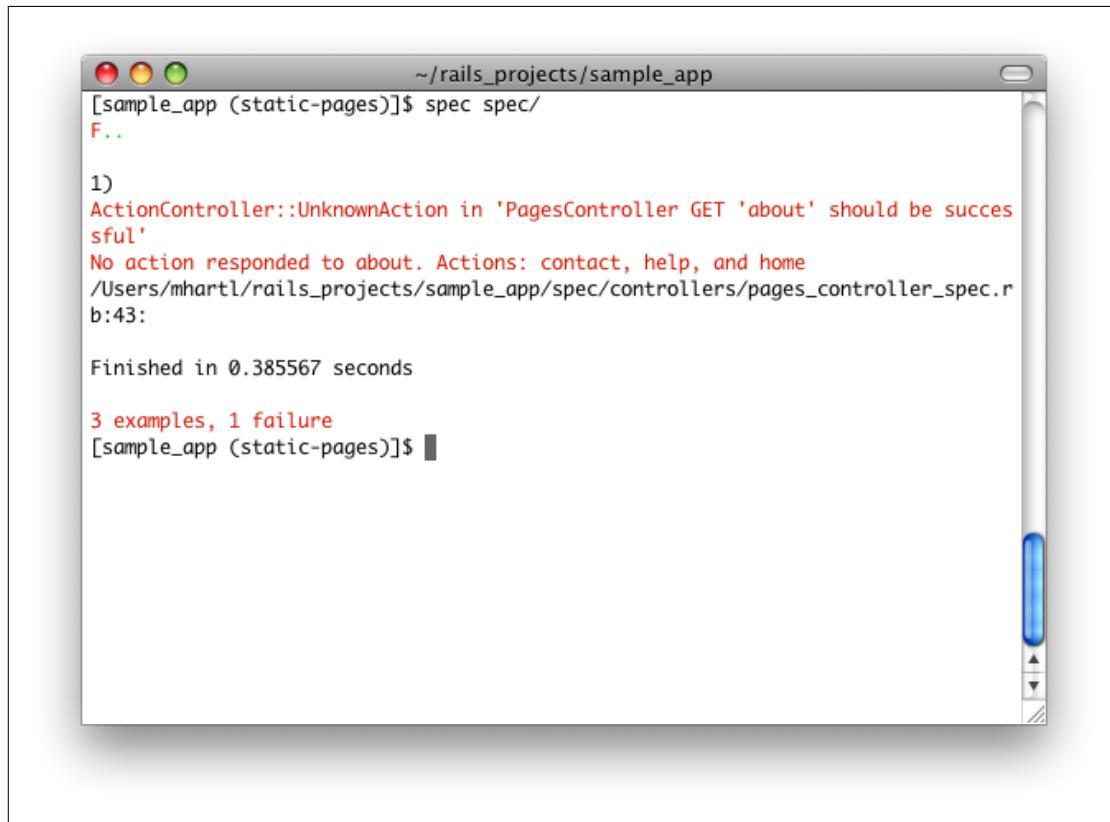


Figure 3.5: Autotest (via `autospec`) in action, with a Growl notification. ([full size](#))



The screenshot shows a terminal window with the title bar 'rails_projects/sample_app'. The command entered is '[sample_app (static-pages)]\$ spec spec/'. The output shows a failure:

```

~/rails_projects/sample_app
[sample_app (static-pages)]$ spec spec/
F..

1)
ActionController::UnknownAction in 'PagesController GET 'about'' should be successful
No action responded to about. Actions: contact, help, and home
/Users/mhartl/rails_projects/sample_app/spec/controllers/pages_controller_spec.r
b:43:

Finished in 0.385567 seconds

3 examples, 1 failure
[sample_app (static-pages)]$ 

```

Figure 3.6: Failing spec for the About page using `spec spec/`. (full size)

```

    get 'contact'
    response.should be_success
  end
end

describe "GET 'about'" do
  it "should be successful" do
    get 'about'
    response.should be_success
  end
end
end

```

The new test attempts to `get` the `about` action, and indicates that the resulting response should be a success. By design, it fails (with a red error message), as seen in Figure 3.6 (`spec spec/`) and Figure 3.7 (`au-`
`tospec`). (If by some chance the test passes, be sure to remove the file `app/views/pages/about.html.erb` if there is one. In general, if `foo.html.erb` exists, Rails will still respond to the `foo` action, even if it's not defined explicitly in the controller.)

The screenshot shows a terminal window titled "SAMPLE_APP: Some RSpec examples failed." It displays the command run: `~/rails_projects/sample_app`. The output shows a failure in the `PagesController` spec:

```
~/rails_projects/sample_app
/Users/mhartl/.rvm/rubies/ruby-1.8.7-p174/bin/ruby /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/rspec-1.3.0/bin/spec --autospec /Users/mhartl/rails_projects/sample_app/spec/controllers/pages_controller_spec.rb -O spec/spec.opts
F..

1)
ActionController::UnknownAction in 'PagesController' GET 'about' should be successful
No action responded to about. Actions: contact, help, and home
/Users/mhartl/rails_projects/sample_app/spec/controllers/pages_controller_spec.rb:43:

Finished in 0.142924 seconds

3 examples, 1 failure
```

Figure 3.7: Failing spec for the About page using Autotest. ([full size](#))

This is Red. Now let's get to Green.

Green

Recall from [Section 3.1.2](#) that we can generate a static page in Rails by creating an action and corresponding view with the page's name. In our case, the About page will first need an action called `about` in the Pages controller. Having written a failing test, we can now be confident that, in getting it to pass, we will actually have created a working `about` page.

Following the models provided by `home` and `contact` from [Listing 3.3](#), let's first add an `about` action in the Pages controller ([Listing 3.9](#)).

Listing 3.9. The Pages controller with added `about` action.

`app/controllers/pages_controller.rb`

```
class PagesController < ApplicationController
  def home
  end

  def contact
  end

  def about
  end
end
```

Next we add the corresponding view. Eventually we'll fill it with something more informative, but for now we'll just mimic the content from the generated views ([Listing 3.4](#) and [Listing 3.5](#)) for the `about` view ([Listing 3.10](#)).

Listing 3.10. A stub About page.

`app/views/pages/about.html.erb`

```
<h1>Pages#about</h1>
<p>Find me in app/views/pages/about.html.erb</p>
```

Running the specs or watching the update from Autotest ([Figure 3.8](#)) should get us back to Green:

```
$ spec spec/
```

Of course, it's never a bad idea to take a look at the page in a browser to make sure our tests aren't completely crazy ([Figure 3.9](#)).

Refactor

Now that we're at Green, we are free to *refactor* our code by changing its form without changing its function. Oftentimes code will start to "smell", meaning that it gets ugly, bloated, or filled with repetition. The

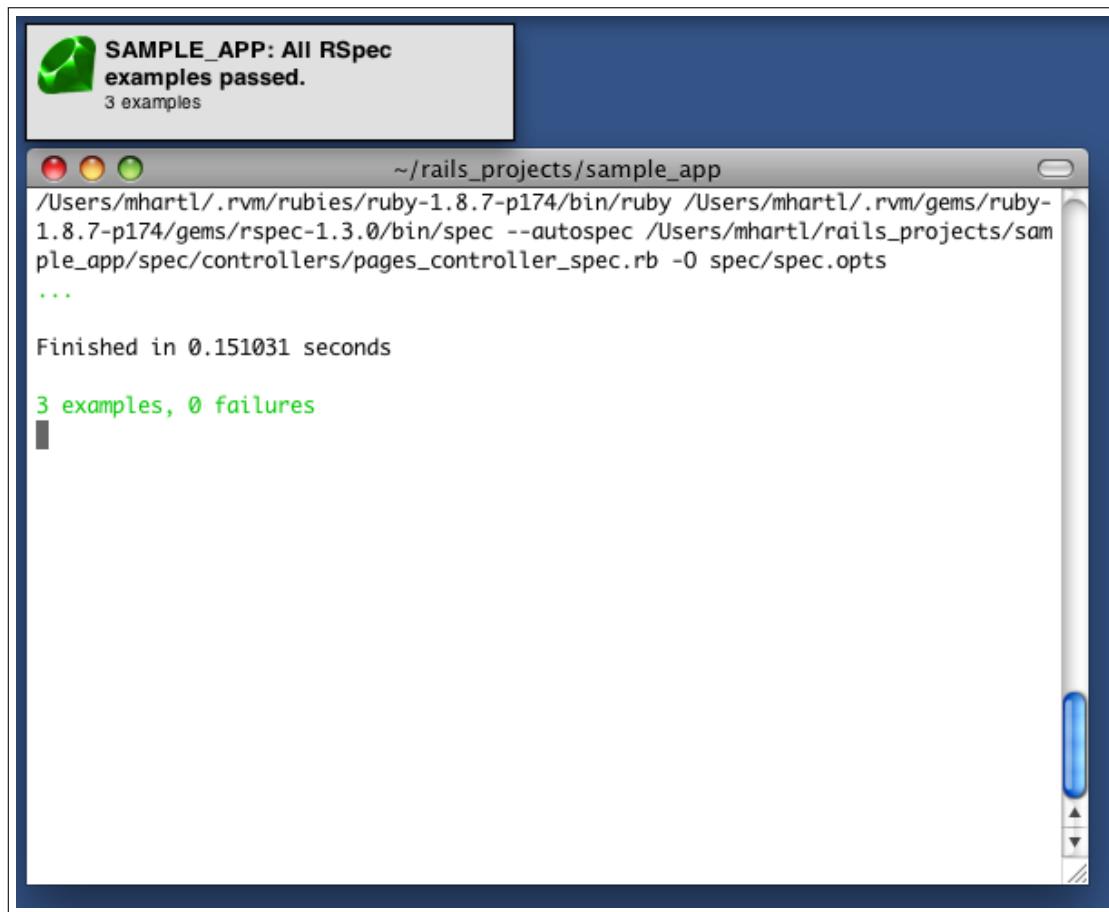


Figure 3.8: Autotest back to Green: all tests passing. ([full size](#))

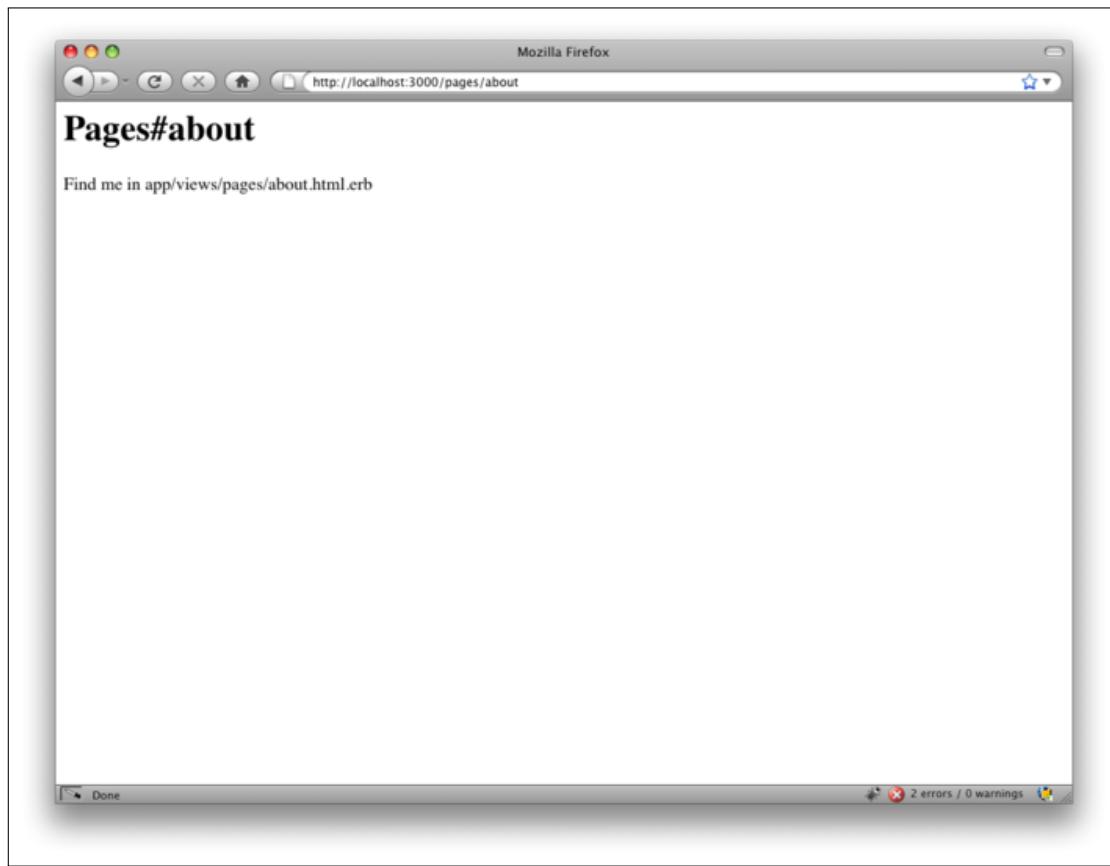


Figure 3.9: The new (and rather raw) About page ([/pages/about](#)). [\(full size\)](#)

Page	URL	Base title	Variable title
Home	/pages/home	"Ruby on Rails Tutorial Sample App"	" Home"
Contact	/pages/contact	"Ruby on Rails Tutorial Sample App"	" Contact"
About	/pages/about	"Ruby on Rails Tutorial Sample App"	" About"

Table 3.1: The (mostly) static pages for the sample app.

computer doesn't care, of course, but humans do, so it is important to keep the code base clean by refactoring frequently. Having a good (passing!) test suite is an invaluable tool in this regard, as it dramatically lowers the probability of introducing bugs while refactoring.

Our sample app is a little too small to refactor right now, but code smell seeps in at every crack, so we won't have to wait long: we'll already get busy refactoring in Section 3.3.3 of this chapter.

3.3 Slightly dynamic pages

Now that we've created the actions and views for some static pages, we'll make them *very slightly* dynamic by adding some content that changes on a per-page basis: we'll have the title of each page change to reflect its content. Whether this represents *truly* dynamic content is debatable, but in any case it lays the necessary foundation for unambiguously dynamic content in Chapter 8.

(If you skipped the TDD material in Section 3.2, be sure to create an About page at this point using the code from Listing 3.9 and Listing 3.10.)

3.3.1 Testing a title change

Our plan is to edit the Home, Contact, and About pages to add the kind of HTML structure we saw in Listing 3.1, including titles that change on each page. It's a delicate matter to decide just which of these changes to test, and in general testing HTML can be quite fragile since content tends to change frequently. We'll keep our tests simple by just testing for the page titles.

By the end of this section, all three of our static pages will have titles of the form "Ruby on Rails Tutorial Sample App — Home", where the last part of the title will vary depending on the page (Table 3.1). We'll build on the tests in Listing 3.8, adding title tests following the model in Listing 3.11.

Listing 3.11. A title test.

```
it "should have the right title" do
  get 'home'
  response.should have_tag("title",
                            "Ruby on Rails Tutorial Sample App | Home")
end
```

This uses the `have_tag` method inside RSpec; the documentation for `have_tag` is surprisingly sparse, but what it does is not surprising at all: the code

```
response.should have_tag("title",
                          "Ruby on Rails Tutorial Sample App | Home")
```

checks to see that the material inside the `<title></title>` tags is "Ruby on Rails Tutorial Sample App | Home". Note that I've broken the material inside `have.tag` into two lines; this tells you something important about Ruby syntax: Ruby doesn't care about newlines.¹² The reason I chose to break the code into pieces is that I prefer to keep lines of source code under 80 characters for legibility.¹³

Adding new tests for each of our three static pages following the model of Listing 3.11 gives us our new Pages controller spec (Listing 3.12).

Listing 3.12. The Pages controller spec with title tests.

`spec/controllers/pages_controller_spec.rb`

```
require 'spec_helper'

describe PagesController do
  integrate_views

  describe "GET 'home'" do
    it "should be successful" do
      get 'home'
      response.should be_success
    end

    it "should have the right title" do
      get 'home'
      response.should have_tag("title",
                                "Ruby on Rails Tutorial Sample App | Home")
    end
  end

  describe "GET 'contact'" do
    it "should be successful" do
      get 'contact'
      response.should be_success
    end

    it "should have the right title" do
      get 'contact'
      response.should have_tag("title",
                                "Ruby on Rails Tutorial Sample App | Contact")
    end
  end

  describe "GET 'about'" do
    it "should be successful" do
      get 'about'
      response.should be_success
    end
  end
end
```

¹²A newline is what comes at the end of a line, starting a, well, new line. In code, it is represented by the character `\n`.

¹³Actually *counting* columns could drive you crazy, which is why many text editors have a visual aid to help you. Consider TextMate, for example; if you take a look back at Figure 1.1, you'll see a small vertical line on the right to help keep code under 80 characters. (It's actually at 78 columns, which gives you a little margin for error.) If you use TextMate, you can find this feature under **View > Wrap Column > 78**.

```
it "should have the right title" do
  get 'about'
  response.should have_tag("title",
    "Ruby on Rails Tutorial Sample App | About")
end
end
end
```

There is one crucial line here that needs emphasis: we are now testing our views inside a controller test, which means we must tell RSpec to *integrate the views*. In other words, by default RSpec just tests actions inside a controller test; if we want it also to render the views, we have to tell it explicitly via the second line:

```
describe PagesController do
  integrate_views

  .
  .
  .
```

With these tests in place, you should run

```
$ spec spec/
```

or use Autotest to verify that our code is now Red (failing tests).

3.3.2 Passing title tests

Now we'll get our title tests to pass, and at the same time add the full HTML structure needed to make valid web pages. (If you were very observant, you might have noticed the red X in the bottom corner of default Home page shown in [Figure 3.9](#); this is the HTML Validator telling us we have a problem, which we'll fix in this section.) Let's start with the Home page ([Listing 3.13](#)), using the same basic HTML skeleton as in the “hello” page from [Listing 3.1](#); we've just added one more element, a series of arguments inside the `html` tag giving more information about the language on the page (`en` for “English”) and the `XML namespace` (`xmlns`).¹⁴

Listing 3.13. The view for the Home page with full HTML structure.

`app/views/pages/home.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
```

¹⁴Don't worry about where these arguments come from or exactly what they mean; I don't know either. I'm just in the habit of looking at the source of well-written web apps, and I'm not shy about borrowing their code.

```

<title>Ruby on Rails Tutorial Sample App | Home</title>
</head>
<body>
  <h1>Sample App Home</h1>
  <p>
    This is the home page for the
    <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
    sample application.
  </p>
</body>
</html>

```

Listing 3.13 uses the title tested for in [Listing 3.12](#):

```

<title>Ruby on Rails Tutorial Sample App | Home</title>

```

As a result, the tests for the Home page should now pass. We're still Red because of the failing Contact and About tests, and we can get to Green with the code in [Listing 3.14](#) and [Listing 3.15](#).

Listing 3.14. The view for the Contact page with full HTML structure.

`app/views/pages/contact.html.erb`

```

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title>Ruby on Rails Tutorial Sample App | Contact</title>
  </head>
  <body>
    <h1>Contact</h1>
    <p>
      Contact Ruby on Rails Tutorial about the sample app at the
      <a href="http://www.railstutorial.org/feedback">feedback page</a>.
    </p>
  </body>
</html>

```

Listing 3.15. The view for the About page with full HTML structure.

`app/views/pages/about.html.erb`

```

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title>Ruby on Rails Tutorial Sample App | About</title>

```

```
</head>
<body>
  <h1>About Us</h1>
  <p>
    <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
    is a project to make a book and screencasts to teach web development
    with <a href="http://rubyonrails.org/">Ruby on Rails</a>. This
    is the sample application for the tutorial.
  </p>
</body>
</html>
```

These example pages introduce the *anchor* tag `a`, which creates links to the given URL (called an “`href`”, or “hypertext reference”, in the context of an anchor tag):

```
<a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
```

You can see the results in [Figure 3.10](#). Note the green checkmark at the lower right of the figure, indicating that the pages are now valid HTML.

3.3.3 Instance variables and Embedded Ruby

We’ve achieved a lot already in this section, generating three valid pages using Rails controllers and actions, but they are purely static HTML and hence don’t show off the power of Rails. Moreover, they suffer from terrible duplication:

- The page titles are almost (but not quite) exactly the same.
- “Ruby on Rails Tutorial Sample App” is common to all three titles.
- The entire HTML skeleton structure is repeated on each page.

Paradoxically, we’ll take the first step toward eliminating duplication by first adding some more: we’ll make the titles of the pages, which are currently quite similar, match *exactly*. This will make it much simpler to remove all the repetition at a stroke.

The technique involves creating *instance variables* inside our actions. Since the Home, Contact, and About page titles have a variable component, we’ll set the variable `@title` (pronounced “at title”) to the appropriate title for each action ([Listing 3.16](#)).

Listing 3.16. The Pages controller with per-page titles.

`app/controllers/pages_controller.rb`

```
class PagesController < ApplicationController
  def home
    @title = "Home"
  end
```

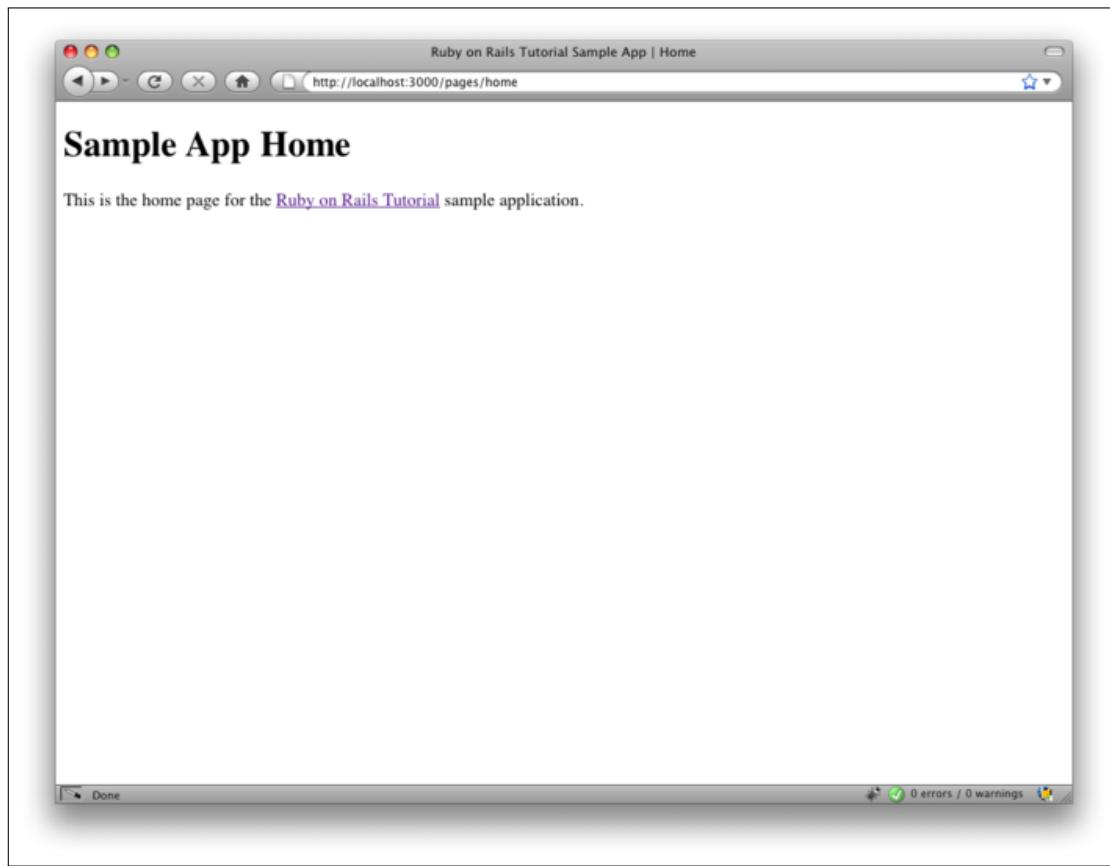


Figure 3.10: A minimal Home page for the sample app ([/pages/home](#)). ([full size](#))

```
def contact
  @title = "Contact"
end

def about
  @title = "About"
end
end
```

A statement such as

```
@title = "Home"
```

is an *assignment*, in this case creating a new variable `@title` with value `"Home"`. The at sign `@` in `@title` indicates that it is an instance variable. Instance variables have a more general meaning in Ruby (see Section 4.2.3), but in Rails their role is primarily to link actions and views: any instance variable defined in the `home` action is automatically available in the `home.html.erb` view, and so on for other action/view pairs.¹⁵

We can see how this works by replacing the literal title “Home” with the contents of the `@title` variable in the `home.html.erb` view (Listing 3.17).

Listing 3.17. The view for the Home page with an Embedded Ruby title.

`app/views/pages/home.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
<head>
  <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
</head>
<body>
  <h1>Sample App Home</h1>
  <p>
    This is the home page for the
    <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
    sample application.
  </p>
</body>
</html>
```

Listing 3.17 is our first example of *Embedded Ruby*, also called *ERb*. (Now you know why HTML views have with the file extension `.html.erb`.) ERb is the primary mechanism in Rails for including dynamic content in web pages.¹⁶ The code

¹⁵In fact, the instance variable is actually visible in *any* view, a fact we’ll make use of in Section 8.2.2.

¹⁶There is a second popular template system called [Haml](#), which I personally love, but it’s not *quite* standard enough yet for use in an introductory tutorial. If there is sufficient interest, I might produce a Rails Tutorial screencast series using Haml for the views. This would also allow for an introduction to [Sass](#), Haml’s sister technology, which if anything is even more awesome than Haml.

```
<%= @title %>
```

indicates using `<%= ... %>` that Rails should insert the contents of the `@title` variable, whatever it may be. When we visit `/pages/home`, Rails executes the body of the `home` action, which makes the assignment `@title = "Home"`, so in the present case

```
<%= @title %>
```

gets replaced with “Home”. Rails then renders the view, using ERb to insert the value of `@title` into the template, which the web server then sends to your browser as HTML. The result is exactly the same as before, only now the variable part of the title is generated dynamically by ERb.

We can verify that all this works by running the tests from [Section 3.3.1](#) and see that they still pass. Then we can make the corresponding replacements for the Contact and About pages ([Listing 3.18](#) and [Listing 3.19](#)).

Listing 3.18. The view for the Contact page with an Embedded Ruby title.

`app/views/pages/contact.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
  </head>
  <body>
    <h1>Contact</h1>
    <p>
      Contact Ruby on Rails Tutorial about the sample app at the
      <a href="http://www.railstutorial.org/feedback">feedback page</a>.
    </p>
  </body>
</html>
```

Listing 3.19. The view for the About page with an Embedded Ruby title.

`app/views/pages/about.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
  </head>
  <body>
    <h1>About Us</h1>
    <p>
```

```
<a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
is a project to make a book and screencasts to teach web development
with <a href="http://rubyonrails.org/">Ruby on Rails</a>. This
is the sample application for the tutorial.
</p>
</body>
</html>
```

As before, the tests still pass.

3.3.4 Eliminating duplication with layouts

Now that we've replaced the variable part of the page titles with instance variables and ERb, each of our pages looks something like this:

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
<head>
  <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
</head>
<body>
  Contents
</body>
</html>
```

In other words, *all* our pages are identical in structure, including even the title (because of Embedded Ruby), with the sole exception of the contents of each page.

Wouldn't it be nice if there were a way to factor out the common elements into some sort of global layout, with the body contents inserted on a per-page basis? Indeed, it would be nice, and Rails happily obliges using a special file called `application.html.erb`, which lives in the `layouts` directory. To capture the structural skeleton, create `application.html.erb` and fill it with the contents of Listing 3.20.

Listing 3.20. The sample application site layout.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
<head>
  <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
</head>
<body>
  <%= yield %>
</body>
</html>
```

Note here the special line

```
<%= yield %>
```

This code is responsible for inserting the contents of each page into the layout. As with `<%= @title %>`, the `<% ... %>` tags indicate Embedded Ruby, and the equals sign in `<%= ... %>` ensures that the results of evaluating the expression are inserted at that exact point in the template. (Don't worry about the meaning of the word "yield" in this context;¹⁷ what matters is that using this layout ensures that visiting the page `/pages/home` converts the contents of `home.html.erb` to HTML and then inserts it in place of `<%= yield %>`.

Of course, our views are still filled with all the HTML structure we just hoisted into the layout, so we have to rip it out, leaving only the interior contents ([Listing 3.21](#), [Listing 3.22](#), and [Listing 3.23](#)).

Listing 3.21. The Home view with HTML structure removed.

`app/views/pages/home.html.erb`

```
<h1>Sample App Home</h1>
<p>
  This is the home page for the
  <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
  sample application.
</p>
```

Listing 3.22. The Contact view with HTML structure removed.

`app/views/pages/contact.html.erb`

```
<h1>Contact</h1>
<p>
  Contact Ruby on Rails Tutorial about the sample app at the
  <a href="http://www.railstutorial.org/feedback">feedback page</a>.
</p>
```

Listing 3.23. The About view with HTML structure removed.

`app/views/pages/about.html.erb`

```
<h1>About Us</h1>
<p>
  <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
  is a project to make a book and screencasts to teach web development
  with <a href="http://rubyonrails.org/">Ruby on Rails</a>. This
  is the sample application for the tutorial.
</p>
```

¹⁷If you've studied Ruby before, you might suspect that Rails is *yielding* the contents to a block, and your suspicion would be correct. But, as far as developing web applications with Rails, it doesn't matter, and I've honestly never given the meaning of `<%= yield %>` a second thought—or even a first one.

With these views defined, the Home, Contact, and About pages are exactly the same as before—i.e., we have successfully refactored them—but have much less duplication. And, as required, the tests still pass.

3.4 Conclusion

Seen from the outside, this chapter hardly accomplished anything: we started with static pages, and ended with...mostly static pages. But appearances are deceiving: by developing in terms of Rails controllers, actions, and views, we are now in a position to add arbitrary amounts of dynamic content to our site. Seeing exactly how this plays out is the task for the rest of this tutorial.

Before moving on, let's take a minute to commit our changes and merge them into the master branch. Back in [Section 3.1.2](#) we created a Git branch for the development of static pages. If you haven't been making commits as we've been moving along, first make a commit indicating that we've reached a stopping point:

```
$ git add .
$ git commit -am "Done with static pages"
```

Then merge the changes back into the master branch using the same technique as in [Section 1.3.5](#):

```
$ git checkout master
$ git merge static-pages
```

Once you reach a stopping point like this, it's usually a good idea to push your code up to a remote repository (which, if you followed the steps in [Section 1.3.4](#), will be GitHub):

```
$ spec spec/
$ git push
```

If you like, at this point you can even deploy the updated application to Heroku:

```
$ spec spec/
$ git push heroku
```

Note that in both cases I've run **spec spec/**, just to be sure that all the tests still pass. Running your tests before pushing or deploying is a good habit to cultivate.

3.5 Exercises

1. Make a Help page for the sample app. First write a test for the existence of a page at the URL `/pages/help`. Then write a second test for the title “Ruby on Rails Tutorial Sample App — Help”. Get your tests to pass, and then fill in the Help page with the content from [Listing 3.24](#).

2. You may have noticed some repetition in the Pages controller spec ([Listing 3.12](#)). In particular, the base title, “Ruby on Rails Tutorial Sample App”, is the same for every title test. Using the RSpec `before(:each)` facility, which executes a block of code before each test case, fill in [Listing 3.25](#) to define a `@base_title` instance variable that eliminates this duplication. (This code uses two new elements: a *symbol*, `:each`, and the string concatenation operator `+`. We’ll learn more about both in [Chapter 4](#), and we’ll see `before(:each)` again in [Section 6.2.1](#).)

Listing 3.24. Code for a proposed Help page.

`app/views/pages/help.html.erb`

```
<h1>Help</h1>
<p>
  Get help on Ruby on Rails Tutorial at the
  <a href="http://www.railstutorial.org/help">Rails Tutorial help page</a>.
  To get help on this sample app, see the
  <a href="http://www.railstutorial.org/book">Rails Tutorial book</a>.
</p>
```

Listing 3.25. The Pages controller spec with a base title.

`spec/controllers/pages_controller_spec.rb`

```
require 'spec_helper'

describe PagesController do
  integrate_views

  before(:each) do
    #
    # Define @base_title here.
    #
  end

  describe "GET 'home'" do
    it "should be successful" do
      get 'home'
      response.should be_success
    end

    it "should have the right title" do
      get 'home'
      response.should have_tag("title", @base_title + " | Home")
    end
  end

  describe "GET 'contact'" do
    it "should be successful" do
      get 'contact'
      response.should be_success
    end
  end
end
```

```
it "should have the right title" do
  get 'contact'
  response.should have_tag("title", @base_title + " | Contact")
end
end

describe "GET 'about'" do
  it "should be successful" do
    get 'about'
    response.should be_success
  end

  it "should have the right title" do
    get 'about'
    response.should have_tag("title", @base_title + " | About")
  end
end
end
```


Chapter 4

Rails-flavored Ruby

Grounded in examples from [Chapter 3](#), this chapter explores some elements of Ruby important for Rails. Ruby is a big language, but fortunately the subset needed to be productive as a Rails developer is relatively small. Moreover, this subset is *different* from the usual approaches to learning Ruby, which is why, if your goal is making dynamic web applications, I recommend learning Rails first, picking up bits of Ruby along the way. To be a Rails *expert*, you need to understand Ruby more deeply, and this book gives you a good foundation for developing that expertise. As noted in [Section 1.1.1](#), after finishing *Rails Tutorial* I suggest reading a pure Ruby book such as *Beginning Ruby*, *The Well-Grounded Rubyist*, or *The Ruby Way*.

This chapter covers a lot of material, and it's OK not to get it all on the first pass. I'll refer back to it frequently in future chapters.

4.1 Motivation

As we saw in the last chapter, it's possible to develop the skeleton of a Rails application, and even start testing it, with essentially no knowledge of the underlying Ruby language. We did this by relying on the generated controller and test code and following the examples we saw there. This situation can't last forever, though, and we'll open this chapter with a couple of additions to the site that bring us face-to-face with our Ruby limitations.

4.1.1 A `title` helper

When we last saw our new application, we had just updated our mostly static pages to use Rails layouts to eliminate duplication in our views ([Listing 4.1](#)).

Listing 4.1. The sample application site layout.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
 "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
 <head>
   <title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
 </head>
 <body>
```

```
<%= yield %>
</body>
</html>
```

This layout works well, but there's one part that could use a little polish. Recall that the title line

```
Ruby on Rails Tutorial Sample App | <%= @title %>
```

relies on the definition of `@title` in the actions, such as

```
class PagesController < ApplicationController

  def home
    @title = "Home"
  end

  .
  .
  .
```

But what if we don't define an `@title` variable? It's a good convention to have a *base title* we use on every page, with an optional variable title if we want to be more specific. We've *almost* achieved that with our current layout, with one wrinkle: as you can see if you delete the `@title` assignment in one of the actions, in the absence of an `@title` variable the title appears as follows:

```
Ruby on Rails Tutorial Sample App |
```

In other words, there's a suitable base title, but there's also a trailing vertical bar character `|` at the end of the title.

One common way to handle this case is to define a *helper*, which is a function designed for use in views. Let's define a `title` helper that returns a base title, "Ruby on Rails Tutorial Sample App", if no `@title` variable is defined, and adds a vertical bar followed by the variable title if `@title` is defined (Listing 4.2).¹

Listing 4.2. Defining a `title` helper.

`app/helpers/application_helper.rb`

```
module ApplicationHelper

  # Return a title on a per-page basis.
  def title
```

¹If a helper is specific to a particular controller, you should put it in the corresponding helper file; for example, helpers for the Pages controller generally go in `app/helpers/pages_helper.rb`. In our case, we expect the `title` helper to be used on all the site's pages, and Rails has a special helper file for this case: `app/helpers/application_helper.rb`.

```
base_title = "Ruby on Rails Tutorial Sample App"
if @title.nil?
  base_title
else
  "#{base_title} | #{@title}"
end
end
end
```

This may look fairly simple to the eyes of an experienced Rails developer, but it's *full* of new Ruby ideas: modules, comments, local variable assignment, booleans, control flow, string interpolation, and return values. We'll cover each of these ideas in this chapter.

Now that we have a helper, we can use it to simplify our layout by replacing

```
<title>Ruby on Rails Tutorial Sample App | <%= @title %></title>
```

with

```
<title><%= title %></title>
```

as seen in Listing 4.3. Note in particular the switch from the instance variable `@title` to the helper method `title` (without the `@` sign). Using Autotest or `spec spec/`, you should verify that the tests from Chapter 3 still pass.

Listing 4.3. The sample application site layout.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
  </head>
  <body>
    <%= yield %>
  </body>
</html>
```

4.1.2 Cascading Style Sheets

There's a second addition to our site that seems simple but adds several new Ruby concepts: including style sheets into our site layout. Though this is a book in web development, not web design, we'll be using cascading style sheets (CSS) to give the sample application some minimal styling, and we'll use the Blueprint CSS framework as a foundation for that styling.

To get started, [download the latest Blueprint CSS](#). (For simplicity, I'll assume you download Blueprint to a **Downloads** directory, but use whichever directory is most convenient.) Using either the command line or a graphical tool, copy the Blueprint CSS directory **blueprint** into the **public/stylesheets** directory, a special directory where Rails keeps stylesheets. On my Mac, the commands looked like this, but your details may differ:

```
$ cp -r ~/Downloads/joshuaclayton-blueprint-css-016c911/blueprint \
> public/stylesheets/
```

Here **cp** is the Unix copy command, and the **-r** flag copies recursively (needed for copying directories). (As mentioned briefly in [Section 3.2.1](#), the tilde `~` means “home directory” in Unix.)

Once you have the stylesheets in the proper directory, Rails provides a helper for including them on our pages using Embedded Ruby ([Listing 4.4](#)).

Listing 4.4. Adding stylesheets to the sample application layout.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
    <%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
    <%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
  </head>
  <body>
    <%= yield %>
  </body>
</html>
```

Let's focus on the new lines:

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
```

These use the built-in Rails helper **stylesheet_link_tag**, which you can read more about at the [Rails API](#).² The first **stylesheet_link_tag** line includes the stylesheet **blueprint/screen.css** for screens (e.g., computer monitors), and the second includes **blueprint/print.css** for printing. (The helper automatically appends the **.css** extension to the filenames if absent, so I've left it off for brevity.) As with the **title** helper, to an experienced Rails developer these lines look simple, but there are at least four new Ruby ideas: built-in Rails methods, method invocation with missing parentheses, symbols, and hashes.

²I've linked my favorite Rails API source, [RailsBrain](#). Its API version is currently a little out of date (2.3.2 instead of 2.3.8), but it doesn't usually matter much. If you think it might matter, you can use the [official Rails API](#), but I prefer the RailsBrain interface. Incidentally, “API” stands for [application programming interface](#).

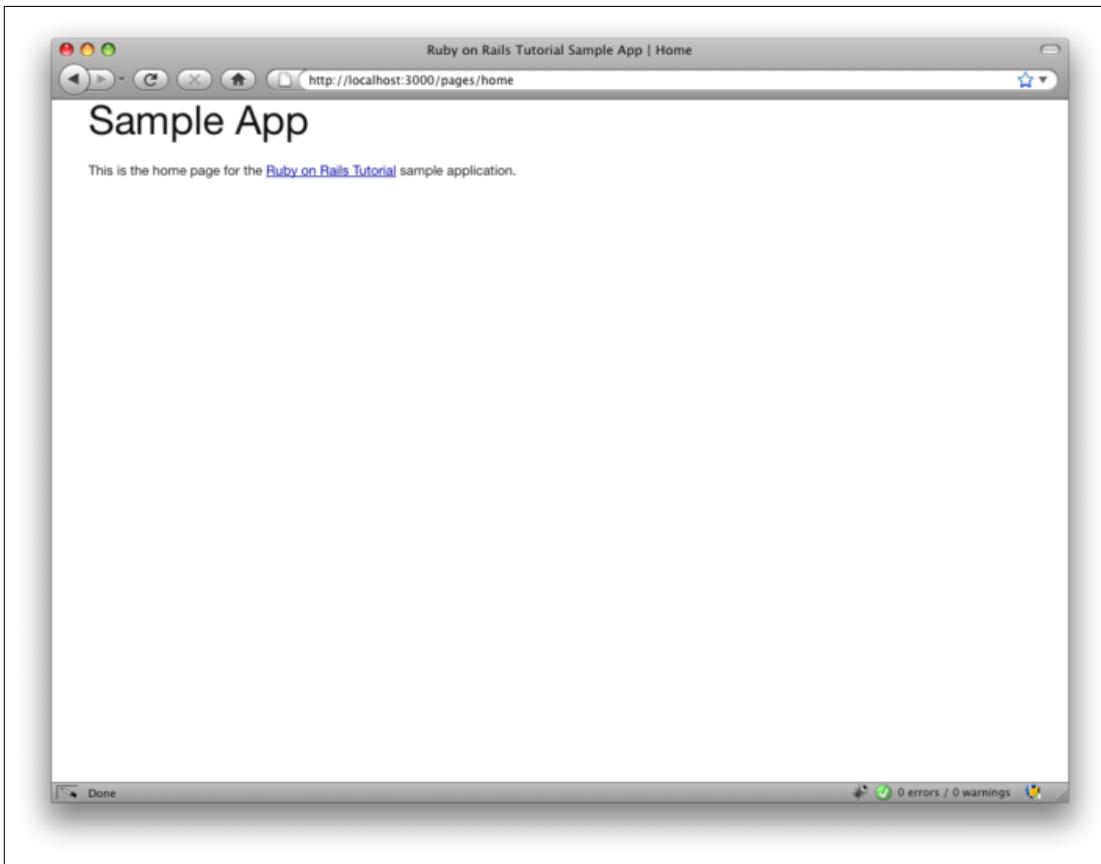


Figure 4.1: The Home page with the new Blueprint stylesheets. ([full size](#))

In this chapter we'll cover these new ideas as well. (We'll see the HTML produced by these stylesheet includes in [Listing 4.5 of Section 4.3.4](#).)

By the way, with the new stylesheets the site doesn't look much different than before, but it's a start (Figure 4.1). We'll build on this foundation starting in [Chapter 5](#).³

4.2 Strings and methods

Our principal tool for learning Ruby will be the *Rails console*, which is a command-line tool for interacting with Rails applications. The console itself is built on top of interactive Ruby (irb), and thus has access to the full power of Ruby. (As we'll see in [Section 4.4.4](#), the console also has access to the Rails environment.) Start the console at the command line as follows:

```
$ script/console  
Loading development environment (Rails 2.3.8)
```

³If you're impatient, feel free to check out the [Blueprint CSS Quickstart tutorial](#).

```
>>
```

By default, the console starts in a *development environment*, which is one of three separate environments defined by Rails (the others are *test* and *production*). This distinction won't be important in this chapter; we'll learn more about environments in [Section 6.3.1](#).

The console is a great learning tool, and you should feel free to explore—don't worry, you (probably) won't break anything. When using the console, type Ctrl-C if you get stuck, or Ctrl-D to exit the console altogether.

Throughout the rest of this chapter, you might find it helpful to consult the [Ruby API](#).⁴ It's packed (perhaps even *too* packed) with information; for example, to learn more about Ruby strings you can look at the [Ruby API entry for the `String` class](#).

4.2.1 Comments

Ruby *comments* start with the pound sign `#` and extend to the end of the line. Ruby (and hence Rails) ignores comments, but they are useful for human readers (including, often, the original author!). In the code

```
# Return a title on a per-page basis.  
def title  
  .  
  .  
  .
```

the first line is a comment indicating the purpose of the subsequent function definition.

You don't ordinarily include comments in console sessions, but for instructional purposes I'll include some comments in what follows, like this:

```
>> 17 + 42    # Integer addition  
=> 59
```

If you follow along in this section typing or copying-and-pasting commands into your own console, you can of course omit the comments if you like; the console will ignore them in any case.

4.2.2 Strings

Strings are probably the most important data structure for web applications, since web pages ultimately consist of strings of characters sent from the server to the browser. Let's get started exploring strings with the console:

```
>> ""          # An empty string  
=> ""  
>> "foo"       # A nonempty string  
=> "foo"
```

⁴I've linked to [RubyBrain](#), sister site to [RailsBrain](#).

These are *string literals* (also, amusingly, called *literal strings*), created using the double quote character `"`. The console prints the result of evaluating each line, which in the case of a string literal is just the string itself.

We can also concatenate strings with the `+` operator:

```
>> "foo" + "bar"      # String concatenation
=> "foobar"
```

Here the result of evaluating `"foo"` plus `"bar"` in the string `"foobar"`.⁵

Another way to build up strings is via *interpolation* using the special syntax `#{}:`⁶

```
>> first_name = "Michael"      # Variable assignment
=> "Michael"
>> "#{first_name} Hartl"      # Variable interpolation
=> "Michael Hartl"
```

Here we've *assigned* the value `"Michael"` to the variable `first_name` and then interpolated it into the string `#{first_name} Hartl`. We could also assign both strings a variable name:

```
>> first_name = "Michael"
=> "Michael"
>> last_name = "Hartl"
=> "Hartl"
>> first_name + " " + last_name      # Concatenation, with a space in between
=> "Michael Hartl"
>> "#{first_name} #{last_name}"      # The equivalent interpolation
=> "Michael Hartl"
```

Note that the final two expressions are equivalent, but I prefer the interpolated version; having to add the single space `" "` seems a bit awkward.

Printing

To *print* a string, the most commonly used Ruby function is `puts` (pronounced “put ess”, for “put string”):

```
>> puts "foo"      # put string
foo
=> nil
```

The `puts` method operates as a *side-effect*: the expression `puts "foo"` prints the string to the screen and then returns *literally nothing*: `nil` is a special Ruby value for “nothing at all”. (In what follows, I'll sometimes suppress the `=> nil` part for simplicity.)

⁵For more on the origins of “foo” and “bar”—and, in particular, the possible *non*-relation of “foobar” to “FUBAR”—see the [Jargon File entry on “foo”](#).

⁶Programmers familiar with Perl or PHP should compare this to the automatic interpolation of dollar sign variables in expressions like `"foo $bar"`.

Using `puts` automatically appends a newline character `\n` to the output; the related `print` method does not:

```
>> print "foo"      # print string (same as puts, but without the newline)
foo=> nil
>> print "foo\n"   # Same as puts "foo"
foo
=> nil
```

Single-quoted strings

All the examples so far have used *double-quoted strings*, but Ruby also supports *single-quoted strings*. For many uses, the two types of strings are effectively identical:

```
>> 'foo'           # A single-quoted string
=> "foo"
>> 'foo' + 'bar'
=> "foobar"
```

There's an important difference, though; Ruby won't interpolate into single-quoted strings

```
>> '#{foo} bar'    # Single-quoted strings don't allow interpolation
"\#{foo} bar"
```

Note how the console returns values using double-quoted strings, which requires a backslash to *escape* characters like `#`.

If double-quoted strings can do everything that single-quoted strings can do, and interpolate to boot, what's the point of single-quoted strings? They are often useful because they are truly literal, and contain exactly the characters you type. For example, the “backslash” character is special on most systems, as in the literal newline `\n`. If you want a variable to contain a literal backslash, single quotes make it easier:

```
>> '\n'           # A literal newline
=> "\\\n"
```

As with the `#` character in our previous example, Ruby needs to escape the backslash with an additional backslash; inside double-quoted strings, a literal backslash is represented with *two* backslashes.

For a small example like this, there's not much savings, but if there are lots of things to escape it can be a real help:

```
>> 'Newlines (\n) and tabs (\t) both use the backslash character \.'
=> "Newlines (\\\n) and tabs (\\\t) both use the backslash character \\".
```

4.2.3 Objects and message passing

Everything in Ruby, including strings and even `nil`, is an *object*. We'll see the technical meaning of this in Section 4.4.2, but I don't think anyone ever understood objects by reading the definition in a book; you have to build up your intuition for objects by seeing lots of examples.

It's easier to describe what objects *do*, which is respond to messages. An object like a string, for example, can respond to the message `length`, which returns the number of characters in the string:

```
>> "foobar".length          # Passing the "length" message to a string
=> 6
```

Typically, the messages that get passed to objects are *methods*, which are functions defined on those objects.⁷ Strings also respond to the `empty?` method:

```
>> "foobar".empty?
=> false
>> "".empty?
=> true
```

Note the question mark at the end of the `empty?` method. This is a Ruby convention indicating that the return value is *boolean*: `true` or `false`. Booleans are especially useful for *control flow*:

```
>> s = "foobar"
>> if s.empty?
>>   "The string is empty"
>> else
>>   "The string is nonempty"
>> end
=> "The string is nonempty"
```

Booleans can also be combined using the `&&` ("and"), `||` ("or"), and `!` ("not") operators:

```
>> x = "foo"
=> "foo"
>> y = ""
=> ""
>> puts "Both strings are empty" if x.empty? && y.empty?
=> nil
>> puts "One of the strings is empty" if x.empty? || y.empty?
=> "One of the strings is empty"
=> nil
>> puts "x is not empty" if !x.empty?
=> "x is not empty"
```

⁷Apologies in advance for switching haphazardly between *function* and *method* throughout this chapter; in Ruby, they're the same thing: all methods are functions, and all functions are methods, because everything is an object.

Since everything in Ruby is an object, it follows that `nil` is an object, so it too can respond to methods. One example is the `to_s` method that can convert virtually any object to a string:

```
>> nil.to_s
""
```

This certainly appears to be an empty string, as we can verify by *chaining* the messages we pass to `nil`:

```
>> nil.empty?
NoMethodError: You have a nil object when you didn't expect it!
You might have expected an instance of Array.
The error occurred while evaluating nil.empty?
>> nil.to_s.empty?      # Message chaining
true
```

We see here that the `nil` object doesn't itself respond to the `empty?` method, but `nil.to_s` does.

There's a special method for testing for `nil`-ness, which you might be able to guess:

```
>> "foo".nil?
=> false
>> "".nil?
=> false
>> nil.nil?
=> true
```

If you look back at Listing 4.2, you'll see that the `title` helper tests to see if `@title` is `nil` using the `nil?` method. This is a hint that there's something special about instance variables (variables with an `@` sign), which can best be understood by contrasting them with ordinary variables. For example, suppose we enter `title` and `@title` variables at the console without defining them first:

```
>> title          # Oops! We haven't defined a title variable.
NameError: undefined local variable or method `title'
>> @title         # An instance variable in the console
=> nil
>> puts "There is no such instance variable." if @title.nil?
There is no such instance variable.
=> nil
>> "#{@title}"   # Interpolating @title when it's nil
""
```

You can see from this example that Ruby complains if we try to evaluate an undefined local variable, but issues no such complaint for an instance variable; instead, instance variables are `nil` if not defined. This code also explains why the code

```
Ruby on Rails Tutorial Sample App | <%= @title %>
```

becomes

```
Ruby on Rails Tutorial Sample App |
```

when `@title` is `nil`: Embedded Ruby inserts the string corresponding to the given variable, and the string corresponding to `nil` is `" "`.

The last example also shows an alternate use of the `if` keyword: Ruby allows you to write a statement that is evaluated only if the statement following `if` is true. There's a complementary `unless` keyword that works the same way:

```
>> string = "foobar"
>> puts "The string '#{string}' is nonempty." unless string.empty?
The string 'foobar' is nonempty.
=> nil
```

It's worth noting that the `nil` object is special, in that it is the *only* Ruby object that is false in a boolean context, apart from `false` itself:

```
>> if nil
>>   true
>> else
>>   false          # nil is false
>> end
=> false
```

In particular, all other Ruby objects are *true*, even 0:

```
>> if 0
>>   true          # 0 (and everything other than nil and false itself) is true
>> else
>>   false
>> end
=> true
```

4.2.4 Method definitions

The console allows us to define methods the same way we did with the `home` action from Listing 3.3 or the `title` helper from Listing 4.2. (Defining methods in the console is a bit cumbersome, and ordinarily you would use a file, but it's convenient for demonstration purposes.) For example, let's define a function

`string_message` that takes a single *argument* and returns a message based on whether the argument is empty or not:

```
>> def string_message(string)
>>   if string.empty?
>>     "It's an empty string!"
>>   else
>>     "The string is nonempty."
>>   end
>> end
=> nil
>> puts string_message("")
It's an empty string!
>> puts string_message("foobar")
The string is nonempty.
```

Note that Ruby functions have an *implicit return*, meaning they return the last statement evaluated—in this case, one of the two message strings, depending on whether the method’s argument `string` is empty or not. Ruby also has an explicit return option; the following function is equivalent to the one above:

```
>> def string_message(string)
>>   return "It's an empty string!" if string.empty?
>>   return "The string is nonempty."
>> end
```

The alert reader might notice at this point that the second `return` here is actually unnecessary—being the last expression in the function, the string "**The string is nonempty.**" will be returned regardless of the `return` keyword, but using `return` in both places has a pleasing symmetry to it.

4.2.5 Back to the `title` helper

We are now in a position to understand the `title` helper from Listing 4.2:⁸

```
module ApplicationHelper
  # Return a title on a per-page basis.
  def title
    base_title = "Ruby on Rails Tutorial Sample App"      # Documentation comment
    if @title.nil?                                       # Method definition
      base_title                                         # Variable assignment
    else                                                 # Boolean test for nil
      "#{base_title} | #{@title}"                         # Implicit return
    end
  end
end
```

⁸Well, there will still be *one* thing left that we don’t understand, which is how Rails ties this all together: mapping URLs to actions, making the `title` helper available in views, etc. This is an interesting subject, and I encourage you to investigate it further, but knowing exactly *how* Rails works is not necessary to *using* Rails. (For a deeper understanding, I recommend *The Rails Way* by Obie Fernandez.)

These elements—function definition, variable assignment, boolean tests, control flow, and string interpolation—come together to make a compact helper method for use in our site layout. The final element is **module ApplicationHelper**: code in Ruby modules can be *mixed in* to Ruby classes. When writing ordinary Ruby, you often write modules and include them explicitly yourself, but in this case Rails handles the inclusion automatically for us. The result is that the **title** method is **automagically** available in all our views.

4.3 Other data structures

Though web apps are ultimately about strings, actually *making* those strings requires using other data structures as well. In this section we'll learn about some Ruby data structures important for writing Rails applications.

4.3.1 Arrays and ranges

An array is just a list of elements in a particular order. We haven't discussed arrays yet in *Rails Tutorial*, but understanding them gives a good foundation for understanding hashes (Section 4.3.3) and for aspects of Rails data modeling (such as the **has_many** association seen in Section 2.3.3 and covered more in Section 11.1.2).

So far we've spent a lot of time understanding strings, and there's a natural way to get from string to arrays using the **split** method:

```
>> "foo bar     baz".split      # Split a string into a three-element array
=> ["foo", "bar", "baz"]
```

The result of this operation is an array of three strings. By default, **split** divides a string into an array by splitting on whitespace, but you can split on nearly anything else:

```
>> "fooxbarxbazx".split('x')
=> ["foo", "bar", "baz"]
```

As is conventional in most computer languages, Ruby arrays are *zero-offset*, which means that the first element in the array has index 0, the second has index 1, and so on:

```
>> a = [42, 8, 17]
=> [42, 8, 17]
>> a[0]                  # Ruby uses square brackets for array access.
=> 42
>> a[1]
=> 8
>> a[2]
=> 17
>> a[-1]                 # Indices can even be negative!
=> 17
```

We see here that Ruby uses square brackets to access array elements. In addition to this bracket notation, Ruby offers synonyms for some commonly accessed elements:

```
>> a                      # Just a reminder of what 'a' is
=> [42, 8, 17]
>> a.first
=> 42
>> a.second
=> 8
>> a.last
=> 17
>> a.last == a[-1]      # Comparison using ==
=> true
```

This last line introduces the equality comparison operator `==`, which Ruby shares with many other languages, along with the associated `!=` (“not equal”), etc.:

```
>> x = a.length          # Like strings, arrays respond to the 'length' method.
=> 3
>> x == 3
=> true
>> x == 1
=> false
>> x != 1
=> true
>> x >= 1
=> true
>> x < 1
=> false
```

In addition to `length` (seen in the first line above), arrays respond to a wealth of other methods:

```
>> a.sort
=> [8, 17, 42]
>> a.reverse
=> [17, 8, 42]
>> a.shuffle
=> [17, 42, 8]
```

By the way, the `shuffle` method is available only on Ruby 1.8.7 or later; if you’re using 1.8.6, you can use

```
>> a.sort_by { rand }
```

instead.

You can also add to arrays with the “push” operator, `<<`:

```
>> a << 7                      # Pushing 7 onto an array
[42, 8, 17, 7]
>> a << "foo" << "bar"       # Chaining array pushes
[42, 8, 17, 7, "foo", "bar"]
```

This last example shows that you can chain pushes together, and also that, unlike arrays in many other languages, Ruby arrays can contain a mixture of different types (in this case, integers and strings).

Before we saw `split` convert a string to an array. We can also go the other way with the `join` method:

```
>> a
[42, 8, 17, 7, "foo", "bar"]
>> a.join                      # Join on nothing
=> "428177foobar"
>> a.join(',')
# Join on comma-space
=> "42, 8, 17, 7, foo, bar"
```

Closely related to arrays are *ranges*, which can probably most easily be understood by converting them to arrays using the `to_a` method:

```
>> 0..9
=> 0..9
>> 0..9.to_a                  # Oops, call to_a on 9
ArgumentError: bad value for range
>> (0..9).to_a                # Use parentheses to call to_a on the range
=> [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Though `0..9` is a valid range, the second expression above shows that we need to add parentheses to call a method on it.

Ranges are useful for pulling out array elements:

```
>> a = %w[foo bar baz quux]      # Use %w to make a string array
["foo", "bar", "baz", "quux"]
>> a[0..2]
=> ["foo", "bar", "baz"]
```

Ranges also work with characters:

```
>> ('a'..'e').to_a
=> ["a", "b", "c", "d", "e"]
```

4.3.2 Blocks

Both arrays and ranges respond to a host of methods that accept *blocks*, which are simultaneously one of Ruby's most powerful and most confusing features:

```
>> (1..5).each { |i| puts 2 * i }
2
4
6
8
10
=> 1..5
```

This code calls the `each` method on the range `(1..5)` and passes it the block `{ |i| puts 2 * i }`. The vertical bars around the variable name in `|i|` are Ruby syntax for a block variable, and it's up to the method to know what to do with the block; in this case, the range's `each` method can handle a block with a single local variable, which we've called `i`, and it just executes the block for each value in the range.

Curly braces are one way to indicate a block, but there is a second way as well:

```
>> (1..5).each do |i|
?>   puts 2 * i
>> end
2
4
6
8
10
=> 1..5
```

Blocks can be more than one line, and often are. In *Rails Tutorial* we'll follow the common convention of using curly braces only for short one-line blocks and the `do..end` syntax for longer one-liners and for multi-line blocks:

```
>> (1..5).each do |number|
?>   puts 2 * number
?>   puts '--'
>> end
2
--
4
--
6
--
8
--
10
--
=> 1..5
```

Here I've used **number** in place of **i** just to emphasize that any variable name will do.

Unless you already have a substantial programming background, there is no shortcut to understanding blocks; you just have to see them a lot, and eventually you'll get used to them.⁹ Luckily, humans are quite good at making generalizations from concrete examples; here are a few more, including a couple using the **map** method:

```
>> 3.times { puts "Betelgeuse!" }      # 3.times takes a block with no variables.
"Betelgeuse!"
"Betelgeuse!"
"Betelgeuse!"
=> 3
>> (1..5).map { |i| i**2 }           # The ** notation is for 'power'.
=> [1, 4, 9, 16, 25]
>> %w[a b c]                      # Recall that %w makes string arrays.
=> ["a", "b", "c"]
>> %w[a b c].map { |char| char.upcase }
=> ["A", "B", "C"]
```

As you can see, the **map** method returns the result of applying the given block to each element in the array or range.

By the way, we're now in a position to understand the line of Ruby I threw into Section 1.4.4 to generate random subdomains:¹⁰

```
('a'..'z').to_a.shuffle[0..7].join
```

Let's build it up step-by-step:

```
>> ('a'..'z').to_a                  # An alphabet array
=> ["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l", "m", "n", "o",
    "p", "q", "r", "s", "t", "u", "v", "w", "x", "y", "z"]
>> ('a'..'z').to_a.shuffle        # Shuffle it!
=> ["c", "g", "l", "k", "h", "z", "s", "i", "n", "d", "y", "u", "t", "j", "q",
    "b", "r", "o", "f", "e", "w", "v", "m", "a", "x", "p"]
>> ('a'..'z').to_a.shuffle[0..7]   # Pull out the first eight elements.
=> ["f", "w", "i", "a", "h", "p", "c", "x"]
>> ('a'..'z').to_a.shuffle[0..7].join # Join them together to make one string.
=> "mznpwybj"
```

4.3.3 Hashes and symbols

Hashes are essentially a generalization of arrays: you can think of hashes as basically like arrays, but not limited to integer indices. (In fact, some languages, especially Perl, call hashes *associative arrays* for this reason.) Instead, hash indices, or *keys*, can be almost any object. For example, we can use strings as keys:

⁹Programming experts, on the other hand, might benefit from knowing that blocks are *closures*, which are one-shot anonymous functions with data attached.

¹⁰Recall that **shuffle** isn't available in Ruby 1.8.6; you have to use **sort_by { rand }** instead.

```

>> user = {}                                # {} is an empty hash
=> {}
>> user["first_name"] = "Michael"          # Key "first_name", value "Michael"
=> "Michael"
>> user["last_name"] = "Hartl"              # Key "last_name", value "Hartl"
=> "Hartl"
>> user["first_name"]                      # Element access is like arrays
=> "Michael"
>> user                                    # A literal representation of the hash
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}

```

Hashes are indicated with curly braces containing key-value pairs; a pair of braces with no key-value pairs—i.e., `{}`—is an empty hash. It's important to note that the curly braces for hashes have nothing to do with the curly braces for blocks. (Yes, this can be confusing.) Though hashes resemble arrays, one important difference is that hashes don't generally guarantee keeping their elements in a particular order.¹¹ If order matters, use an array.

Instead of defining hashes one item at a time using square brackets, it's easy to use their literal representation:

```

>> user = { "first_name" => "Michael", "last_name" => "Hartl" }
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}

```

Here I've used the usual Ruby convention of putting an extra space at the two ends of the hash—a convention ignored by the console output. (Don't ask me why the spaces are conventional; probably some early influential Ruby programmer liked the look of the extra spaces, and the convention stuck.)

So far we've used strings as hash keys, but in Rails it is much more common to use *symbols* instead. Symbols look kind of like strings, but prefixed with a colon instead of surrounded by quotes. For example, `:name` is a symbol. You can think of symbols as basically strings without all the extra baggage:¹²

```

>> "name".length
4
>> :name.length
NoMethodError: undefined method `length' for :name:Symbol
>> "foobar".reverse
=> "rabooft"
>> :foobar.reverse
NoMethodError: undefined method `reverse' for :foobar:Symbol

```

Symbols are a special Ruby data type shared with very few other languages, so they may seem weird at first, but Rails uses them a lot, so you'll get used to them fast.

In terms of symbols as hash keys, we can define a `user` hash as follows:

¹¹ Apparently Ruby 1.9 guarantees that hashes keep their elements in the same order entered, but it would be unwise ever to count on a particular ordering.

¹² As a result of having less baggage, symbols are easier to compare to each other; strings need to be compared character by character, while symbols can be compared all in one go. This makes them ideal for use as hash keys.

```
>> user = { :name => "Michael Hartl", :email => "michael@example.com" }
=> { :name=>"Michael Hartl", :email=>"michael@example.com" }
>> user[:name]                      # Access the value corresponding to :name.
=> "Michael Hartl"
>> user[:password]                 # Access the value of an undefined key.
=> nil
```

We see here from the last example that the hash value for an undefined key is simply `nil`.

Hash values can be virtually anything, even other hashes:

```
>> params = {}           # Define a hash called 'params' (short for 'parameters').
=> {}
>> params[:user] = { :name => "Michael Hartl", :email => "mhartl@example.com" }
=> { :name=>"Michael Hartl", :email=>"mhartl@example.com" }
>> params
=> { :user=>{ :name=>"Michael Hartl", :email=>"mhartl@example.com" } }
>> params[:user][:email]
=> "mhartl@example.com"
```

These sorts of hashes-of-hashes are heavily used by Rails, as we'll see starting in [Section 8.2](#).

As with arrays and ranges, hashes respond to the `each` method. For example, consider a hash named `flash` with keys for two conditions, `:success` and `:error`:

```
>> flash = { :success => "It worked!", :error => "It failed. :-(" }
=> { :success=>"It worked!", :error=>"It failed. :-(" }
>> flash.each do |key, value|
?>   puts "Key #{key.inspect} has value #{value.inspect}"
>> end
Key :success has value "It worked!"
Key :error has value "It failed. :-("
```

Note that, while the `each` method for arrays takes a block with only one variable, `each` for hashes takes two, a `key` and a `value`. Thus, the `each` method for a hash iterates through the hash one key-value *pair* at a time.

The last example uses the useful `inspect` method, which returns a string with a literal representation of the object it's called on:

```
>> puts flash          # Put the flash hash as a string (with ugly results).
successIt worked!errorIt failed. :-(
>> puts flash.inspect    # Put the flash hash as a pretty string
{:success=>"It worked!", :error=>"It failed. :-("}
>> puts :name, :name.inspect
name
:name
```

```
>> puts "It worked!", "It worked!".inspect
It worked!
"It worked!"
```

By the way, using `inspect` to print an object is common enough that there's a shortcut for it, the `p` function:

```
>> p flash          # Same as 'puts flash.inspect'
{:success=>"It worked!", :error=>"It failed. :-("}
```

4.3.4 CSS revisited

It's time now to revisit the lines from Listing 4.4 used in the layout to include the cascading style sheets:

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print',   :media => 'print' %>
```

We are now nearly in a position to understand this. As mentioned briefly in Section 4.1.2, Rails defines a special function to include stylesheets, and

```
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

is a call to this function. But there are two mysteries. First, where are the parentheses? In Ruby, they are optional; these two lines are equivalent:

```
# Parentheses on function calls are optional.
stylesheet_link_tag('blueprint/screen', :media => 'screen')
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

Second, the `:media` argument sure looks like a hash, but where are the curly braces? When hashes are the *last* argument in a function call, the curly braces are optional; these two lines are equivalent:

```
# Curly braces on final hash arguments are optional.
stylesheet_link_tag 'blueprint/screen', { :media => 'screen' }
stylesheet_link_tag 'blueprint/screen', :media => 'screen'
```

So, we see now that each of the lines

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print',   :media => 'print' %>
```

calls the `stylesheet_link_tag` function with two arguments: a string, indicating the path to the stylesheet, and a hash, indicating the media type ('`screen`' for the computer screen and '`print`' for a printed version). Because of the `<%= %>` brackets, the results are inserted into the template by ERb, and if you view the source of the page in your browser you should see the HTML needed to include a stylesheet (Listing 4.5).¹³

Listing 4.5. The HTML source produced by the CSS includes.

```
<link href="/stylesheets/blueprint/screen.css" media="screen" rel="stylesheet"
      type="text/css" />
<link href="/stylesheets/blueprint/print.css" media="print" rel="stylesheet"
      type="text/css" />
```

4.4 Ruby classes

We've said before that everything in Ruby is an object, and in this section we'll finally get to define some of our own. Ruby, like many object-oriented languages, uses *classes* to organize methods; these classes are then *instantiated* to create objects. If you're new to object-oriented programming, this may sound like gibberish, so let's look at some concrete examples.

4.4.1 Constructors

We've seen lots of examples of using classes to instantiate objects, but we have yet to do so explicitly. For example, we instantiated a string using the double quote characters, which is a *literal constructor* for strings:

```
>> s = "foobar"          # A literal constructor for strings using double quotes
=> "foobar"
>> s.class
=> String
```

We see here that strings respond to the method `class`, and simply return the class they belong to.

Instead of using a literal constructor, we can use the equivalent *named constructor*, which involves calling the `new` method on the class name:

```
>> s = String.new("foobar")    # A named constructor for a string
=> "foobar"
>> s.class
=> String
>> s == "foobar"
=> true
```

This is equivalent to the literal constructor, but it's more explicit about what we're doing.

Arrays work the same way as strings:

¹³You may see some funky numbers, like `?1257465942`, after the CSS filenames. These are inserted by Rails to ensure that browsers reload the CSS when it changes on the server.

```
>> a = Array.new([1, 3, 2])
=> [1, 3, 2]
```

Hashes, in contrast, are different. While the array constructor `Array.new` takes an initial value for the array, `Hash.new` takes a *default* value for the hash, which is the value of the hash for a nonexistent key:

```
>> h = Hash.new
=> {}
>> h[:foo]           # Try to access the value for the nonexistent key :foo.
=> nil
>> h = Hash.new(0)    # Arrange for nonexistent keys to return 0 instead of nil.
=> {}
>> h[:foo]
=> 0
```

4.4.2 Class inheritance

When learning about classes, it's useful to find out the *class hierarchy* using the `superclass` method:

```
>> s = String.new("foobar")
=> "foobar"
>> s.class           # Find the class of s.
=> String
>> s.class.superclass # Find the superclass of String.
=> Object
>> s.class.superclass.superclass # Find the superclass of Object (it's nil!).
=> nil
```

We see here that the superclass of `String` is `Object`, but `Object` has no superclass. This is the technical meaning of “everything in Ruby is an object”.

To understand classes a little more deeply, there's no substitute for making one of our own. Let's make a `Word` class with a `palindrome?` method that returns `true` if the word is the same spelled forward and backward:

```
>> class Word
>>   def palindrome?(string)
>>     string == string.reverse
>>   end
>> end
=> nil
```

We can use it as follows:

```
>> w = Word.new                      # Make a new Word object
=> #<Word:0x22d0b20>
>> w.palindrome?("foobar")
=> false
>> w.palindrome?("level")
=> true
```

If this example strikes you as a bit contrived, good; this is by design. It's odd to create a new class just to create a method that takes a string as an argument. Since a word *is a* string, it's more natural to have our **Word** class *inherit* from **String** (you should exit the console and re-enter it to clear out the old definition of **Word**):

```
>> class Word < String                # Word inherits from String.
>>   # Return true if the string is its own reverse.
>>   def palindrome?
>>     self == self.reverse          # self is the string itself.
>>   end
>> end
=> nil
```

Here **Word < String** is the Ruby syntax for inheritance (discussed briefly in [Section 3.1.2](#)), which ensures that, in addition to the new **palindrome?** method, words also have all the same methods as strings:

```
>> s = Word.new("level")      # Make a new Word, initialized with "level".
=> "level"
>> s.palindrome?            # Words have the palindrome? method.
=> true
>> s.length                 # Words also inherit all the normal string methods.
=> 5
```

Since the **Word** class inherits from **String**, we can use the console to see the class hierarchy explicitly:

```
>> s.class
=> Word
>> s.class.superclass
=> String
```

Note that checking that the word is its own reverse involves accessing the word inside the Word class. Ruby allows us to do this using the **self** keyword: inside the Word class, **self** is the object itself, which means we can use

```
self == self.reverse
```

to check if the word is a palindrome.¹⁴

4.4.3 Modifying built-in classes

While inheritance is a powerful idea, in the case of palindromes it might be even more natural to add the `palindrome?` method to the `String` class itself, so that (among other things) we can call `palindrome?` on a string literal, which we currently can't do:

```
>> "level".palindrome?
NoMethodError: undefined method `palindrome?' for "level":String
```

Somewhat amazingly, Ruby lets you do just this; Ruby classes can be *opened* and modified, allowing ordinary mortals such as ourselves to add methods to them:¹⁵

```
>> class String
>>   # Return true if the string is its own reverse.
>>   def palindrome?
>>     self == self.reverse
>>   end
>> end
=> nil
>> "deified".palindrome?
=> true
```

(I don't know which is cooler: that Ruby lets you add methods to built-in classes, or that "`deified`" is a palindrome.)

Modifying built-in classes is a powerful technique, but with great power comes great responsibility, and it's considered bad form to add methods to built-in classes without having a *really* good reason for doing so. Rails does have some good reasons; for example, in web applications we often want to prevent variables from being `blank`—e.g., a user's name should be something other than spaces and other `whitespace`—so Rails adds a `blank?` method to Ruby. Since the Rails console automatically includes the Rails extensions, we can see an example here (this won't work in plain `irb`):

```
>> "".blank?
=> true
>> " ".empty?
=> false
>> " ".blank?
=> true
>> nil.blank?
=> true
```

¹⁴For more on Ruby classes and the `self` keyword, see the [RailsTips](#) post on [Class and Instance Variables in Ruby](#).

¹⁵For those familiar with JavaScript, this functionality is comparable to using a built-in class prototype object to augment the class. (Thanks to reader [Erik Eldridge](#) for pointing this out.)

We see that a string of spaces is not *empty*, but it is *blank*. Note also that `nil` is blank; since `nil` isn't a string, this is a hint that Rails actually adds `blank?` to `String`'s base class, which (as we saw at the beginning of this section) is `Object` itself. We'll see some other examples of Rails additions to Ruby classes in Section 9.3.3.

4.4.4 A controller class

All this talk about classes and inheritance may have triggered a flash of recognition, because we have seen both before, in the Pages controller ([Listing 3.16](#)):

```
class PagesController < ApplicationController

  def home
    @title = "Home"
  end

  def contact
    @title = "Contact"
  end

  def about
    @title = "About"
  end
end
```

You're now in a position to appreciate, at least vaguely, what this code means: `PagesController` is a class that inherits from `ApplicationController`, and comes equipped with `home`, `contact`, and `about` methods, each of which defines the instance variable `@title`. Since each Rails console session loads the local Rails environment, we can even create a controller explicitly and examine its class hierarchy:¹⁶

```
>> controller = PagesController.new
=> #<PagesController:0x22855d0>
>> controller.class
=> PagesController
>> controller.class.superclass
=> ApplicationController
>> controller.class.superclass.superclass
=> ActionController::Base
>> controller.class.superclass.superclass.superclass
=> Object
```

We can even call the actions, which are just methods:

¹⁶You don't have to know what each class in this hierarchy does. I don't know what they all do, and I've been programming in Ruby on Rails since 2005. This means either that (a) I'm grossly incompetent or (b) you can be a skilled Rails developer without knowing all its innards. I hope for both our sakes that it's the latter. :-)

```
>> controller.home
=> "Home"
```

This return value of "`Home`" comes from the assignment `@title = "Home"` in the `home` action.

But wait—actions don’t have return values, at least not ones that matter. The point of the `home` action, as we saw in [Chapter 3](#), is to render a web page. And I sure don’t remember ever calling `PagesController.new` anywhere. What’s going on?

What’s going on is that Rails is *written in Ruby*, but Rails isn’t Ruby. Some Rails classes are used like ordinary Ruby objects, but some are just *grist* for Rails’ magic mill. Rails is *sui generis*, and should be studied and understood separately from Ruby. This is why, if your principal programming interest is writing web applications, I recommend learning Rails first, then learning Ruby, then looping back to Rails.

4.4.5 A user class

We end our tour of Ruby with a complete class of our own, a `User` class that anticipates the User model coming up in [Chapter 6](#).

So far we’ve entered class definitions at the console, but this quickly becomes tiresome; instead, create the file `example_user.rb` in your Rails root directory and fill it with the contents of [Listing 4.6](#).

Listing 4.6. Code for an example user.

`example_user.rb`

```
class User
  attr_accessor :name, :email

  def initialize(attributes = {})
    @name = attributes[:name]
    @email = attributes[:email]
  end

  def formatted_email
    "#{@name} <#{@email}>"
  end
end
```

There’s quite a bit going on here, so let’s take it step by step. The first line,

```
attr_accessor :name, :email
```

creates *attribute accessors* corresponding to a user’s name and email address. This creates “getter” and “setter” methods that allow us to retrieve (get) and assign (set) `@name` and `@email` instance variables.

The first method, `initialize`, is special in Ruby: it’s the method called when we execute `User.new`. This particular `initialize` takes one argument, `attributes`:

```
def initialize(attributes = {})
  @name = attributes[:name]
  @email = attributes[:email]
end
```

Here the `attributes` variable has a *default value* equal to the empty hash, so that we can define a user with no name or email address (recall from Section 4.3.3 that hashes return `nil` for nonexistent keys, so `attributes[:name]` will be `nil` if there is no `:name` key, and similarly for `attributes[:email]`).

Finally, our class defines a method called `formatted_email` method that uses the values of the assigned `@name` and `@email` variables to build up a nicely formatted version of the user's email address using string interpolation (Section 4.2.2):

```
def formatted_email
  "#{@name} <#{@email}>"
end
```

Let's fire up the console, `require` the example user code, and take our User class out for a spin:

```
>> require 'example_user'          # This is how you load the example_user code.
=> ["User"]
=> example = User.new
=> #<User:0x224ceec @email=nil, @name=nil>
=> example.name                  # nil since attributes[:name] is nil
=> nil
=> example.name = "Example User"      # Assign a non-nil name
=> "Example User"
=> example.email = "user@example.com"    # and a non-nil email address
=> "user@example.com"
=> example.formatted_email
=> "Example User <user@example.com>"
```

This code creates an empty example user and then fills in the name and email address by assigning directly to the corresponding attributes (assignments made possible by the `attr_accessor` line in Listing 4.6). When we write

```
example.name = "Example User"
```

Ruby is setting the `@name` variable to `"Example User"` (and similarly for the `email` attribute), which we then use in the `formatted_email` method.

Recalling from Section 4.3.4 we can omit the curly braces for final hash arguments, we can create another user by passing a hash to the `initialize` method to create a user with pre-defined attributes:

```
>> user = User.new(:name => "Michael Hartl", :email => "mhartl@example.com")
=> #<User:0x225167c @email="mhartl@example.com", @name="Michael Hartl">
>> user.formatted_email
=> "Michael Hartl <mhartl@example.com>"
```

We will see starting in [Chapter 8](#) that initializing objects using a hash argument is common in Rails applications.

4.5 Exercises

1. Using [Listing 4.7](#) as a guide, combine the `split`, `shuffle`, and `join` methods to write a function that shuffles the letters in a given string.
2. Using [Listing 4.8](#) as a guide, add a `shuffle` method to the `String` class.
3. Create three hashes called `person1`, `person2`, and `person3`, with first and last names under the keys `:first` and `:last`. Then create a `params` hash so that `params[:father]` is `person1`, `params[:mother]` is `person2`, and `params[:child]` is `person3`. Verify that, for example, `params[:father][:first]` has the right value.
4. Learn about the `Hash` method `merge`.

Listing 4.7. Skeleton for a string shuffle function.

```
>> def string_shuffle(s)
>>   s.split('') .?.
>> end
=> nil
>> string_shuffle("foobar")
```

Listing 4.8. Skeleton for a `shuffle` method attached to the `String` class.

```
>> class String
>>   def shuffle
>>     self.split('') .?.
>>   end
>> end
=> nil
>> "foobar".shuffle
```

Chapter 5

Filling in the layout

In the process of taking a brief tour of Ruby in [Chapter 4](#), we added some basic cascading style sheets to our site layout ([Section 4.1.2](#)). In this chapter, we'll add some custom styles of our own, as we fill in the layout with links to the pages (such as Home and About) that we've created so far. Along the way, we'll learn about partials, Rails routes, and integration tests. We'll end by taking a first important step toward letting users sign up to our site.

5.1 Adding some structure

Rails Tutorial is a book on web development, not web design, but it would be depressing to work on an application that looks like *complete crap*, so in this section we'll add some structure to the layout and give it some minimal styling with CSS. We'll also give our *code* some styling, so to speak, using *partials* to tidy up the layout once it gets a little cluttered.

When building web applications, it is often useful to getting a high-level overview of the user interface even as early as possible. Throughout the rest of this book, I will thus often include *mockups* (in a web context often called *wireframes*), which are rough sketches of what the eventual application will look like.¹ In this chapter, we will principally be developing the static pages introduced in [Section 3.1](#), including a site logo, a navigation header, and a site footer. A mockup for the most important of these pages, the Home page, appears in [Figure 5.1](#). (You can see the final result in [Figure 5.8](#). You'll note that it differs in some details—for example, the footer has four links instead of three—but that's fine, since a mockup need not be exact.)

As usual, if you're using Git for version control, now would be a good time to make a new branch:

```
$ git checkout -b filling-in-layout
```

You might still have the `example_user.rb` file from [Chapter 4](#) in your project directory; if so, you should probably just remove it.

¹The mockups in *Ruby on Rails Tutorial* are made with an excellent free online mockup application called [Mockingbird](#).

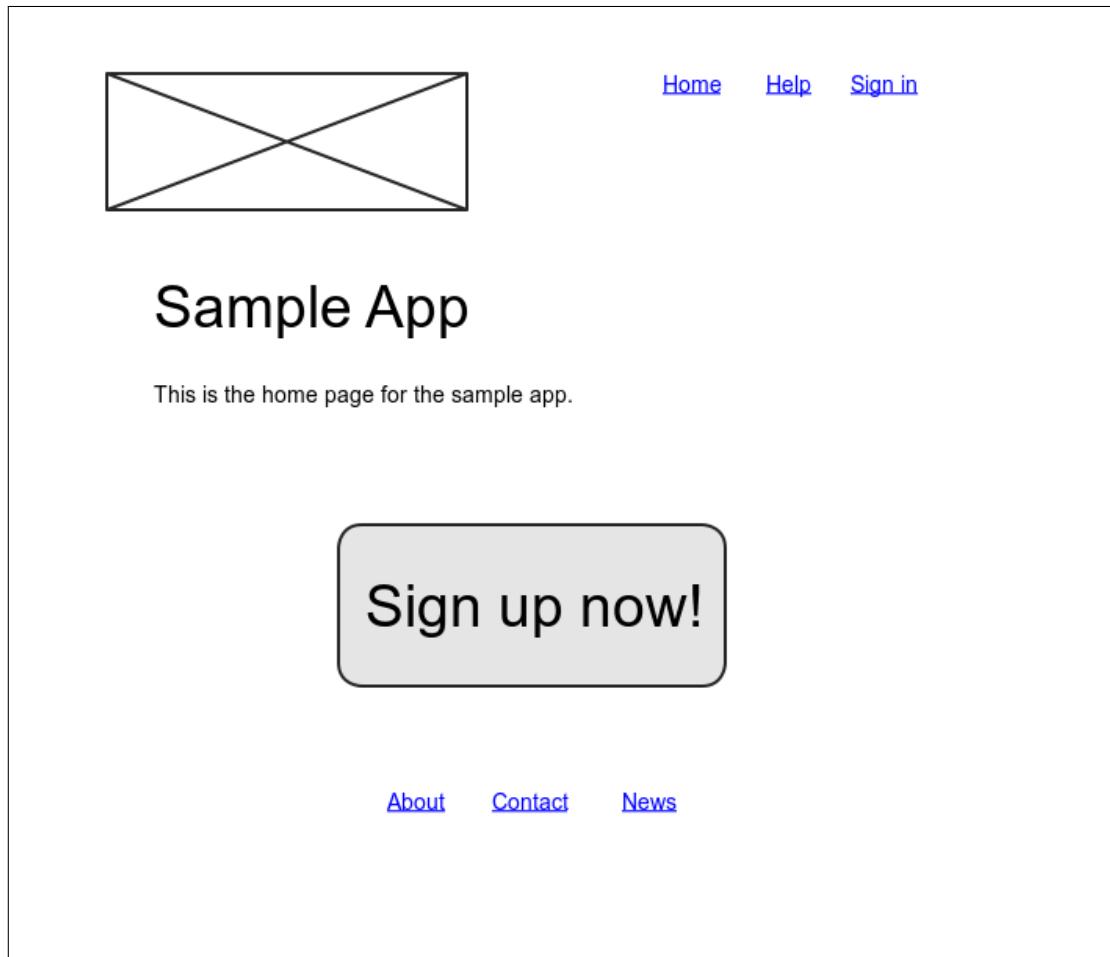


Figure 5.1: A mockup of the sample application's Home page. ([full size](#))

5.1.1 Site navigation

When we last saw the site layout file `application.html.erb` in Listing 4.3, we had just added Blueprint stylesheets using the Rails `stylesheet_link_tag` helper. It's time to add a couple more stylesheets, one specifically for Internet Explorer browsers and one for our (soon-to-be-added) custom CSS. We'll also add some additional divisions (`divs`), some `ids` and `classes`, and the start of our site navigation. The full file is in Listing 5.1; explanations for the various pieces follow immediately thereafter. If you'd rather not delay gratification, you can see the results in Figure 5.2. (Note: it's not (yet) very gratifying.)

Listing 5.1. The site layout with added structure.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
    <%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
    <%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
    <!--[if lt IE 8]><%= stylesheet_link_tag 'blueprint/ie' %><![endif]-->
    <%= stylesheet_link_tag 'custom', :media => 'screen' %>
  </head>
  <body>
    <div class="container">
      <div id="header">
        <%= image_tag("logo.png", :alt => "Sample App", :class => "round") %>
        <ul class="navigation round">
          <li><%= link_to "Home", '#' %></li>
          <li><%= link_to "Help", '#' %></li>
          <li><%= link_to "Sign in", '#' %></li>
        </ul>
      </div>
      <div id="content" class="round">
        <%= yield %>
      </div>
    </div>
  </body>
</html>
```

Let's look at the new elements from top to bottom. First we have a stylesheet that only gets included if the browser is Microsoft Internet Explorer (IE), any version less than (`lt`) 8:

```
<!--[if lt IE 8]><%= stylesheet_link_tag 'blueprint/ie' %><![endif]-->
```

IE has a large number of idiosyncrasies (especially before version 8), and Blueprint comes with a special `ie.css` file that fixes a bunch of them. The weird `[if lt IE 8]` syntax is *not* part of Rails; it's actually just HTML, and it seems (at least to me) rather surprising that HTML lets you do this sort of conditional inclusion. It's a good thing, though, because it means we can include the extra stylesheet *only* for IE browsers less than version 8, leaving other browsers such as Firefox and Safari unaffected.

After the IE stylesheet comes a stylesheet link for a file that doesn't exist yet, `custom.css`, where we'll put our custom CSS:

```
<%= stylesheet_link_tag 'custom', :media => 'screen' %>
```

CSS is very forgiving, and even though the file doesn't exist our page will still work just fine. (We'll create `custom.css` in Section 5.1.2.)

The next section places a *container div* around our site navigation and content, which is a `div` tag with class `container`. This container div is needed by Blueprint (see the [Blueprint tutorial](#) for more information). Then there are divs with id `header` and `content`; the header div contains the site navigation:

```
<div class="container">
  <div id="header">
    <%= image_tag("logo.png", :alt => "Sample App", :class => "round") %>
    <ul class="navigation round">
      <li><%= link_to "Home", '#' %></li>
      <li><%= link_to "Help", '#' %></li>
      <li><%= link_to "Sign in", '#' %></li>
    </ul>
  </div>
  <div id="content">
    <%= yield %>
  </div>
</div>
```

The `div` tag in HTML is a generic division; it doesn't do anything apart from divide the document into distinct parts. All HTML elements, including divs, can be assigned *classes*² and *ids*; these are merely labels, and are useful for styling with CSS (Section 5.1.2). The main difference between classes and ids is that classes can be used multiple times on a page, but ids can only be used once.

Inside the header div is a Rails helper called `image_tag` ([API link](#)):

```
<%= image_tag("logo.png", :alt => "Sample App", :class => "round") %>
```

Note that, as with `stylesheet_link_tag` (Section 4.3.4), we pass a hash of options, in this case setting the `alt` and `class` attributes of the image tag using symbols `:alt` and `:class`. To make this clearer, let's look at the HTML this tag produces:³

```

```

²These are completely unrelated to Ruby classes.

³You might notice that the `img` tag, rather than looking like `...`, instead looks like ``. Tags that follow this form are known as *self-closing* tags.

The `alt` attribute is what will be displayed if there is no image,⁴ and the class will be used for styling in Section 5.1.2. (Rails helpers often take options hashes in this way, giving us the flexibility to add arbitrary HTML options without ever leaving Rails.) You can see the result in Figure 5.2; we'll add the logo image at the end of this section.

The second element inside the layout header is a list of navigation links, made using the *unordered list* tag `ul`, together with the *list item* tag `li`:

```
<ul class="navigation round">
  <li><%= link_to "Home", '#' %></li>
  <li><%= link_to "Help", '#' %></li>
  <li><%= link_to "Sign in", '#' %></li>
</ul>
```

This list uses the Rails helper `link_to` (API link) to create links (which we created directly with the anchor tag `a` in Section 3.3.2); the first argument is the link text, while the second is the URL. We'll fill in the URLs with *named routes* in Section 5.2.3, but for now we use the stub URL '`#`' commonly used in web design. Once Rails has processed this layout and evaluated the Embedded Ruby, the list looks like this:

```
<ul class="navigation round">
  <li><a href="#">Home</a></li>
  <li><a href="#">Help</a></li>
  <li><a href="#">Sign in</a></li>
</ul>
```

One thing you might notice is that the `ul` tag has *two* classes, separated by a space:

```
<ul class="navigation round">
```

This is convenient for the common case of an element with two different kinds of styles.

Our layout is now complete, and we can look at the results by visiting, e.g., the Home page. In anticipation of adding users to our site in Chapter 8, let's add a signup link to the `home.html.erb` view (Listing 5.2).

Listing 5.2. The Home page with a link to the signup page.

`app/views/pages/home.html.erb`

```
<h1>Sample App</h1>

<p>
  This is the home page for the
  <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
```

⁴The `alt` text is also what will be displayed by screen readers for the visually impaired. Though people are sometimes sloppy about including the `alt` attribute for images, it is in fact required by the HTML standard. Luckily, Rails includes a default `alt` attribute; if you don't specify the attribute in the call to `image.tag`, Rails just uses the image filename (minus extension). In this case, though, `Sample App` is more descriptive than `logo`, so I've elected to set the `alt` text explicitly.

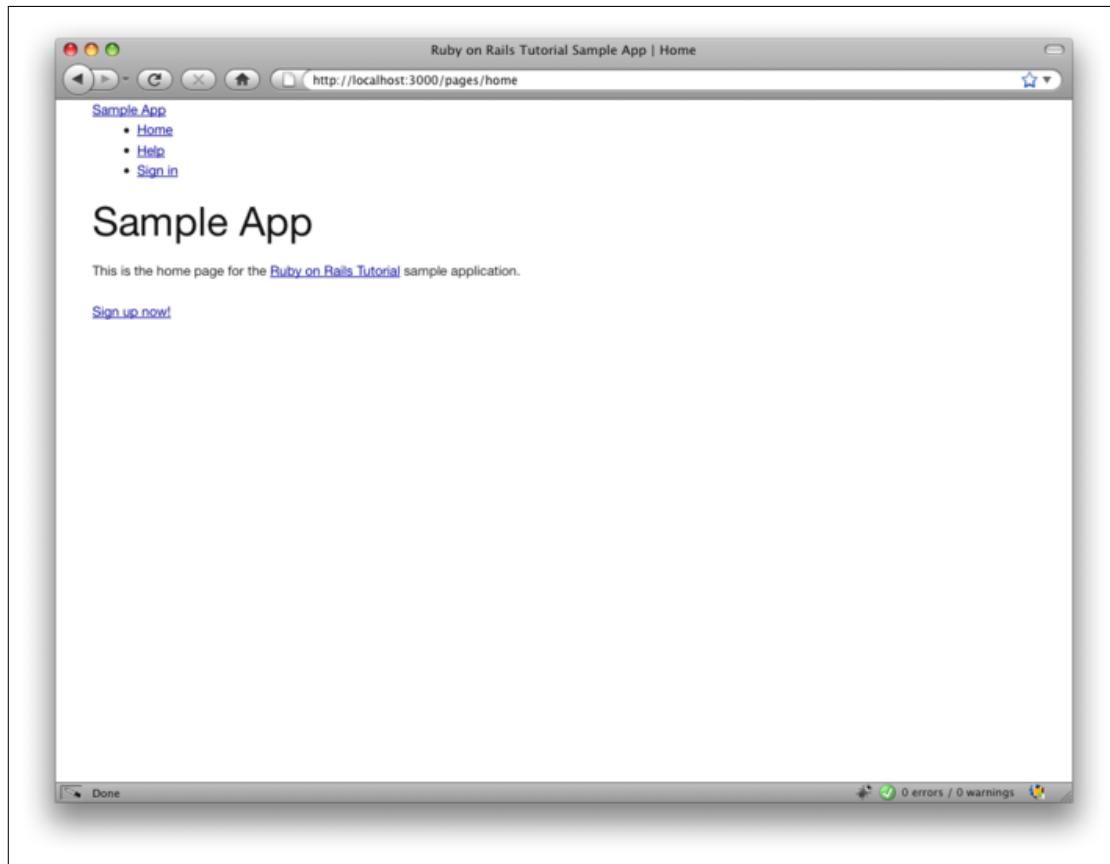


Figure 5.2: The Home page (</pages/home>) with no logo image or custom CSS. ([full size](#))

```
sample application.  
</p>  
  
<%= link_to "Sign up now!", '#', :class => "signup_button round" %>
```

As with the previous uses of `link_to`, this just creates a stub link of the form

```
<a href="#" class="signup_button round">Sign up now!</a>
```

Note again the recurring theme of options hashes, in this case used to add a couple CSS classes to the anchor tag.

Now we're finally ready to see the fruits of our labors (Figure 5.2). Pretty underwhelming, you say? Perhaps so. Happily, though, we've done a good job of giving our HTML elements sensible classes and ids, which puts us in a great position to add style to the site with CSS.

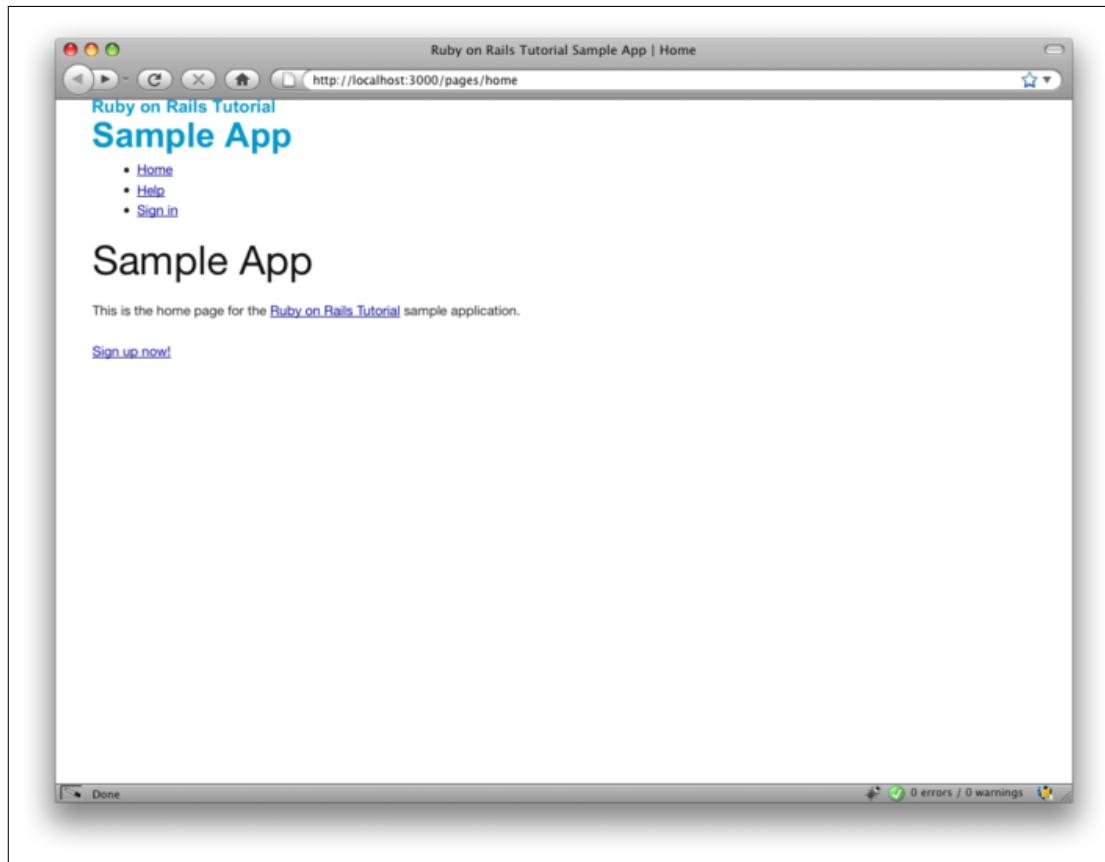


Figure 5.3: The Home page ([/pages/home](#)) with a logo image but no custom CSS. ([full size](#))

Before we move on to CSS styling, let's replace the logo alt text with a logo image; you can download the sample application logo at

http://www.railstutorial.org/images/sample_app/logo.png

Put the logo in **public/images** so that Rails can find it. The result appears in Figure 5.3.

5.1.2 Custom CSS

In Section 5.1.1, you may have noticed that the CSS classes and headers are *semantic*, that is, they have meaning in English beyond the structure of the page. For example, instead of writing that the navigation menu was “right-top” we gave it the class “navigation”. This gives us considerable flexibility in constructing a layout based on CSS.

Let's get started by filling in the **custom.css** file with Listing 5.3. (There are quite a few rules in Listing 5.3. To get a sense of what a CSS rule does, it's often helpful to comment it out using CSS comments, i.e., by putting it inside `/* ... */`, and seeing what changes.)

Listing 5.3. CSS for the container, body, and links.

`public/stylesheets/custom.css`

```
.container {
  width: 710px;
}

body {
  background: #cff;
}

#header {
  padding-top: 20px;
}

#header img {
  padding: 1em;
  background: #fff;
}

#content {
  margin-top: 1em;
  font-size: 120%;
  padding: 20px;
  background: #fff;
}

/* Links */

a {
  color: #09c;
  text-decoration: none;
}

a:hover {
  color: #069;
  text-decoration: underline;
}

a:visited {
  color: #069;
}
```

You can see the results of this CSS in [Figure 5.4](#). There's a lot of CSS here, but it has a consistent form. Each rule refers either to a class, an id, an HTML tag, or some combination thereof, followed by a list of styling commands. For example,

```
body {
  background: #cff;
}
```

changes the background color of the `body` tag to baby blue, while

```
#header img {  
  padding: 1em;  
  background: #ffff;  
}
```

puts a padding layer of one `em` (roughly the width of the letter M) around the image (`img`) inside an element with id `header`. (The hash mark at the beginning of `#header` refers to a CSS id.) This rule also makes the background color `#ffff`, which is white.⁵ Similarly,

```
.container {  
  width: 710px;  
}
```

styles an element with class `container`, in this case giving it a width of 710 pixels (corresponding to 18 Blueprint columns).⁶ The dot `.` in `.container` indicates that the rule styles a *class* called “container” in the same way that the pound sign `#` in `#header` indicates that the rule styles an *id* called “header”.

Changing colors is nice, but the navigation links are still hanging down on the left side of the page. Let’s move them to a better location and give them a nicer appearance with the navigation rules in Listing 5.4. As before, we style the `navigation` class using a dot preceding the name, i.e., `.navigation`. The results appear in Figure 5.5.

Listing 5.4. Navigation CSS.

`public/stylesheets/custom.css`

```
.  
./* Navigation */  
  
.navigation {  
  float: right;  
}  
  
.navigation {  
  background-color: white;  
  padding: 0 0.7em;  
  white-space: nowrap;  
}  
  
ul.navigation {
```

⁵HTML colors can be coded with three base-16 (hexadecimal) numbers, one each for the primary colors red, green, and blue. `#ffff` maxes out all three colors, yielding pure white. See this [HTML colors](#) site for more information. By the way, the `#` in `#ffff` has nothing to do with the one in `#header`.

⁶Blueprint CSS uses a grid of columns 40 pixels across, 30 pixels for the column itself and 10 pixels of padding. The rightmost column doesn’t need padding, so 18 columns is 710 pixels: $18 * 40 - 10 = 710$.

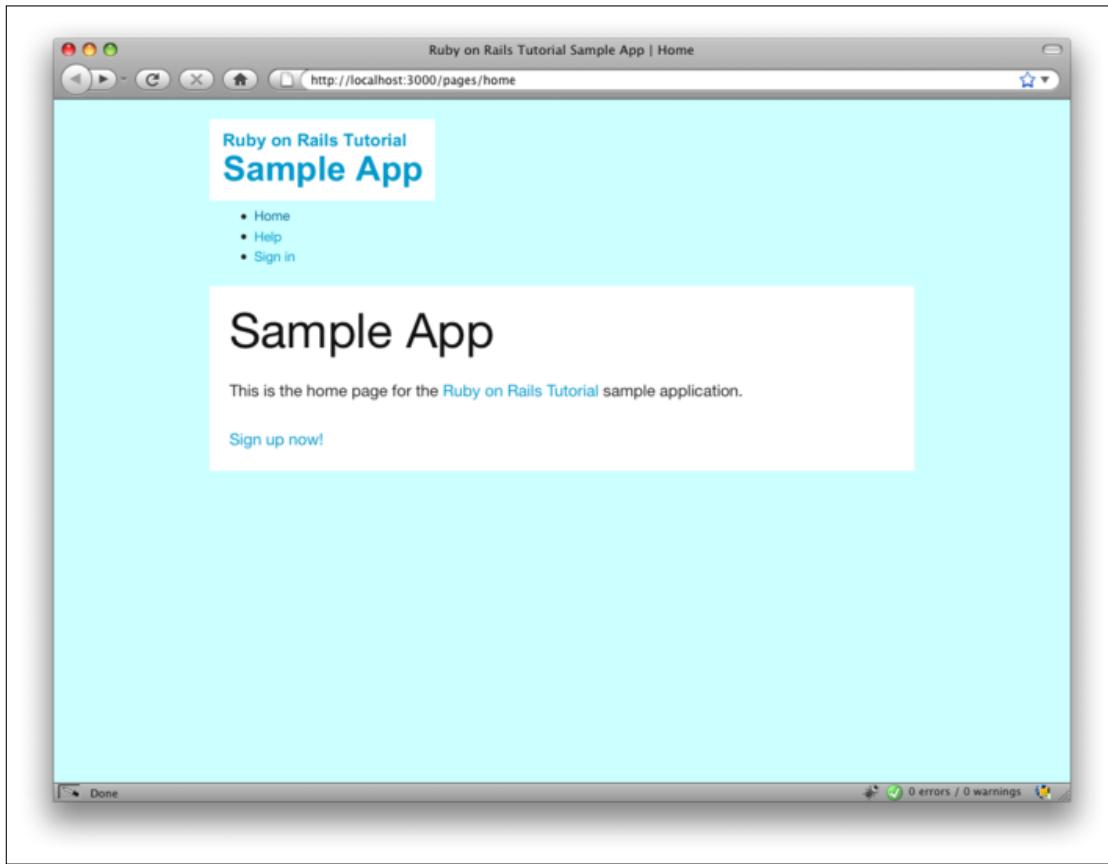


Figure 5.4: The Home page (</pages/home>) with custom colors. ([full size](#))

```
    margin: 0;
}

.navigation li {
  list-style-type: none;
  display: inline-block;
  padding: 0.2em 0;
}

.navigation li a {
  padding: 0 5px;
  font-weight: bold;
}

.navigation li a:visited {
  color: #09c;
}

.navigation li a:hover {
  text-decoration: underline;
}
```

Note that `ul.navigation` styles a `ul` tag with class `navigation`, whereas `.navigation li` styles an `li` tag *inside* an element with class `navigation`.

As the `penultimate` step, we'll make the link to our site's signup page a little more obvious. (Though for the sample app we don't care, on any real site it's naturally quite important to make the signup link *very* prominent.) Listing 5.5 shows CSS to make the signup link big, green, and clickable (so a click anywhere inside the box will follow the link).

Listing 5.5. CSS to make the signup button big, green, and clickable (`/pages/home`).

`public/stylesheets/custom.css`

```
.

.

.

/* Sign up button */

.a.signup_button {
  margin-left: auto;
  margin-right: auto;
  display: block;
  text-align: center;
  width: 190px;
  color: #fff;
  background: #006400;
  font-size: 150%;
  font-weight: bold;
  padding: 20px;
}
```

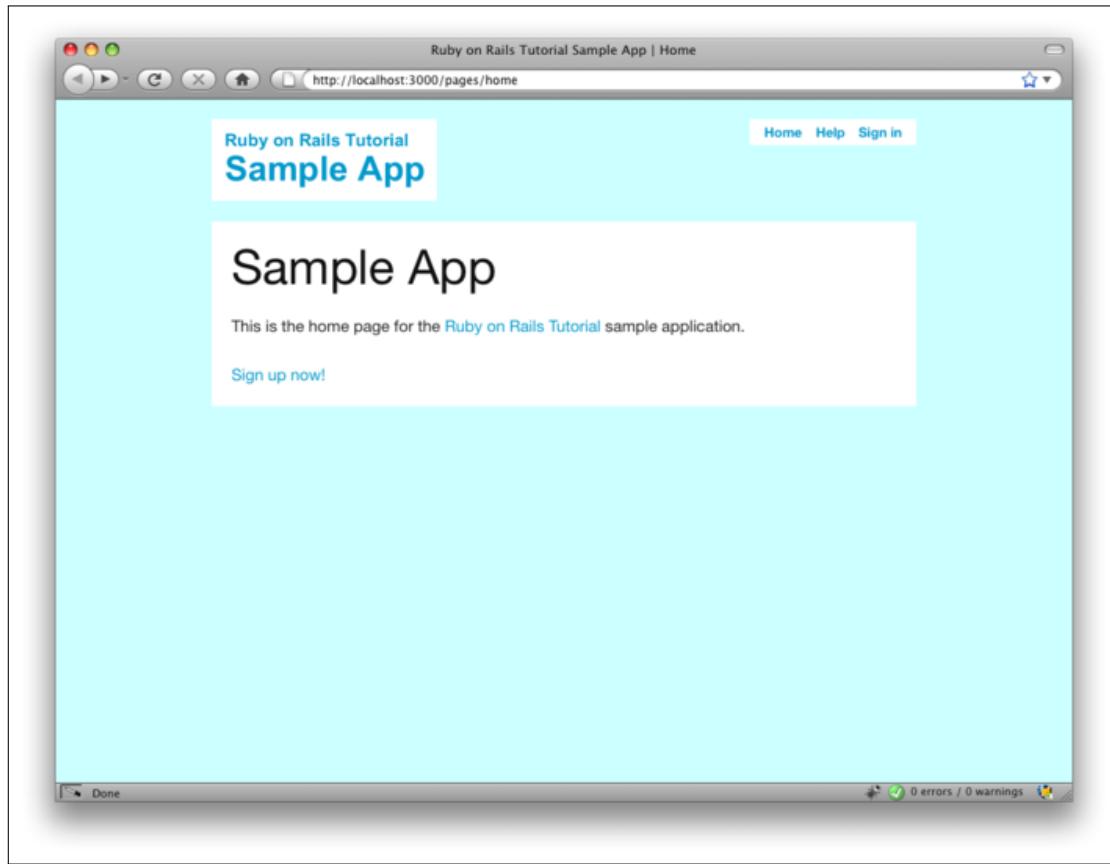


Figure 5.5: The Home page (</pages/home>) with navigation styling. ([full size](#))

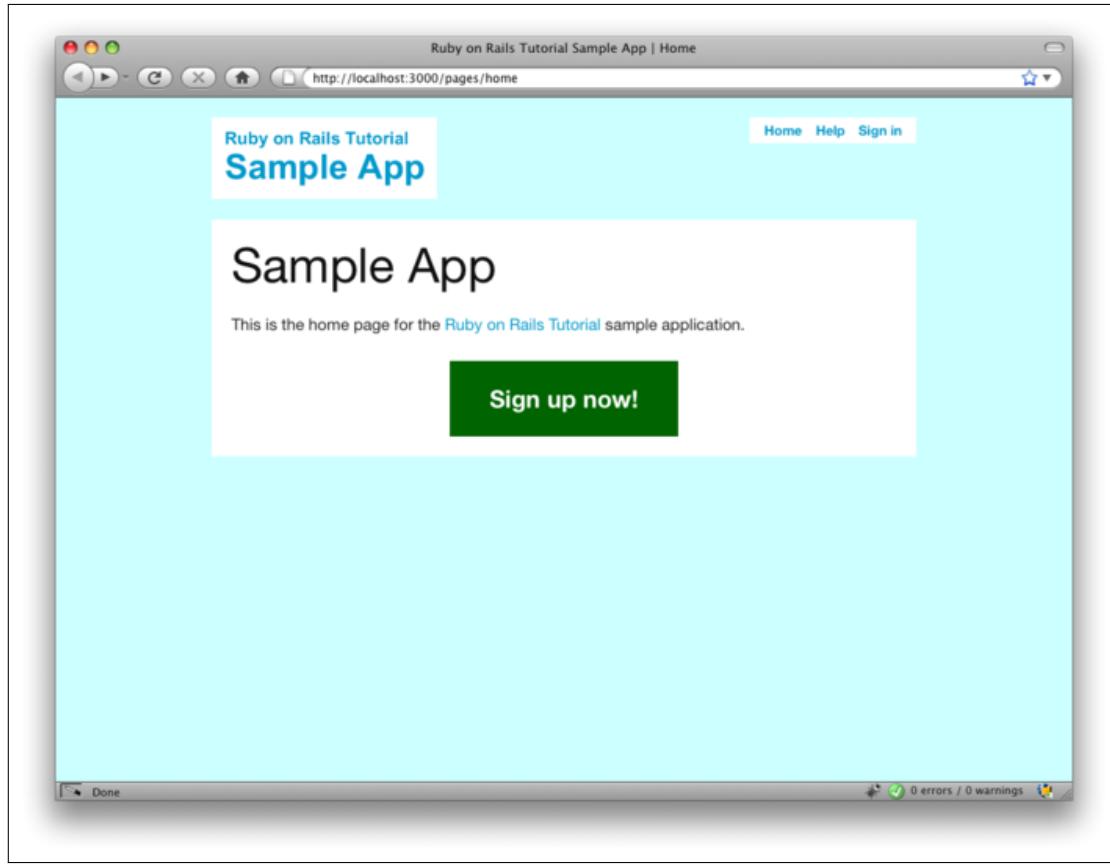


Figure 5.6: The Home page ([/pages/home](#)) with a signup button. ([full size](#))

There are a bunch of rules here; as usual, comment a line out and reload the page if you want to see what each one does. The end result is a signup link that's hard to miss (Figure 5.6).

As a final touch, we'll make use of the `round` class we've placed on many of our site elements. Though the current sharp-cornered boxes aren't terrible, it's a little friendlier to soften the corners so they won't slice up our users. We can accomplish this using the CSS code in Listing 5.6, with the results shown in Figure 5.7.

Listing 5.6. Stylesheet rules for round corners.

`public/stylesheets/custom.css`

```
.  
. .  
. .  
/* Round corners */  
  
.round {  
  -moz-border-radius: 10px;  
  -webkit-border-radius: 10px;  
}
```

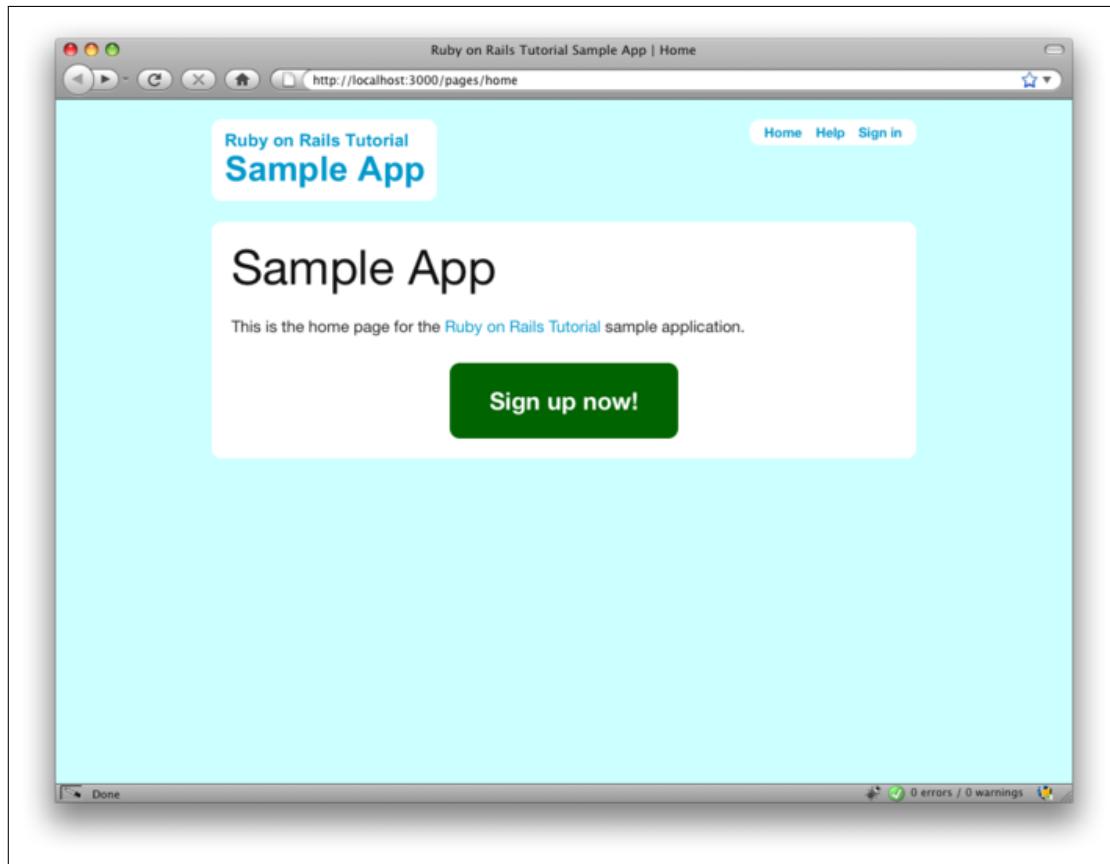


Figure 5.7: The Home page (`/pages/home`) with round corners. ([full size](#))

It's worth noting that this trick works on Firefox, Safari, and many other browsers, but it doesn't work on Internet Explorer. There *are* ways of getting round corners that work on all browsers, but there's no other technique that's even *close* to this easy, so we'll just risk leaving our IE users with a few tiny cuts.

5.1.3 Partials

Though the layout in Listing 5.1 serves its purpose, it's getting a little cluttered: there are four lines of CSS includes and eight lines of header for what are logically only two ideas. We can tuck these sections away using a convenient Rails facility called *partials*. Let's first take a look at what the layout looks like after the partials are defined (Listing 5.7).

Listing 5.7. The site layout with partials for the stylesheets and header.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"  
  "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">  
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
```

```
<head>
  <title><%= title %></title>
  <%= render 'layouts/stylesheets' %>
</head>
<body>
  <div class="container">
    <%= render 'layouts/header' %>
    <div id="content" class="round">
      <%= yield %>
    </div>
  </div>
</body>
</html>
```

In Listing 5.7, we've replaced the stylesheet lines with a single call to a Rails helper called `render`:

```
<%= render 'layouts/stylesheets' %>
```

The effect of this line is to look for a file called `app/views/layouts/_stylesheets.html.erb`, evaluate its contents, and insert the results into the view. (Recall that `<%= ... %>` is the Embedded Ruby syntax needed to evaluate a Ruby expression and then insert the results into the template.) Note the leading underscore on the filename `_stylesheets.html.erb`; this underscore is the universal convention for naming partials, and among other things makes it possible to identify all the partials in a directory at a glance.

Of course, to get the partial to work, we have to fill it with some contents; in the case of the stylesheet partial, this is just the four lines of stylesheet includes from Listing 5.1; the result appears in Listing 5.8.

Listing 5.8. A partial for stylesheet includes.

`app/views/layouts/_stylesheets.html.erb`

```
<%= stylesheet_link_tag 'blueprint/screen', :media => 'screen' %>
<%= stylesheet_link_tag 'blueprint/print', :media => 'print' %>
<!--[if lt IE 8]><%= stylesheet_link_tag 'blueprint/ie' %><![endif]-->
<%= stylesheet_link_tag 'custom', :media => 'screen' %>
```

Similarly, we can move the header material into the partial shown in Listing 5.9 and insert it into the layout with another call to `render`:

```
<%= render 'layouts/header' %>
```

Listing 5.9. A partial for the site header.

`app/views/layouts/_header.html.erb`

```
<div id="header" class="round">
  <%= image_tag("logo.png", :alt => "Sample App", :class => "round") %>
  <ul class="navigation round">
    <li><%= link_to "Home", '#' %></li>
    <li><%= link_to "Help", '#' %></li>
    <li><%= link_to "Sign in", '#' %></li>
  </ul>
</div>
```

Now that we know how to make partials, let's add a site footer to go along with the header. By now you can probably guess that we'll call it `_footer.html.erb` and put it in the layouts directory (Listing 5.10).

Listing 5.10. A partial for the site footer.

`app/views/layouts/_footer.html.erb`

```
<div id="footer">
  <ul class="navigation round">
    <li><%= link_to "About", '#' %></li>
    <li><%= link_to "Contact", '#' %></li>
    <li><a href="http://news.railstutorial.org/">News</a></li>
    <li><a href="http://www.railstutorial.org/">Rails Tutorial</a></li>
  </ul>
</div>
```

As with the header, in the footer we've used `link_to` for the internal links to the About and Contact pages and stubbed out the URLs with '`#`' for now.

We can render the footer partial in the layout by following the same pattern as the stylesheets and header partials (Listing 5.11).

Listing 5.11. The site layout with a footer partial.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html lang="en" xml:lang="en" xmlns="http://www.w3.org/1999/xhtml">
  <head>
    <title><%= title %></title>
    <%= render 'layouts/stylesheets' %>
  </head>
  <body>
    <div class="container">
      <%= render 'layouts/header' %>
      <div id="content" class="round">
        <%= yield %>
      </div>
      <%= render 'layouts/footer' %>
    </div>
```

```
</body>
</html>
```

Of course, the footer will be ugly without some styling (Listing 5.12). The results appear in Figure 5.8.

Listing 5.12. Adding the CSS for the site footer.

`public/stylesheets/custom.css`

```
.
.
.

#footer {
  text-align: center;
  margin-top: 10px;
  width: 710px;
  margin-left: auto;
  margin-right: auto;
}

#footer .navigation {
  float: none;
}
```

Note here the rule

```
#footer .navigation {
  float: none;
}
```

that overrides the previous rule

```
.navigation {
  float: right;
}
```

so that the footer is centered in the bottom of the page rather than pushed off to the right like the navigation in the header. This convention of having a succession of rules, with subsequent rules possibly overriding previous ones, is what puts the *cascading* in cascading style sheets.

5.2 Layout links

Now that we've finished a site layout with decent styling, it's time to start filling in the links we've stubbed out with '`#`'. Of course, we could hard-code links like

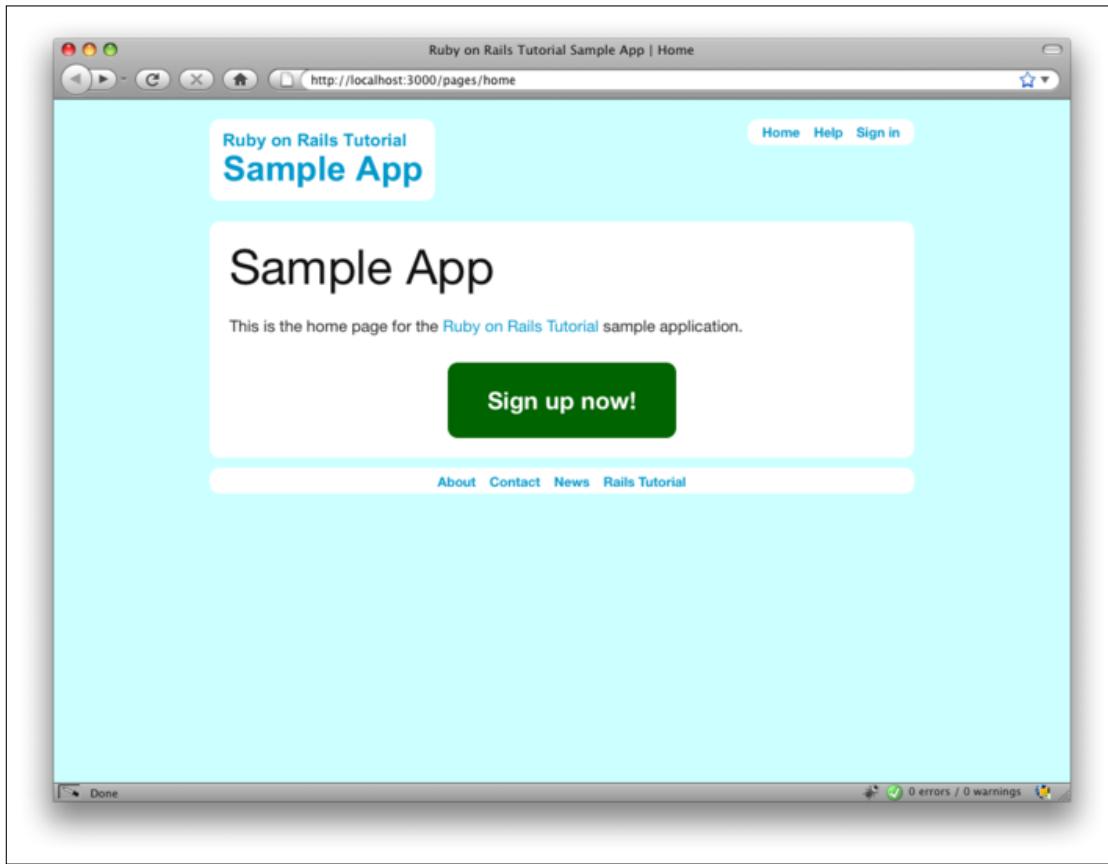


Figure 5.8: The Home page ([/pages/home](#)) with an added footer. (full size)

Page	URL	Named route
Home	/	<code>root_path</code>
About	/about	<code>about_path</code>
Contact	/contact	<code>contact_path</code>
Help	/help	<code>help_path</code>
Sign up	/signup	<code>signup_path</code>
Sign in	/signin	<code>signin_path</code>

Table 5.1: Route and URL mapping for site links.

```
<a href="/pages/about">About</a>
```

but that isn't the Rails Way. For one, it would be nice if the URL for the about page were /about rather than /pages/about; moreover, Rails conventionally uses *named routes*, which involves code like

```
<%= link_to "About", about_path %>
```

This way the code has a more transparent meaning, and it's also more flexible since we can change the definition of `about_path` and have the URL change everywhere `about_path` is used.

The full list of our planned links appears in [Table 5.1](#), along with their mapping to URLs and routes. We'll implement all but the last one by the end of this chapter. (We'll make the last one in [Chapter 9](#).)

5.2.1 Integration tests

Before writing the routes for our application, we'll continue with our test-driven development by writing some tests for them. There are several ways to test routes, and I'm going to take this opportunity to introduce *integration tests*, which I think are rather under-used in Rails land. Integration tests can help us test routes and a lot more besides.

We start with a generator for a layout links integration spec:

```
$ mkdir spec/integration
$ script/generate integration_spec layout_links
      create spec/integration/layout_links_spec.rb
```

Note that the generator automatically appends `.spec.rb` to the name of our test file.

Our integration test will use the same `get` function we used in [Section 3.2](#) in the Pages controller spec, with code like this:

```
describe "GET 'home'" do
  it "should be successful" do
    get 'home'
```

```
    response.should be_success
  end
end
```

In this section, we want to test URLs like / and /about, but you can't `get` these URLs inside a controller test—controller tests only know about URLs defined for that exact controller. In contrast, integration tests are bound by no such restriction, since they are designed as *integrated* tests for the whole application and hence can `get` any page they want.

Following the model of the Pages controller spec, we can write an integration spec for each of the pages in [Table 5.1](#) that we've already created, namely, Home, About, Contact, and Help.⁷ To make sure the right template (i.e., view) is rendered in each case, we'll use the `render_template` function inside RSpec. The test definitions appear in [Listing 5.13](#). Of course, at this point they should fail (Red); we'll get them to Green in [Section 5.2.2](#).

Listing 5.13. Integration test for routes.

`spec/integration/layout_links_spec.rb`

```
require 'spec_helper'

describe "Layout links" do
  it "should have a Home page at '/'" do
    get '/'
    response.should render_template('pages/home')
  end

  it "should have a Contact page at '/contact'" do
    get '/contact'
    response.should render_template('pages/contact')
  end

  it "should have an About page at '/about'" do
    get '/about'
    response.should render_template('pages/about')
  end

  it "should have a Help page at '/help'" do
    get '/help'
    response.should render_template('pages/help')
  end
end
```

By the way, if you're running Autotest (via the `autospec` command), you might notice that it doesn't run the integration test. Apparently this is by design, since integration tests can be slow and hence can disrupt the red-green-refactor cycle. Nevertheless, I still find it preferable to include the integration tests in Autotest. To get this to work, you just have to tell Autotest to run tests in the `spec/integration` directory ([Listing 5.14](#) or [Listing 5.15](#)).

⁷If you don't have a Help page, now would be a good time to look back at the [Chapter 3](#) exercises in Section 3.5.

Listing 5.14. Additions to `.autotest` needed to run integration tests with Autotest on OS X.

```
Autotest.add_hook :initialize do |autotest|
  autotest.add_mapping(/spec/integration/.*_spec\.rb$/) do
    autotest.files_matching(/spec/integration/.*_spec\.rb$/)
  end
end
```

Listing 5.15. Additions to `.autotest` needed to run integration tests with Autotest on Ubuntu Linux.

```
Autotest.add_hook :initialize do |autotest|
  autotest.add_mapping(%r%^spec/(integration)/.*rb$%) do |filename, _|
    filename
  end
end
```

(Don’t worry about where this code comes from; I don’t know the Autotest API either. At some point I Googled around with search terms like “rspec autotest integration” and found it, and when I dropped it into my `.autotest` file, it worked.)

5.2.2 Rails routes

Now that we have tests for the URLs we want, it’s time get them to work. The file Rails uses for URL mappings is `config/routes.rb`. This is our first encounter with the `config` directory, so it’s helpful to take a quick look at it (Figure 5.9). This directory is where Rails collects files needed for the application configuration (hence the name); probably the two most commonly used files are `routes.rb` and `environment.rb`.

If you take a look at the default `routes.rb` file, you’ll see that it’s quite a mess, but it’s a useful mess—full of commented-out example route mappings. I suggest reading through it at some point, and also taking a look at the [Rails Guides article “Rails Routing from the outside in”](#) for a much more in-depth treatment of routes. For now, though, we’ll stick with the examples in Listing 5.16.⁸

Listing 5.16. Routes for static pages.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  map.contact '/contact', :controller => 'pages', :action => 'contact'
  map.about   '/about',   :controller => 'pages', :action => 'about'
  map.help    '/help',    :controller => 'pages', :action => 'help'
  .
  .
  .
end
```

⁸In the line `ActionController::Routing::Routes.draw do |map|` you might recognize that the `draw` method takes a *block*, a construction we last saw in Section 4.3.2. The `draw` method takes a block with one variable, which is nearly always called `map`.

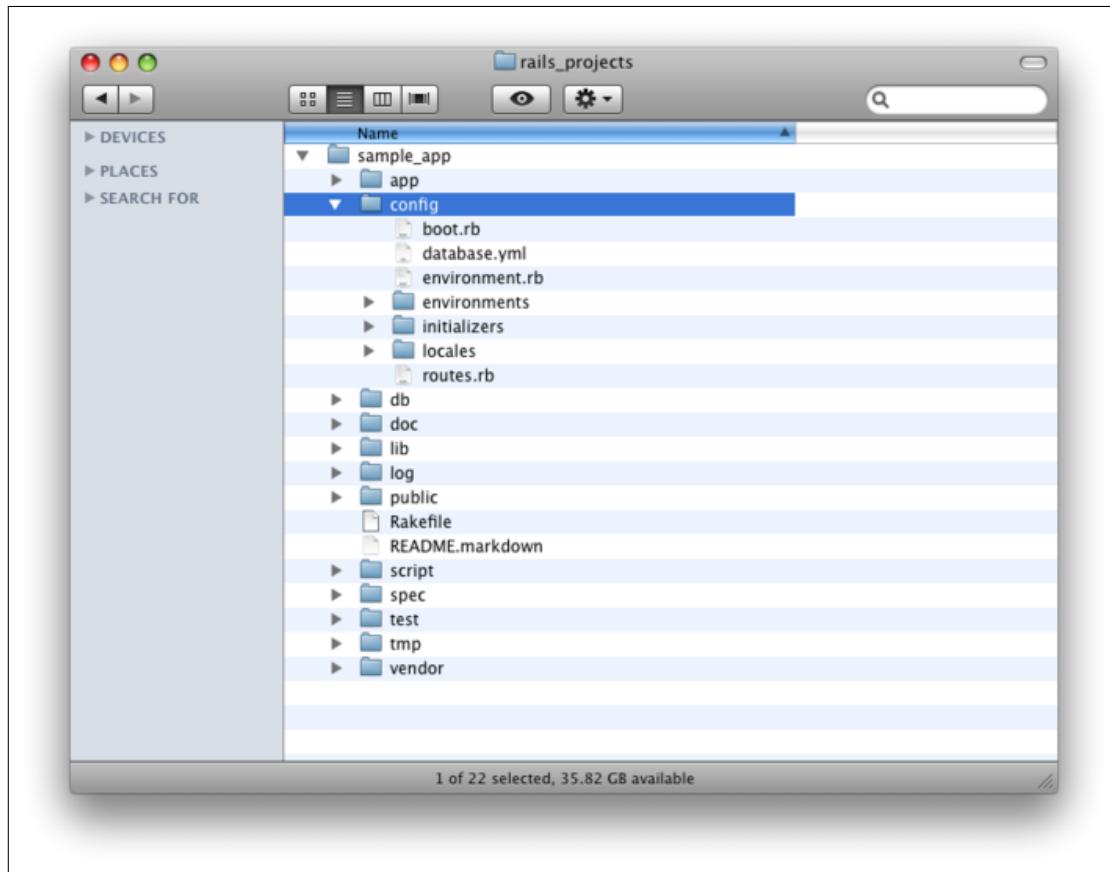


Figure 5.9: Contents of the sample app's `config` directory. ([full size](#))

If you read this code carefully, you can probably figure out basically what it does; for example, you can see that

```
map.about '/about', :controller => 'pages', :action => 'about'
```

maps the URL `/about` to the `about` action in the Pages controller. (Before this was more explicit: we used `/pages/about` to get to the same place, but `/about` is more succinct.) What isn't obvious is that `map.about` also automatically creates *named routes* for use in the controllers and views:

```
about_path => '/about'
about_url => 'http://localhost:3000/about'
```

Note that `about.url` is the *full* URL `http://localhost:3000/about` (with `localhost:3000` being replaced with the domain name, such as `example.com`, for a fully deployed site). As discussed in Section 5.2, to get just `/about`, you use `about.path`. (*Rails Tutorial* uses the `path` form for consistency, but the difference rarely matters in practice.)

With these routes now defined, the tests for the About, Contact, and Help pages should pass. (As usual, use Autotest or `spec spec/` to check.) This leaves the test for the Home page.

To establish the route mapping for the Home page, we could use code like

```
map.home '/', :controller => 'pages', :action => 'home'
```

but this is unnecessary; Rails has a special route for the root URL `/` (“slash”) located lower down in the file (Listing 5.17).

Listing 5.17. The mapping for the root route.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  .
  .
  .
  # You can have the root of your site routed with map.root
  # -- just remember to delete public/index.html.
  map.root :controller => 'pages', :action => 'home'
  .
  .
  .
end
```

This code maps the root URL `/` to `/pages/home`, and also gives URL helpers as follows:

```
root_path => '/'
root_url  => 'localhost:3000/'
```

We should also heed the comment and delete `public/index.html` to prevent Rails from rendering the default page ([Figure 1.5](#)) when we visit `/`.⁹ You can of course simply remove the file by trashing it, but if you're using Git for version control there's a way to tell Git about the removal at the same time using `git rm`:

```
$ git rm public/index.html
$ git commit -am "Removed default Rails page"
```

With that, all of the routes for static pages are working, and the tests should pass. Now we just have to fill in the links in the layout.

5.2.3 Named routes

Let's put the named routes created in [Section 5.2.2](#) to work in our layout. This will entail filling in the second arguments of the `link_to` functions with the proper named routes. For example, we'll convert

```
<%= link_to "About", '#' %>
```

to

```
<%= link_to "About", about_path %>
```

and so on.

We'll start in the header partial, `_header.html.erb` ([Listing 5.18](#)), which has links to the Home and Help pages. While we're at it, we'll follow a common web convention and link the logo image to the Home page as well.

Listing 5.18. Header partial with links.

`app/views/layouts/_header.html.erb`

```
<div id="header" class="round">
  <% logo = image_tag("logo.png", :alt => "Sample App", :class => "round") %>
  <%= link_to logo, root_path %>
  <ul class="navigation round">
```

⁹As of this writing, unfortunately the integration test for the Home page passes even if we *don't* remove `public/index.html`, even though actually visiting the root directory `/` yields the default Rails page from [Figure 1.5](#). Evidently integration tests aren't a perfect simulation of web browsers. If your application ever *requires* a perfect simulation, take a look at the [Selenium web application testing system](#).

```
<li><%= link_to "Home", root_path %></li>
<li><%= link_to "Help", help_path %></li>
<li><%= link_to "Sign in", '#' %></li>
</ul>
</div>
```

We won't have a named route for the "Sign in" link until [Chapter 9](#), so we've left it as '`#'` for now. Note that this code defines the local variable `logo` for the logo image tag, and then links to it in the next line:

```
<% logo = image_tag("logo.png", :alt => "Sample App", :class => "round") %>
<%= link_to logo, root_path %>
```

This is a little cleaner than stuffing it all into one line. It's especially important to notice that the ERb for the variable assignment doesn't have an equals sign; it's just `<% ... %>`, because we don't want that line inserted into the template. (Using a local variable in this manner is only one way to do it. An even cleaner way might be to define a `logo` helper; see [Section 5.5](#).)

The other place with links is the footer partial, `_footer.html.erb`, which has links for the About and Contact pages ([Listing 5.19](#)).

Listing 5.19. Footer partial with links.

`app/views/layouts/_footer.html.erb`

```
<div id="footer">
  <ul class="navigation round">
    <li><%= link_to "About", about_path %></li>
    <li><%= link_to "Contact", contact_path %></li>
    <li><a href="http://news.railstutorial.org/">News</a></li>
    <li><a href="http://www.railstutorial.org/">Rails Tutorial</a></li>
  </ul>
</div>
```

With that, our layout has links to all the static pages created in [Chapter 3](#), so that, for example, `/about` goes to the About page ([Figure 5.10](#)).

By the way, it's worth noting that, although we haven't actually tested for the presence of the links on the layout, our tests will fail if the routes aren't defined. You can check this by commenting out the routes in [Listing 5.16](#) and running your test suite. For a testing method that actually makes sure the links go to the right places, see [Section 5.5](#).

5.3 User signup: A first step

As a capstone to our work on the layout and routing, in this section we'll make a route for the signup page, which will mean creating a second controller along the way. This is a first important step toward allowing users to register for our site; we'll take the next step, modeling users, in [Chapter 6](#), and we'll finish the job in [Chapter 8](#).

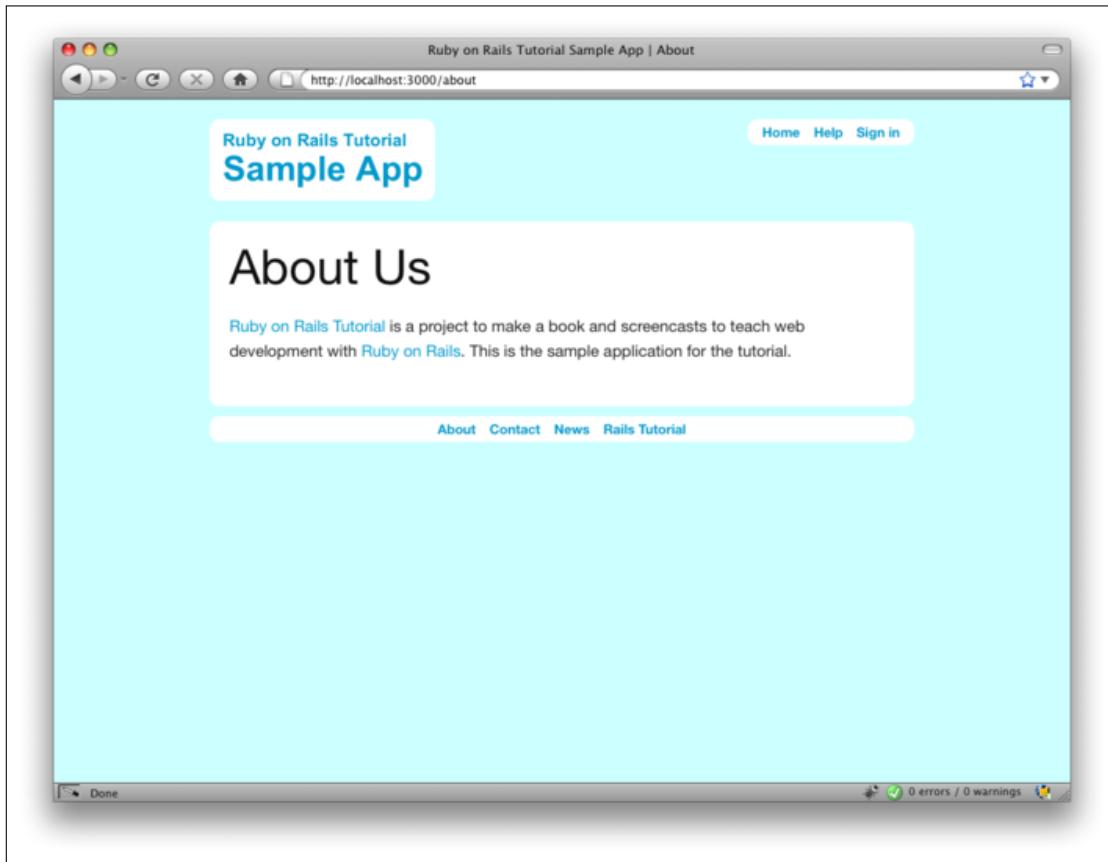


Figure 5.10: The About page at [/about](#). (full size)

5.3.1 Users controller

It's been a while since we created our first controller, the Pages controller, way back in [Section 3.1.2](#). It's time to create a second one, the Users controller. As before, we'll use `generate` to make the simplest controller that meets our present needs, namely, one with a stub signup page for new users. Following the conventional [REST architecture](#) favored by Rails, we'll call the action for new users `new` and pass it as an argument to the `rspec_controller` generator to create it automatically ([Listing 5.20](#)).

Listing 5.20. Generating a Users controller.

```
$ script/generate rspec_controller Users new
exists  app/controllers/
exists  app/helpers/
create  app/views/users
create  spec/controllers/
create  spec/helpers/
create  spec/views/users
create  spec/controllers/users_controller_spec.rb
create  spec/helpers/users_helper_spec.rb
create  app/controllers/users_controller.rb
create  app/helpers/users_helper.rb
create  spec/views/users/new.html.erb_spec.rb
create  app/views/users/new.html.erb
```

As with the Pages controller, this generates view specs that we won't need, so remove them:

```
$ rm -rf spec/views
```

The controller generator makes both the Users controller a useful default test that verifies that the `new` action responds properly to a GET request ([Box 3.1](#)); the code appears in [Listing 5.21](#). This code should look familiar; it follows the exact same form as the Pages controller spec last seen in [Section 3.3.1](#) ([Listing 3.12](#)).

Listing 5.21. Testing the signup page.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do

  describe "GET 'new'" do
    it "should be successful" do
      get 'new'
      response.should be_success
    end
  end
end
```

By construction, the Users controller already has the proper `new` action and `new.html.erb` template to get this test to pass (Listing 5.22). (To view the page at `/users/new`, you might have to restart the server.)

Listing 5.22. Action for the new user (signup) page.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def new
  end

end
```

To get back in the spirit of test-driven development, let's add a second (failing) test of our own by testing for a title that contains the string "`Sign up`" (Listing 5.23). Be sure to add `integrate_views` as we did in the Pages controller spec (Listing 3.12); otherwise, the test won't pass even after we add the proper title.

Listing 5.23. A test for the signup page title.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  describe "GET 'new'" do
    it "should be successful" do
      get 'new'
      response.should be_success
    end

    it "should have the right title" do
      get 'new'
      response.should have_tag("title", /Sign up/)
    end
  end
end
```

This test uses the `have.tag` method we've seen before (Section 3.3.1),¹⁰ along with a new construction, a *regular expression*. Rather than testing for the full title, which will be "Ruby on Rails Tutorial Sample App — Sign up", we just test to make sure that "Sign up" is present; the Pages controller spec already deals with the general case of the base title, so we wouldn't gain much by including it here. If you've programmed much before, chances are you're already familiar with regular expressions, but here's a quick `script/console` session in case you need a refresher.¹¹

¹⁰Note that, as in Section 3.3.1, `have.tag` needs the `integrate_views` line since it tests the view along with the action.

¹¹We'll talk more about regular expressions in Section 6.2.3.

```
>> "Sign" =~ /Sign up/      # Compare a string to a regular expression with =~
=> nil
>> "Sign up" =~ /Sign up/
=> 0
>> "Ruby on Rails Tutorial Sample App | Sign up" =~ /Sign up/
=> 36
```

The equals-tilde operator compares a string to a *regular expression* (or *regex*), which is a pattern designed to match only particular strings. In the case of the regex `/Sign up/`, any string that contains those exact characters will match. The console's return values may look odd: for no match, the regex comparison returns `nil`; for a match, it returns the *index* (position) in the string where the match starts.¹² Usually, though, the exact index doesn't matter, since the comparison is usually used in a boolean context: recall from [Section 4.2.3](#) that `nil` is `false` in a boolean context and that anything else (even `0`) is true. Thus, we can write code like this:

```
>> title = "Ruby on Rails Tutorial Sample App | Sign up"
=> "Ruby on Rails Tutorial Sample App | Sign up"
>> "It's a match!" if title =~ /Sign up/
=> "It's a match!"
```

Implicitly, the RSpec code

```
it "should have the right title" do
  get 'new'
  response.should have_tag("title", /Sign up/)
end
```

performs this exact comparison, comparing the value of the page's `title` tag with the regex `/Sign up/`. This way, we don't have to worry about what else is in the title string; what we really care about is the "Sign up" part.

Of course, by design this test currently fails (Red). To get a custom title, we need to make an `@title` instance variable as in [Section 3.3.3](#). We can thus get to Green with the code in [Listing 5.24](#).

Listing 5.24. Setting the custom title for the new user page.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def new
    @title = "Sign up"
  end
end
```

¹²The indices are zero-offset, as with arrays ([Section 4.3.1](#)), so a return value of `0` means the string matches the regular expression starting with the first character.

5.3.2 Signup URL

With the code from Section 5.3.1, we already have a working page for new users at `/users/new`, but recall from Table 5.1 that we want the URL to be `/signup` instead. As in Section 5.2, we'll first write a test (Listing 5.25).

Listing 5.25. Simple integration test for user signup link.

`spec/integration/layout_links_spec.rb`

```
require 'spec_helper'

describe "Layout links" do
  .
  .
  .
  it "should have a signup page at '/signup'" do
    get '/signup'
    response.should render_template('users/new')
  end
end
```

Note that this is the *same* file as the one used for the other layout links, even though the Signup page is in a different controller. Being able to hit pages in multiple controllers is one of the advantages of using integration tests.

The last step is to make a named route for signups. We'll follow the examples from Listing 5.16 and add a `map.signup` rule for the signup URL (Listing 5.26).

Listing 5.26. A route for the signup page.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  .
  .
  .
  map.help      '/help',      :controller => 'pages', :action => 'help'
  map.signup   '/signup',     :controller => 'users', :action => 'new'
  .
  .
  .
end
```

The test now passes. All that's left is to add the right link to the button on the Home page. As with the other routes, `map.signup '/signup'` gives us the named route `signup_path`, which we put to use in Listing 5.27.

Listing 5.27. Linking the button to the Signup page.

`app/views/pages/home.html.erb`

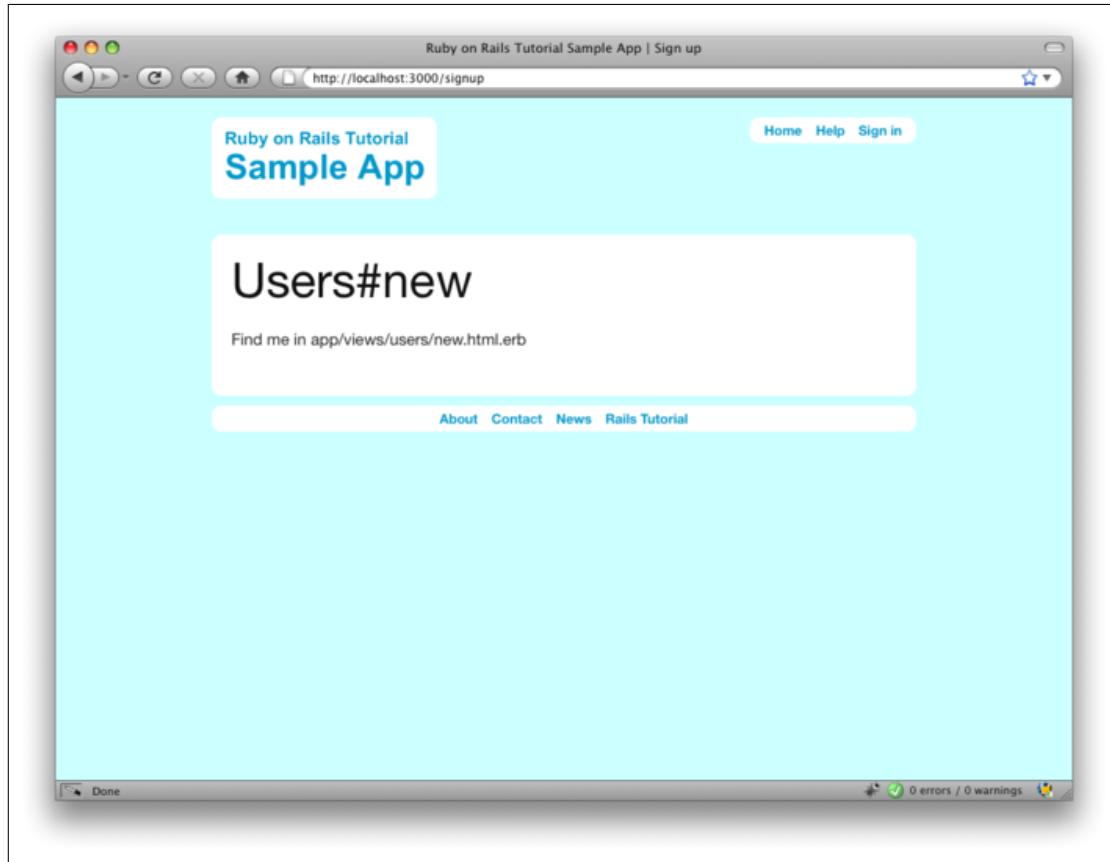


Figure 5.11: The new signup page at `/signup`. ([full size](#))

```
<h1>Sample App</h1>

<p>
  This is the home page for the
  <a href="http://www.railstutorial.org/">Ruby on Rails Tutorial</a>
  sample application.
</p>

<%= link_to "Sign up now!", signup_path, :class => "signup_button round" %>
```

With that, we're done with the links and named routes, at least until we add a route for signing in (Chapter 9). The resulting new user page (with URL `/signup`) appears in Figure 5.11.

5.4 Conclusion

We've hammered our application layout into shape and polished up the routes. If you're following along with Git, be sure to commit and merge (and, just to be paranoid, run your tests first):

```
$ spec spec/
$ git add .
$ git commit -am "Finished layout and routes"
$ git checkout master
$ git merge filling-in-layout
```

You might also want to push up to GitHub, or deploy to Heroku:

```
$ git push
$ git push heroku
```

5.5 Exercises

1. Replace the local variable `logo` in [Listing 5.18](#) with a helper method of the same name, so that the new partial looks like [Listing 5.28](#). Use the code in [Listing 5.29](#) to help you get started.
2. You may have noticed that our tests for the layout links test the routing but don't actually check that the links on the layout go to the right pages. One way to implement these tests is to use *Webrat* in an integration test. Install it by following the instructions at the [Webrat home page](#) and then fill in the code in [Listing 5.30](#).¹³

Listing 5.28. Header partial with the logo helper from [Listing 5.29](#).

`app/views/layouts/_header.html.erb`

```
<div id="header" class="round">
  <%= link_to logo, root_path %>
  <ul class="navigation round">
    <li><%= link_to "Home", root_path %></li>
    <li><%= link_to "Help", help_path %></li>
    <li><%= link_to "Sign in", '#' %></li>
  </ul>
</div>
```

Listing 5.29. A template for the `logo` helper.

`app/helpers/application_helper.rb`

¹³As noted by a helpful reader, the Webrat installation instructions specify some code to be added to `test/test_helper.rb`—which we deleted in Section 3.2.2 (since we're using RSpec instead of Test::Unit). As our intrepid reader correctly guessed, you should put the specified code in `spec/spec_helper.rb` instead.

```
module ApplicationHelper

  def logo
    # Fill in.
  end
end
```

Listing 5.30. A tests for the links on the layout.

`spec/integration/layout_links_spec.rb`

```
require 'spec_helper'

describe "Layout links" do
  .
  .
  .
  it "should have the right links on the layout" do
    visit root_path
    click_link "About"
    response.should render_template('pages/about')
    click_link "Help"
    response.should # fill in
    click_link "Contact"
    response.should # fill in
    click_link "Home"
    response.should # fill in
    click_link "Sign up now!"
    response.should # fill in
  end
end
```


Chapter 6

Modeling and viewing users, part I

In [Chapter 5](#), we ended with a stub page for creating new users ([Section 5.3](#)); over the course of the next three chapters, we'll fulfill the promise implicit in this incipient signup page. The first critical step is to create a *data model* for users of our site, together with a way to store that data. Completing this task is the goal for this chapter and the next ([Chapter 7](#)).

In parallel with our data modeling, we'll also develop a web page for *showing* users, which will give us a first taste of the REST architecture for structuring web applications. Though we won't get very far in this chapter, our eventual goal for the user profile pages is to show the user's profile image, basic user data, and a list of microposts, as mocked up in [Figure 6.1](#).¹ ([Figure 6.1](#) has our first example of *lorem ipsum* text, which has a [fascinating story](#) that you should definitely read about some time.) In this chapter, we'll lay the essential foundation for the user show page, and will start filling in the details starting in [Chapter 7](#).

As usual, if you're following along using Git for version control, now would be a good time to make a topic branch for modeling users:

```
$ git checkout master
$ git checkout -b modeling-users
```

(The first line here is just to make sure that you start on the master branch, so that the `modeling-users` topic branch is based on `master`. You can skip that command if you're already on the master branch.)

6.1 User model

Although the ultimate goal of the next three chapters is to make a signup page for our site, it would do little good to accept signup information now, since we don't currently have any place to put it. Thus, the first step in signing up users is to make a data structure to capture and store their information. In Rails, the default data structure for a data model is called, naturally enough, a *model* (the M in MVC from [Section 1.2.4](#)). The default Rails solution to the problem of persistence is to use a *database* for long-term data storage, and the default library for interacting with the database is called *Active Record*.²

¹ [Mockingbird](#) doesn't support custom images like the profile photo in [Figure 6.1](#); I put that in by hand using [Adobe Fireworks](#). The hippo here is from <http://www.flickr.com/photos/43803060@N00/24308857/>.

²The name comes from the "active record pattern", identified and named in *Patterns of Enterprise Application Architecture* by Martin Fowler.

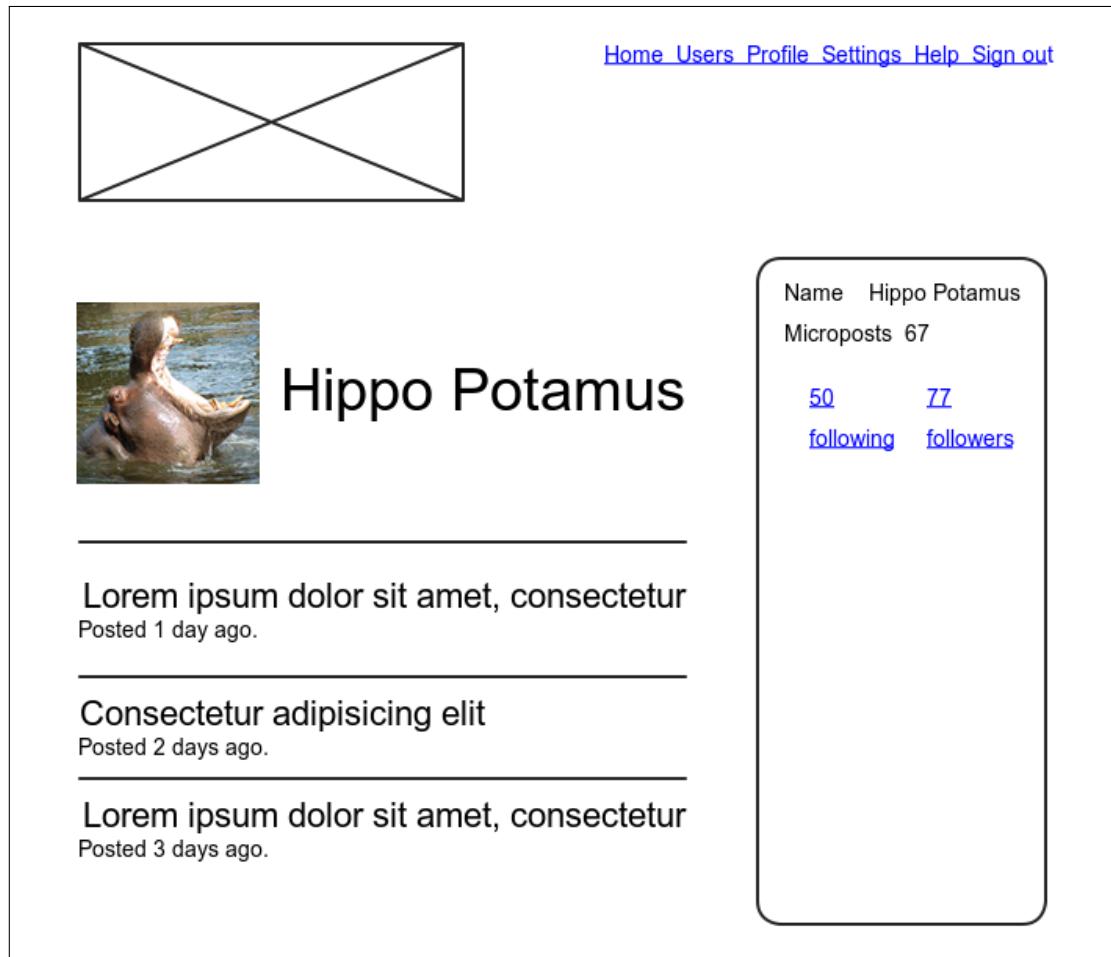


Figure 6.1: A mockup of our best guess at the user show page. ([full size](#))

Active Record comes with a host of methods for creating, saving, and finding data objects, all without having to use the structured query language (SQL)³ used by relational databases. Moreover, Rails has a feature called *migrations* to allow data definitions to be written in pure Ruby, without having to learn an SQL data definition language (DDL).⁴ The effect is that Rails insulates you almost entirely from the details of the data store. In this book, by using SQLite for development and Heroku for deployment (Section 1.4), we have developed this theme even further, to the point where we barely ever have to think about how Rails stores data, even for production applications.⁵

6.1.1 Database migrations

You may recall from Section 4.4.5 that we have already encountered, via a custom-built `User` class, user objects with `name` and `email` attributes. That class served as a useful example, but it lacked the critical property of *persistence*: when we created a User object at the Rails console, it disappeared as soon as we exited. Our goal in this section is to create a model for users that won't disappear quite so easily.

As with the User class in Section 4.4.5, we'll start by modeling a user with two attributes, a `name` and an `email` address, the latter of which we'll use as a unique username.⁶ (We'll add a password attribute in Section 7.1.) In Listing 4.6, we did this with Ruby's `attr_accessor` keyword:

```
class User
  attr_accessor :name, :email
  .
  .
  .
end
```

In contrast, when using Rails to model users we don't need to identify the attributes explicitly. As noted briefly above, to store data Rails uses a relational database by default, which consists of *tables* composed of data *rows*, where each row has *columns* of data attributes. For example, to store users with names and email addresses, we'll create a `users` table with `name` and `email` columns (with each row corresponding to one user). By naming the columns in this way, we'll let Active Record figure out the User object attributes for us.

Let's see how this works. (If this discussion gets too abstract for your taste, be patient; the console examples starting in Section 6.1.3 and the database browser screenshots in Figure 6.3 and Figure 6.8 should make things clearer.) In Section 5.3.1, recall that we created a Users controller (along with a `new` action with the command

```
$ script/generate rspec_controller Users new
$ rm -rf spec/views
```

³Pronounced “ess-cue-ell”, though the alternate pronunciation “sequel” is also common.

⁴In its earliest incarnations, Rails did require knowledge of an SQL DDL. Even after Rails added migrations, setting up the old default database (MySQL) was quite involved. Happily, as noted in Section 1.2.5, Rails now uses SQLite by default, which stores its data as a simple file—no setup required.

⁵Occasionally, it is necessary to pierce this abstraction layer, but one design goal of this tutorial is to make all the code database-independent. (Indeed, this is a worthy goal in general.) In case you ever do need to write database-specific code to deploy on Heroku, you should know that they use the excellent PostgreSQL (“post-gres-cue-ell”) database. PostgreSQL is free, open-source, and cross-platform; if you develop PostgreSQL-specific applications, you can install it locally, and configure Rails to use it in development by editing the `config/database.yml` file. Such configuration is beyond the scope of this tutorial, but there are lots of resources on the web; use a search engine to find the most up-to-date information for your platform.

⁶By using an email address as the username, we open the theoretical possibility of communicating with our users at a future date.

There is an analogous command for making a User model ([Listing 6.1](#)).

Listing 6.1. Generating a User model.

```
$ script/generate rspec_model User name:string email:string
```

(Note that, in contrast to the plural convention for controller names, model names are singular: a Users controller, but a User model.) By passing the optional parameters `name:string` and `email:string`, we tell Rails about the two attributes we want, along with what types those attributes have; compare this with including the action names in [Listing 3.2](#) and [Listing 5.20](#).

One of the results of the `generate` command in [Listing 6.1](#) is a new file called a *migration*. Migrations provide a way to alter the structure of the database incrementally, so that our data model can adapt to changing requirements. In the case of the User model, the migration is created automatically by the model generation script; it creates a `users` table with two columns, `name` and `email`, as shown in Listing 6.2. (We'll see in [Section 6.2.4](#) how to make a migration from scratch.)

Listing 6.2. Migration for the User model (to create a `users` table).

```
db/migrate/<timestamp>.create_users.rb
```

```
class CreateUsers < ActiveRecord::Migration
  def self.up
    create_table :users do |t|
      t.string :name
      t.string :email

      t.timestamps
    end
  end

  def self.down
    drop_table :users
  end
end
```

Note that the name of the migration is prefixed by a *timestamp* based on when the migration was generated. In the early days of migrations, the filenames were prefixed with incrementing integers, which caused conflicts for collaborating teams when two programmers had migrations with the same number. Using timestamps conveniently avoids such collisions.

Let's focus on the `self.up` method, which uses a Rails method called `create_table` to create a *table* in the database for storing users. (The use of `self` in `self.up` identifies it as a *class method*. This doesn't matter now, but we'll learn about class methods when we make one of our own in [Section 7.2.4](#).) The `create_table` method accepts a block ([Section 4.3.2](#)) with one block variable, in this case called `t` (for "table"). Inside the block, the `create_table` method uses the `t` object to create `name` and `email` columns in the database, both of type `string`.⁷ Here the table name is plural (`users`) even though the model name is singular (User), which reflects a linguistic convention followed by Rails: a model represents

⁷Don't worry about exactly how the `t` object manages to do this; the beauty of *abstraction layers* is that we don't have to know. We can just trust the `t` object to do its job.

users	
id	integer
name	string
email	string
created_at	datetime
updated_at	datetime

Figure 6.2: The users data model produced by Listing 6.2.

a single user, whereas a database table consists of many users. The final line in the block, `t.timestamps`, is a special command that creates two *magic columns* called `created_at` and `updated_at`, which are timestamps that automatically record when a given user is created and updated. (We’ll see concrete examples of the magic columns starting in Section 6.1.3.) The full data model represented by this migration is shown in Figure 6.2.

We can run the migration, known as “migrating up”, using the `rake` command (Box 1.2) as follows:⁸

```
$ rake db:migrate
```

(You may recall that we have run this command before, in Section 1.2.5 and again in Chapter 2.) The first time `db:migrate` is run, it creates a file called `db/development.sqlite3`, which is an `SQLite`⁹ database. We can see the structure of the database using the excellent `SQLite Database Browser` to open the `db/development.sqlite3` file (Figure 6.3); compare with the diagram in Figure 6.2. You might note that there’s one column in Figure 6.3 not accounted for in the migration: the `id` column. As noted briefly in Section 2.2, this column is created automatically, and is used by Rails to identify each row uniquely.

You’ve probably inferred that running `db:migrate` executes the `self.up` command in the migration file. What, then, of `self.down`? As you might guess, `down` migrates *down*, reversing the effects of migrating up. In our case, this means *dropping* the `users` table from the database:

```
class CreateUsers < ActiveRecord::Migration
  .
  .
  .
  def self.down
    drop_table :users
  end
end
```

You can execute `down` with `rake` using the argument `db:rollback`:

⁸We’ll see how to migrate up on a remote Heroku server in Section 7.4.2.

⁹Officially pronounced “ess-cue-ell-ite”, although the (mis)pronunciation “sequel-ite” is also common.

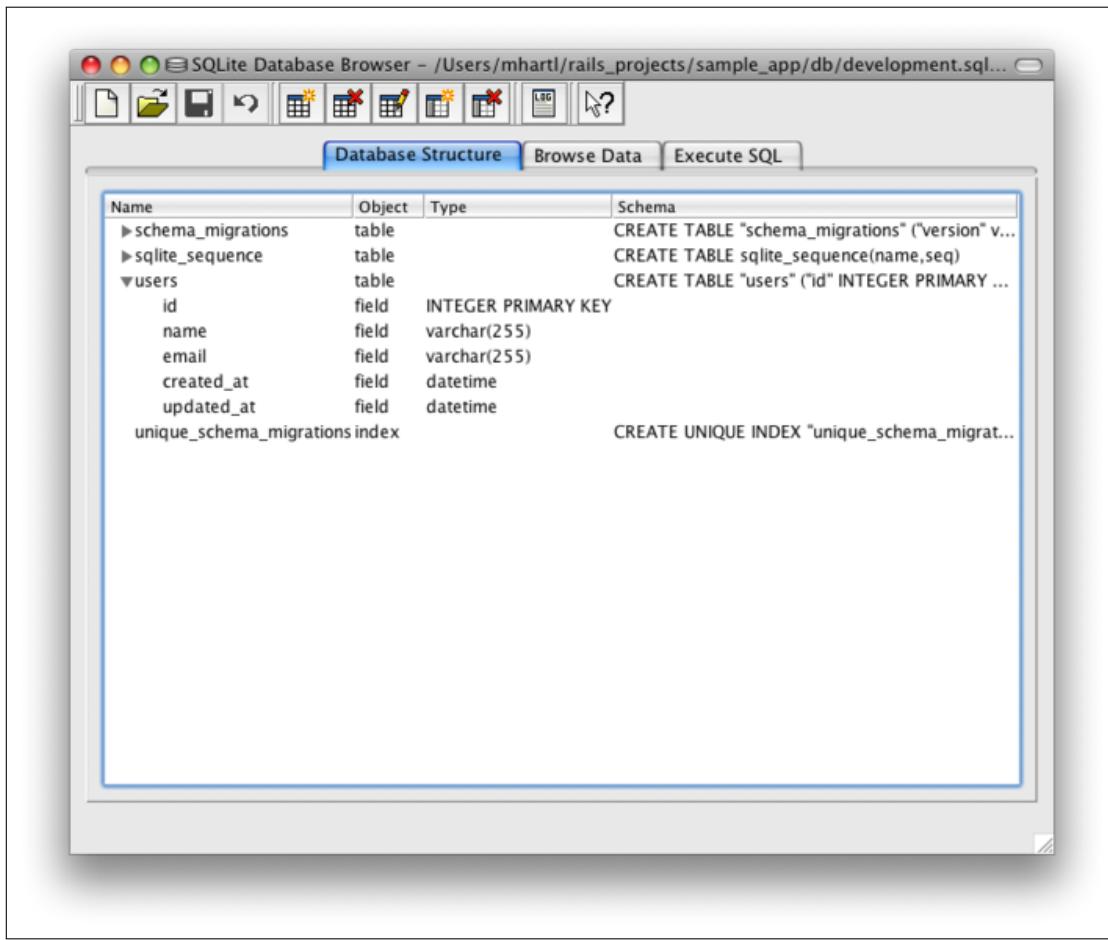


Figure 6.3: The SQLite Database Browser with our new `users` table. ([full size](#))

```
$ rake db:rollback
```

This is often useful if you realize there's another column you want to add but don't want the trouble of making a new migration: you can roll back the migration, add the desired column, and then migrate back up. (This isn't always convenient, and we'll learn how to add columns to an existing table in [Section 7.1.2](#).)

If you rolled back the database, migrate up again before proceeding:

```
$ rake db:migrate
```

6.1.2 The model file

We've seen how the User model generation in [Listing 6.1](#) generated a migration file ([Listing 6.2](#)), and we saw in [Figure 6.3](#) the results of running this migration: it updated a file called `development.sqlite3` by creating a table `users` with columns `id`, `name`, `email`, `created_at`, and `updated_at`. Listing 6.1 also created the model itself; the rest of this section is dedicated to understanding it.

We begin by looking at the code for the User model, which lives in the file `user.rb` inside the `app/models/` directory; it is, to put it mildly, very compact ([Listing 6.3](#)).

Listing 6.3. The brand new User model.

`app/models/user.rb`

```
class User < ActiveRecord::Base
end
```

Recall from [Section 4.4.2](#) that the syntax `class User < ActiveRecord::Base` means that the `User` class *inherits* from `ActiveRecord::Base`, so that the User model automatically has all the functionality of the `ActiveRecord::Base` class. Of course, knowledge of this inheritance doesn't do any good unless we know what `ActiveRecord::Base` contains, and we'll get a first taste starting momentarily. Before we move on, though, there are two tasks to complete.

Model annotation

Though it's not strictly necessary, I like to *annotate* my Rails models using the `annotate_models` plugin:

```
$ script/plugin install \
> http://repo.pragprog.com/svn/Public/plugins/annotate_models
$ rake annotate_models
```

This is our first example of a Rails *plugin*, which is a method for distributing additions to Rails without having to touch the core framework. In the present case, the plugin gives us a Rake task called `annotate_models`, which simply adds the comments containing the data model to the model file ([Listing 6.4](#)).

Listing 6.4. The annotated User model.`app/models/user.rb`

```
# == Schema Information
# Schema version: <timestamp>
#
# Table name: users
#
# id          :integer        not null, primary key
# name        :string(255)
# email       :string(255)
# created_at  :datetime
# updated_at  :datetime
#
class User < ActiveRecord::Base
end
```

I find that having the data model visible in the model files is useful for reminding me which attributes the model has. (Future code listings will usually omit the annotations for brevity.)

After adding a plugin, you should commit it before you forget:

```
$ git add .
$ git commit -am "Annotated models"
```

I find model annotation so useful that I usually add it to the Git *pre-commit hook*, which is a script that gets executed automatically before every Git commit. Throughout the rest of this tutorial, we will be adding new models and adding columns to existing models, and this way we don't have to remember to run `rake annotate_models` after every change. To get it to work, first fill the `.git/hooks/pre-commit` file with the contents of Listing 6.5.¹⁰

Listing 6.5. A Git pre-commit hook for annotating models.`.git/hooks/pre-commit`

```
#!/bin/sh
rake annotate_models > /dev/null
```

Then make the hook file executable:

```
$ chmod +x .git/hooks/pre-commit
```

Now the hook will execute every time we run `git commit`.

¹⁰I get tired of seeing the rather verbose annotation output on every commit, so in Listing 6.5 I redirect the output to `/dev/null`, which is basically a Unix black hole that swallows anything sent to it.

Accessible attributes

Another step that isn't strictly necessary but is a really good idea is to tell Rails which attributes of the model are accessible, i.e., which attributes can be modified by outside users (such as users submitting requests with web browsers). We do this with the `attr_accessible` method ([Listing 6.6](#)). We'll see in [Chapter 10](#) that using `attr_accessible` is important for preventing a *mass assignment* vulnerability, a distressingly common and often serious security hole in many Rails applications.

Listing 6.6. Making the `name` and `email` attributes accessible.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessible :name, :email
end
```

6.1.3 Creating user objects

We've done some good prep work, and now it's time to cash in and learn about Active Record by playing with our newly created User model. As in [Chapter 4](#), our tool of choice is the Rails console. Since we don't (yet) want to make any changes to our database, we'll start the console in a *sandbox*:

```
$ script/console --sandbox
Loading development environment in sandbox (Rails 2.3.8)
Any modifications you make will be rolled back on exit
>>
```

As indicated by the helpful message "Any modifications you make will be rolled back on exit", when started in a sandbox the console will "roll back" (i.e., undo) any database changes introduced during the session.

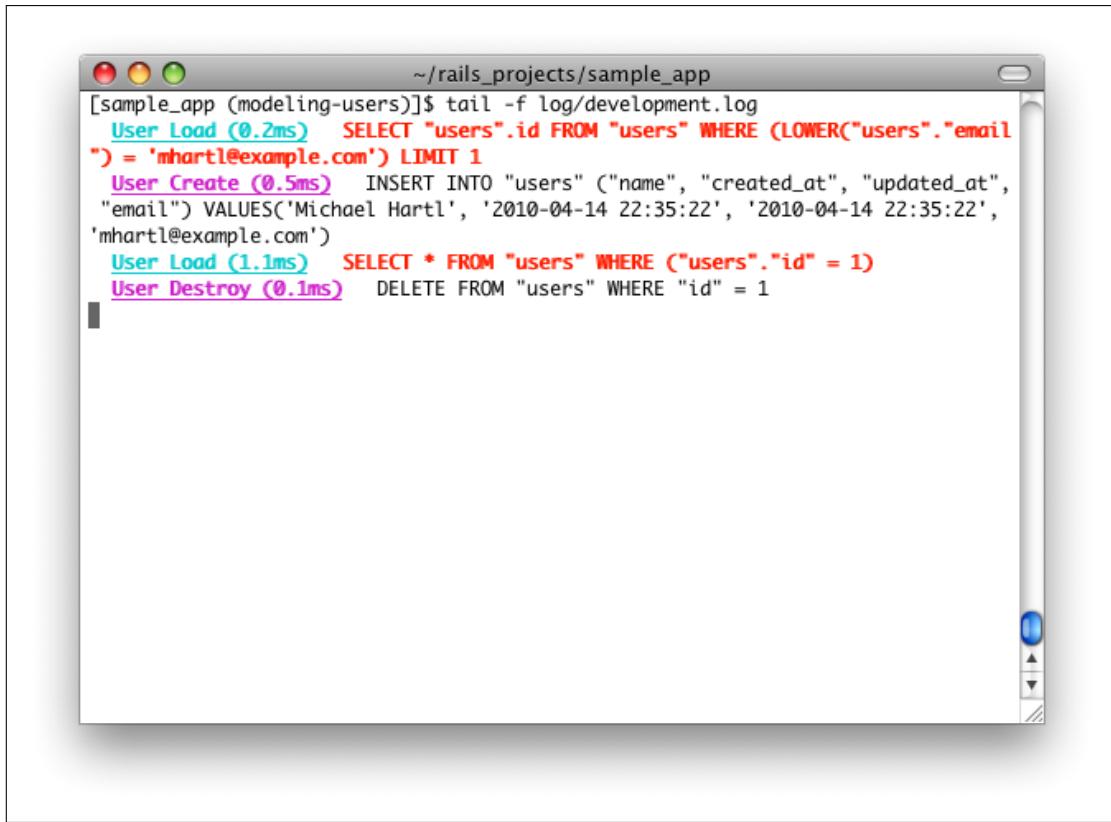
When working at the console, it's useful to keep an eye on the *development log*, which records the actual low-level SQL statements being issued by Active Record, as shown in [Figure 6.4](#). The way to get this output at a Unix command line is to `tail` the log:

```
$ tail -f log/development.log
```

The `-f` flag ensures that `tail` will display additional lines as they are written. I recommend keeping a terminal window for tailing the log open whenever working at the console.

In the console session in [Section 4.4.5](#), we created a new user object with `User.new`, which we had access to only after requiring the example user file in [Listing 4.6](#). With models, the situation is different; as you may recall from [Section 4.4.4](#), the Rails console automatically loads the Rails environment, which includes the models. This means that we can make a new user object without any further work:

```
>> User.new
=> #<User id: nil, name: nil, email: nil, created_at: nil, updated_at: nil>
```



```
~/rails_projects/sample_app
[sample_app (modeling-users)]$ tail -f log/development.log
User Load (0.2ms)  SELECT "users".id FROM "users" WHERE (LOWER("users"."email")
") = 'mhartl@example.com') LIMIT 1
User Create (0.5ms)  INSERT INTO "users" ("name", "created_at", "updated_at",
"email") VALUES('Michael Hartl', '2010-04-14 22:35:22', '2010-04-14 22:35:22',
'mhartl@example.com')
User Load (1.1ms)  SELECT * FROM "users" WHERE ("users"."id" = 1)
User Destroy (0.1ms)  DELETE FROM "users" WHERE "id" = 1
```

Figure 6.4: Tailing the development log. ([full size](#))

We see here the default console representation of a user object, which prints out the same attributes shown in Figure 6.3 and Listing 6.4.

When called with no arguments, `User.new` returns an object with all `nil` attributes. In Section 4.4.5, we designed the example User class to take an *initialization hash* to set the object attributes; that design choice was motivated by Active Record, which allows objects to be initialized in the same way:

```
>> user = User.new(:name => "Michael Hartl", :email => "mhartl@example.com")
#<User id: nil, name: "Michael Hartl", email: "mhartl@example.com",
created_at: nil, updated_at: nil>
```

Here we see that the name and email attributes have been set as expected.

If you've been tailing the development log, you may have noticed that no new lines have shown up yet. This is because calling `User.new` doesn't touch the database; it simply creates a new Ruby object in memory. To save the user object to the database, we call the `save` method on the `user` variable:

```
>> user.save
=> true
```

The `save` method returns `true` if it succeeds and `false` otherwise. (Currently, all saves should succeed; we'll see cases in Section 6.2 when some will fail.) As soon as you save, you should see a line in the development log with the SQL command to `INSERT INTO "users"`. Because of the many methods supplied by Active Record, we won't ever need raw SQL in this book, and I'll omit discussion of the SQL commands from now on. But you can learn a lot by watching the log.

You may have noticed that the new user object had `nil` values for the `id` and the magic columns `created_at` and `updated_at` attributes. Let's see if our `save` changed anything:

```
>> user
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">
```

We see that the `id` has been assigned a value of `1`, while the magic columns have been assigned the current time and date.¹¹ Currently the created and updated timestamps are identical; we'll see them differ in Section 6.1.5.

As with the User class in Section 4.4.5, instances of the User model allow access to their attributes using a dot notation:¹²

¹¹In case you're curious about `"2010-01-05 00:57:46"`, I'm not writing this after midnight; the timestamps are recorded in **Coordinated Universal Time** (UTC), which for most practical purposes is the same as **Greenwich Mean Time**. From the **NIST Time and Frequency FAQ**: Q: Why is UTC used as the acronym for Coordinated Universal Time instead of CUT? A: In 1970 the Coordinated Universal Time system was devised by an international advisory group of technical experts within the International Telecommunication Union (ITU). The ITU felt it was best to designate a single abbreviation for use in all languages in order to minimize confusion. Since unanimous agreement could not be achieved on using either the English word order, CUT, or the French word order, TUC, the acronym UTC was chosen as a compromise.

¹²Note the value of `user.updated_at`. Told you the timestamp was in UTC.

```
>> user.name
=> "Michael Hartl"
>> user.email
=> "mhartl@example.com"
>> user.updated_at
=> Tue, 05 Jan 2010 00:57:46 UTC +00:00
```

As we'll see in [Chapter 8](#), it's often convenient to make and save a model in two steps as we have above, but Active Record also lets you combine them into one step with `User.create`:

```
>> User.create(:name => "A Nother", :email => "another@example.org")
=> #<User id: 2, name: "A Nother", email: "another@example.org", created_at:
"2010-01-05 01:05:24", updated_at: "2010-01-05 01:05:24">
>> foo = User.create(:name => "Foo", :email => "foo@bar.com")
=> #<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2010-01-05
01:05:42", updated_at: "2010-01-05 01:05:42">
```

Note that `User.create`, rather than returning `true` or `false`, returns the User object itself, which we can optionally assign to a variable (such as `foo` in the second command above).

The inverse of `create` is `destroy`:

```
>> foo.destroy
=> #<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2010-01-05
01:05:42", updated_at: "2010-01-05 01:05:42">
```

Oddly, `destroy`, like `create`, returns the object in question, though I can't recall ever having used the return value of `destroy`. Even odder, perhaps, is that the `destroyed` object still exists in memory:

```
>> foo
=> #<User id: 3, name: "Foo", email: "foo@bar.com", created_at: "2010-01-05
01:05:42", updated_at: "2010-01-05 01:05:42">
```

How do we know if we really destroyed an object? And for saved and non-destroyed objects, how can we retrieve users from the database? It's time to learn how to use Active Record to find user objects.

6.1.4 Finding user objects

Active Record provides several options for finding objects. Let's use them to find the first user we created while verifying that the third user (`foo`) has been destroyed. We'll start with the existing user:

```
>> User.find(1)
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">
```

Here we've passed the id of the user to `User.find`; Active Record returns the user with that `id` attribute.

Let's see if the user with an `id` of `3` still exists in the database:

```
>> User.find(3)
ActiveRecord::RecordNotFound: Couldn't find User with ID=3
```

Since we destroyed our third user in [Section 6.1.3](#), Active Record can't find it in the database. Instead, `find` raises an *exception*, which is a way of indicating an exceptional event in the execution of a program—in this case, a nonexistent Active Record id, which causes `find` to raise an `ActiveRecord::RecordNotFound` exception.¹³

In addition to the generic `find`, Active Record also allows us to find users by specific attributes:

```
>> User.find_by_email("mhartl@example.com")
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
  created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">
```

Since we will be using email addresses as usernames, this sort of find will be useful when we learn how to let users sign in to our site ([Chapter 8](#)).¹⁴

We'll end with a couple of more general ways of finding users. First, there's `first`:

```
>> User.first
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
  created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">
```

Naturally, `first` just returns the first user in the database. There's also `all`:

```
>> User.all
=> [#<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
  created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">,
 #<User id: 2, name: "A Nother", email: "another@example.org", created_at:
  "2010-01-05 01:05:24", updated_at: "2010-01-05 01:05:24">]
```

No prizes for inferring that `all` returns an array ([Section 4.3.1](#)) of all users in the database.

6.1.5 Updating user objects

Once we've created objects, we often want to update them. There are two basic ways to do this. First, we can assign attributes individually, as we did in [Section 4.4.5](#):

¹³Exceptions and exception handling are somewhat advanced Ruby subjects, and we won't need them much in this book. They are important, though, and I suggest learning about them using one of the Ruby books recommended in [Section 1.1.1](#).

¹⁴To those worried that `find_by_email` will be inefficient if there are a large number of users, you're ahead of the game. We'll cover this issue, and its solution via database indices, in [Section 6.2.4](#).

```
>> user          # Just a reminder about our user's attributes
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
   created_at: "2010-01-05 00:57:46", updated_at: "2010-01-05 00:57:46">
>> user.email = "mhartl@example.net"
=> "mhartl@example.net"
>> user.save
=> true
```

Note that the final step is necessary to write the changes to the database. We can see what happens without a save by using `reload`, which reloads the object based on the database information:

```
>> user.email
=> "mhartl@example.net"
>> user.email = "foo@bar.com"
=> "foo@bar.com"
>> user.reload.email
=> "mhartl@example.net"
```

Now that we've updated the user, the magic columns differ, as promised in [Section 6.1.3](#):

```
>> user.created_at
=> "2010-01-05 00:57:46"
>> user.updated_at
=> "2010-01-05 01:37:32"
```

The second way to update attributes is to use `update_attributes`:

```
>> user.update_attributes(:name => "The Dude", :email => "dude@abides.org")
=> true
>> user.name
=> "The Dude"
>> user.email
=> "dude@abides.org"
```

The `update_attributes` method accepts a hash of attributes, and on success performs both the update and the save in one step (returning `true` to indicate that the save went through). It's worth noting that, once you have defined some attributes as accessible using `attr_accessible` ([Section 6.1.2](#)), *only* those attributes can be modified using `update_attributes`. If you ever find that your models mysteriously start refusing to update certain columns, check to make sure that those columns are included in the call to `attr_accessible`.

6.2 User validations

The User model we created in Section 6.1 now has working `name` and `email` attributes, but they are completely generic: any string (including an empty one) is currently valid in either case. And yet, names and email addresses are more specific than this. For example, `name` should be non-blank, and `email` should match the specific format characteristic of email addresses. Moreover, since we'll be using email addresses as unique usernames when users sign in, we shouldn't allow email duplicates in the database.

In short, we shouldn't allow `name` and `email` to be just any strings; we should enforce certain constraints on their values. Active Record allows us to impose such constraints using *validations*. In this section we'll cover several of the most common cases, validating *presence*, *length*, *format* and *uniqueness*. In Section 7.1.1 we'll add a final common validation, *confirmation*. And we'll see in Section 8.2 how validations give us convenient error messages when users make submissions that violate them.

As with the other features of our sample app, we'll add User model validations using test-driven development. Since we've changed the data model, it's a good idea to prepare the test database before proceeding:

```
$ rake db:test:prepare
```

This just ensures that the data model from the development database, `db/development.sqlite3`, is reflected in the test database, `db/test.sqlite3`.

6.2.1 Validating presence

We'll start with a test for the presence of a `name` attribute. Although the first step in TDD is to write a *failing* test (Section 3.2.2), in this case we don't yet know enough about validations to write the proper test, so we'll write the validation first, using the console to understand it. Then we'll comment out the validation, write a failing test, and verify that uncommenting the validation gets the test to pass. This procedure may seem pedantic for such a simple test, but I have seen¹⁵ many “simple” tests that test the wrong thing; being meticulous about TDD is simply the *only* way to be confident that we're testing the right thing. (This comment-out technique is also useful when rescuing an application whose application code is already written but—*quelle horreur!*—has no tests.)

The way to validate the presence of particular attribute is to use `validates_presence_of`, whose arguments are symbols indicating which attributes to validate—in this case, `name` (Listing 6.7).

Listing 6.7. Validating the presence of a `name` attribute.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessible :name, :email

  validates_presence_of :name
end
```

This may look like magic, but `validates_presence_of` is just a method, as indeed is `attr_accessible`; an equivalent formulation of Listing 6.7 is as follows:

¹⁵and written

```
class User < ActiveRecord::Base
  attr_accessible(:name, :email)

  validates_presence_of(:name)
end
```

Let's drop into the console to see the effects of adding a validation to our User model:¹⁶

```
$ script/console --sandbox
>> user = User.new(:name => "", :email => "mhartl@example.com")
>> user.save
=> false
>> user.valid?
=> false
```

Here `user.save` returns `false`, indicating a failed save. In the final command, we use the `valid?` method, which returns `false` when the object fails one or more validations, and `true` when all validations pass. (Recall from [Section 4.2.3](#) that Ruby uses a question mark to indicate such true/false *boolean* methods.) In this case, we only have one validation, so we know which one failed, but it can still be helpful to check using the `errors` object generated on failure:

```
>> user.errors.full_messages
=> ["Name can't be blank"]
```

(The error message is a hint that Rails validates the presence of an attributing the `blank?` method, which we saw at the end of [Section 4.4.2](#).)

Now for the failing test. To ensure that our incipient test will fail, let's comment out the validation at this point ([Listing 6.8](#)).

Listing 6.8. Commenting out a validation to ensure a failing test.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessible :name, :email

  # validates_presence_of :name
end
```

As in the case of controller generation (e.g., [Listing 5.20](#)), the model generate command in [Listing 6.1](#) produces an initial spec for testing users. I've modified the default a bit to produce the initial tests shown in [Listing 6.9](#).

¹⁶I'll omit the output of console commands when they are not particularly instructive—for example, the results of `User.new`.

Listing 6.9. The initial user spec.`spec/models/user_spec.rb`

```
# == Schema Information
#
# Table name: users
#
# id          :integer        not null, primary key
# name        :string(255)
# email       :string(255)
# created_at  :datetime
# updated_at  :datetime
#
# require 'spec_helper'

describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end

  it "should create a new instance given valid attributes" do
    User.create!(@attr)
  end

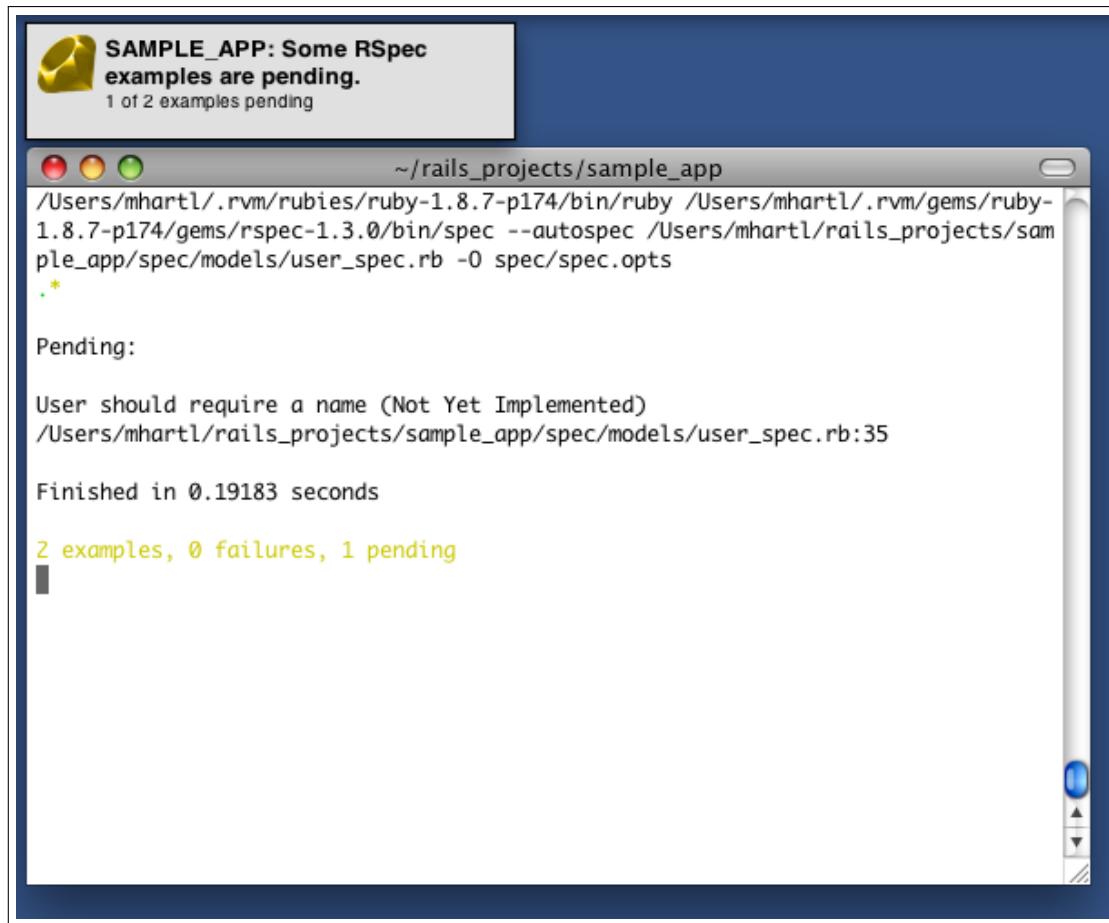
  it "should require a name"
end
```

Let's break this down. At the top is the same information about the data model that we saw in Listing 6.4, which (as in `user.rb`) was generated automatically by the `annotate_models` plugin. We've seen `require` and `describe` before, most recently in Listing 5.25. The next line is a `before(:each)` block; this was covered briefly in an exercise (Listing 3.25), and all it does is run the block before each test, in this case setting the `@attr` instance variable to an initialization hash.

The first test is just a sanity check, verifying that the `User` model is basically working. It uses `User.create!` (read “create bang”), which works just like the `create` method we saw in Section 6.1.3 except that it raises an `ActiveRecord::RecordInvalid` exception if the creation fails (similar to the `ActiveRecord::RecordNotFound` exception we saw in Section 6.1.4). As long as the attributes are valid, it won't raise any exceptions, and the test will pass.

The final line is the test for the presence of the `name` attribute—or rather, it *would* be the actual test, if it had anything in it. Instead, the test is just a stub, but a useful stub it is: it's a *pending spec*, which is a way to write a description of the application's behavior without worrying yet about the implementation. Pending specs are handled well by programs for running specs, as seen for Autotest in Figure 6.5. (The output of `spec spec/` is similarly useful.)

In order to fill in the pending spec, we need a way to make an attributes hash with an invalid name. (The `@attr` hash is valid by construction, with a non-blank `name` attribute.) The `Hash` method `merge` does the trick, as we can see with `script/console`:



The screenshot shows a terminal window titled "SAMPLE_APP: Some RSpec examples are pending." It displays the command used to run the tests and the resulting output. The output indicates 1 of 2 examples are pending, specifically a pending requirement for a user name.

```
~/rails_projects/sample_app
/Users/mhartl/.rvm/rubies/ruby-1.8.7-p174/bin/ruby /Users/mhartl/.rvm/gems/ruby-1.8.7-p174/gems/rspec-1.3.0/bin/spec --autospec /Users/mhartl/rails_projects/sample_app/spec/models/user_spec.rb -O spec/spec.opts
.*Pending:
User should require a name (Not Yet Implemented)
/Users/mhartl/rails_projects/sample_app/spec/models/user_spec.rb:35
Finished in 0.19183 seconds
2 examples, 0 failures, 1 pending
```

Figure 6.5: Autotest (via `autospec`) with a pending User spec. ([full size](#))

```
>> @attr = { :name => "Example User", :email => "user@example.com" }
=> { :name => "Example User", :email => "user@example.com" }
>> @attr.merge(:name => "")
=> { :name => "", :email => "user@example.com"}
```

With `merge` in hand, we're ready to make the new spec (using a trick I'll explain momentarily), as seen in Listing 6.10.

Listing 6.10. A failing test for validation of the `name` attribute.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end

  .

  .

  .

  it "should require a name" do
    no_name_user = User.new(@attr.merge(:name => ""))
    no_name_user.should_not be_valid
  end
end
```

Here we use `merge` to make a new user called `no_name_user` with a blank name. The second line then uses the RSpec `should_not` method to verify that the resulting user is *not* valid. The trick I alluded to above is related to `be_valid`: we know from earlier in this section that a `User` object responds to the `valid?` boolean method. RSpec adopts the useful convention of allowing us to test *any* boolean method by dropping the question mark and prepending `be_`. In other words,

```
no_name_user.should_not be_valid
```

is equivalent to

```
no_name_user.valid?.should_not == true
```

Since it sounds more like natural language, writing `should_not be_valid` is definitely more idiomatically correct RSpec.

With that, our new test should fail, which we can verify with Autotest or by running the `user_spec.rb` file using the `spec` script:

```
$ spec spec/models/user_spec.rb
.F

1)
'User should require a name' FAILED
expected valid? to return false, got true
./spec/models/user_spec.rb:14:

2 examples, 1 failure
```

Now uncomment the validation (i.e., revert [Listing 6.8](#) back to [Listing 6.7](#)) to get the test to pass:

```
$ spec spec/models/user_spec.rb
..
2 examples, 0 failures
```

Of course, we also want to validate the presence of email addresses. The test ([Listing 6.11](#)) is analogous to the one for the `name` attribute.

Listing 6.11. A test for presence of the `email` attribute.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end
  .

  it "should require an email address" do
    no_email_user = User.new(@attr.merge(:email => ""))
    no_email_user.should_not be_valid
  end
end
```

The only trick to the implementation is knowing that `validates_presence_of` (like `attr_accessible`) can take multiple arguments, as seen in [Listing 6.12](#).

Listing 6.12. Validating the presence of the `name` and `email` attributes.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessible :name, :email
```

```
validates_presence_of :name, :email
end
```

Now all the tests should pass, and the “presence of” validations are complete.

6.2.2 Length validation

We’ve constrained our User model to require a name for each user, but we should go further: the user’s names will be displayed on the sample site, so we should enforce some limit on their length. With all the work we did in [Section 6.2.1](#), this step is easy.

We start with a test. There’s no science to picking a maximum length; we’ll just pull **50** out of thin air as a reasonable upper bound, which means verifying that names of **51** characters are too long ([Listing 6.13](#)).

Listing 6.13. A test for `name` length validation.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end

  .

  .

  it "should reject names that are too long" do
    long_name = "a" * 51
    long_name_user = User.new(@attr.merge(:name => long_name))
    long_name_user.should_not be_valid
  end
end
```

For convenience, we’ve used “string multiplication” in [Listing 6.13](#) to make a string 51 characters long. We can see how this works using the console:

```
>> s = "a" * 51
=> "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa"
>> s.length
=> 51
```

The test in [Listing 6.13](#) should fail. To get it to pass, we need to know about the validation to constrain length, `validates_length_of`; in this case, we use a `:maximum` parameter to enforce the upper bound ([Listing 6.14](#)). (This is a one-element *options hash*; recall from [Section 4.3.4](#) that curly braces are optional when passing hashes as the final argument in a method. As noted in [Section 5.1.1](#), the use of options hashes is a recurring theme in Rails.)

Listing 6.14. Adding length validation for the `name` attribute.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessible :name, :email

  validates_presence_of :name, :email
  validates_length_of    :name, :maximum => 50
end
```

With our test suite passing again, we can move on to a more challenging validation: email format.

6.2.3 Format validation

Our validations for the `name` attribute enforce only minimal constraints—any non-blank name under 51 characters will do—but of course the `email` attribute must satisfy more stringent requirements. So far we’ve only rejected blank email addresses; in this section, we’ll require email addresses to conform to the familiar pattern `user@example.com`.

Neither the tests nor the validation will be exhaustive, just good enough to accept most valid email addresses and reject most invalid ones. We’ll start with a couple tests involving collections of valid and invalid addresses. To make these collections, it’s worth knowing about a useful method for making arrays of strings, as seen in this console session:

```
>> %w[foo bar baz]
=> ["foo", "bar", "baz"]
>> addresses = %w[user@foo.com THE_USER@foo.bar.org first.last@foo.jp]
=> ["user@foo.com", "THE_USER@foo.bar.org", "first.last@foo.jp"]
>> addresses.each do |address|
?>   puts address
>> end
user@foo.com
THE_USER@foo.bar.org
first.last@foo.jp
```

Here we’ve iterated over the elements of the `addresses` array using the `each` method (Section 4.3.2). With this technique in hand, we’re ready to write some basic email format validation tests (Listing 6.15).

Listing 6.15. Tests for email format validation.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end

  .
  .
  .

  it "should accept valid email addresses" do
```

```

addresses = %w[user@foo.com THE_USER@foo.bar.org first.last@foo.jp]
addresses.each do |address|
  valid_email_user = User.new(@attr.merge(:email => address))
  valid_email_user.should be_valid
end
end

it "should reject invalid email addresses" do
  addresses = %w[user@foo.com user_at_foo.org example.user@foo.]
  addresses.each do |address|
    invalid_email_user = User.new(@attr.merge(:email => address))
    invalid_email_user.should_not be_valid
  end
end
end

```

As noted above, these are far from exhaustive, but we do check the common valid email forms `user@foo.com`, `THE_USER@foo.bar.org` (uppercase, underscores, and compound domains), and `first.last@foo.jp` (the standard corporate username `first.last`, with a two-letter top-level domain `jp`), along with several invalid forms.

The application code for email format validation uses a *regular expression* (covered briefly in Section 5.3.1) to define the format, along with the Rails `validates_format_of` method (Listing 6.16).

Listing 6.16. Validating the email format with a regular expression.

`app/models/user.rb`

```

class User < ActiveRecord::Base
  attr_accessible :name, :email

  EmailRegex = /\A[\w+\-\.]+\@[a-z\d\-\.]+\.[a-z]+\z/i

  validates_presence_of :name, :email
  validates_length_of   :name, :maximum => 50
  validates_format_of   :email, :with => EmailRegex
end

```

Here `EmailRegex` is a Ruby *constant*, as indicated by the leading capital letter.¹⁷ (Using a constant is simply conventional in this context; we could have used a variable, too.) The value of `EmailRegex` is a *regular expression*, also known as a *regex*. The code

```

EmailRegex = /\A[\w+\-\.]+\@[a-z\d\-\.]+\.[a-z]+\z/i
.
.
.
validates_format_of :email, :with => EmailRegex

```

¹⁷It's not a coincidence that it looks like a `ClassName`—class names are constants, too.

Expression	Meaning
/\A[\w+\-\.]+@[a-z\d\-\.]+\.[a-z]+\z/i	full regex
/	start of regex
\A	match start of a string
[\w+\-\.]+	at least one word character, plus, hyphen, or dot
@	literal “at sign”
[a-z\d\-\.]+	at least one letter, digit, hyphen, or dot
\.	literal dot
[a-z]+	at least one letter
\z	match end of a string
/	end of regex
i	case insensitive

Table 6.1: Breaking down the email regex from Listing 6.16.

ensures that only email addresses that match the pattern will be valid.

So, what of the pattern? Regular expressions consist of a terse (some would say [unreadable](#)) language for matching text patterns; learning to construct regexes is an art, and to get you started I’ve broken [Email-Regex](#) into bite-sized pieces ([Table 6.1](#)).¹⁸ To really learn about regular expressions, though, I consider the amazing [Rubular](#) regular expression editor ([Figure 6.6](#)) to be simply essential.¹⁹ The Rubular website has a beautiful interactive interface for making regular expressions, along with a handy regex quick reference. I encourage you to study [Table 6.1](#) with a browser window open to Rubular—no amount of reading about regular expressions can replace a couple of hours playing with Rubular. (Video is a better format than print for understanding regexes, and I plan to cover them in more depth in the [Rails Tutorial screencasts](#).)

By the way, there actually exists a full regex for matching email addresses according to the official standard, but it’s really not worth the trouble. The one in [Listing 6.16](#) is fine, maybe even better than the official one.²⁰

The tests should all be passing now. (In fact, the tests for valid email addresses should have been passing all along; since regexes are notoriously error-prone, the valid email tests are there mainly as a sanity check on [EmailRegex](#).) This means that there’s only one constraint left: enforcing the email addresses to be unique.

6.2.4 Uniqueness validation

To enforce uniqueness of email addresses (so that we can use them as usernames), we’ll be using the method [validates_uniqueness_of](#). But be warned: there’s a *major* caveat, so don’t just skim this section—read it carefully.

We’ll start, as usual, with our tests. In our previous model tests, we’ve mainly used [User.new](#), which just creates a Ruby object in memory, but for uniqueness tests we actually need to put a record into the database.²¹ The (first) duplicate email test appears in [Listing 6.17](#).

¹⁸Note that, in [Table 6.1](#), “letter” really means “lower-case letter”, but the `i` at the end of the regex enforces case-insensitive matching.

¹⁹If you find it as useful as I do, I encourage you to [donate to Rubular](#) to reward developer [Michael Lovitt](#) for his wonderful work.

²⁰Did you know that `"Michael Hartl"@example.com`, with quotation marks and a space in the middle, is a valid email address according to the standard? Incredibly, it is—but it’s absurd. If you don’t have an email address that contains only letters, numbers, underscores, and dots, then get one. N.B. The regex in [Listing 6.16](#) allows plus signs, too, because Gmail (and possibly other email services) does something useful with them: for example, to filter orders from Amazon, you can use `username+amazon@gmail.com`, which will go to the Gmail address `username@gmail.com`, allowing you to filter on the string `amazon`.

²¹As noted briefly in the introduction to this section, there is a dedicated test database, `db/test.sqlite3`, for this purpose.

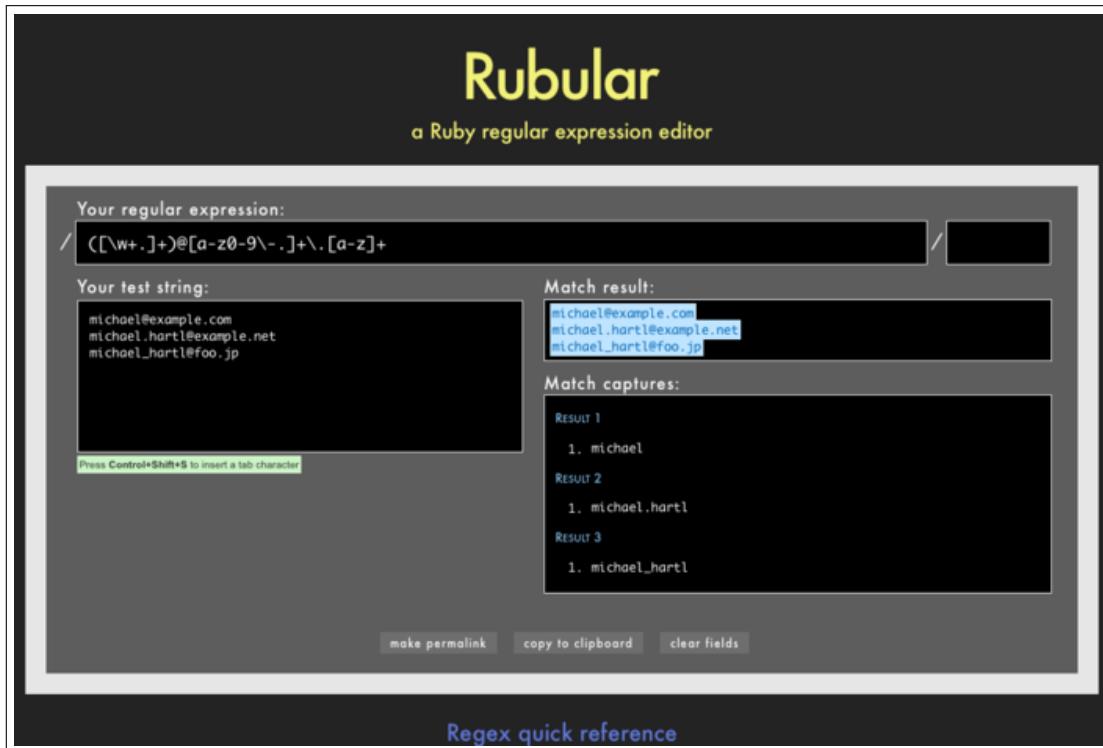


Figure 6.6: The awesome [Rubular](#) regular expression editor. ([full size](#))

Listing 6.17. A test for the rejection of duplicate email addresses.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = { :name => "Example User", :email => "user@example.com" }
  end

  .

  .

  it "should reject duplicate email addresses" do
    # Put a user with given email address into the database.
    User.create!(@attr)
    user_with_duplicate_email = User.new(@attr)
    user_with_duplicate_email.should_not be_valid
  end
end
```

The method here is to create a user and then try to make another one with the same email address. (We use the noisy method `create!`, first seen in Listing 6.9, so that it will raise an exception if anything goes wrong. Using `create`, without the bang `!`, risks having a silent error in our test, a potential source of elusive bugs.) We can get this test to pass with the code in Listing 6.18.²²

Listing 6.18. Validating the uniqueness of email addresses.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .

  .

  validates_format_of      :email, :with => EmailRegex
  validates_uniqueness_of :email
end
```

We're not quite done, though. Email addresses are case-insensitive—`foo@bar.com` goes to the same place as `FOO@BAR.COM` or `FoO@Bar.coM`—so our validation should cover this case as well. We test for this with the code in Listing 6.19.

Listing 6.19. A test for the rejection of duplicate email addresses, insensitive to case.

`spec/models/user_spec.rb`

```
describe User do
```

²²If you're wondering why the `create!` line in Listing 6.9 doesn't cause this to fail by creating a duplicate user, it's because Rails tests are *transactional*: each test is wrapped in a `transaction`, which *rolls back* the database after the test executes. This way, each test runs against a fresh database.

```

before(:each) do
  @attr = { :name => "Example User", :email => "user@example.com" }
end

.
.

it "should reject email addresses identical up to case" do
  upcased_email = @attr[:email].upcase
  User.create!(@attr.merge(:email => upcased_email))
  user_with_duplicate_email = User.new(@attr)
  user_with_duplicate_email.should_not be_valid
end
end

```

Here we are using the `upcase` method on strings (seen briefly in [Section 4.3.2](#)). This test does the same thing as the first duplicate email test, but with an upper-case email address instead. If this test feels a little abstract, go ahead and fire up the console:

```

$ script/console --sandbox
>> @attr = { :name => "Example User", :email => "user@example.com" }
=> {:name => "Example User", :email => "user@example.com"}
>> upcased_email = @attr[:email].upcase
=> "USER@EXAMPLE.COM"
>> User.create!(@attr.merge(:email => upcased_email))
>> user_with_duplicate_email = User.new(@attr)
>> user_with_duplicate_email.valid?
=> true

```

Of course, currently `user_with_duplicate_email.valid?` is `true`, since this is a failing test, but we want it to be `false`. Fortunately, `validates_uniqueness_of` accepts an option, `:case_insensitive`, for just this purpose ([Listing 6.20](#)).

Listing 6.20. Validating the uniqueness of email addresses, ignoring case.

`app/models/user.rb`

```

class User < ActiveRecord::Base
  .
  .
  .
  validates_format_of :email, :with => EmailRegex
  validates_uniqueness_of :email, :case_sensitive => false
end

```

With that, our tests pass, and email uniqueness is validated.

The uniqueness caveat

There's just one small problem, the caveat alluded to above:

Using `validates_uniqueness_of` does not guarantee uniqueness.

D'oh! But what can go wrong? Here's what:

1. Alice signs up for the sample app, with address `alice@wonderland.com`.
2. Alice accidentally clicks on “Submit” *twice*, sending two requests in quick succession.
3. The following sequence occurs: request 1 creates a user in memory that passes validation, request 2 does the same, request 1’s user gets saved, request 2’s user gets saved.
4. Result: two user records with the exact same email address, despite the uniqueness validation.

If the above sequence seems implausible, believe me, it isn’t: it happens on any Rails website with significant traffic.²³ Luckily, the solution is straightforward to implement; we just need to enforce uniqueness at the database level as well. Our method is to create a database *index* on the email column, and then require that the index be unique.

The email index represents an update to our data modeling requirements, which (as discussed in [Section 6.1.1](#)) is handled in Rails using migrations. We saw in [Section 6.1.1](#) that generating the User model automatically created a new migration ([Listing 6.2](#)); in the present case, we are adding structure to an existing model, so we need to create a migration directly using the `migration` generator:

```
$ script/generate migration add_email_uniqueness_index
```

Unlike the migration for users, the email uniqueness migration is not pre-defined, so we need to fill in its contents with [Listing 6.21](#).²⁴

Listing 6.21. The migration for enforcing email uniqueness.

`db/migrate/<timestampe>.add_email_uniqueness_index.rb`

```
class AddEmailUniquenessIndex < ActiveRecord::Migration
  def self.up
    add_index :users, :email, :unique => true
  end

  def self.down
    remove_index :users, :email
  end
end
```

This uses a Rails method called `add_index` to add an index on the `email` column of the `users` table. The index by itself doesn’t enforce uniqueness, but the option `:unique => true` does.

The final step is to migrate the database:

²³Yes, it happened to me. How do you think I found out about this issue. ;-)

²⁴Of course, we could just edit the migration file for the `users` table in [Listing 6.2](#) but that would require rolling back and then migrating back up. The Rails Way is to use migrations every time we discover that our data model needs to change.

```
$ rake db:migrate
```

Now the Alice scenario above will work fine: the database will save a user record based on the first request, and will reject the second save for violating the uniqueness constraint. (An error will appear in the Rails log, but that doesn't do any harm. You can actually catch the `ActiveRecord::StatementInvalid` exception that gets raised—see [here](#) for an example—but in this tutorial we won't bother with this step.) Adding this index on the email attribute accomplishes a second goal, alluded to briefly in [Section 6.1.4](#): it fixes an efficiency problem in `find_by_email` ([Box 6.1](#)).

Box 6.1. Database indices

When creating a column in a database, it is important to consider if we will need to *find* records by that column. Consider, for example, the `email` attribute created by the migration in [Listing 6.2](#). When we allow users to sign in to the sample app starting in [Chapter 8](#), we will need to find the user record corresponding to the submitted email address; unfortunately, based on the naïve data model, the only way to find a user by email address is to look through *each* user row in the database and compare its email attribute to the given email. This is known in the database business as a *full-table scan*, and for a real site with thousands of users it is a [Bad Thing](#).

Putting an index on the email column fixes the problem. To understand a database index, it's helpful to consider the analogy of a book index. In a book, to find all the occurrences of a given string, say “foobar”, you would have to scan each page for “foobar”. With a book index, on the other hand, you can just look up “foobar” in the index to see all the pages containing “foobar”. (Of course, with *this* book you can always just [view Rails Tutorial as a single page](#) or [get the PDF](#) and search it electronically.) A database index works essentially the same way.

6.3 Viewing users

We're not quite done with the basic user model—we need to add passwords, a task for [Chapter 7](#)—but we do have enough in place to make a minimalist page for showing user information. This will allow a gentle introduction to the REST style of organizing the actions for our site's users. Since this is just a rough demonstration for now, there are no tests in this section; we'll add tests when we flesh out the user view in [Section 7.3](#).

6.3.1 Debug and Rails environments

As preparation for adding dynamic pages to our sample application, now is a good time to add some debug information to our site layout ([Listing 6.22](#)). This displays some useful information about each page using the built-in `debug` method and `params` variable (which we'll learn more about in [Section 6.3.2](#)), as seen in [Figure 6.7](#).

Listing 6.22. Adding some debug information to the site layout.

```
app/views/layouts/application.html.erb
```

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html>
.
.
.
<body>
<div class="container">
.
.
.
<%= render 'layouts/footer' %>
<%= debug(params) if Rails.env.development? %>
</div>
</body>
</html>
```

Since we don't want to display debug information to users of a deployed application, we use

```
if Rails.env.development?
```

to restrict the debug information to the *development environment*. Though we've seen evidence of Rails environments before (most recently in Section 6.1.3), this is the first time it has mattered to us.

Rails comes equipped with three environments: `test`, `development`, and `production`.²⁵ The default environment for the Rails console is `development`:

```
$ script/console
Loading development environment (Rails 2.3.8)
>> Rails.env
=> "development"
>> Rails.env.development?
=> true
>> Rails.env.test?
=> false
```

As you can see, Rails provides a `Rails` object with an `env` attribute and associated environment boolean methods. In particular, `Rails.env.development?` is `true` only in a development environment, so the Embedded Ruby

```
<%= debug(params) if Rails.env.development? %>
```

²⁵You can define your own custom environments as well; see the [Railscast on adding an environment](#) for details.

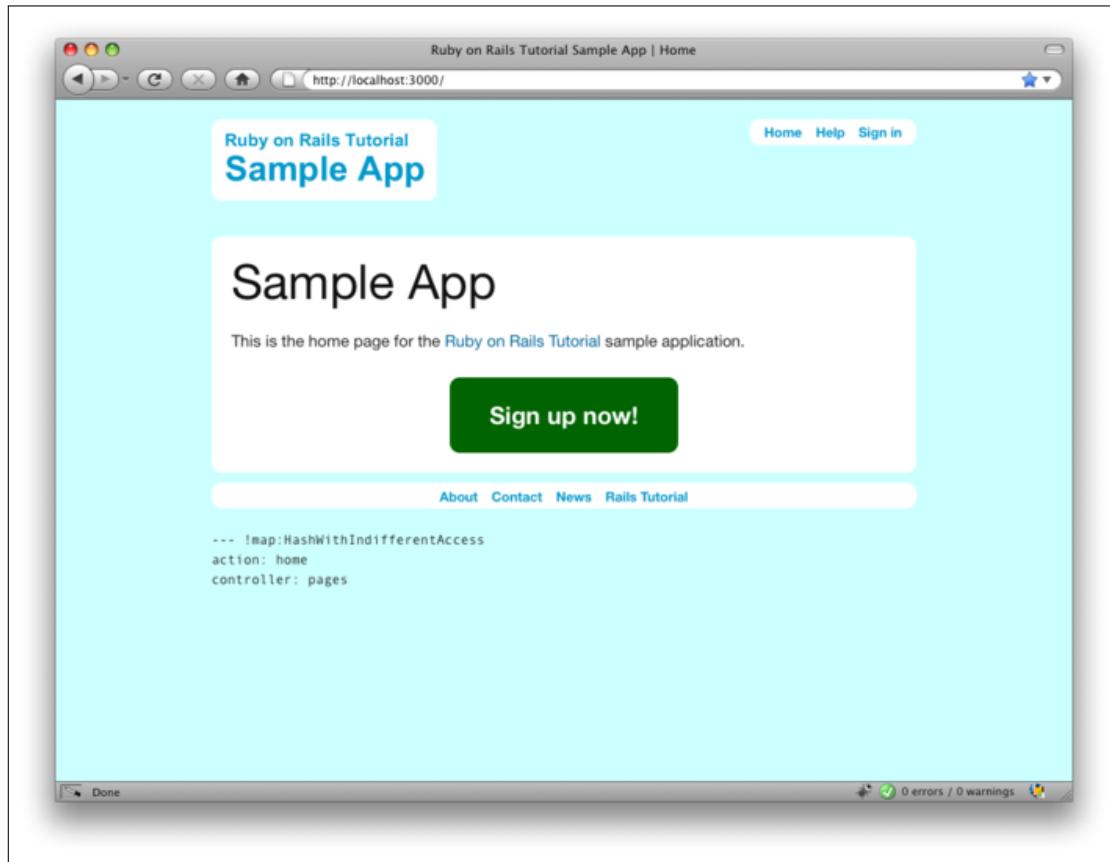


Figure 6.7: The sample application Home page ([/](#)) with debug information at the bottom. ([full size](#))

won't be inserted into production applications or tests. (Inserting the debug information into tests probably doesn't do any harm, but it probably doesn't do any good, either, so it's best to restrict the debug display to development only.)

If you ever need to run a console in a different environment (to debug a test, for example), you can pass the environment as a parameter to the `console` script:

```
$ script/console test
Loading test environment (Rails 2.3.8)
>> Rails.env
=> "test"
```

As with the console, `development` is the default environment for the local Rails server, but you can also run it in a different environment:

```
$ script/server --environment production
```

If you view your app running in production, it won't work without a production database, which we can create by running `rake db:migrate` in production:

```
$ rake db:migrate RAILS_ENV=production
```

(I find it confusing that the console, server, and migrate commands specify non-default environments in three mutually incompatible ways, which is why I bothered showing all three.)

By the way, if you have deployed your sample app to Heroku, you can see its environment using the `heroku` gem, which provides its own (remote) console:

```
$ heroku console
Ruby console for yourapp.herokuapp.com
>> Rails.env
=> "production"
>> Rails.env.production?
=> true
```

Naturally, since Heroku is a platform for production sites, it runs each application in a production environment.

6.3.2 User model, view, controller

In order to make a page to view a user, we'll use the *User model* to put a user into the database, make a *view* to display some user information, and then add an action to the *Users controller* to handle the browser request. In other words, for the first time in this tutorial, we'll see in one place all three elements of the model-view-controller architecture first discussed in [Section 1.2.4](#).

Our first step is to create a user using the console, which we'll take care *not* to start in a sandbox since this time the whole point is to save a record to the database:

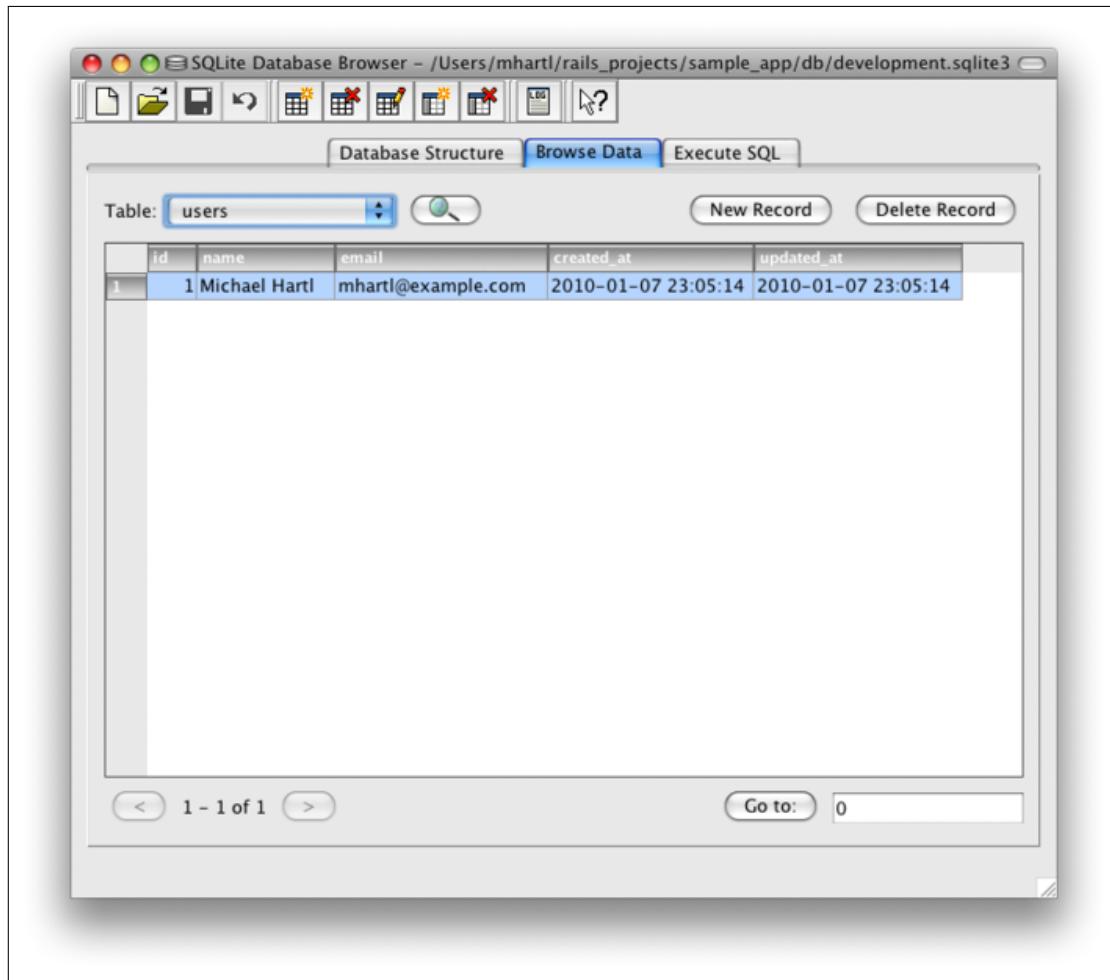


Figure 6.8: A user row in the SQLite database `db/development.sqlite3`. ([full size](#))

```
$ script/console
Loading development environment (Rails 2.3.8)
>> User.create!(:name => "Michael Hartl", :email => "mhartl@example.com")
=> #<User id: 1, name: "Michael Hartl", email: "mhartl@example.com",
  created_at: "2010-01-07 23:05:14", updated_at: "2010-01-07 23:05:14">
```

To double-check that this worked, let's look at the row in the development database using the SQLite Database Browser (Figure 6.8). Note that the columns correspond to the attributes of the data model defined in Section 6.1.

Next comes the view, which is minimalist to emphasize that this is just a demonstration (Listing 6.23). We use the standard Rails location for showing a user, `app/views/users/show.html.erb`; unlike the `new.html.erb` view, which we created with the generator in Listing 5.20, the `show.html.erb` file

doesn't currently exist, so you'll have to create it by hand.

Listing 6.23. A stub view for showing user information.

`app/views/users/show.html.erb`

```
<%= @user.name %>, <%= @user.email %>
```

This view uses Embedded Ruby to display the user's name and email address, assuming the existence of an instance variable called `@user`. Of course, eventually the real user show page will look very different, and won't display the email address publicly.

Finally, we'll add the `show` action to the Users controller (corresponding to the `show.html.erb` view) with the code in Listing 6.24, which defines the `@user` instance variable needed by the view.

Listing 6.24. The Users controller with a `show` action.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def show
    @user = User.find(params[:id])
  end

  def new
    @title = "Sign up"
  end
end
```

Here we've gotten a little ahead of ourselves by using the standard Rails `params` object to retrieve the user id. When we make the appropriate request to the Users controller, `params[:id]` will be the user id `1`, so the effect is the same as the `find` command

```
User.find(1)
```

we saw in Section 6.1.4.

With the `show` action defined as in Listing 6.24, we can view the user show page at the url `/users/show/1` (Figure 6.9), which adds `id` to the traditional `/controller/action/id` pattern we first encountered in Section 3.1.2.

If you look closely at the debug information in Figure 6.9, you'll see the following:²⁶

```
--- !map:HashWithIndifferentAccess
action: show
id: "1"
controller: users
```

²⁶The Rails `debug` information is shown as `YAML` (a recursive acronym standing for “YAML Ain’t Markup Language”), which is a friendly data format designed to be both machine- *and* human-readable.

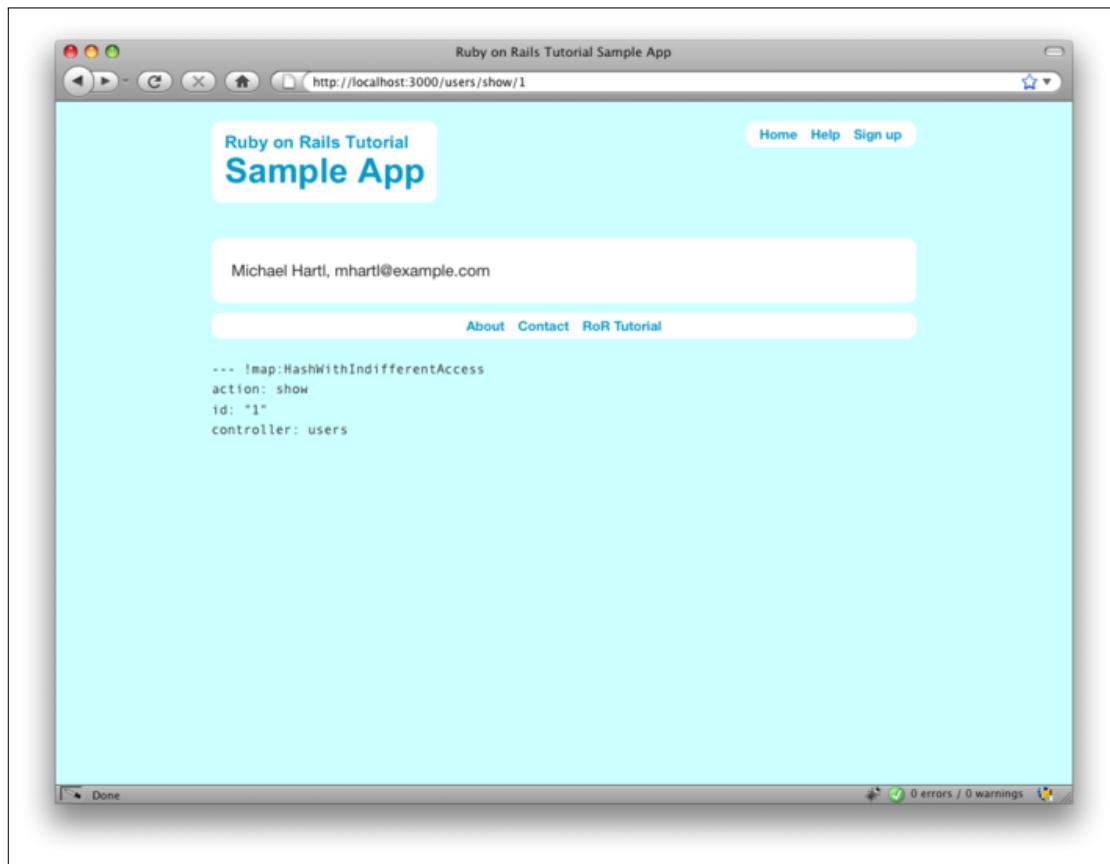


Figure 6.9: The user show page at </users/show/1>. (full size)

This is a representation of `params` (which is basically a hash), and we see that its `id` attribute is "`1`". This is why

```
User.find(params[:id])
```

finds the user with id `1`. (The `find` method knows how to convert the string "`1`" into the integer `1`.)

6.3.3 A Users resource

Although the user show URL in [Section 6.3.2](#) works fine for now, it doesn't fit in well with the REST architecture favored in Rails applications. This style is based on the ideas of *representational state transfer* identified and named by computer scientist [Roy Fielding](#) in his doctoral dissertation *Architectural Styles and the Design of Network-based Software Architectures*.²⁷ The REST design style emphasizes representing data as *resources* that can be created, shown, updated, or destroyed—four actions corresponding to the four fundamental operations GET, POST, PUT, and DELETE defined by the [HTTP standard](#) ([Box 3.1](#)).

When following REST principles, resources are typically referenced using the resource name and a unique identifier. What this means in the context of users—which we're now thinking of as a *Users resource*—is that we should view the user with id `1` by issuing a GET request to the URL `/users/1`. Here the `show` action is *implicit* in the type of request—when Rails' REST features are activated, GET requests are *automatically* handled by the `show` action.

Unfortunately, the URL `/users/1` doesn't work quite yet ([Figure 6.10](#)). By default, Rails only recognizes the `/controller/action/id` pattern, but recall that Rails URLs are determined by a routes file ([Section 5.2.2](#)). We can get the REST-style uses URL to work by adding users as a resource to `config/routes.rb` ([Listing 6.25](#)).

Listing 6.25. Adding a Users resource to the routes file.

```
config/routes.rb
```

```
ActionController::Routing::Routes.draw do |map|
  map.resources :users
  .
  .
  .
end
```

Now that we've added the Users resource, the URL `/users/1` works perfectly ([Figure 6.11](#)).

The one additional resource line in [Listing 6.25](#) doesn't just add a working `/users/1` URL; it endows our sample application with all the actions needed for a RESTful Users resource,²⁸ along with a large number of named routes ([Section 5.2.3](#)) for generating user URLs. The resulting correspondence of URLs, actions, and named routes is shown in [Table 6.2](#). (Compare to [Table 2.2](#).) Over the course of the next three chapters, we'll cover all of the other entries in [Table 6.2](#) as we fill in all the actions necessary to make users a fully RESTful resource.²⁹

²⁷Fielding, Roy Thomas. *Architectural Styles and the Design of Network-based Software Architectures*. Doctoral dissertation, University of California, Irvine, 2000.

²⁸This means that the *routing* works, but the corresponding pages don't necessarily work at this point. For example, `/users/1/edit` gets routed properly to the `edit` action of the Users controller, but since the `edit` action doesn't exist yet ac-

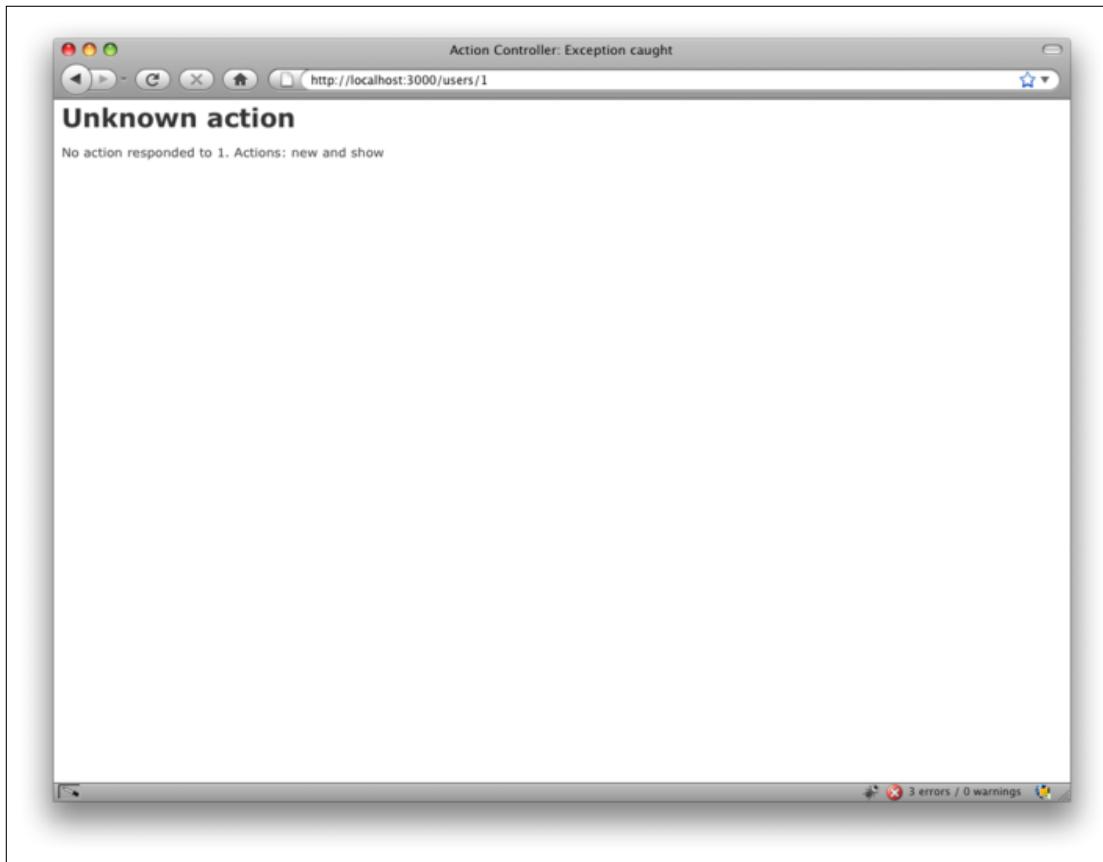


Figure 6.10: The initial effect of hitting `/users/1.` ([full size](#))

HTTP request	URL	Action	Named route	Purpose
GET	/users	<code>index</code>	<code>users_path</code>	page to list all users
GET	/users/1	<code>show</code>	<code>user_path(1)</code>	page to show user with id 1
GET	/users/new	<code>new</code>	<code>new_user_path</code>	page to make a new user (signup)
POST	/users	<code>create</code>	<code>users_path</code>	create a new user
GET	/users/1/edit	<code>edit</code>	<code>edit_user_path(1)</code>	page to edit user with id 1
PUT	/users/1	<code>update</code>	<code>user_path(1)</code>	update user with id 1
DELETE	/users/1	<code>destroy</code>	<code>user_path(1)</code>	delete user with id 1

Table 6.2: RESTful routes provided by the Users resource in [Listing 6.25](#).

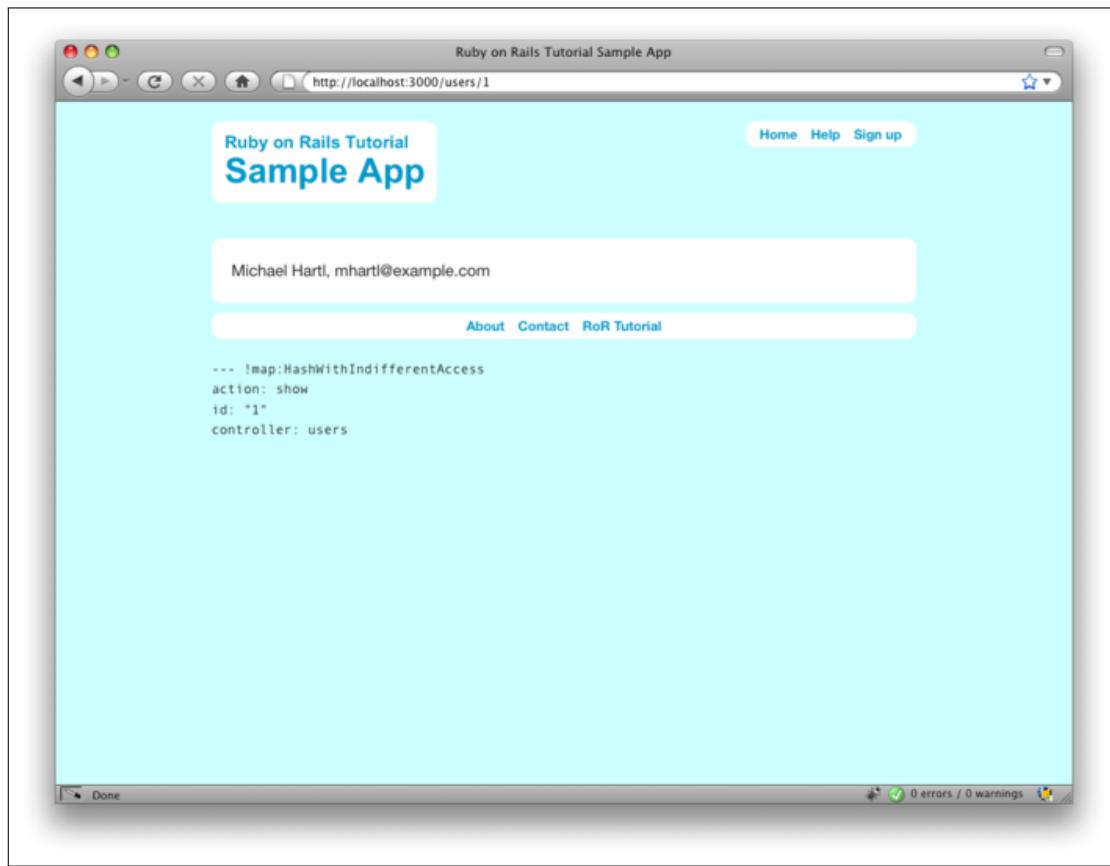


Figure 6.11: The user show page at `/users/1` after adding a Users resource. ([full size](#))

6.4 Conclusion

This chapter is the first half of the two-step process of creating a working User model. Our users now have `name` and `email` attributes, together with validations enforcing several important constraints on their values. We've also taken a first small steps toward a working user show page and a Users resource based on the principles of representational state transfer (REST). In [Chapter 7](#), we'll complete the process by adding user passwords and a more useful user view.

If you're using Git, now would be a good time to commit if you haven't done so in a while:

```
$ git add .
$ git commit -am "Finished first cut of the User model"
```

6.5 Exercises

1. Read through the [Rails API entry for ActiveRecord::Base](#) to get a sense of its capabilities.
2. Study the entries in the Rails API for each of the `validates` methods in this chapter to learn more about their capabilities and options.
3. Spend a couple hours playing with [Rubular](#).

tually hitting that URL will return an error.

²⁹Note that [Table 6.2](#) includes the URL `/users/new`, which actually worked fine in [Section 5.3.1](#). This is because `/users/new` matches both the `/controller/action` pattern and the RESTful routes pattern.

Chapter 7

Modeling and viewing users, part II

In [Chapter 6](#), we created the first iteration of a User model to represent users of our application, but the job is only half-done. Virtually any website with users, including ours, needs *authentication* as well, but currently any user signing up for the site would only have a name and email address, with no way to verify their identity. In this chapter, we'll add the `password` attribute needed for an initial user signup ([Chapter 8](#)) and for signing in with an email/password combination ([Chapter 9](#)). In the process, we'll re-use several of the ideas from [Chapter 6](#), including migrations and validations, and also introduce some new ideas such as virtual attributes, private methods, and Active Record callbacks.

Once we have a working `password` attribute, we'll make a working action and view for showing user profiles ([Section 7.3](#)). By the end of the chapter, our user profiles will display names and profile photos (as indicated by the mockup in [Figure 7.1](#)), and they will be nicely tested with user factories and RSpec stubs.

7.1 Insecure passwords

Making industrial-strength passwords requires a lot of machinery, so we'll break the process into two main steps. In this section, we'll make a `password` attribute and add validations. The resulting User model will be functionally complete but badly insecure, with the passwords stored as plain text in the database. In [Section 7.2](#), we'll fix this problem by encrypting the passwords before saving them, thereby protecting our site against potential attackers.

7.1.1 Password validations

Even though we have yet even to add a column for passwords to our database, we're already going to start writing tests for them. Our initial plan is to have tests to validate the presence, length, and confirmation of passwords. This is our biggest single block of tests so far, so see if you can read it all in one go. If you get stuck, it might help to review the analogous validations from [Section 6.2](#) or skip ahead to the application code in [Listing 7.2](#).

In order to minimize typos in passwords, when making a user signup page in [Chapter 8](#) we'll adopt the common convention of requiring that users *confirm* their passwords. To get started, let's review the user attributes hash last seen in [Listing 6.19](#):

```
describe User do
```

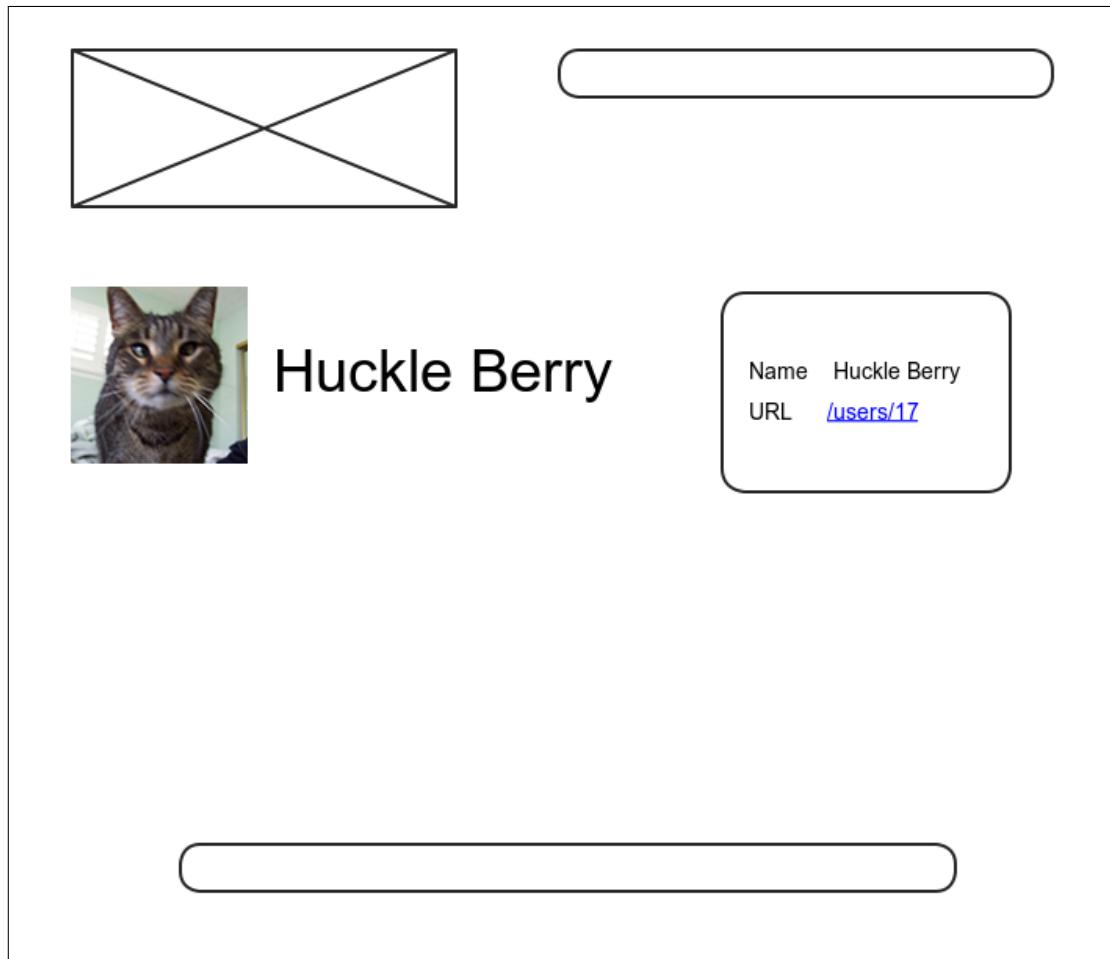


Figure 7.1: A mockup of the user profile made in Section 7.3. ([full size](#))

```
before(:each) do
  @attr = { :name => "Example User", :email => "user@example.com" }
end

.
.
.

end
```

To write tests for passwords, we'll need to add *two* new attributes to the `@attr` hash, `password` and `password_confirmation`. As you can probably guess, the `password_confirmation` attribute will be used for the password confirmation step.

Let's write tests for the presence of the password and its confirmation, together with tests confirming that the password is a valid length (restricted somewhat arbitrarily to be between 6 and 40 characters long). The results appear in Listing 7.1.

Listing 7.1. Tests for password validations.

`spec/models/user_spec.rb`

```
describe User do

  before(:each) do
    @attr = {
      :name => "Example User",
      :email => "user@example.com",
      :password => "foobar",
      :password_confirmation => "foobar"
    }
  end

  it "should create a new instance given valid attributes" do
    User.create!(@attr)
  end

  .
  .
  .

  describe "password validations" do

    it "should require a password" do
      User.new(@attr.merge(:password => "", :password_confirmation => "")).
        should_not be_valid
    end

    it "should require a matching password confirmation" do
      User.new(@attr.merge(:password_confirmation => "invalid")).
        should_not be_valid
    end

    it "should reject short passwords" do
      short = "a" * 5
      hash = @attr.merge(:password => short, :password_confirmation => short)
    end
  end
end
```

```
User.new(hash).should_not be_valid
end

it "should reject long passwords" do
  long = "a" * 41
  hash = @attr.merge(:password => long, :password_confirmation => long)
  User.new(hash).should_not be_valid
end
end
end
```

Note in Listing 7.1 how we first collect a set of valid user attributes in `@attr`. If for some reason those attributes aren't valid—as would be the case, for example, if we didn't implement password confirmations properly—then the first test

```
it "should create a new instance given valid attributes" do
  User.create!(@attr)
end
```

would catch the error. The subsequent tests then check each validation in turn, using the same `@attr.merge` technique first introduced in Listing 6.10.

Now for the application code, which contains a trick. Actually, it contains *two* tricks. First, you might expect at this point that we would run a migration to add a `password` attribute to the User model, as we did with the `name` and `email` attributes in Listing 6.1. But this is not the case: we will store only an *encrypted* password in the database; for the password, we will introduce a *virtual attribute* (that is, an attribute not corresponding to a column in the database) using the `attr_accessor` keyword, much as we did with the original `name` and `email` attributes for the example user in Section 4.4.5. The `password` attribute will not ever be written to the database, but will exist only in memory for use in performing the password confirmation step (implemented below) and the encryption step (implemented in Section 7.1.2 and Section 7.2).

The second trick is that we will *not* introduce a `password_confirmation` attribute, not even a virtual one. Instead, we will use the special validation

```
validates_confirmation_of :password
```

which will *automatically* create a virtual attribute called `password_confirmation`, while confirming that it matches the `password` attribute at the same time.

Thus prepared to understand the implementation, let's take a look at the code itself (Listing 7.2).

Listing 7.2. Validations for the `password` attribute.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessor :password
  attr_accessible :name, :email, :password, :password_confirmation
  .
```

```
.
.
# Automatically create the virtual attribute 'password_confirmation'.
validates_confirmation_of :password

# Password validations.
validates_presence_of :password
validates_length_of    :password, :within => 6..40
end
```

As promised, we use `attr_accessor :password` to create a virtual `password` attribute (as in Section 4.4.5), and include `validates_confirmation_of` to reject users whose password and password confirmations don't match. We also see `validates_presence_of`, which we first saw in Listing 6.7. Finally, we see a second application of `validates_length_of`; previously, we constrained the `name` attribute to be 50 characters or less, using the `:maximum` option:

```
validates_length_of :name, :maximum => 50
```

For the password length validation, instead we've used the `:within` option, passing it the range¹ `6..40` to enforce the desired length constraints.

Finally, since we'll be accepting passwords and password confirmations as part of the signup process in Chapter 8, we need to add the password and its confirmation to the list of accessible attributes (first mentioned in Section 6.1.2), which we've done in the line

```
attr_accessible :name, :email, :password, :password_confirmation
```

7.1.2 A password migration

At this point, you may be concerned that we're not storing user passwords anywhere; since we've elected to use a virtual password, rather than storing it in the database, it exists only in memory. How can we use this password for authentication? The solution is to create a separate attribute dedicated to password storage, and our strategy will be to use the virtual password as raw material for an *encrypted password*, which we *will* store in the database upon user signup (Chapter 8) and retrieve later for use in user authentication (Chapter 9).

Let's plan to store the encrypted password using an `encrypted_password` attribute in our User model. We'll discuss the implementation details in Section 7.2, but we can get started with our encrypted password tests by noting that the encrypted password should at the least *exist*. We can test this using the Ruby method `respond_to?`, which accepts a symbol and returns `true` if the object responds to the given method or attribute and `false` otherwise:

```
$ script/console --sandbox
>> user = User.new
>> user.respond_to?(:password)
```

¹We saw ranges before in Section 4.3.1.

```
=> true
>> user.respond_to?(:encrypted_password)
=> false
```

We can test the existence of an `encrypted_password` attribute with the code in [Listing 7.3](#), which uses RSpec's `respond_to` helper method.

Listing 7.3. Testing for the existence of an `encrypted_password` attribute.

`spec/models/user_spec.rb`

```
describe User
  .
  .
  .
  describe "password encryption" do
    before(:each) do
      @user = User.create!(@attr)
    end

    it "should have an encrypted password attribute" do
      @user.should respond_to?(:encrypted_password)
    end
  end
end
```

Note that in the `before(:each)` block we *create* a user, rather than just calling `User.new`. We could actually get this test to pass using `User.new`, but (as we'll see momentarily) *setting* the encrypted password will require that the user be saved to the database. Using `create!` in this first case does no harm, and putting it in `before(:each)` will allow us to keep all the encrypted password tests in one `describe` block.

To get this test to pass, we'll need a migration to add the `encrypted_password` attribute to the `users` table:

```
$ script/generate migration add_password_to_users encrypted_password:string
```

Here the first argument is the migration name, and we've also supplied a second argument with the name and type of attribute we want to create. (Compare this to the original generation of the `users` table in [Listing 6.1](#).) We can choose any migration name we want, but it's convenient to end the name with `_to_` `users`, since in this case Rails can automatically construct a migration to add columns to the `users` table. Moreover, by including the second argument, we've given Rails enough information to construct the entire migration for us, as seen in [Listing 7.4](#).

Listing 7.4. The migration to add an `encrypted_password` column to the `users` table.

`db/migrate/<timestampl>.add_password_to_users.rb`

users	
id	integer
name	string
email	string
encrypted_password	string
created_at	datetime
updated_at	datetime

Figure 7.2: The User model with an added (encrypted) password attribute.

```
class AddPasswordToUsers < ActiveRecord::Migration
  def self.up
    add_column :users, :encrypted_password, :string
  end

  def self.down
    remove_column :users, :encrypted_password
  end
end
```

This code uses the `add_column` method to add an `encrypted_password` column to the `users` table (and the complementary `remove_column` method to remove it when migrating down). The result is the data model shown in Figure 7.2.

Now if we run the migration and prepare the test database, the test should pass, since the User model will respond to the `encrypted_password` attribute:

```
$ rake db:migrate
$ rake db:test:prepare
```

We can of course run the full test suite with `spec spec/`, but sometimes it's convenient to run just *one* RSpec example, which we can do with the `-e` ("example") flag:

```
$ spec spec/models/user_spec.rb \
> -e "should have an encrypted password attribute"
.

1 example, 0 failures
```

7.1.3 An Active Record callback

Now that our User model has an attribute for storing the password, we need to arrange to generate and save the encrypted password when Active Record saves the user to the database. We'll do this with a technique called a *callback*, which is a method that gets invoked at a particular point in the lifetime of an Active Record object. In the present case, we'll use a `before_save` callback to create `encrypted_password` just before the user is saved.²

We start with a test for the encrypted password attribute. Since we're deferring the implementation details—and, in particular, the method of encryption—to Section 7.2, in this section we'll just make sure that a saved user's `encrypted_password` attribute is not the empty string `" "`. We do this by combining the `empty?` method on strings (Section 4.2.3) with the RSpec convention for boolean methods (first seen in the context of `valid?/be_valid` in Listing 6.10), yielding the test in Listing 7.5.

Listing 7.5. Testing that the `encrypted_password` attribute is nonempty.

`spec/models/user_spec.rb`

```
describe User
  .
  .
  .
  describe "password encryption" do
    before(:each) do
      @user = User.create!(@attr)
    end
    .
    .
    .
    it "should set the encrypted password" do
      @user.encrypted_password.should_not be_blank
    end
  end
end
```

This uses the RSpec boolean convention first mentioned just after Listing 6.10 in Section 6.2.1, in this case verifying that `encrypted_password.blank?` is not true using the construction `should_not be_blank`. (We first saw the `blank?` method at the end of Section 4.4.2.)

To get this test to pass, we *register* a callback called `encrypt_password` by passing a symbol of that name to the `before_save` method, and then define an `encrypt_password` method to perform the encryption. With the `before_save` in place, Active Record will automatically call the corresponding method before saving the record. The result appears in Listing 7.6.

Listing 7.6. A `before_save` callback to create the `encrypted_password` attribute.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
```

²For more details on the kind of callbacks supported by Active Record, see the [discussion of callbacks at the Rails Guides](#).

```
.  
. .  
before_save :encrypt_password  
  
private  
  
def encrypt_password  
  self.encrypted_password = encrypt(password)  
end  
  
def encrypt(string)  
  string # Only a temporary implementation!  
end  
end
```

Here the `encrypt_password` callback delegates the actual encryption to an `encrypt` method; as noted in the comment, this is only a temporary implementation—as currently constructed, Listing 7.6 simply sets the encrypted to the *unencrypted* password, which kind of defeats the purpose. But it’s enough to get our test to pass, and we’ll make the `encrypt` method do some actual encryption in Section 7.2.

Before trying to understand the implementation, first note that the encryption methods appear after the `private` keyword; `private` introduces a section of code that is used internally by the object and is not intended for public use.³ For an example, we can look at a User object in the console:

```
>> user = User.new  
>> user.encrypt_password  
NoMethodError: Attempt to call private method
```

Here Ruby raises a `NoMethodError` exception and issues a warning that the method is private.

In the present context, making the `encrypt_password` and `encrypt` methods private isn’t strictly necessary, but it’s a good practice to make them private unless they are needed for the public interface.⁴

Now that we understand the `private` keyword, let’s take another look at the `encrypt_password` method:

```
def encrypt_password  
  self.encrypted_password = encrypt(password)  
end
```

This is a one-line method (the best kind!), but it contains not one but *two* subtleties. First, the left-hand side of the statement explicitly assigns the `encrypted_password` attribute using the `self` keyword. (Recall

³The extra level of indentation is a typographical reminder that we’re in a private section; otherwise, it’s easy to miss the `private` keyword and be confused when trying to access a private method that you think is public. I thought the extra indentation was a stupid convention until I burned an hour on just this problem a couple years back. Now I add the extra indentation.

⁴Ruby has a closely related keyword called `protected` that differs subtly from `private`. As far as I can tell, the only reason to learn the difference is so that you can ace a job interview that asks you “In Ruby, what is the difference between `private` and `protected`?”. But do you really want to work at a company that asks you such a lame interview question? At his keynote at RubyConf in 2008, Dave Thomas (author of *Programming Ruby*) suggested eliminating `protected` from future versions of Ruby, and I agree with the sentiment. Just use `private` and you’ll be fine.

from [Section 4.4.2](#) that inside the class `self` refers to the object itself, which for the User model is just the user.) The use of `self` is *required* in this context; if we omitted `self` and wrote

```
def encrypt_password
  encrypted_password = encrypt(password)
end
```

Ruby would create a *local variable* called `encrypted_password`, which isn't what we want at all.

Second, the right hand side of the assignment calls `encrypt` on `password`, but there is no `password` in sight. In the console, we would access the password attribute through a user object

```
>> user = User.new(:password => "foobar")
>> user.password
=> "foobar"
```

Inside the User class, the user object is just `self`, and we could write

```
def encrypt_password
  self.encrypted_password = encrypt(self.password)
end
```

in analogy with the console example, just replacing `user` with `self`. But the `self` is optional, so for brevity we can write simply

```
def encrypt_password
  self.encrypted_password = encrypt(password)
end
```

as in [Listing 7.6](#) above. (Of course, as we've noted, the `self` is *not* optional when assigning to an attribute, so we have to write `self.encrypted_password` in this case.)

7.2 Secure passwords

With the code from [Section 7.1](#), in principle we are done: although the “encrypted” password is the same as the unencrypted password, as long as we are willing to store unencrypted passwords in the database we have the necessary foundation for user login and authentication.⁵ Our standards in *Rails Tutorial* are much loftier, though: any web developer worth his salt should know how to implement a password system with *secure one-way hashing*. In this section we will build on the material from [Section 7.1](#) to implement just such an industrial-strength password system.

⁵I am ashamed to admit that this is how we implemented passwords in *RailsSpace*. Consider this section my penance.

7.2.1 A secure password test

As hinted at in [Section 7.1.3](#), all of the machinery for password encryption will be tucked away in the `private` regions of the User model, which presents a challenge for testing it. What we need is some sort of *public interface* that we can expose to the rest of the application. One useful aspect of test-driven development is that, by acting as a client for our application code, the tests motivate us to design a useful interface right from the start.

Authenticating users involves comparing a submitted password to the (encrypted) password of a given user. This means we need to define some method to perform the comparison, which we'll call `has_password?`; this will be our public interface to the encryption machinery.⁶ The `has_password?` method will test whether a user has the same password as one submitted on a sign-in form (to be written in [Chapter 9](#)); a skeleton method for `has_password?` appears in [Listing 7.7](#).

Listing 7.7. A `has_password?` method for users.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
  .

  # Return true if the user's password matches the submitted password.
  def has_password?(submitted_password)
    # Compare encrypted_password with the encrypted version of
    # submitted_password.
  end
  .
  .
  .
end
```

With this method, we can write tests as in [Listing 7.8](#), which uses the RSpec methods `be_true` and `be_false` to test that `has_password?` returns `true` or `false` in the proper cases.

Listing 7.8. Tests for the `has_password?` method.

`spec/models/user_spec.rb`

```
describe User
  .
  .
  .

  describe "password encryption" do
    before(:each) do
      @user = User.create!(@attr)
```

⁶The alert reader may notice that none of what we do in this section *requires* encryption, but, once we develop some of the theory of secure passwords and write a basic implementation ([Section 7.2.2](#)), the only way for the `has_password?` method to work properly is for all the encryption machinery to work properly as well.

```

end
.
.
.
describe "has_password? method" do
  it "should be true if the passwords match" do
    @user.has_password?(@attr[:password]).should be_true
  end

  it "should be false if the passwords don't match" do
    @user.has_password?("invalid").should be_false
  end
end
end

```

In Section 7.2.3, we'll complete the implementation of `has_password?` (and get the test to pass in the process). But first we need to learn a little more about secure passwords.

7.2.2 Some secure password theory

The basic idea of encrypted passwords is simple: rather than storing a raw password in the database (known as “cleartext”), we store a string generated using a [cryptographic hash function](#), which is essentially irreversible, so that even an attacker in possession of the hashed password will be unable to infer the original. To verify that a submitted password matches the user’s password, we first encrypt the submitted string and then compare the hashes. Let’s drop into a console session to see how this works:

```

$ script/console
>> require 'digest'
>> def secure_hash(string)
>>   Digest::SHA2.hexdigest(string)
>> end
=> nil
>> password = "secret"
=> "secret"
>> encrypted_password = secure_hash(password)
=> "2bb80d537b1da3e38bd30361aa855686bde0eacd7162fef6a25fe97bf527a25b"
>> submitted_password = "secret"
=> "secret"
>> encrypted_password == secure_hash(submitted_password)
=> true

```

Here we’ve defined a function called `secure_hash` that uses a cryptographic hash function called SHA2, part of the [SHA family of hash functions](#), which we include into Ruby through the `digest` library.⁷ It’s not

⁷In my setup, the `require 'digest'` line is unnecessary, but several readers have reported getting a `NameError` exception if they don’t include it explicitly. It does no harm in any case, so I’ve included the explicit `require` just to be safe.

important to know exactly how these hash functions work; for our purposes what's important is that they are one-way: there is no computationally tractable way to discover that

```
2bb80d537b1da3e38bd30361aa855686bde0eacd7162fef6a25fe97bf527a25b
```

is the SHA2 hash of the string "**secret**".

If you think about it, though, we still have a problem: if an attacker ever got hold of the hashed passwords, he would still have a chance at discovering the originals. For example, he could guess that we used SHA2, and so write a program to compare a given hash to the hashed values of potential passwords:

```
>> hash = "2bb80d537b1da3e38bd30361aa855686bde0eacd7162fef6a25fe97bf527a25b"
>> secure_hash("secede") == hash
=> false
>> secure_hash("second") == hash
=> false
>> secure_hash("secret") == hash
=> true
```

So our attacker has a match—bad news for any users with password "**secret**". This technique is known, for obvious reasons, as a *dictionary attack*.

To foil a potential dictionary attack, we can use a *salt*, which is a different unique string for each user. One common way to (nearly) ensure uniqueness is to hash the current time (in UTC to be time-zone-independent) along with the password, so that two users will have the same salt only if they are created at exactly the same time *and* have the same password. Let's see how this works using the **secure_hash** function defined in the console above:

```
>> Time.now.utc
=> Fri Jan 29 18:11:27 UTC 2010
>> password = "secret"
=> "secret"
>> salt = secure_hash("#{Time.now.utc}#{password}")
=> "d1c2f8e1a4f5caed8818fcfa72117e65631e3168af4adcfda5bcfa2c5a36d311"
>> encrypted_password = secure_hash("#{salt}#{password}")
=> "918574661e80aacd3d81a380f5d22489eeb1e2f0bd33bb6575e8a64ac5a5aecc"
```

In the the last line, we've hashed the salt with the password, yielding an encrypted password that is virtually impossible to crack.

7.2.3 Implementing `has_password?`

Having finished with the theory, we're now ready for the implementation. Let's look ahead a little to see where we're going. Each user object knows its own encrypted password, so to check for a match with a submitted password we can define `has_password?` as follows:

```
def has_password?(submitted_password)
  encrypted_password == encrypt(submitted_password)
end
```

As long as we encrypt the submitted password using the same salt used to encrypt the original password, this function will be true if and only if the submitted password matches.

Since comparing a user password with a submitted password will involve encrypting the submitted password with the salt, we need to store the salt somewhere, so the first step is to add a `salt` column to the `users` table:

```
$ script/generate migration add_salt_to_users salt:string
```

As with the `encrypted_password` migration (Section 7.1.2), this migration has a name that ends in `_to_users` and passes a second argument containing the attribute name and type, so Rails automatically constructs the right migration (Listing 7.9).

Listing 7.9. The migration to add a `salt` column to the `users` table.

```
db/migrate/<timestamp>.add_salt_to_users.rb
```

```
class AddSaltToUsers < ActiveRecord::Migration
  def self.up
    add_column :users, :salt, :string
  end

  def self.down
    remove_column :users, :salt
  end
end
```

Then we migrate the database and prepare the test database as usual:

```
$ rake db:migrate
$ rake db:test:prepare
```

The result is a database with the data model shown in Figure 7.3.

Finally, we're ready for the full implementation. When last we saw the `encrypt` function (Listing 7.6), it did nothing, simply returning the string in its argument. With the ideas from Section 7.2.2, we're now in a position to use a secure hash instead (Listing 7.10).⁸

Listing 7.10. The `has_password?` method with secure encryption.

```
app/models/user.rb
```

```
require 'digest'
class User < ActiveRecord::Base
  .
  .
```

⁸As noted in Section 7.2.2, the explicit `require 'digest'` line is unnecessary on some systems, but it does no harm to include it.

users	
id	integer
name	string
email	string
encrypted_password	string
salt	string
created_at	datetime
updated_at	datetime

Figure 7.3: The User model with an added salt.

```

  ..
before_save :encrypt_password

def has_password?(submitted_password)
  encrypted_password == encrypt(submitted_password)
end

private

def encrypt_password
  self.salt = make_salt
  self.encrypted_password = encrypt(password)
end

def encrypt(string)
  secure_hash("#{salt}#{string}")
end

def make_salt
  secure_hash("#{Time.now.utc}#{password}")
end

def secure_hash(string)
  Digest::SHA2.hexdigest(string)
end
end

```

This code contains the same two subtleties mentioned in Section 7.1.3, namely, the assignment to an Active Record attribute in the line

```
self.salt = make_salt
```

and the omission of the `self` keyword in the `encrypt` method:

```
def encrypt(string)
  secure_hash("#{salt}#{string}")
end
```

Since we're inside the User class, Ruby knows that `salt` refers to the user's `salt` attribute.

At this point, the tests from Listing 7.8 should pass:⁹

```
$ spec spec/models/user_spec.rb -e "should be true if the passwords match"
.
1 example, 0 failures

$ spec spec/models/user_spec.rb -e "should be false if the passwords don't match"
.
1 example, 0 failures
```

7.2.4 An authenticate method

Having a `has_password?` method for each user is nice, but by itself it isn't very useful. We'll end our discussion of passwords by using `has_password?` to write a method to authenticate a user based on an email/password combination. In Chapter 9, we'll use this `authenticate` method when signing users in to our site.

We can get a hint of how this will work by using the console. First, we'll create a user, and then retrieve that user by email address to verify that it has a given password:¹⁰

```
$ script/console --sandbox
>> User.create(:name => "Michael Hartl", :email => "mhartl@example.com",
?>                 :password => "foobar", :password_confirmation => "foobar")
>> user = User.find_by_email("mhartl@example.com")
>> user.has_password?("foobar")
=> true
```

Using these ideas, let's write a method that will return an authenticated user on password match, and `nil` otherwise. We should be able to use the resulting `authenticate` class method as follows:

⁹It would be nice to be able to run all the examples in a particular `describe` block (in this case, "`has_password? method`"), but the `spec` script doesn't have this option. You can, however, pass `spec` the *line number* of the `describe` block using the `-l` option. Though hard-coding the line number into the tutorial would be brittle and error-prone, feel free to try using `-l` with the appropriate line number from your current version of `user_spec.rb` to see how it works.

¹⁰Recall from Box 6.1 that the `index` on the `email` column ensures that this retrieval is efficient.

```
User.authenticate(email, submitted_password)
```

We start with the tests, which we'll use to specify the behavior we expect from `User.authenticate`. There are three cases to check: `authenticate` (1) should return `nil` when the email/password combination is invalid or (2) when no user exists with the given email address, and (3) should return the user object itself on success. With this information, we can write the tests for `authenticate` as in Listing 7.11.

Listing 7.11. Tests for the `User.authenticate` method.
`spec/models/user_spec.rb`

```
describe User
  .
  .
  .
  describe "password encryption" do
    .
    .
    .
  describe "authenticate method" do
    it "should return nil on email/password mismatch" do
      wrong_password_user = User.authenticate(@attr[:email], "wrongpass")
      wrong_password_user.should be_nil
    end

    it "should return nil for an email address with no user" do
      nonexistent_user = User.authenticate("bar@foo.com", @attr[:password])
      nonexistent_user.should be_nil
    end

    it "should return the user on email/password match" do
      matching_user = User.authenticate(@attr[:email], @attr[:password])
      matching_user.should == @user
    end
  end
end
end
```

Now we're ready for the implementation, which will get our tests to pass and show how to define a *class method* as a bonus. We've mentioned class methods several times before, most recently in Section 6.1.1; a class method is simply a method attached to a class, rather than an instance of that class. For example, `new`, `find`, and `find_by_email` are all class methods on the `User` class. Outside of the class, they are invoked using the class name, as in `User.find`, but inside the class we can omit the class name.

The way to define a class method is to use the `self` keyword in the method definition. Listing 7.12 shows this construction in the context of the `authenticate` method. Note the call to `find_by_email`, in which we omit the explicit `User` class name since this method is already inside the `User` class.

Listing 7.12. The `User.authenticate` method.`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
  .
  def has_password?(submitted_password)
    encrypted_password == encrypt(submitted_password)
  end

  def self.authenticate(email, submitted_password)
    user = find_by_email(email)
    return nil if user.nil?
    return user if user.has_password?(submitted_password)
  end

  private
  .
  .
  .
end
```

There are several equivalent ways to write the `authenticate` method, but I find the implementation above the clearest. It handles two cases (invalid email and a successful match) with explicit `return` keywords, and handles the third case (password mismatch) implicitly, since in that case we reach the end of the method, which automatically returns `nil`. See [Section 7.5](#) for some of the other possible ways to implement this method.

7.3 Better user views

Now that User model is now effectively complete,¹¹ we are in a position to add a sample user to the development database and make a `show` page to show some of that user's information. Along the way, we'll add some tests to the Users controller spec started in [Section 5.3.1](#).

Before continuing, it's helpful to see where we left off by recalling what the Users controller spec looks like right now ([Listing 7.13](#)). Our tests for the user show page will follow this example, but we'll find that, unlike the tests for the `new` action, the tests for the `show` action will require the use of an instance of the User model. We'll meet this challenge using *factories* and *stubs*.

Listing 7.13. The current Users controller spec.`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
```

¹¹We'll plan to add a couple more attributes (one to identify administrative users and one to allow a "remember me" feature), but they are not strictly necessary. All the *essential* user attributes have now been defined.

```

integrate_views

describe "GET 'new'" do

  it "should be successful" do
    get 'new'
    response.should be_success
  end

  it "should have the right title" do
    get 'new'
    response.should have_tag("title", /Sign up/)
  end
end
end

```

7.3.1 Testing the user show page (with factories)

One goal of automated testing is to test each component separately if possible, so that, for example, model tests query the database (since that's one of the things models do) but controller/view tests do not.¹² As seen in Listing 7.14, the Users controller `show` action needs an instance of the User model, so the tests for this action will require that we create an `@user` variable somehow. We'll accomplish this with a user *factory*, which is an object that simulates a user object without ever touching the database.¹³

Listing 7.14. The user `show` action from Listing 6.24.

`app/controllers/users_controller.rb`

```

class UsersController < ApplicationController

  def show
    @user = User.find(params[:id])
  end
  .
  .
  .
end

```

We'll be using the factories generated by [Factory Girl](#),¹⁴ a Ruby gem produced by the good people at [thoughtbot](#). As with other Ruby gems, we can install it using the `gem` command:¹⁵

¹² Some programmers take this to the next level, and attempt to test views separately from controllers. I think this goes too far, since views and controllers are necessarily coupled. In fact, if you take the default tests generated by RSpec and change the name of an instance variable in a controller but not in the view, all the tests still pass, even though the application is completely broken. In *Rails Tutorial*, we combine controller and views specs using `integrate_views` to avoid this problem.

¹³ Many experienced Rails programmers find that this factory approach is much more flexible than *fixtures*, which Rails uses by default but can be quite brittle and difficult to maintain.

¹⁴ Presumably “Factory Girl” is a reference to the [movie of the same name](#).

¹⁵ As noted briefly in Section 1.2.2, when installing gems I prefer to use the version flag `-v` to install a gem version known to work with the tutorial—in this case, version 1.2.3 of Factory Girl. Feel free omit this flag to use the latest version, but you’re on your own if it breaks. :-)

```
$ [sudo] gem install factory_girl -v 1.2.3
```

Next, we need to tell Rails about Factory Girl, which we can do by using a special gem configuration command in the environment file for our tests, as seen in [Listing 7.15](#).¹⁶

Listing 7.15. Including the Factory Girl gem into the test environment file.

`config/environments/test.rb`

```
# Settings specified here will take precedence over those in
# config/environment.rb

config.gem 'factory_girl'
```

Now we're ready to create the file `spec/factories.rb` and define a User factory, as shown in [Listing 7.16](#). By putting the `factories.rb` file in the `spec/` directory, we arrange for RSpec to load our factories automatically whenever the tests run.

Listing 7.16. A factory to simulate User model objects.

`spec/factories.rb`

```
# By using the symbol ':user', we get Factory Girl to simulate the User model.
Factory.define :user do |user|
  user.name          "Michael Hartl"
  user.email         "mhartl@example.com"
  user.password      "foobar"
  user.password_confirmation "foobar"
end
```

With the definition in [Listing 7.16](#), we can create a User factory in the tests like this:

```
@user = Factory(:user)
```

As noted in the comment in the first line of [Listing 7.16](#), by using the symbol `:user` we ensure that Factory Girl will guess that we want to use the User model, so in this case `@user` will simulate an instance of `User`.

To use our new User factory in the Users controller spec, we'll create an `@user` variable in a `before (:each)` block, stub it out (details below), and then `get` the show page and verify success (just as we did with the `new` page in [Listing 7.13](#)). The result appears in [Listing 7.17](#).

Listing 7.17. A test for getting the user `show` page, with a factory and a stub.

`spec/controllers/users_controller_spec.rb`

¹⁶You can also put the `config.gem` line in `config/environment.rb`, but we only need factories in the test environment, so there's no need to include it in development or production. Also, when a gem is loaded with `config.gem` in the `test.rb` file, you can have Rails *automatically* install it using `rake gems:install RAILS_ENV=test`. That's a bit complicated at this point, though, and gem dependencies are completely different in Rails 3, so I'm not bothering with this detail now.

```
require 'spec_helper'

describe UsersController do
  integrate_views

  describe "GET 'show'" do
    before(:each) do
      @user = Factory(:user)
      # Arrange for User.find(params[:id]) to find the right user.
      User.stub!(:find, @user.id).and_return(@user)
    end

    it "should be successful" do
      get :show, :id => @user
      response.should be_success
    end
  end
  .
  .
  .
end
```

Apart from the first use of a factory, the real novelty here is the use of a *stub*, which is a facility provided by RSpec to avoid hits to the database. The code

```
User.stub!(:find, @user.id).and_return(@user)
```

ensures that any call to `User.find` with the given `id` will return `@user`. Since this is just what we have in the application code ([Listing 7.14](#)), the stub will cause RSpec to intercept the call to `User.find` and, instead of hitting the database, returning `@user` instead. Since `@user` is a factory, the database never gets touched.

There are two other details worth noting. First, in the call to `get`, the test uses the *symbol* `:show` instead of the string `'show'`, which is different from the convention in the other tests (for example, in [Listing 3.7](#) we wrote `get 'home'`). Both

```
get :show
```

and

```
get 'show'
```

do the same thing, but when testing the canonical REST actions ([Table 6.2](#)) I prefer to use symbols, which

for some reason feel more natural in this context.¹⁷ Second, note that the value of the hash key `:id`, instead of being the user's `id` attribute `@user.id`, is the user object itself:

```
get :show, :id => @user
```

We could use the code

```
get :show, :id => @user.id
```

to accomplish the same thing, but in this context Rails automatically converts the user object to the corresponding id.¹⁸ Using the more succinct construction

```
get :show, :id => @user
```

is a very common Rails idiom.

Because of the code we added in Listing 6.24, the test in this section already passes. If you're feeling paranoid, you can comment out the line

```
@user = User.find(params[:id])
```

and verify that the test fails, then uncomment it to get it to pass. (We went through this same process once before, in Section 6.2.1.)

7.3.2 A name and a Gravatar

In this section, we'll improve the look of the user show page by adding a heading with the user's name and profile image. This is one of those situations where I can go either way on test-driven development, and often when making views I'll experiment with the HTML before bothering with tests. Let's stick with the TDD theme for now, and test for a second-level heading (`h2` tag) containing the user's name and an `img` tag of class `gravatar`. (We'll talk momentarily about what this second part means.)

To view a working user show page in a browser, we'll need to create a sample user in the development database. To do this, first reset the database with `rake db:reset`, which will clear out any old sample users from previous sessions, and then start the console (*not* in a sandbox this time) and create the user:

```
$ rake db:reset
$ script/console
>> User.create(:name => "Example User", :email => "user@example.com",
?>                 :password => "foobar", :password_confirmation => "foobar")
```

¹⁷I used `get 'new'` in Listing 5.21 and subsequent tests for the `new` action because at that point we had yet to encounter the idea of standard REST actions. I'll switch to `get :new` in future tests.

¹⁸It does this by calling the `to_param` method on the `@user` variable.

The tests in this section are similar to the tests for the `new` page seen in Listing 5.23. In particular, we use the `have_tag` method to check the title and the content of the `h2` tag, as seen in Listing 7.18.

Listing 7.18. Tests for the user show page.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  describe "GET 'show'" do
    .
    .
    .
    it "should have the right title" do
      get :show, :id => @user
      response.should have_tag("title", /#{@user.name}/)
    end

    it "should include the user's name" do
      get :show, :id => @user
      response.should have_tag("h2", /#{@user.name}/)
    end

    it "should have a profile image" do
      get :show, :id => @user
      response.should have_tag("h2>img", :class => "gravatar")
    end
  end
  .
  .
  .
end
```

The first new construction is

```
/#{@user.name}/
```

to interpolate the user name into a regular expression. This works much like

```
"#{@name} <#{@email}>"
```

from Listing 4.6, which interpolates the name and email into a string. RSpec's `have_tag` method uses the resulting regular expression to test for a match, in this case verifying the presence of an `h2` tag containing the

user's name. The other new element is the code `h2>img`, which makes sure that the `img` tag is *inside* the `h2` tag.¹⁹

We can get the first test to pass by setting the `@title` variable for use in the `title` helper (Section 4.1.1), in this case setting it to the user's name (Listing 7.19).

Listing 7.19. A title for the user show page.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController

  def show
    @user = User.find(params[:id])
    @title = @user.name
  end
end
```

This code introduces a potential problem: a user could enter a name with malicious code—called a [cross-site scripting attack](#)—which would be injected into our application by the `title` helper defined in Listing 4.2. The solution is to *escape* potentially problematic code using the `h` method (short for `html_escape`), as seen in Listing 7.20.²⁰

Listing 7.20. The title helper using `h` to escape the HTML.

`app/helpers/application_helper.rb`

```
module ApplicationHelper

  # Return a title on a per-page basis.
  def title
    base_title = "Ruby on Rails Tutorial Sample App"
    if @title.nil?
      base_title
    else
      "#{base_title} | #{h(@title)}"
    end
  end
end
```

The `h` method converts, e.g., `<script>` to `<script>`, rendering any malicious scripts completely harmless.

Now for the other tests. Creating an `h2` with the (escaped) user name is easy (Listing 7.21).

¹⁹It's not necessarily always a good idea to make HTML tests this specific, since we don't always want to constrain the HTML layout this tightly. Feel free to experiment and find the right level of detail for your projects and tastes.

²⁰I'm not covering the details on HTML escaping because in Rails 3 all HTML inserted by ERb will be escaped by default, so application developers won't have to worry about it any more.

Listing 7.21. The user show view with the user's name.

`app/views/users/show.html.erb`

```
<h2>
  <%= h @user.name %>
</h2>
```

Getting the `img` test to pass is trickier. The first step is to install a plugin to handle each user's Gravatar,²¹ which is a "globally recognized avatar":²²

```
$ script/plugin install \
> svn://rubyforge.org/var/svn/gravatarplugin/plugins/gravatar
```

(If you get an error complaining that the `svn` command isn't found, you may need to [install Subversion](#) for your system.)

Gravatars are a convenient way to include user profile images without going through the trouble of managing image upload, cropping, and storage.²³ The Gravatar plugin comes with a helper method called `gravatar_for`, which we'll use in our user show view, as seen in [Listing 7.22](#). The result (after stopping and restarting the development server to load the plugin) appears in [Figure 7.4](#), which shows our example user with the default Gravatar image. And since the `gravatar_for` helper automatically assigns its `img` tag the class `gravatar`, the tests from [Listing 7.18](#) should now pass.

Listing 7.22. The user show view with name and Gravatar.

`app/views/users/show.html.erb`

```
<h2>
  <%= gravatar_for @user %>
  <%= h @user.name %>
</h2>
```

If this Gravatar business seems like magic, worry not. Let's fire up the console to get a little more insight into what's going on:

```
>> user = User.first
>> user.update_attributes(:email => "example@railstutorial.org",
?>                         :password => "foobar",
?>                         :password_confirmation => "foobar")
=> true
```

²¹Gravatar was originally created by Tom Preston-Werner, cofounder of [GitHub](#), and was acquired and scaled by [Automattic](#) (best known as the makers of [WordPress](#)).

²²In Hinduism, an avatar is the manifestation of a deity in human or animal form. By extension, the term *avatar* is commonly used to mean some kind of personal representation, especially in a virtual environment. But you've seen the movie by now, so you already knew this.

²³If your application does need to handle images or other file uploads, [Paperclip](#) is the way to go. Like Factory Girl, Paperclip is brought to you by [thoughtbot](#). (Though I do know several people there, I have no vested interest in promoting thoughtbot; they just make good software.)

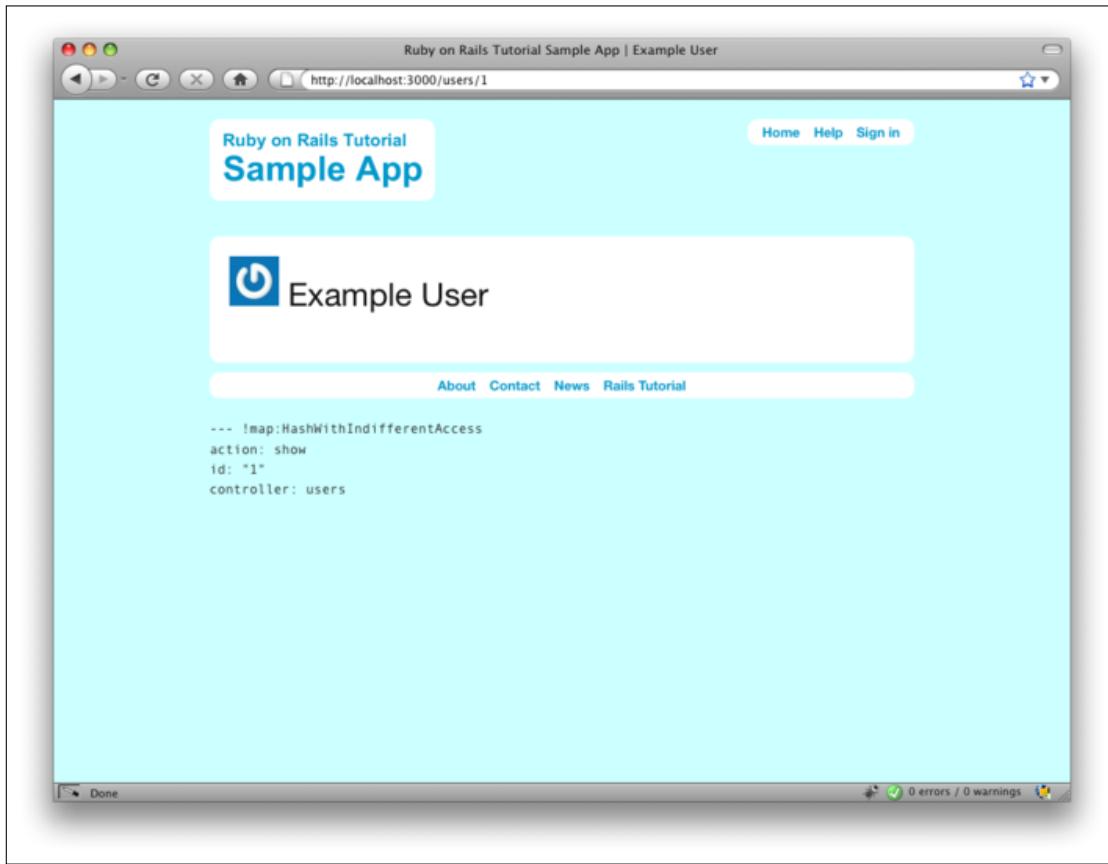


Figure 7.4: The initial user show page [/users/1](#) with the default Gravatar. ([full size](#))

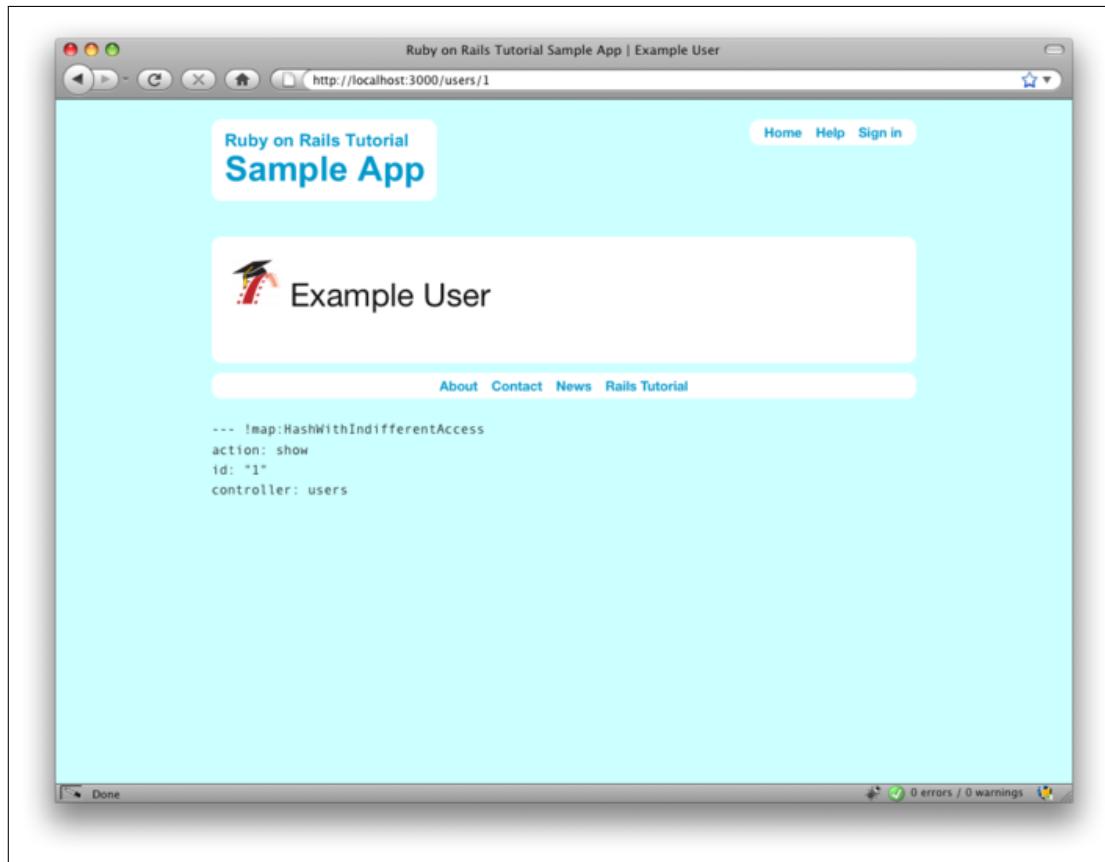


Figure 7.5: The user show page `/users/1` with the Rails Tutorial Gravatar. ([full size](#))

Note that we can pull out the first (and, at this point, only) user in the database with the handy `User.first` method. In the `update_attributes` step we've reassigned the user's email address, changing it to `example@railstutorial.org`. As you can see from Figure 7.5, this change results in a new Gravatar being displayed: the Rails Tutorial logo. What's going on is that Gravatar works by associating images with email addresses; since `user@example.com` is an invalid email address (the `example.com` domain is reserved for examples), there is no Gravatar for that email address. But at my Gravatar account I've associated the address `example@railstutorial.org` with the Rails Tutorial logo, so when updating the example user with that email address the Gravatar changes automatically.

If you're wondering how the `gravatar_for` helper works, you can get a hint from Listing 7.23, which is an excerpt from the Gravatar plugin. You can see that `gravatar_for` calls the `gravatar` method (not shown) with the argument `user.email`. In other words, since our User model already *has* an `email` attribute, the code `gravatar_for(@user)` Just Works.TM(What would you do if there were no `email` attribute on the User model? That's a good question, and it's left as an exercise (Section 7.5).)

Listing 7.23. The definition of `gravatar_for` from the Gravatar plugin source code.
`vendor/plugins/gravatar/lib/gravatar.rb`

```

.
.
.

# Return the HTML img tag for the given user's gravatar. Presumes that
# the given user object will respond_to "email", and return the user's
# email address.
def gravatar_for(user, options={})
  gravatar(user.email, options)
end
.
.
.

```

7.3.3 A user sidebar

Even though our tests are passing, and the user show page is much-improved, it's still nice to polish it up just a bit more. In Listing 7.24, we have a `table` tag with one table row (`tr`) and two table data cells (`td`).²⁴

Listing 7.24. The user show page `/users/1` with a sidebar.

`app/views/users/show.html.erb`

```


|                                                                      |                                                                                                             |
|----------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <h2> &lt;%= gravatar_for @user %&gt; &lt;%= h @user.name %&gt; </h2> | <strong>Name</strong> <%= h @user.name %><br /> <strong>URL</strong> <%= link_to user_path(@user), @user %> |
|----------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|


```

Here we've used an HTML break tag `
` to put a break between the user's name and URL. Also note the use of `user_path` to make a clickable link so that users can easily share their profile URLs. This is only the first of many named routes (Section 5.2.2) associated with the User resource (Listing 6.25); we'll see many more in the next few chapters. The code

```
user_path(@user)
```

²⁴If anyone gives you grief for using, horror of horrors, *tables for layout*, have them point their Firebug inspector at Twitter's profile sidebar and tell you what they see. In fact, you'll find that, while "semantic markup" using `divs` and `spans` is increasingly common, virtually all sites resort to `tables` for layout on occasion. In the present case, getting the vertical alignment just right is *much* easier with tables.

Named route	Path
<code>users_path</code>	/users
<code>user_path(@user)</code>	/users/1
<code>new_user_path</code>	/users/new
<code>edit_user_path(@user)</code>	/users/1/edit
<code>users_url</code>	http://localhost:3000/users
<code>user_url(@user)</code>	http://localhost:3000/users/1
<code>new_user_url</code>	http://localhost:3000/users/new
<code>edit_user_url(@user)</code>	http://localhost:3000/users/1/edit

Table 7.1: Named routes provided by the users resource in Listing 6.25.

returns the path to the user, in this case /users/1. The related code

```
user_url(@user)
```

just returns the entire URL, http://localhost:3000/users/1. (Compare to the routes created in Section 5.2.2.) Both are examples of the named routes created by the users resource in Listing 6.25; a list of all the named routes appears in Table 7.1.

Note that in

```
<%= link_to user_path(@user), @user %>
```

`user_path(@user)` is the link *text*, while the address is just `@user`. In the context of a `link_to`, Rails converts `@user` to the appropriate URL; in other words, the code above is equivalent to the code

```
<%= link_to user_path(@user), user_path(@user) %>
```

Either way works fine, but, as in the `:id => @user` idiom from Listing 7.17, using just `@user` is a common Rails convention. In both cases, the Embedded Ruby produces the HTML

```
<a href="/users/1">/users/1</a>
```

With the HTML elements and CSS classes in place, we can style the show page with the CSS shown in Listing 7.25. The resulting page is shown in Figure 7.6.

Listing 7.25. CSS for styling the user show page, including the sidebar.

```
public/stylesheets/custom.css
```

```
.  
./* User show page */  
  
table.profile {  
  width: 100%;  
  margin-bottom: 0;  
}  
  
td.main {  
  width: 70%;  
  padding: 1em;  
}  
  
td.sidebar {  
  width: 30%;  
  padding: 1em;  
  vertical-align: top;  
  background: #ffc;  
}  
  
.profile img.gravatar {  
  border: 1px solid #999;  
  margin-bottom: -15px;  
}
```

7.4 Conclusion

In this chapter, we've effectively finished the User model, so we're now fully prepared to sign up new users and to let them sign in securely with an email/password combination. Moreover, we have a nice first cut of the user profile page, so after signing in users will have a place to go.

7.4.1 Git commit

Before moving on, we should close the Git loop opened in the introduction to [Chapter 6](#) by making a final commit to the `modeling-users` branch and then merging into `master`.²⁵ First, verify that we are on the `modeling-users` branch:

```
$ git branch  
  master  
* modeling-users
```

²⁵Ordinarily, I recommend making more frequent, smaller commits, but frequent Git commits throughout the tutorial would be hard to maintain and would break up the flow of the discussion.

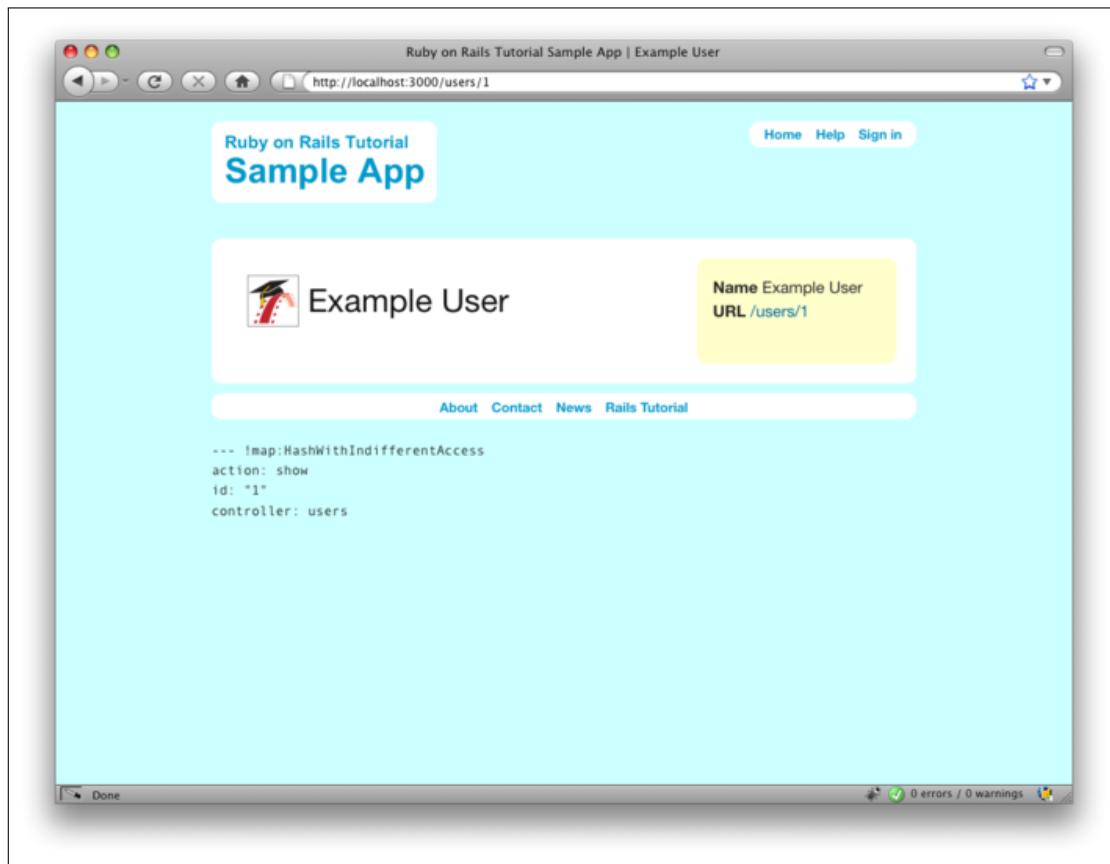


Figure 7.6: The user show page [/users/1](#) with a sidebar and CSS. ([full size](#))

As noted in [Section 1.3.5](#), the asterisk here identifies the present branch, so we are indeed ready to commit and merge.²⁶

```
$ git add .
$ git commit -am "User model with passwords, and user show page"
$ git checkout master
$ git merge modeling-users
```

7.4.2 Heroku deploy

If you've deployed your sample application to Heroku, you can push it up at this point:

```
$ git push heroku
```

Then migrate the database on the remote server using the `heroku` command:

```
$ heroku rake db:migrate
```

Now if you want to create a sample user on Heroku, you can use the Heroku console:

```
$ heroku console
>> User.create(:name => "Example User", :email => "user@example.com",
?>                 :password => "foobar", :password_confirmation => "foobar")
```

7.5 Exercises

1. Copy each of the variants of the `authenticate` method from [Listing 7.26](#) through [Listing 7.30](#) into your User model, and verify that they are correct by running your test suite.
2. The final `authenticate` example ([Listing 7.30](#)) is particularly challenging. Experiment with the console to see if you can understand how it works.
3. How could you get the Gravatar helper `gravatar_for` to work if our User model used `email_address` instead of `email` to represent email addresses? *Hint:* The best answer does *not* involve editing the Gravatar plugin.

Listing 7.26. The `authenticate` method with `User` in place of `self`.

²⁶If you're *not* on the right branch, run `git checkout modeling-users` before proceeding.

```
def User.authenticate(email, submitted_password)
  user = find_by_email(email)
  return nil if user.nil?
  return user if user.has_password?(submitted_password)
end
```

Listing 7.27. The `authenticate` method with an explicit third `return`.

```
def self.authenticate(email, submitted_password)
  user = find_by_email(email)
  return nil if user.nil?
  return user if user.has_password?(submitted_password)
  return nil
end
```

Listing 7.28. The `authenticate` method using an `if` statement.

```
def self.authenticate(email, submitted_password)
  user = find_by_email(email)
  if user.nil?
    nil
  elsif user.has_password?(submitted_password)
    user
  else
    nil
  end
end
```

Listing 7.29. The `authenticate` method using an `if` statement and an implicit return.

```
def self.authenticate(email, submitted_password)
  user = find_by_email(email)
  if user.nil?
    nil
  elsif user.has_password?(submitted_password)
    user
  end
end
```

Listing 7.30. The `authenticate` method using the ternary operator.

```
def self.authenticate(email, submitted_password)
  user = find_by_email(email)
  user && user.has_password?(submitted_password) ? user : nil
end
```


Chapter 8

Sign up

Now that we have a working User model, it's time to add an ability few websites can live with out: letting users sign up for the site—thus fulfilling the promise implicit in [Section 5.3](#), “User signup: A first step”. We'll use an HTML *form* to submit user signup information to our application in [Section 8.1](#), which will then be used to create a new user and save its attributes to the database in [Section 8.3](#). As usual, we'll write tests as we develop, and in [Section 8.4](#) we'll combine RSpec with the remarkable *Webrat* utility to write succinct and expressive integration tests.

Since we'll be creating a new user in this chapter, you might want to reset the database to clear out any users created at the console (e.g., in [Section 7.3.2](#)), so that your results will match those shown in the tutorial. You can do this as follows:

```
$ rake db:reset
```

If you're following along with version control, make a topic branch as usual:

```
$ git checkout master
$ git checkout -b signing-up
```

8.1 Signup form

Recall from [Section 5.3.1](#) that we already have tests for the new users (signup) page, originally seen in [Listing 5.23](#) and reproduced in [Listing 8.1](#). (As promised in [Section 7.3.1](#), we've switched from `get 'new'` to `get :new` because that's what my fingers want to type.) In addition, we saw in [Figure 5.11](#) (shown again in [Figure 8.1](#)) that this signup page is currently blank: useless for signing up new users. The goal of this section is to start changing this sad state of affairs by producing the signup form mocked up in [Figure 8.2](#).

Listing 8.1. The tests for the new users page (first seen in [Listing 5.23](#)).

```
spec/controllers/users_controller_spec.rb
```

```
require 'spec_helper'

describe UsersController do
  integrate_views

  describe "GET 'new'" do
    it "should be successful" do
      get :new
      response.should be_success
    end

    it "should have the right title" do
      get :new
      response.should have_tag("title", /Sign up/)
    end
  end

  .
  .
  .

end
```

8.1.1 Using `form_for`

The HTML element needed for submitting information to a remote website is a *form*, which suggests a good first step toward registering users is to make a form to accept their signup information. We can accomplish this in Rails with the `form_for` helper method; the result appears in Listing 8.2.

Listing 8.2. A form to sign up new users.
`app/views/users/new.html.erb`

```
<h2>Sign up</h2>

<% form_for(@user) do |f| %>
  <div class="field">
    <%= f.label :name %><br />
    <%= f.text_field :name %>
  </div>
  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :password %><br />
    <%= f.password_field :password %>
  </div>
  <div class="field">
    <%= f.label :password_confirmation, "Confirmation" %><br />
```

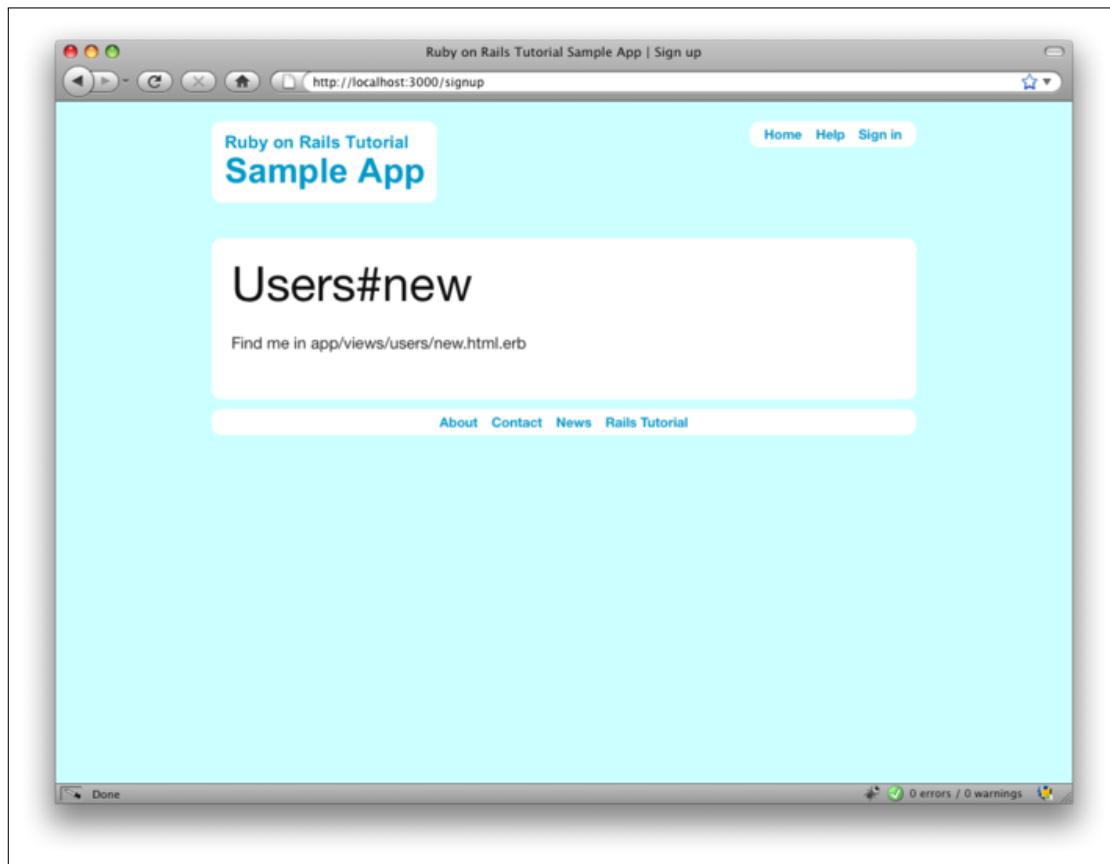


Figure 8.1: The current state of the signup page [/signup](#). ([full size](#))

Sign up

Name

Email

Password

Confirmation

(Sign up)

Figure 8.2: A mockup of the user signup page. (full size)

```
<%= f.password_field :password_confirmation %>
</div>
<div class="actions">
  <%= f.submit "Sign up" %>
</div>
<% end %>
```

Let's break this down into pieces. The presence of the `do` keyword indicates that `form_for` takes a block (Section 4.3.2), which has one variable, which we've called `f` for "form". Inside of the `form_for` helper, `f` is an object that represents a form; as is usually the case with Rails helpers, we don't need to know any details about the implementation, but what we *do* need to know is what the `f` object does: when called with a method corresponding to an [HTML form element](#)—such as a text field, radio button, or password field—it returns code for that element specifically designed to set an attribute of the `@user` object. In other words,

```
<div class="field">
  <%= f.label :name %><br />
  <%= f.text_field :name %>
</div>
```

creates the HTML needed to make a labeled text field element appropriate for setting the `name` attribute of a User model.

To see this in action, we need to drill down and look at the actual HTML produced by this form, but here we have a problem: the page currently breaks, because we have not set the `@user` variable—like all undefined instance variables (Section 4.2.3), `@user` is currently `nil`. Appropriately, if you run your test suite at this point, you'll see that the signup page tests fail. To get them to pass and get our form to render, we must define an `@user` variable in the controller action corresponding to `new.html.erb`, i.e., the `new` action in the Users controller. The `form_for` helper expects `@user` to be a User object, and since we're creating a *new* user we simply use `User.new`, as seen in Listing 8.3.

Listing 8.3. Adding an `@user` variable to the `new` action.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
  def new
    @user = User.new
    @title = "Sign up"
  end
end
```

With the `@user` variable so defined, the tests should be passing again,¹ and now the form (with the tiny bit of styling from Listing 8.4) appears as in Figure 8.3.

¹If you get an error like `views/users/new.html.erb.spec.rb` fails, remove those accursed view specs with `$ rm -rf spec/views`.

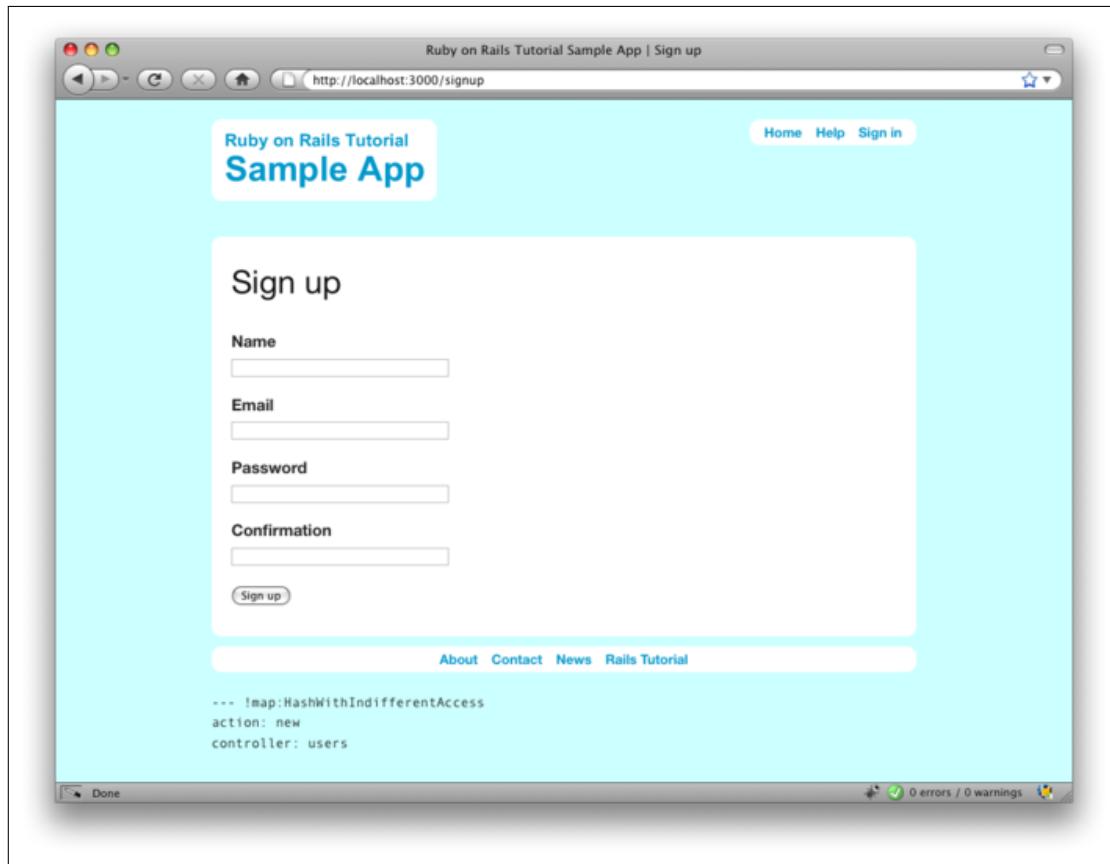


Figure 8.3: The signup form [/signup](#) for new users. ([full size](#))

Listing 8.4. A [wafer-thin](#) amount of CSS for the signup form.

`public/stylesheets/custom.css`

```
.  
. .  
. .  
div.field, div.actions {  
  margin-bottom: 10px;  
}
```

8.1.2 The form HTML

As indicated by Figure 8.3, the signup page now renders properly, indicating that the `form_for` code in Listing 8.2 is producing valid HTML. If you look at the HTML for the generated form (using either [Firebug](#) or the “view page source” feature of your browser), you should see markup as in Listing 8.5. Although many

of the details are irrelevant for our purposes, let's take a moment to highlight the most important parts of its structure.

Listing 8.5. The HTML for the form in Figure 8.3.

```
<form action="/users" class="new_user" id="new_user" method="post">
<div style="margin:0;padding:0;display:inline">
<input name="authenticity_token" type="hidden"
       value="rB82sI7Qw5J9J1UMILG/VQL411vH5putR+JwlxLScMQ=" />
</div>

<div class="field">
  <label for="user_name">Name</label><br />
  <input id="user_name" name="user[name]" size="30" type="text" />
</div>

<div class="field">
  <label for="user_email">Email</label><br />
  <input id="user_email" name="user[email]" size="30" type="text" />
</div>
<div class="field">
  <label for="user_password">Password</label><br />
  <input id="user_password" name="user[password]" size="30" type="password" />
</div>

<div class="field">
  <label for="user_password_confirmation">Confirmation</label><br />
  <input id="user_password_confirmation" name="user[password_confirmation]"
         size="30" type="password" />
</div>
<div class="actions">
  <input id="user_submit" name="commit" type="submit" value="Sign up" />
</div>
</form>
```

We'll start with the internal structure. Comparing Listing 8.2 with Listing 8.5, we see that the Embedded Ruby

```
<div class="field">
  <%= f.label :name %><br />
  <%= f.text_field :name %>
</div>
```

produces the HTML

```
<div class="field">
  <label for="user_name">Name</label><br />
```

```
<input id="user_name" name="user[name]" size="30" type="text" />
</div>
```

and

```
<div class="field">
  <%= f.label :password %><br />
  <%= f.password_field :password %>
</div>
```

produces the HTML

```
<div class="field">
  <label for="user_password">Password</label><br />
  <input id="user_password" name="user[password]" size="30" type="password" />
</div>
```

As seen in Figure 8.4, text fields (`type="text"`) simply display their contents, whereas password fields (`type="password"`) obscure the input for security purposes, as seen in Figure 8.4.

As we'll see in Section 8.3, the key to creating a user is the special `name` attribute in each `input`:

```
<input id="user_name" name="user[name]" - - - />
.
.
.
<input id="user_password" name="user[password]" - - - />
```

These `name` values allow Rails to construct an initialization hash (via the `params` variable first seen in Section 6.3.2) for creating users using the values entered by the user, as we'll see in Section 8.2.

The second important element is the `form` tag itself. Rails creates the `form` tag using the `@user` object: because every Ruby object knows its own class (Section 4.4.1), Rails figures out that `@user` is of class `User`; moreover, since `@user` is a *new* user, Rails knows to construct a form with the `post` method, which is the proper verb for creating a new object (Box 3.1):

```
<form action="/users" class="new_user" id="new_user" method="post">
```

Here the `class` and `id` attributes are largely irrelevant; what's important is `action="/users"` and `method="post"`. Together, these constitute instructions to issue an HTML POST request to the `/users` URL. We'll see in the next two sections what effects this has.

Finally, note the rather obscure code for the “authenticity token”:

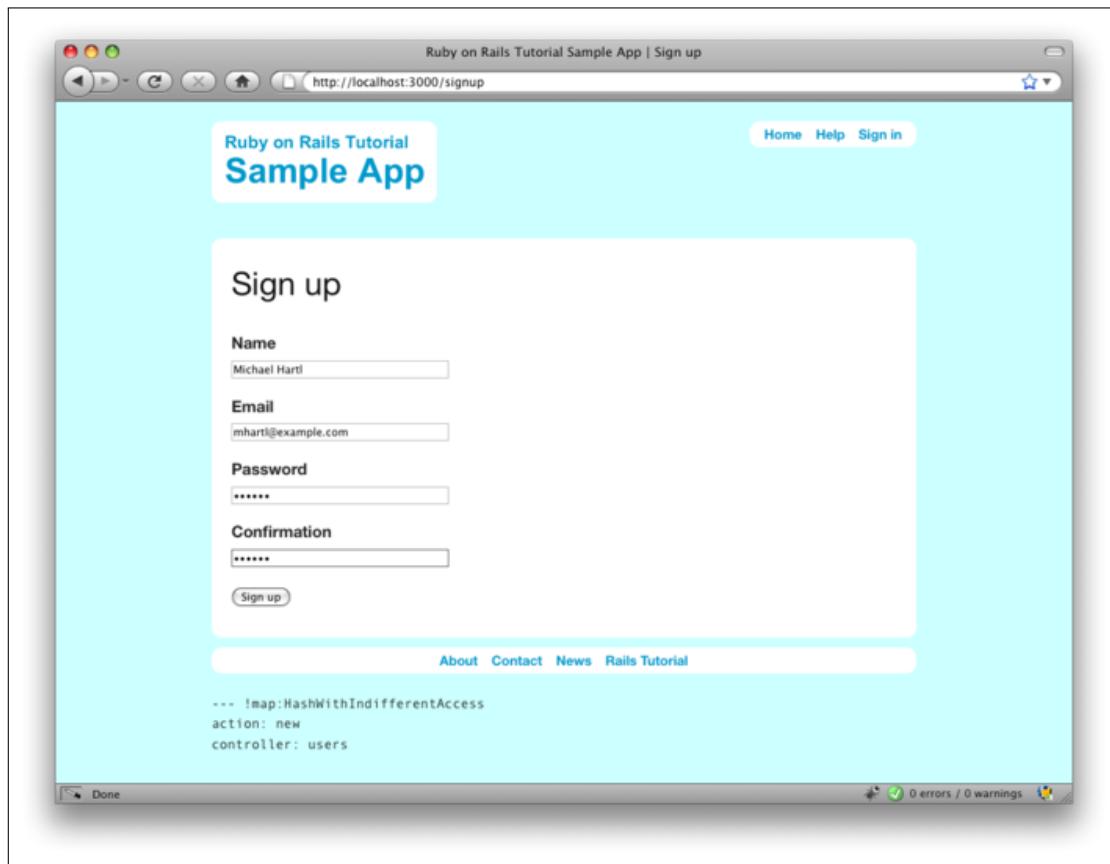


Figure 8.4: A filled-in form, showing the difference between `text` and `password` fields. ([full size](#))

```
<div style="margin:0;padding:0;display:inline">
<input name="authenticity_token" type="hidden"
       value="rB82sI7Qw5J9J1UMILG/VQL411vH5putR+JwlxLScMQ=" />
</div>
```

Here Rails uses a special unique value to thwart a particular kind of cross-site scripting attack called a *forgery*; see [here](#) if you’re interested in the details of how this works and why it’s important. Happily, Rails takes care of the problem for you, and the input tag is **hidden** so you don’t really have to give it a second thought, but it shows up when you view the form source so I wanted at least to address it.

8.2 Signup failure

Though we’ve briefly examined the HTML for the form in [Figure 8.3](#) (shown in [Listing 8.5](#)), it’s best understood in the context of *signup failure*. Just getting a signup form that accepts an invalid submission and re-renders the signup page (as mocked up in [Figure 8.5](#)) is a significant accomplishment, and it’s the goal of this section.

8.2.1 Testing failure

Recall from [Section 6.3.3](#) that adding `map.resources :users` to the `routes.rb` file ([Listing 6.25](#)) automatically ensures that our Rails application responds to the RESTful URLs from [Table 6.2](#). In particular, it ensures that a POST request to `/users` is handled by the `create` action. Our strategy for the `create` action is to use the form submission make a new user object using `User.new`, try (and fail) to save that user, and then render the signup page for possible resubmission. Our task is to write tests for this action, and then add `create` to the Users controller to get it to pass.

Let’s get started by reviewing the code for the signup form:

```
<form action="/users" class="new_user" id="new_user" method="post">
```

As noted in [Section 8.1.2](#), this HTML issues a POST request to the `/users` URL. In analogy with the `get` method, which issues a GET request inside of tests, we use the `post` method to issue a POST request to the `create` action. As we’ll see shortly, `create` takes in a hash corresponding to the object type being created; since this is a test for *signup failure*, we’ll just pass an `@attr` hash with blank entries, as seen in [Listing 8.6](#). This is essentially equivalent to visiting the signup page and clicking on the button without filling in any of the fields.

Listing 8.6. Tests for failed user signup.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views
  .
```

Sign up

- Name can't be blank
- Email is invalid
- Password is too short

Name

Email

Password

Confirmation

Sign up

Figure 8.5: A mockup of the signup failure page. ([full size](#))

```

.
.

describe "POST 'create'" do
  describe "failure" do
    before(:each) do
      @attr = { :name => "", :email => "", :password => "",
                :password_confirmation => "" }
      @user = Factory.build(:user, @attr)
      User.stub!(:new).and_return(@user)
      @user.should_receive(:save).and_return(false)
    end

    it "should have the right title" do
      post :create, :user => @attr
      response.should have_tag("title", /sign up/i)
    end

    it "should render the 'new' page" do
      post :create, :user => @attr
      response.should render_template('new')
    end
  end
end
end

```

By themselves, these two tests are relatively straightforward: we make sure that the title is correct, and then we check that a failed signup attempt just re-renders the new user page. On the other hand, the `before(:each)` block is a little tricky—its purpose is to verify that the `create` action attempts to `save` the user. To understand it, let's break it down into pieces:

```

before(:each) do
  @attr = { :name => "", :email => "", :password => "",
            :password_confirmation => "" }
  @user = Factory.build(:user, @attr)
  User.stub!(:new).and_return(@user)
  @user.should_receive(:save).and_return(false)
end

```

First, note that we use Factory Girl not to simulate an *existing* user (as in, e.g., Listing 7.17), but rather to `build` a *new* user based on a particular attributes hash:

```

@attr = { :name => "", :email => "", :password => "",
          :password_confirmation => "" }
@user = Factory.build(:user, @attr)

```

We then use an RSpec stub (first seen in Listing 7.17) to ensure that a call to `User.new` returns our factory user:

```
User.stub! (:new).and_return(@user)
```

The final line sets up a *message expectation*, indicating that the `@user` object should receive the `save` method in each test:

```
@user.should_receive(:save).and_return(false)
```

It's essential to understand that, unlike stubs, which merely return values, message expectations are *requirements*. In other words,

```
User.stub! (:new).and_return(@user)
```

ensures that a call to `User.new` returns `@user`, but it doesn't *require* it, whereas

```
@user.should_receive(:save).and_return(false)
```

leads to an *error* if `@user.save` is never called. This expectation error ensures that the tests initially fail.

It's also important to understand that the expectation

```
@user.should_receive(:save).and_return(false)
```

arranges for `@user.save` to return `false`. (Recall from Section 6.1.3 that `@user.save` returns `true` if the save succeeds and `false` if it fails.) In particular, it doesn't *require* the method to return `false`; the expectation *enforces* it (Box 8.1).

Box 8.1. An expectation subtlety

Naïvely, it seems that the code

```
@user.should_receive(:save).and_return(false)
```

requires that `@user.save` gets called *and* requires that it return `false`—but this is not the case. Instead, though it *does* require `@user.save` to be called, it intercepts the call and *arranges* for it to return `false`. The first part is an expectation; the second part is a guarantee.

Why does a message expectation not verify the return value? Doesn't that need to be tested as well? It does—and it has already been tested in the model layer (Chapter 6). Inside a controller test, we *assume* that the model layer does its job correctly. The purpose of a message expectation is to verify that the right methods get called, and then (like stubs) to supply a return value without ever hitting the database.

8.2.2 A working form

We can get the tests from [Section 8.2.1](#) to pass with the code in [Listing 8.7](#). This listing includes a second use of the `render` method, which we first saw in the context of partials ([Section 5.1.3](#)); as you can see, `render` works in controller actions as well. Note that we've taken this opportunity to introduce an `if-else` branching structure, which allows us to handle the cases of failure and success separately based on the value of `@user.save`.

Listing 8.7. A `create` action that can handle signup failure (but not success).

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
  def create
    @user = User.new(params[:user])
    if @user.save
      # Handle a successful save.
    else
      @title = "Sign up"
      render 'new'
    end
  end
end
```

The best way understand how the code in [Listing 8.7](#) works is to *submit* the form with some invalid signup data; the results appear in [Figure 8.6](#).

To get a clearer picture of how Rails handles the submission, let's take a closer look at the `params` hash in the debug information at the bottom of [Figure 8.6](#):

```
--- !map:HashWithIndifferentAccess
commit: Sign up
authenticity_token: rB82sI7Qw5J9J1UMILG/VQL411vH5puR+Jw1xL5cMQ=
action: create
controller: users
user: !map:HashWithIndifferentAccess
  name: Foo Bar
  password_confirmation: dude
  password: dude
  email: foo@invalid
```

We saw starting in [Section 6.3.2](#) that the `params` hash contains information about each request; in the case of a URL like `/users/1`, the value of `params[:id]` is the `id` of the corresponding user (`1` in this example). In the case of posting to the signup form, `params` instead contains a hash of hashes, a construction we first saw in [Section 4.3.3](#), which introduced the strategically named `params` variable in a console session. This debug information above shows that submitting the form results in a `user` hash with attributes corresponding to the submitted values, where the keys come from the `name` attributes of the `input` tags seen in [Listing 8.2](#); for example, the value of

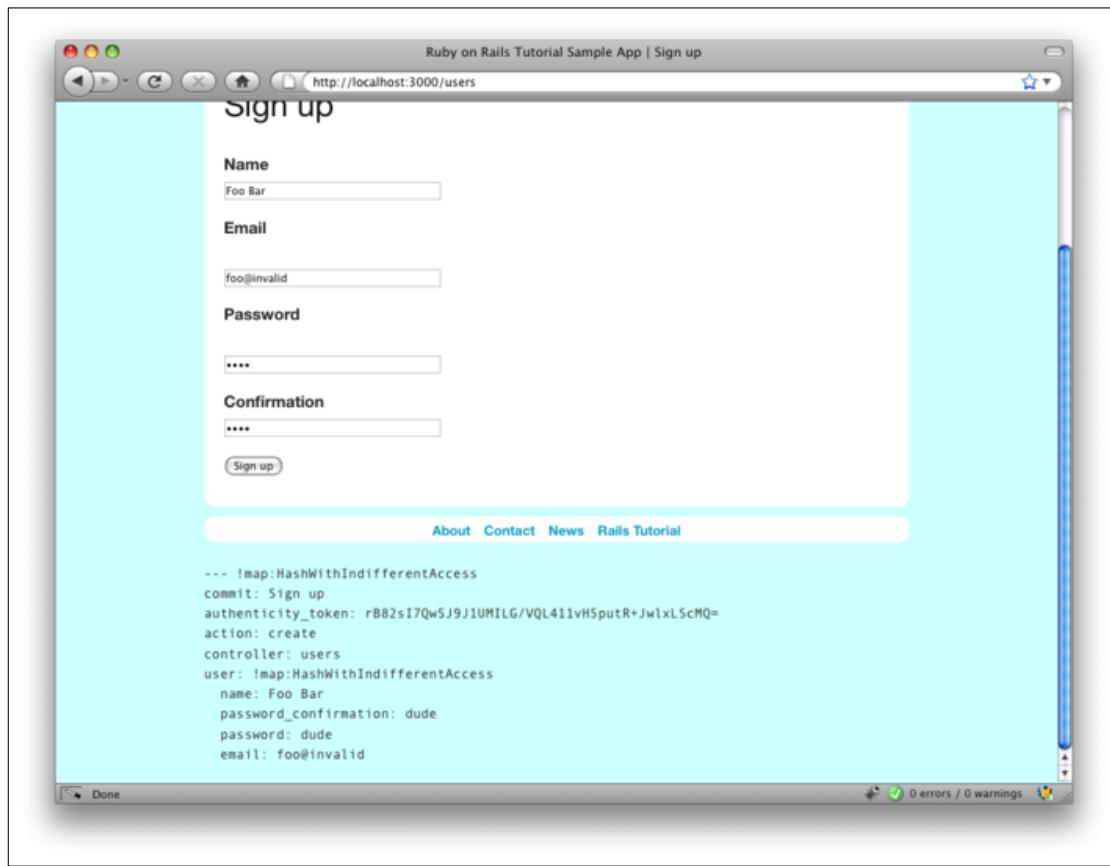


Figure 8.6: Signup failure with a `params` hash. (full size)

```
<input id="user_email" name="user[email]" size="30" type="text" />
```

with name "`user[email]`" is precisely the `email` attribute of the `user` hash.

Though the hash keys appears as strings in the debug output, internally Rails uses symbols, so that `params[:user]` is the hash of user attributes—in fact, exactly the attributes needed as an argument to `User.new`, as first seen in [Section 4.4.5](#) and appearing in [Listing 8.7](#). This means that the line

```
@user = User.new(params[:user])
```

is equivalent to

```
@user = User.new(:name => "Foo Bar", :email => "foo@invalid",
                  :password => "dude", :password_confirmation => "dude")
```

This is exactly the format needed to initialize a `User` model object with the given attributes.

Of course, instantiating such a variable has implications for successful signup—as we'll see in [Section 8.3](#), once `@user` is defined properly, calling `@user.save` is all that's needed to complete the registration—but it has consequences even in the failed signup considered here. Note in [Figure 8.6](#) that the fields are *pre-filled* with the data from the failed submission. This is because `form_for` automatically fills in the fields with the attributes of the `@user` object, so that, for example, if `@user.name` is "Foo" then

```
<% form_for(@user) do |f| %>
  <div class="field">
    <%= f.label :name %><br />
    <%= f.text_field :name %>
  </div>
  .
  .
  .
```

will produce the HTML

```
<form action="/users" class="new_user" id="new_user" method="post">

  <div class="field">
    <label for="user_name">Name</label><br />
    <input id="user_name" name="user[name]" size="30" type="text" value="Foo"/>
  </div>
  .
  .
  .
```

Here the `value` of the `input` tag is "Foo", so that's what appears in the text field.

8.2.3 Signup error messages

Though not strictly necessary, it's helpful to output error messages on failed signup to indicate the problems that prevented successful user registration. Rails provides just such messages based on the User model validations. For example, consider trying to save a user with invalid email and with a short password:

```
$ script/console
>> user = User.new(:name => "Foo Bar", :email => "foo@invalid",
?>                      :password => "dude", :password_confirmation => "dude")
>> user.save
=> false
>> user.errors.full_messages
=> ["Email is invalid", "Password is too short (minimum is 6 characters)"]
```

Here the `errors.full_messages` object (which we saw briefly in Section 6.2.1) contains an array of error messages.

As in the console session above, the failed save in Listing 8.7 generates a list of error messages associated with the `@user` object. To display the messages in the browser, we use the special `error_messages` object on the form variable `f`, as seen in Listing 8.8.

Listing 8.8. Code to add error message display to the signup form.

`app/views/users/new.html.erb`

```
<h2>Sign up</h2>
<% form_for(@user) do |f| %>
  <%= f.error_messages %>
  .
  .
  .
<% end %>
```

The `error_messages` method inserts the HTML needed for a list of form submission errors, including CSS ids and classes for use in styling the response.² After adding the CSS in Listing 8.9, on failed submission the error messages appear as in Figure 8.7. Because the messages are generated by the model validations, they will automatically change if you ever change your mind about, say, the format of email addresses, or the minimum length on passwords.

Listing 8.9. CSS for styling error messages.

`public/stylsheets/custom.css`

```
.
.
.
```

²Exactly how this happens is a bit obscure; the form variable `f` has access to the `@user` variable inside the block, and must at some point iterate through the user's error messages in order to display them as HTML. Discovering exactly how `error_messages` works is left as an exercise (Section 8.6).

```
.fieldWithErrors {
  padding: 2px;
  background-color: red;
  display: table;
}

.fieldWithErrors label {
  color: #fff;
}

#errorExplanation {
  width: 400px;
  border: 2px solid red;
  padding: 7px;
  padding-bottom: 12px;
  margin-bottom: 20px;
  background-color: #f0f0f0;
}

#errorExplanation h2 {
  text-align: left;
  font-weight: bold;
  padding: 5px 5px 5px 15px;
  font-size: 12px;
  margin: -7px;
  background-color: #c00;
  color: #fff;
}

#errorExplanation p {
  color: #333;
  margin-bottom: 0;
  padding: 5px;
}

#errorExplanation ul li {
  font-size: 12px;
  list-style: square;
}
```

8.2.4 Filtering parameter logging

Before moving on to successful signup, there's one loose end to tie off. You might have noticed that, even though we went to great pains to encrypt the password in [Chapter 7](#), both the password and its confirmation appear as `cleartext` in the debug information. By itself this is no problem—recall from [Listing 6.22](#) that this information only appears for applications running in `development` mode, so actual users would never see it—but it does hint at a real problem: the passwords also appear unencrypted in the *log file* that Rails uses to record information about the running application. Indeed, if you look in the development log file you should see lines like those shown in [Listing 8.10](#). (This information also appears in the server log, which we saw in

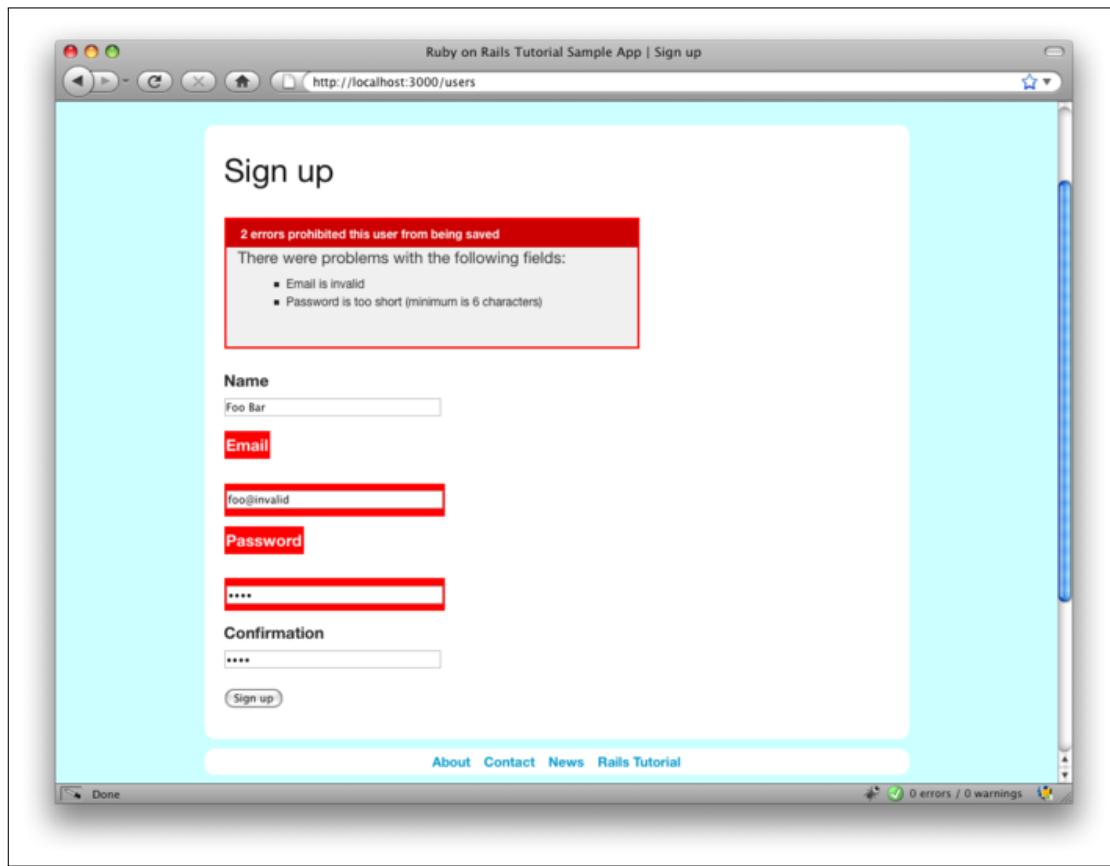


Figure 8.7: Failed signup with error messages. (full size)

Figure 1.7.)

Listing 8.10. The development log with visible passwords.

`log/development.log`

```
Parameters: {"commit"=>"Sign up", "action"=>"create",
"authenticity_token"=>"K1HchFF8uYE8ZaQKz5DVG9vF2KGoXJu4JGp/VE3NMjA=",
"controller"=>"users",
"user"=>{ "name"=>"Foo Bar", "password_confirmation"=>"dude",
"password"=>"dude", "email"=>"foo@invalid" } }
```

It would be a terrible security breach to store unencrypted passwords in the log files—if anyone ever got a hold of the file, they would potentially obtain the passwords for every user on the system. (Of course, here the signup fails, but the problem is exactly the same for successful submissions.) Happily, Rails offers a simple solution: just *filter* the password parameter so that it doesn’t appear in the log (Listing 8.11).³ After you add this line, the log appears as in Listing 8.12. (Of course, in production the log file will be `log/production.log`, and the filtering will work the same way.)

Listing 8.11. Code to filter `password` and `password_confirmation` from the log.

`app/controllers/application_controller.rb`

```
class ApplicationController < ActionController::Base
.
.
.
# Scrub sensitive parameters from your log
filter_parameter_logging :password
end
```

Listing 8.12. The development log with filtered passwords.

`log/development.log`

```
Parameters: {"commit"=>"Sign up", "action"=>"create",
"authenticity_token"=>"K1HchFF8uYE8ZaQKz5DVG9vF2KGoXJu4JGp/VE3NMjA=",
"controller"=>"users",
"user"=>{ "name"=>"Foo Bar", "password_confirmation"=>"[FILTERED]",
"password"=>"[FILTERED]", "email"=>"foo@invalid" } }
```

8.3 Signup success

Having handled invalid form submissions, now it’s time to complete the signup form by actually saving a new user (if valid) to the database. First, we try to save the user; if the save succeeds, the user’s information gets written to the database automatically, and we then *redirect* the browser to show the user’s profile (together with a friendly greeting), as mocked up in Figure 8.8. If it fails, we simply fall back on the behavior developed in Section 8.2.

³Password parameter filtering is enabled by default in Rails 3.

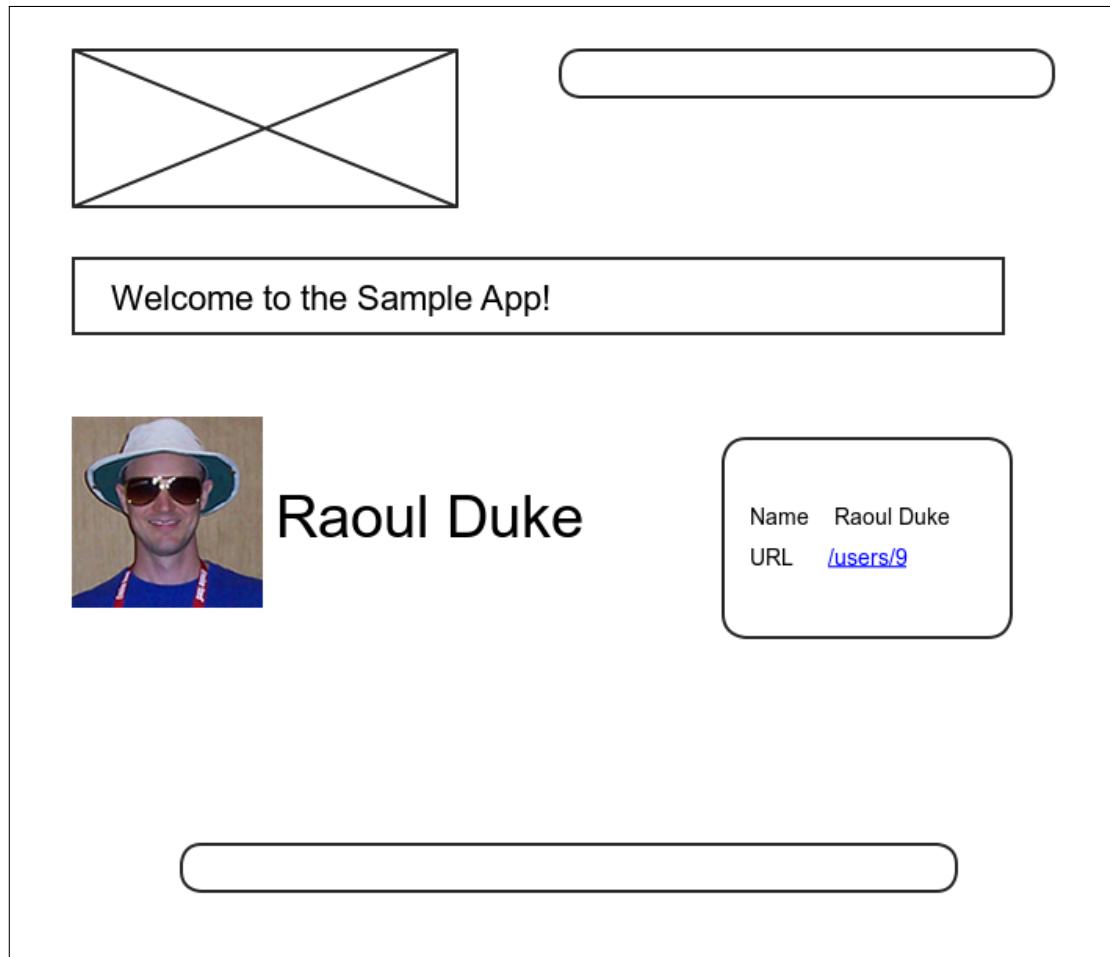


Figure 8.8: A mockup of successful signup. ([full size](#))

8.3.1 Testing success

The tests for a successful signup follow the lead of the failed signup tests from Listing 8.6. Let's take a look at the result in Listing 8.13.

Listing 8.13. Tests for signup success.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  .
  .
  .

  describe "POST 'create'" do
    .
    .
    .

    describe "success" do
      before(:each) do
        @attr = { :name => "New User", :email => "user@example.com",
                  :password => "foobar", :password_confirmation => "foobar" }
        @user = Factory(:user, @attr)
        User.stub!(:new).and_return(@user)
        @user.should_receive(:save).and_return(true)
      end

      it "should redirect to the user show page" do
        post :create, :user => @attr
        response.should redirect_to(user_path(@user))
      end
    end
  end
end
end
```

As with the signup failure tests (Listing 8.6), here we use `post :create` to hit the `create` action with an HTTP POST request. Unlike in Listing 8.6, which used `Factory.build` to simulate a new user, Listing 8.13 uses `Factory(:user)` to simulate a *saved* user. As before, our expectation is that `@user.save` will be called, but since this is a test for a successful save we arrange for it to return `true`.⁴

```
@user.should_receive(:save).and_return(true)
```

After a successful save, we check for a redirect to the user show page using this code:

⁴Recall from Box 8.1 that the `and_return` part of the statement is not an expectation, but rather a guarantee.

```
response.should redirect_to(user_path(@user))
```

This is the kind of redirect that happens on nearly every successful form submission on the web, and with RSpec's helpful syntax you don't have to know anything about the underlying HTTP response code.⁵ The URL itself is generated using the named route `user_path` shown in [Table 7.1](#).

8.3.2 The finished signup form

To get these tests to pass and thereby complete a working signup form, fill in the commented-out section in [Listing 8.7](#) with a redirect, as shown in [Listing 8.14](#).

Listing 8.14. The user `create` action with a save and a redirect.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
  def create
    @user = User.new(params[:user])
    if @user.save
      redirect_to @user
    else
      @title = "Sign up"
      render 'new'
    end
  end
end
```

Note that we can omit the `user_path` in the redirect, writing simply `redirect_to @user` to redirect to the user show page, a convention we saw before with `link_to` in [Listing 7.24](#). This syntax is nicely succinct, but unfortunately RSpec doesn't understand it, so we have to use the more verbose `user_path(@user)` in that case.

8.3.3 The flash

Before submitting a valid registration in a browser, we're going to add a bit of polish common in web applications: a message that appears temporarily and then disappears upon page reload. (If this is unclear now, be patient; a concrete example appears shortly.) The Rails way to accomplish this is to use a special variable called the *flash*, which operates like `flash memory` in that it stores its data temporarily. The `flash` variable is effectively a hash; you may even recall the console example in [Section 4.3.3](#), where we saw how to iterate through a hash using a strategically named `flash` hash. To recap, try this console session:

⁵In case you're curious, the response code is 302, in contrast to the "permanent" 301 redirect discussed briefly in Box 3.1.

```
$ script/console
>> flash = { :success => "It worked!", :error => "It failed. :-(" }
=> { :success => "It worked!", :error => "It failed. :-(" }
>> flash.each do |key, value|
?>   puts "#{key}"
?>   puts "#{value}"
>> end
success
It worked!
error
It failed. :-(

```

We can arrange to display the contents of the flash site-wide by including it in our application layout, as in Listing 8.15.

Listing 8.15. Adding the contents of the `flash` variable to the site layout.
`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html>
  .
  .
  .
  <div id="content" class="round">
    <% flash.each do |key, value| %>
      <div class="flash <%= key %>"><%= value %></div>
    <% end %>
    <%= yield %>
  </div>
  .
  .
  .
</html>
```

This code arranges to insert a `div` tag for each element in the flash, with a CSS class indicating the type of message. For example, if `flash[:success] = "Welcome to the Sample App!"`, then the code

```
<% flash.each do |key, value| %>
  <div class="flash <%= key %>"><%= value %></div>
<% end %>
```

will produce this HTML:⁶

⁶Note that the key `:success` is a symbol, but Embedded Ruby automatically converts it to the string "`success`" before inserting it into the template.

```
<div class="flash success">Welcome to the Sample App!</div>
```

The reason we iterate through all possible key/value pairs is so that we can include other kinds of flash messages; for example, in [Listing 9.8](#) we'll see `flash[:error]` used to indicate a failed signin attempt.⁷

Let's test for the right flash message by making sure the right message appears under the key `:success` ([Listing 8.16](#)).

Listing 8.16. A test for a flash message on successful user signup.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  .

  .

  describe "POST 'create'" do
    .

    .

    describe "success" do
      .

      .

      it "should have a welcome message" do
        post :create, :user => @attr
        flash[:success].should =~ /welcome to the sample app/i
      end
    end
  end
end
```

This uses the case-insensitive regular expression `/.../i`, a construction we last saw in the [EmailRegex](#) of [Listing 6.16](#). Note that we don't yet test for the appearance of the actual flash message HTML; we'll fix this by testing for the actual `div` tag in [Section 8.4.3](#).

We get the flash test to pass by assigning to `flash[:success]` in the `create` action as in [Listing 8.17](#). The message uses different capitalization from the one in the test, but the test passes anyway because of the `i` at the end of the regular expression. This way we won't break the test if we write, e.g., `sample app` in place of `Sample App`.

Listing 8.17. Adding a flash message to user signup.

`app/controllers/users_controller.rb`

⁷Actually, we'll use the closely related `flash.now`, but we'll defer that subtlety until we need it.

```

class UsersController < ApplicationController
  .
  .
  .
  def create
    @user = User.new(params[:user])
    if @user.save
      flash[:success] = "Welcome to the Sample App!"
      redirect_to @user
    else
      @title = "Sign up"
      render 'new'
    end
  end
end

```

8.3.4 The first signup

We can see the result of all this work by signing up our first user (under the name “Rails Tutorial” and email address “example@railstutorial.org”), which shows a friendly message upon successful signup, as seen in [Figure 8.9](#). (The nice green styling for the `success` class comes included with the Blueprint CSS framework from [Section 4.1.2](#).) Then, upon reloading the user show page, the flash message disappears as promised ([Figure 8.10](#)).

We can now check our database just to be double-sure that the new user was actually created:

```

$ script/console
>> user = User.first
=> #<User id: 1, name: "Rails Tutorial", email: "example@railstutorial.org",
  created_at: "2010-02-17 03:07:53", updated_at: "2010-02-17 03:07:53",
  encrypted_password: "48aa8f4444b71f3f713d87d051819b0d44cd89f4a963949f201...", 
  salt: "f52924ba502d4f92a634d4f9647622ccce26205176cceca2adc...">

```

Success!

8.4 RSpec integration tests

In principle we are done with user signup at this point, but you may have noticed that we haven’t tested the structure of the signup form, nor have we tested that submissions actually work. Of course, we *have* checked these things by viewing the pages in our browser, but the whole point of automated testing is to make sure that once things work they stay that way. Making such tests is the goal of this section—and the results are pretty sweet.

One testing method would be to check the HTML structure of the form (using `integrate_views` and the `have_tag` method), and indeed this is a good way to test-drive views. ([Section 8.6](#) has an exercise to this effect.) But I prefer not to test the detailed HTML structure of views—I don’t see any reason why we should have to know that Rails implements user email submission using `name="user[email]"`, and indeed any

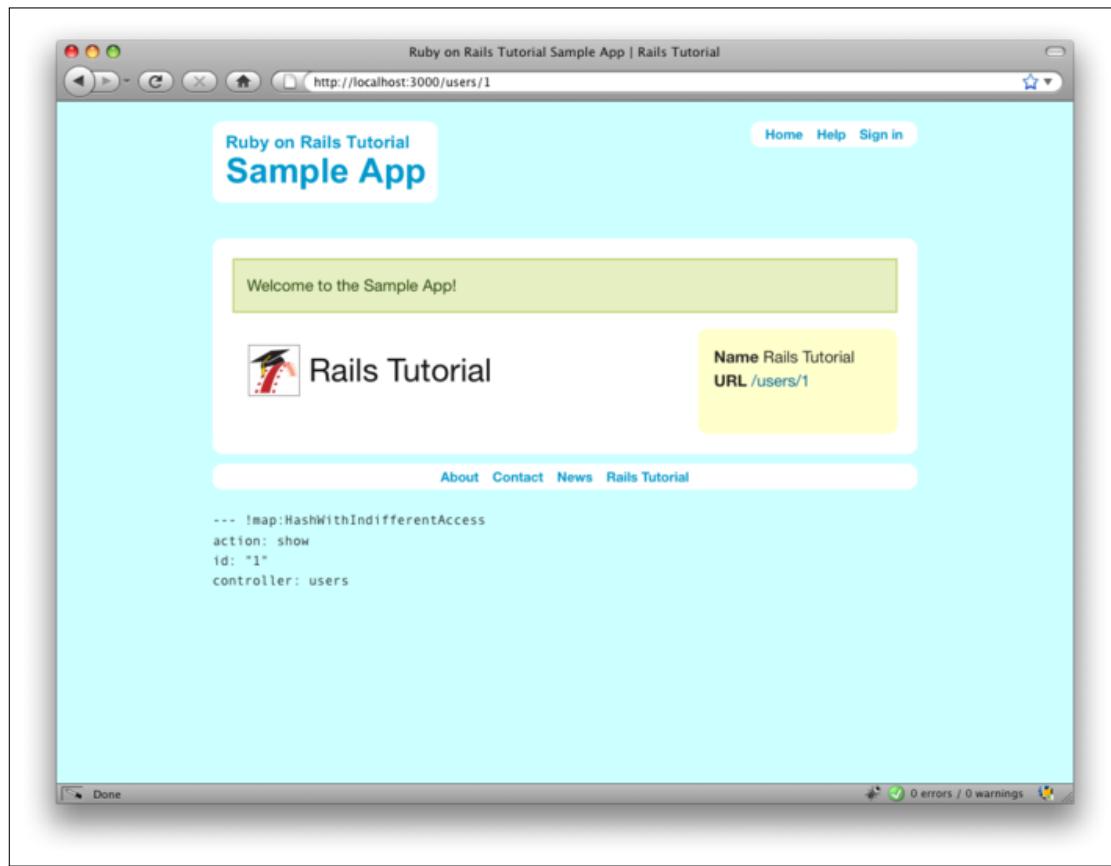


Figure 8.9: The results of a successful user signup, with flash message. ([full size](#))

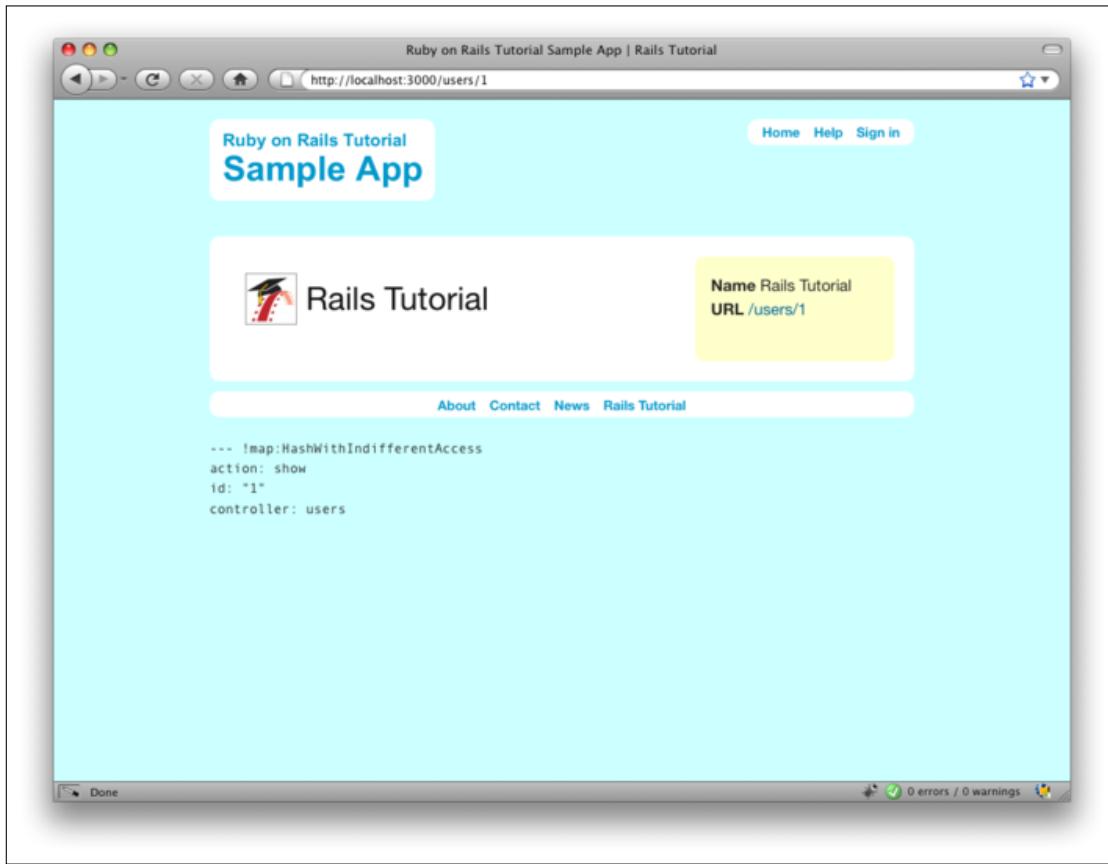


Figure 8.10: The flash-less profile page after a browser reload. ([full size](#))

test of that structure would break if a future Rails version changed this convention. Moreover, it would be nice to have a test for the entire signup process: visiting the signup page, filling in the form values, clicking the button, and making sure (if the submission is valid) that a new user gets created in the (test) database.

Though it's not the only way (see [Box 8.2](#)), my preferred solution to this problem is to use an RSpec integration test, which we first used in [Section 5.2.1](#) to test the custom routes (such as `/about` for the About page). In that section, we saw only a tiny fraction of the power of integration tests; below, we'll see just how amazing they can be.

Box 8.2. Integration alternatives

As we've seen in this and previous chapters, *Ruby on Rails Tutorial* uses RSpec for all its tests, including integration tests. In my view, there is no match for the simplicity and power of RSpec integration tests, especially when combined with the expressiveness of Webrat. There are a couple of viable alternatives, though. One is the Rails default, integration testing with `Test::Unit`. This is fine if you use `Test::Unit` elsewhere, but we're using RSpec in this tutorial, and I prefer not to mix RSpec and `Test::Unit` in a single project. A second option is [Cucumber](#), which works nicely with RSpec and allows the definition of plain-text stories describing application behavior. Cucumber is quite powerful and has a growing following in the Ruby community; although I find that the plain-text stories can be a bit verbose and (cuc)cumbersome, I suggest taking a look at some [Cucumber tutorials](#) to see if it suits you.

8.4.1 Webrat

The first step toward amazing signup integration tests is to install *Webrat*, which allows us to write tests using a highly expressive mini-language for interacting with web pages. (The [Webrat home page](#) has installation instructions, and the [Webrat wiki](#) has further documentation.) If you're using Linux, you'll probably have to install a couple libraries first:

```
$ [sudo] apt-get install libxslt1-dev libxml2-dev
```

Otherwise, you just need to install the `webrat` gem (together with `launchy`, a gem needed to spawn new browser windows when Webrat encounters exceptions):

```
$ [sudo] gem install webrat -v 0.7.0
$ [sudo] gem install launchy -v 0.3.5
```

Then add a couple lines in your RSpec helper file ([Listing 8.18](#)) and you'll be good to go.

Listing 8.18. Code needed to use Webrat with RSpec.

`spec/spec_helper.rb`

```
.
.
.
```

```

require 'spec/rails'
require 'webrat'

Webrat.configure do |config|
  config.mode = :rails
end
.
.
.

```

8.4.2 Users signup failure should not make a new user

Now we're ready to make an integration test for signing up users. RSpec comes with a generator to make such an *integration spec*:

```

$ script/generate integration_spec users
      create  spec/integration/users_spec.rb

```

We start with signup failure. A simple way to arrange a failing signup is to visit the signup URL and just click the button, resulting in a page as in Figure 8.11. Upon failed submission, the response should render the `users/new` template. If you inspect the resulting HTML, you should see something like the markup in Listing 8.19. This means that we can test for the presence of error messages by looking for a `div` tag with the CSS id `"errorExplanation"`. A test for these steps appears in Listing 8.20.

Listing 8.19. The error explanation `div` from the page in Figure 8.11.

```

<div class="errorExplanation" id="errorExplanation">
  <h2>5 errors prohibited this user from being saved</h2>
  <p>There were problems with the following fields:</p>
  <ul>
    <li>Name can't be blank</li>
    <li>Email can't be blank</li>
    <li>Email is invalid</li>
    <li>Password can't be blank</li>
    <li>Password is too short (minimum is 6 characters)</li>
  </ul>
</div>

```

Listing 8.20. Testing signup failure.

`spec/integration/users_spec.rb`

```

require 'spec_helper'

describe "Users" do

```

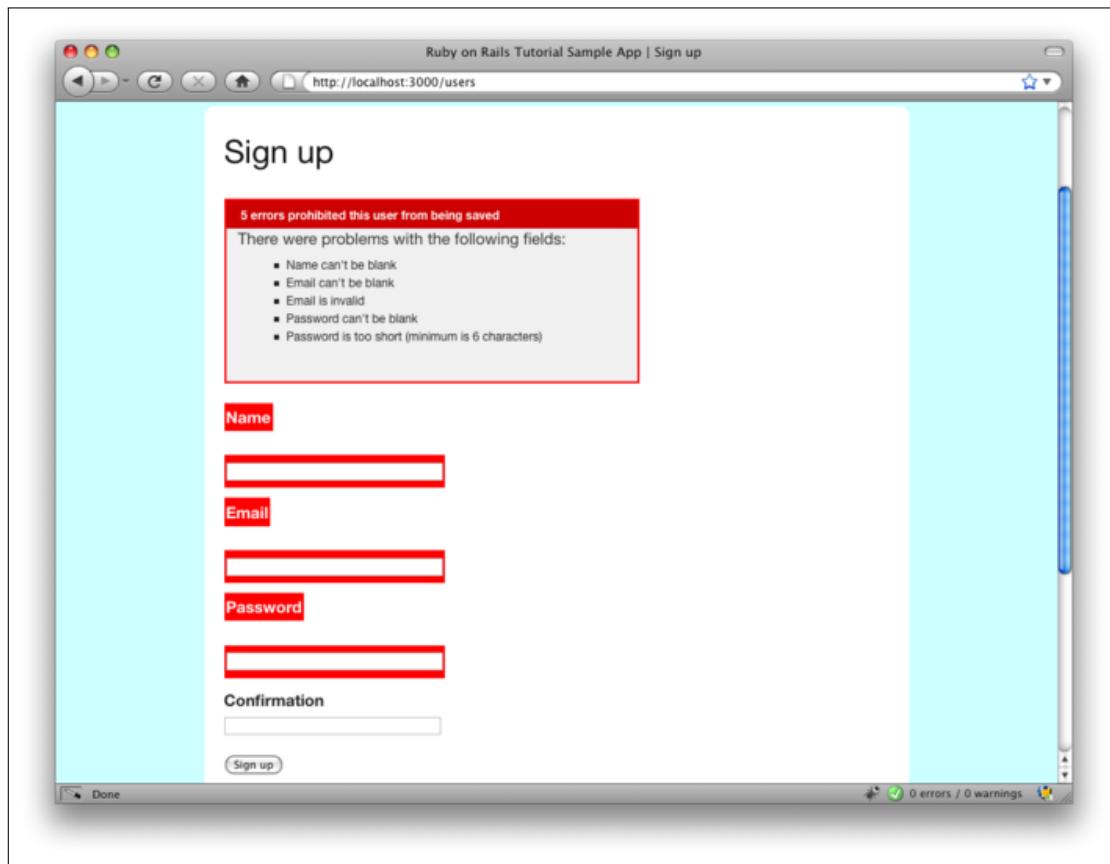


Figure 8.11: The result of visiting `/signup` and just clicking “Sign up”. ([full size](#))

```

describe "signup" do
  describe "failure" do
    it "should not make a new user" do
      visit signup_path
      click_button
      response.should render_template('users/new')
      response.should have_tag("div#errorExplanation")
    end
  end
end

```

Here "`div#errorExplanation`" is CSS-inspired shorthand for

```
<div id="errorExplanation">...</div>
```

(Recall that CSS uses the pound sign `#` to style ids.)

Notice how natural the language is in Listing 8.20. The only problem is that it doesn't *quite* test what we want: we're not actually testing that a failed submission fails to create a new user. To do so, we need two new elements. First, Active Record models all have a `count` method, which simply returns how many records of that type are in the database. If you cleared the database at the beginning of the chapter and then created a new user in Section 8.3, this number should currently be 1:

```

$ script/console
>> User.count
=> 1

```

Second, we need a way to wrap the test steps in a single package, and then check that it doesn't change the `User` count. The way to do this is to use a `lambda`,⁸ as shown in Listing 8.21.

Listing 8.21. Testing signup failure with a `lambda`.

`spec/integration/users_spec.rb`

```

require 'spec_helper'

describe "Users" do
  describe "signup" do
    describe "failure" do
      it "should not make a new user" do

```

⁸The name comes from the [lambda calculus](#), a mathematical system for representing functions and their operations.

```
lambda do
  visit signup_path
  click_button
  response.should render_template('new')
  response.should have_tag("div#errorExplanation")
end.should_not change(User, :count)
end
end
end
```

This wraps the signup failure steps in a `lambda` package and verifies that `User.count` doesn't get changed by the operation. The `change` method is part of RSpec, and

```
should_not change(User, :count)
```

verifies that the code inside the `lambda` block doesn't change the value of `User.count`.

This test should pass, and it provides an end-to-end verification that our signup machinery is working, at least for failed submissions.

8.4.3 Users signup success should make a new user

We come now to the integration test for successful signup. In this case, we need to fill in the signup fields with valid user data. When we do, the result should be the user show page with a “flash success” `div` tag, and it should change the User count by 1. Listing 8.22 shows how to do it.

Listing 8.22. Testing signup success.

`spec/integration/users_spec.rb`

```
require 'spec_helper'

describe "Users" do

  describe "signup"
  .
  .
  .

  describe "success" do

    it "should make a new user" do
      lambda do
        visit signup_path
        fill_in "Name",           :with => "Example User"
        fill_in "Email",          :with => "user@example.com"
        fill_in "Password",       :with => "foobar"
        fill_in "Confirmation",  :with => "foobar"
        click_button
      end
    end
  end
end
```

```

    response.should render_template('users/show')
  end.should change(User, :count).by(1)
end
end
end
end

```

By the way, although it's not obvious from the Webrat documentation, you can use the CSS id of the text box instead of the label, so `fill_in :user_name` also works.⁹ (This is especially nice for forms that don't use labels.)

I hope you agree that the combination of Webrat and RSpec syntax is incredibly natural and succinct. For example, to fill in a field with a value, we just use code like this:

```

fill_in "Name",           :with => "Example User"
fill_in "Email",          :with => "user@example.com"
fill_in "Password",       :with => "foobar"
fill_in "Confirmation",  :with => "foobar"

```

Here the first arguments to `fill_in` are the label values, i.e., exactly the text the user sees in the browser; there's no need to know anything about the underlying HTML structure generated by the Rails `form_for` helper.

Finally, we come to the coup de grâce—testing that successful signup actually creates a user in the database:

```

it "should make a new user" do
  lambda do
    .
    .
    .
  end.should change(User, :count).by(1)

```

As in Listing 8.21, we've wrapped the code for a successful signup in a `lambda` block. In this case, instead of making sure that the User count *doesn't* change, we verify that it increases by 1 due to a User record being created in the test database.

With that, our signup integration tests are complete, and we can be confident that, if users don't join our site, it's not because the signup form is broken.

8.5 Conclusion

Being able to sign up users is a major milestone for our application. Though the sample app has yet to accomplish anything useful, we have laid an essential foundation for all future development. In the next two chapters, we will complete two more major milestones: first, in Chapter 9 we will complete our authentication

⁹You can use Firebug or your browser's "view source" if you need to figure out the id. Or you can note that Rails uses the name of the resource and the name of the attribute separated with an underscore, yielding `user_name`, `user_email`, etc.

machinery by allowing users to sign in and out of the application; second, in [Chapter 10](#) we will allow all users to update their account information and will allow site administrators to delete users, while also adding page protection to enforce a site security model, thereby completing the full suite of the Users resource REST actions from [Table 6.2](#).

As usual, if you’re using Git, you should merge your changes into the `master` branch at this point:

```
$ git commit -am "User signup complete"
$ git checkout master
$ git merge signing-up
```

8.6 Exercises

1. **(advanced)** Dig into the latest Rails API to discover how the `error_messages` method from [Listing 8.8](#) is implemented.
2. Using the model in [Listing 8.23](#), write tests to check for the presence of each field on the signup form. (Don’t forget the `integrate_views` line, which is essential for this to work.)
3. Oftentimes signup forms will clear the password field for failed submissions, as shown in [Figure 8.12](#). Modify the Users controller `create` action to replicate this behavior. *Hint:* Reset `@user.password`.
4. The flash HTML in [Listing 8.15](#) is a particularly ugly combination of HTML and ERB. Verify by running the test suite that the cleaner code in [Listing 8.24](#), which uses the Rails `content_tag` helper, also works.

Listing 8.23. A template for testing for each field on the signup form.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  .

  .

  describe "GET 'new'" do
    .

    .

    it "should have a name field" do
      get :new
      response.should have_tag("input[name=?][type=?]", "user[name]", "text")
    end

    it "should have an email field"
```

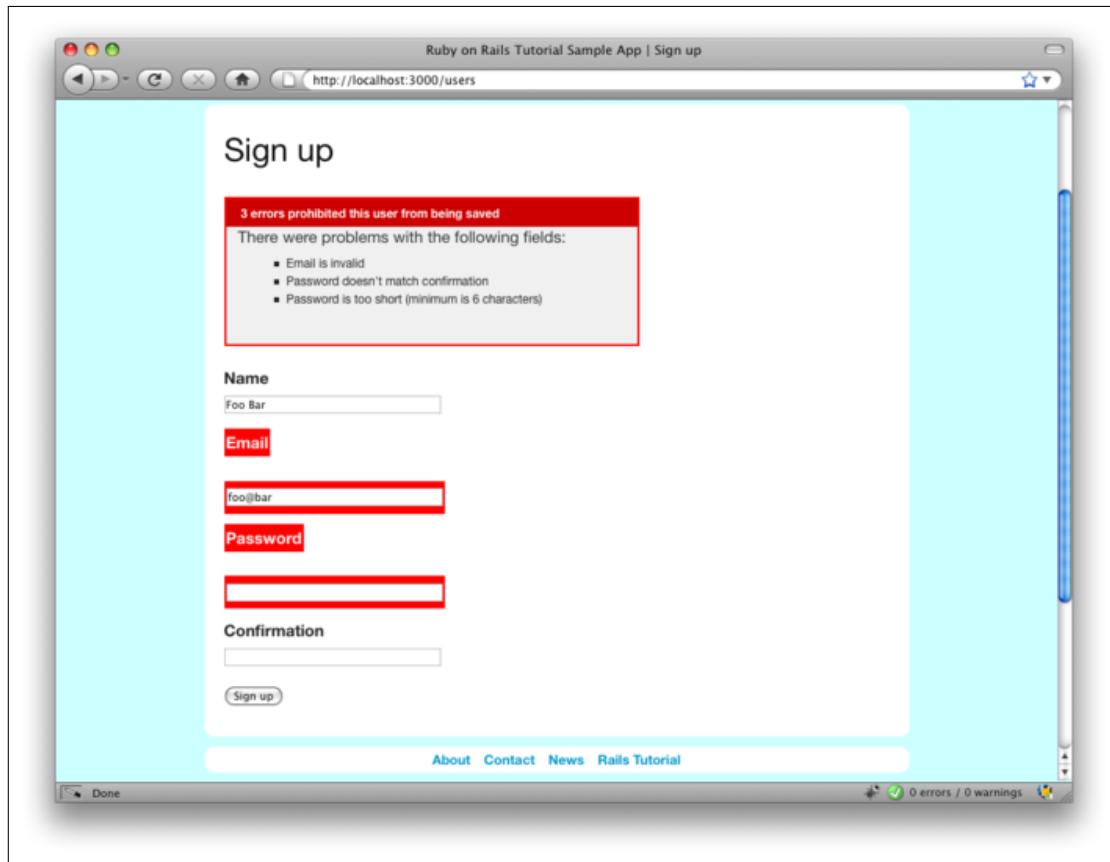


Figure 8.12: A failed signup form submission with the password field cleared. ([full size](#))

```
it "should have a password field"  
  
it "should have a password confirmation field"  
end  
.  
. .  
end
```

Listing 8.24. The `flash` ERb in the site layout using `content_tag`.
`app/views/layouts/application.html.erb`

```
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"  
"http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">  
<html>
```

```
•  
•  
•  


<% flash.each do |key, value| %>  
    <%= content_tag(:div, value, :class => "flash #{key}") %>  
  <% end %>  
  <%= yield %>  
</div>  
•  
•  
•  
</html>


```


Chapter 9

Sign in, sign out

Now that new users can sign up for our site ([Chapter 8](#)), it's time to give registered users the ability to sign in and sign out. This will allow us to add customizations based on signin status and depending on the identity of the current user. For example, in this chapter we'll update the site header with signin/signout links and a profile link; in [Chapter 11](#), we'll use the identity of a signed-in user to create microposts associated with that user, and in [Chapter 12](#) we'll allow the current user to follow other users of the application (thereby receiving a feed of their microposts).

Having users sign in will also allow us to implement a security model, restricting access to particular pages based on the identity of the signed-in user. For instance, as we'll see in [Chapter 10](#), only signed-in users will be able to access the page used to edit user information. The signin system will also make possible special privileges for administrative users, such as the ability (also in [Chapter 10](#)) to delete users from the database.

As in previous chapters, we'll do our work on a topic branch and merge in the changes at the end:

```
$ git checkout -b sign-in-out
```

9.1 Sessions

A [*session*](#) is a semi-permanent connection between two computers, such as a client computer running a web browser and a server running Rails. There are several different models for session behavior common on the web: “forgetting” the session on browser close, using an optional “remember me” checkbox for persistent sessions, and remembering sessions until the user explicitly signs out.¹ We'll opt for the final of these options: when users sign in, we will remember their signin status “forever”,² clearing the session only when the user explicitly signs out.

It's convenient to model sessions as a RESTful resource: we'll have a signin page for [**new**](#) sessions, signing in will [**create**](#) a session, and signing out will [**destroy**](#) it. We will therefore need a Sessions controller with [**new**](#), [**create**](#), and [**destroy**](#) actions. Unlike the case of the Users controller, which uses a database back-end (via the User model) to persist data, the Sessions controller will use a [*cookie*](#), which is a small piece of text placed on the user's browser. Much of the work involved in signin comes from

¹Another common model is to expire the session after a certain amount of time. This is especially appropriate on sites containing sensitive information, such as banking and financial trading accounts.

²We'll see in [Section 9.3.2](#) just how long “forever” is.

building this cookie-based authentication machinery. In this section and the next, we'll prepare for this work by constructing a Sessions controller, a signin form, and the relevant controller actions. (Much of this work parallels user signup from [Chapter 8](#).) We'll then complete user signin with the necessary cookie-manipulation code in [Section 9.3](#).

9.1.1 Sessions controller

The elements of signing in and out correspond to particular REST actions of the Sessions controller: the signin form is handled by the `new` action (covered in his section), actually signing in is handled by sending a POST request to the `create` action ([Section 9.2](#) and [Section 9.3](#)), and signing out is handled by sending a DELETE request to the `destroy` action ([Section 9.4](#)). (Recall the association of HTTP verbs with REST actions from [Table 6.2](#).) Since we know that we'll need a `new` action, we can create it when we generate the Sessions controller (just as with the Users controller in [Listing 5.20](#)):³

```
$ script/generate rspec_controller Sessions new
$ rm -rf spec/views
```

Now, as with the signup form in [Section 8.1](#), we create a new file for the Sessions controller specs and add a couple tests for the `new` action and corresponding view ([Listing 9.1](#)). (This pattern should start to look familiar by now.)

Listing 9.1. Tests for the `new` session action and view.

`spec/controllers/sessions_controller_spec.rb`

```
require 'spec_helper'

describe SessionsController do
  integrate_views

  describe "GET 'new'" do
    it "should be successful" do
      get :new
      response.should be_success
    end

    it "should have the right title" do
      get :new
      response.should have_tag("title", /sign in/i)
    end
  end
end
```

As in the case of `/users/new` from [Section 5.3.1](#), the first test already passes because `/sessions/new` matches the default `/controller/action/id` pattern (with `nil` id), but it's a good practice to add a

³If given the `create` and `destroy` actions as well, the generate script would make `views` for those actions, which we don't need. Of course, we could delete the views, but I've elected to omit them from `generate` and instead define the actions by hand.

HTTP request	URL	Named route	Action	Purpose
GET	/signin	<code>signin_path</code>	<code>new</code>	page for a new session (signin)
POST	/sessions	<code>sessions_path</code>	<code>create</code>	create a new session
DELETE	/signout	<code>signout_path</code>	<code>destroy</code>	delete a session (sign out)

Table 9.1: RESTful routes provided by the sessions rules in Listing 9.2.

RESTful route as well.⁴ While we’re at it, we’ll create all the actions needed throughout the chapter as well. We generally follow the example from Listing 6.25, but in this case we define only the particular actions we need, i.e., `new`, `create`, and `destroy`, and also add named routes for signin and signout (Listing 9.2).

Listing 9.2. Adding a resource to get the standard RESTful actions for sessions.

`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  map.resources :users
  map.resources :sessions, :only => [:new, :create, :destroy]
  map.signin '/signin', :controller => 'sessions', :action => 'new'
  map.signout '/signout', :controller => 'sessions', :action => 'destroy'
  .
  .
  .
end
```

As you can see, the `map.resources` method can take an options hash, which in this case has key `:only` and value equal to an array of the actions the Sessions controller has to respond to. The resources defined in Listing 9.2 provide URLs and actions similar to those for users (Table 6.2), as shown in Table 9.1.

We can get the second test in Listing 9.1 to pass by adding the proper title instance variable to the `new` action, as shown in Listing 9.3 (which also defines the `create` and `destroy` actions for future reference).

Listing 9.3. Adding the title for the signin page.

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController

  def new
    @title = "Sign in"
  end

  def create
  end

  def destroy
  end
end
```

⁴In Rails 3, the `/controller/action/id` pattern is disabled by default, so the RESTful route will be necessary in this case.

Sign up now!'."/>

The form mockup shows a top navigation bar with 'Home', 'Help', and 'Sign in' buttons. The main content area is titled 'Sign in'. It contains two input fields labeled 'Email' and 'Password', each with a corresponding text input box. Below the inputs is a 'Sign in' button. At the bottom left is a link 'New user? [Sign up now!](#)'.

Figure 9.1: A mockup of the signin form. ([full size](#))

With that, we're ready to make the actual signin form.

9.1.2 Signin form

The signin form (or, equivalently, the new session form) is similar in appearance to the signup form, except with two fields (email and password) in place of four. A mockup appears in Figure 9.1.

Recall from Listing 8.2 that the signup form uses the `form_for` helper, taking as an argument the user instance variable `@user`:

```
<% form_for (@user) do |f| %>
  .
  .
  .
<% end %>
```

The main difference between this and the new session form is that we have no Session model, and hence no analogue for the `@user` variable. This means that, in constructing the new session form, we have to give `form_for` slightly more information; in particular, whereas

```
form_for(@user)
```

allows Rails to infer that the `action` of the form should be to POST to the URL `/users`, in the case of sessions we need to indicate both the *name* of the resource and the appropriate URL:

```
form_for(:session, :url => sessions_path)
```

Since we're authenticating users with email address and password, we need a field for each one inside the form; the result appears in Listing 9.4.

Listing 9.4. Code for the signin form.

`app/views/sessions/new.html.erb`

```
<h2>Sign in</h2>

<% form_for(:session, :url => sessions_path) do |f| %>
  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :password %><br />
    <%= f.password_field :password %>
  </div>
  <div class="actions">
    <%= f.submit "Sign in" %>
  </div>
<% end %>

<p>New user? <%= link_to "Sign up now!", signup_path %></p>
```

With the code in Listing 9.4, the signin form appears as in Figure 9.2.

Though you'll soon get out of the habit of looking at the HTML generated by Rails (instead trusting the helpers to do their job), for now let's take a look at it (Listing 9.5).

Listing 9.5. HTML for the signin form produced by Listing 9.4.

```
<form action="/sessions" method="post">
  <div class="field">
    <label for="session_email">Email</label><br />
    <input id="session_email" name="session[email]" size="30" type="text" />
```

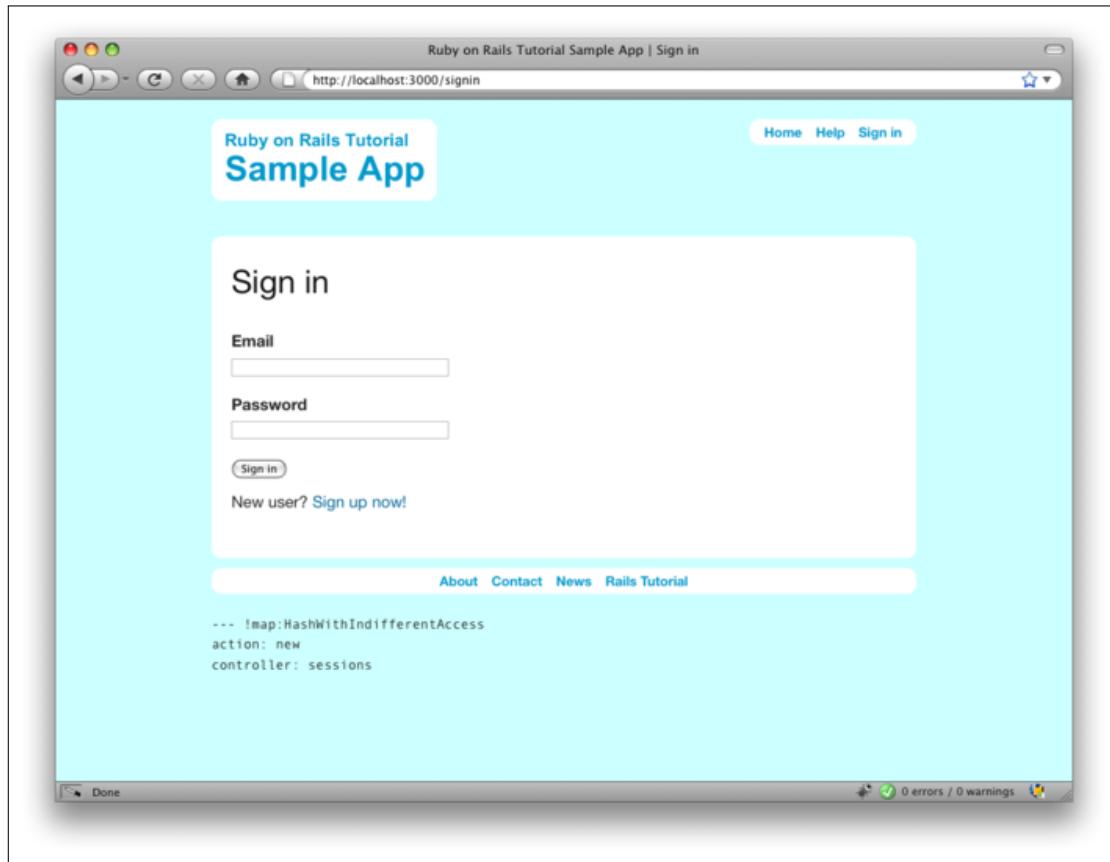


Figure 9.2: The signin form ([/sessions/new](#)). (full size)

```
</div>
<div class="field">
  <label for="session_password">Password</label><br />
  <input id="session_password" name="session[password]" size="30"
         type="password" />
</div>
<div class="actions">
  <input id="session_submit" name="commit" type="submit" value="Sign in" />
</div>
</form>
```

Comparing Listing 9.5 with Listing 8.5, you might be able to guess that submitting this form will result in a `params` hash where `params[:session][:email]` and `params[:session][:password]` correspond to the email and password fields. Handling this submission—and, in particular, authenticating users based on the submitted email and password—is the goal of the next two sections.

9.2 Signin failure

As in the case of creating users (`signup`), the first step in creating sessions (`signin`) is to handle *invalid* input. We'll start by reviewing what happens when a form gets submitted, and then arrange for helpful error messages to appear in the case of signin failure (as mocked up in Figure 9.3.) Finally, we'll lay the foundation for successful signin (Section 9.3) by evaluating each signin submission based on the validity of its email/password combination.

9.2.1 Reviewing form submission

Let's start by defining a minimalist `create` action for the Sessions controller (Listing 9.6), which does nothing but render the `new` view. Submitting the `/sessions/new` form with blank fields then yields the result shown in Figure 9.4.

Listing 9.6. A preliminary version of the Sessions `create` action.

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController
  .
  .
  .
  def create
    render 'new'
  end
  .
  .
  .
end
```

Carefully inspecting the debug information in Figure 9.4 shows that, as hinted at the end of Section 9.1.2, the submission results in a `params` hash containing the email and password under the key `:session`:

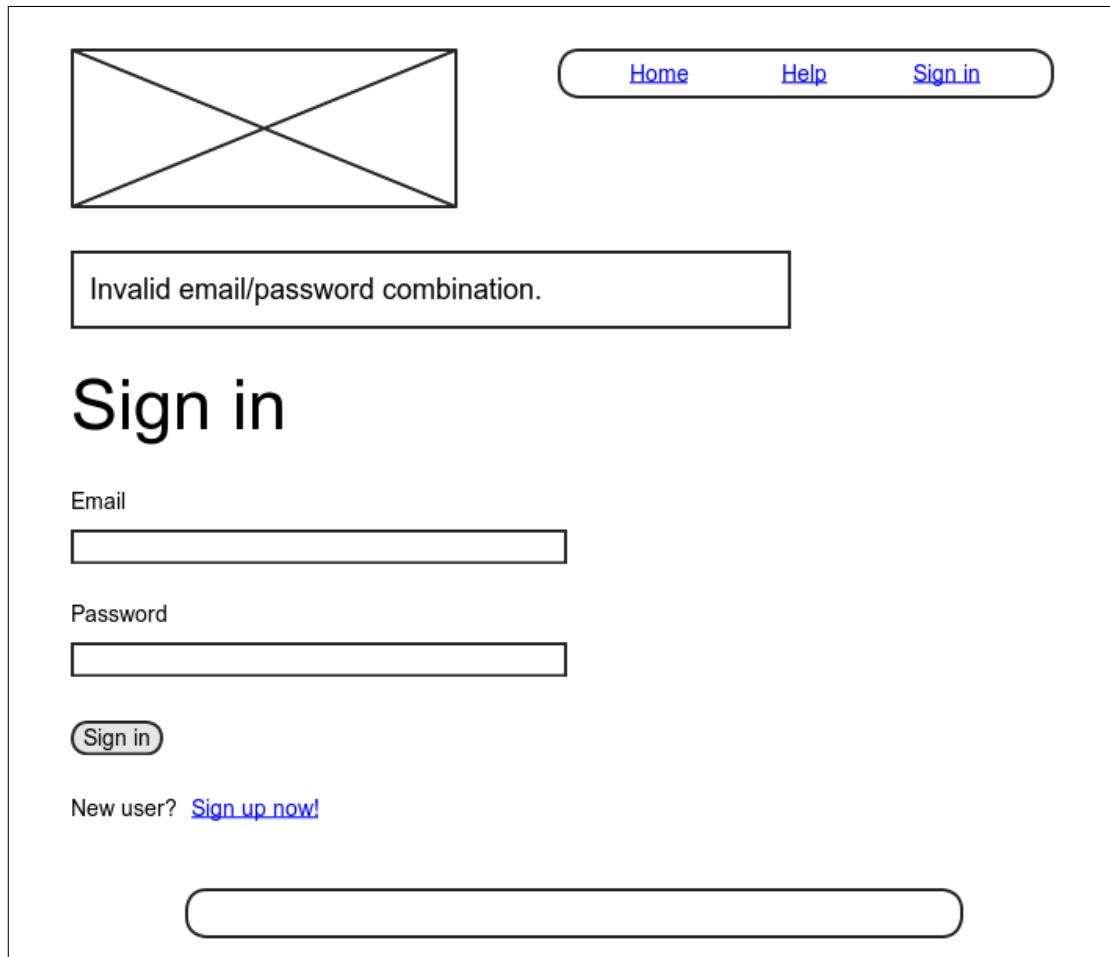


Figure 9.3: A mockup of signin failure. ([full size](#))

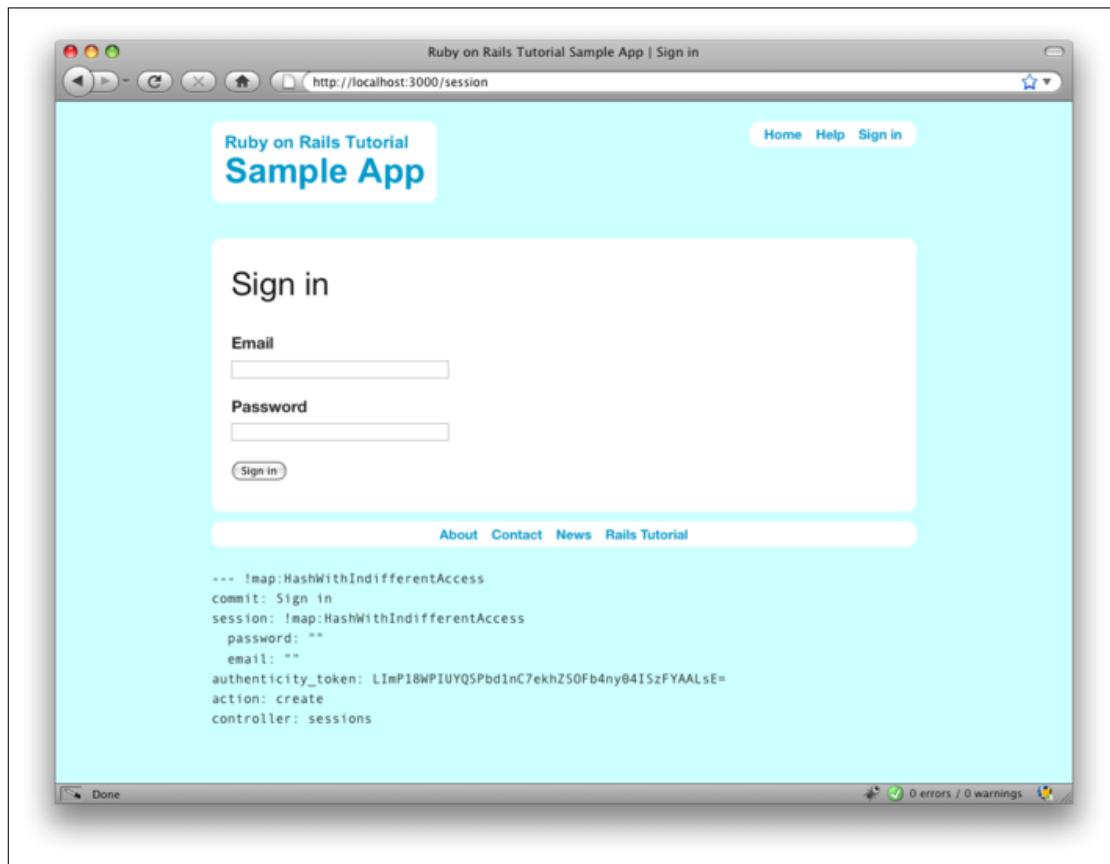


Figure 9.4: The initial failed signin, with `create` as in Listing 9.6. (full size)

```
--- !map:HashWithIndifferentAccess
commit: Sign in
session: !map:HashWithIndifferentAccess
  password: ""
  email: ""
authenticity_token: LImP18WPIUYQSPbd1nC7ekhZSOFb4ny04ISzFYAALsE=
action: create
controller: sessions
```

This means that

```
params[:session][:email]
```

contains the submitted email address, and

```
params[:session][:password]
```

has the password.

In other words, inside the `create` action the `params` hash has all the information needed to authenticate users by email and password. Not coincidentally, we have already developed exactly the method needed: `User.authenticate` from Section 7.2.4 (Listing 7.12). Recalling that `authenticate` returns `nil` for an invalid authentication, our strategy for user signin can be summarized as follows:

```
def create
  user = User.authenticate(params[:session][:email],
                          params[:session][:password])
  if user.nil?
    # Create an error message and re-render the signin form.
  else
    # Sign the user in and redirect to the user's show page.
  end
end
```

9.2.2 Failed signin (test and code)

In order to handle a failed signin attempt, first we need to determine that it's a failure. We thus have an expectation that the `User.authenticate` method will be called in the `create` action; we can express this in test code using the same type of message expectation first seen in Listing 8.6, where we indicated the expectation that `@user.save` would be called:

```
@user.should_receive(:save).and_return(false)
```

To describe invalid signin, we'll use a similar construction, in this case expressing the expectation that the `User` model will receive the `authenticate` method, as shown in Listing 9.7. (As with the user signup specs, Listing 9.7 doesn't actually test that the form has right fields for signin or that form submissions actually work. We'll test these aspects of signin with integration tests in Section 9.4.4.)

Listing 9.7. Tests for a failed signin attempt.

`spec/controllers/sessions_controller_spec.rb`

```
require 'spec_helper'

describe SessionsController do
  integrate_views
  .

  .

  .

  describe "POST 'create'" do
    describe "invalid signin" do
      before(:each) do
        @attr = { :email => "email@example.com", :password => "password" }
        User.should_receive(:authenticate).
          with(@attr[:email], @attr[:password]).
          and_return(nil)
      end

      it "should re-render the new page" do
        post :create, :session => @attr
        response.should render_template('new')
      end

      it "should have the right title" do
        post :create, :session => @attr
        response.should have_tag("title", /sign in/i)
      end
    end
  end
end
```

Here the tests follow the example from the analogous tests for user signup (Listing 8.6). The main difference is in the message expectation, which adds the `with` method:⁵

```
User.should_receive(:authenticate).
  with(@attr[:email], @attr[:password]).
  and_return(nil)
```

In English, this says that the `User` model should receive the `authenticate` method with the email and password attributes from the form submission. Since this is an invalid signin, we arrange for the return value

⁵See the RSpec documentation on message expectations for more about `with`.

be `nil`. (Note that we *don't* have to check that the email/password combination is actually invalid, since that logic is covered by the User model tests in Listing 7.11. Recall the message expectation subtlety from Box 8.1: `and_return(nil)` isn't an expectation; it's a guarantee.)

The application code needed to get these tests to pass appears in Listing 9.8. As promised in Section 9.2.1, we extract the submitted email address and password from the `params` hash, and then pass them to the `User.authenticate` method, thereby fulfilling the message expectation from the test. If the user is not authenticated (i.e., if it's `nil`), we set the title and re-render the signin form.⁶ We'll handle the other branch of the if-else statement in Section 9.3; for now we'll just leave a descriptive comment.

Listing 9.8. Code for a failed signin attempt.

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController
  .
  .
  .
  def create
    user = User.authenticate(params[:session][:email],
                           params[:session][:password])
    if user.nil?
      flash.now[:error] = "Invalid email/password combination."
      @title = "Sign in"
      render 'new'
    else
      # Sign the user in and redirect to the user's show page.
    end
  end
end
```

Recall from Section 8.4.2 that we displayed signup errors using the User model error messages. Since the session isn't an Active Record model, this strategy won't work here, so instead we've put a message in the flash (or, rather, in `flash.now`,⁷ see Box 9.1). Thanks to the flash message display in the site layout (Listing 8.15), the `flash[:error]` message automatically gets displayed; thanks to the Blueprint CSS, it automatically gets nice styling (Figure 9.5).

Box 9.1. Flash dot now

There's a subtle difference between `flash` and `flash.now`. The `flash` variable is designed to be used before a *redirect*, and it persists on the resulting page for one request—that is, it appears once, and disappears when you click on another link. Unfortunately, this means that if we *don't* redirect, and instead simply render a page (as in Listing 9.8), the flash message persists for *two* requests: it appears on the rendered page but is still waiting for a “redirect” (i.e., a second request), and thus appears *again* if you click a link.

To avoid this weird behavior, when *rendering* rather than *redirecting* we use `flash.now` instead of `flash`. The `flash.now` object is specifically designed for displaying flash messages on rendered pages.

⁶If case you're wondering why we use `user` instead of `@user` in Listing 9.8, it's because this user variable is never needed in any view, so there is no reason to use an instance variable here. (Using `@user` still works, though.)

⁷We don't have a controller test for `flash.now` because Rails 2.3 broke `flash.now` support in tests (it worked in Rails 2.2). It's not a big deal; we'll cover it as part of the integration tests in Section 9.4.4.

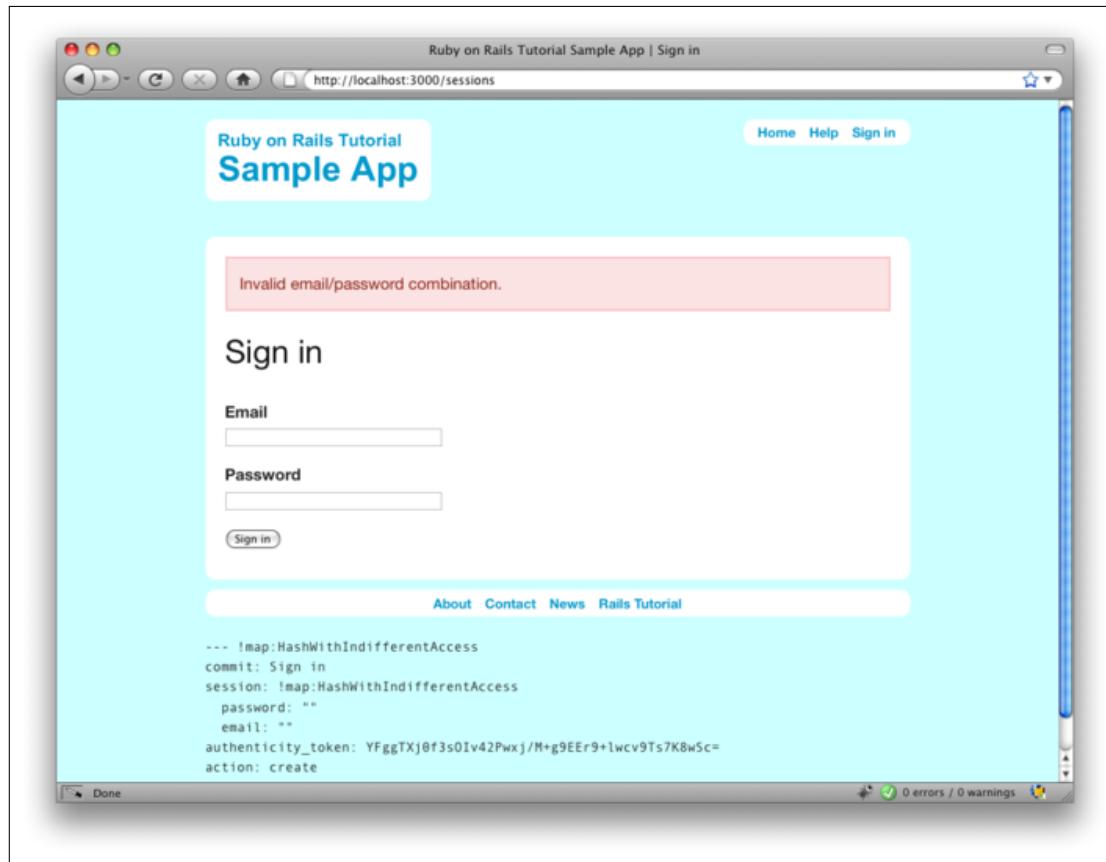


Figure 9.5: A failed signin (with a flash message). ([full size](#))

If you ever find yourself wondering why a flash message is showing up where you don't expect it, the chances are good that you need to replace `flash` with `flash.now`.

9.3 Signin success

Having handled a failed signin, we now need to actually sign a user in. A hint of where we're going—the user profile page, with modified navigation links—is mocked up in [Figure 9.6](#).⁸ Getting there will require some of the most challenging Ruby programming so far in this tutorial, so hang in there through the end and be prepared for a little heavy lifting. Happily, the first step is easy—completing the Sessions controller `create` action is a snap. Unfortunately, it's also a cheat.

⁸Image from <http://www.flickr.com/photos/hermanusbackpackers/3343254977/>.

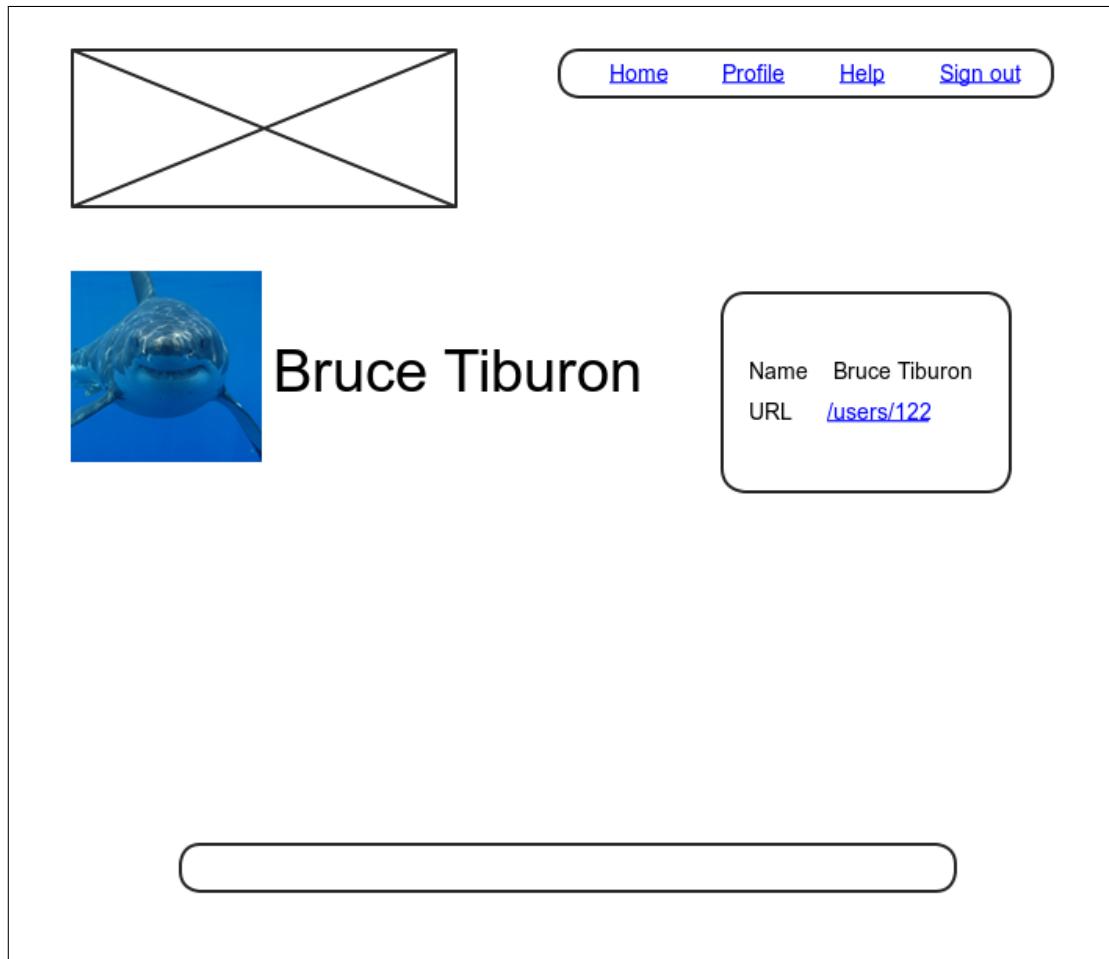


Figure 9.6: A mockup of the user profile after a successful signin (with updated nav links). ([full size](#))

9.3.1 The completed `create` action

Filling in the area now occupied by the signin comment ([Listing 9.8](#)) is simple: upon successful signin, we sign the user in using the `sign_in` function, and then redirect to the profile page ([Listing 9.9](#)). We see now why this is a cheat: alas, `sign_in` doesn't currently exist. Writing it will occupy the rest of this section.

Listing 9.9. The completed Sessions controller `create` action (not yet working).

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController
  .
  .
  .
  def create
    user = User.authenticate(params[:session][:email],
                           params[:session][:password])
    if user.nil?
      flash.now[:error] = "Invalid email/password combination."
      @title = "Sign in"
      render 'new'
    else
      sign_in user
      redirect_to user
    end
  end
end
```

Even though we lack the `sign_in` function, we can still write the tests ([Listing 9.10](#)). (We'll fill in the body of the first test in [Section 9.3.4](#).)

Listing 9.10. Pending tests for user signin (to be completed in [Section 9.3.4](#)).

`spec/controllers/sessions_controller_spec.rb`

```
describe SessionsController do
  .
  .
  .
  describe "POST 'create'" do
    .
    .
    .
    describe "with valid email and password" do
      before(:each) do
        @user = Factory(:user)
        @attr = { :email => @user.email, :password => @user.password }
        User.should_receive(:authenticate).
          with(@user.email, @user.password).
          and_return(@user)
      end
    end
  end
end
```

```

it "should sign the user in" do
  post :create, :session => @attr
  # Fill in with tests for a signed-in user.
end

it "should redirect to the user show page" do
  post :create, :session => @attr
  redirect_to user_path(@user)
end
end
end
end

```

In analogy with the message expectation in Listing 9.7, Listing 9.10 sets up the expectation that `User.authenticate` gets called with the right email/password combination, in this case (for a valid combination) returning the factory user:

```

@user = Factory(:user)
@attr = { :email => @user.email, :password => @user.password }
User.should_receive(:authenticate).
  with(@user.email, @user.password).
  and_return(@user)

```

9.3.2 Remember me

There are many moving parts in the signin machinery. Our strategy is to push a few elements onto the `stack`, and then pop them off as the section unfolds. Have patience if you get bogged down; you'll probably have to get all the way through before you see how it all fits together.

The signin functions will end up crossing the traditional Model-View-Controller lines; in particular, several signin functions will need to be available in both controllers and views. You may recall from Section 4.2.5 that Ruby provides a `module` facility for packaging functions together and including them in multiple places, and that's the plan for the authentication functions. We could make an entirely new module for authentication, but the Sessions controller already comes equipped with a module, namely, `SessionsHelper`. Moreover, helpers are automatically included in Rails views, so all we need to do to use the Sessions helper functions in controllers is to include the module into the Application controller (Listing 9.11).

Listing 9.11. Including the Sessions helper module into the Application controller.

`app/controllers/application_controller.rb`

```

class ApplicationController < ActionController::Base
  helper :all # include all helpers, all the time
  include SessionsHelper
  .
  .
  .
end

```

Here the line `helper :all` includes the helpers in all the `views`, but not in the controllers. We need the methods from the Sessions helper in both places, so we have to include it explicitly.

Now we're ready for the first signin element, the `sign_in` function itself. Our authentication method is to place a *remember token* as a cookie on the user's browser, and then use the token to find the user record in the database as the user moves from page to page (implemented in Section 9.3.4). The result, Listing 9.12, pushes three things onto the stack: the `remember_me!` method, the `cookies` hash, and `current_user`. Let's start popping them off.

Listing 9.12. The complete (but not-yet-working) `sign_in` function.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper

  def sign_in(user)
    user.remember_me!
    cookies[:remember_token] = { :value => user.remember_token,
                                 :expires => 20.years.from_now.utc }
    self.current_user = user
  end
end
```

We begin with the `remember_me!` method. This method's job is to save the user's remember token to the database. We can outline the expected behavior using some tests, which, because `remember_me!` acts on users, are tests on the User model (Listing 9.13).

Listing 9.13. Tests for the `remember_me!` method.

`spec/models/user_spec.rb`

```
describe User do
  .
  .
  .
  describe "remember me" do
    before(:each) do
      @user = User.create! (@attr)
    end

    it "should have a remember_me! method" do
      @user.should respond_to(:remember_me!)
    end

    it "should have a remember token" do
      @user.should respond_to(:remember_token)
    end

    it "should set the remember token" do
      @user.remember_me!
      @user.remember_token.should_not be_nil
    end
  end
end
```

```
    end
  end
end
```

The `remember_me!` method itself then appears as in Listing 9.14 (with the `has_password?` method shown to give you some context).

Listing 9.14. The `remember_me!` method.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
  .
  before_save :encrypt_password
  .

  .
  .

  def has_password?(submitted_password)
    encrypted_password == encrypt(submitted_password)
  end

  def remember_me!
    self.remember_token = encrypt("#{salt}--#{id}")
    save_without_validation
  end

  .
  .
  .

  private

  def encrypt_password
    unless password.nil?
      self.salt = make_salt
      self.encrypted_password = encrypt(password)
    end
  end

  .
  .
  .

end
```

Here we set the remember token for each user using the `encrypt` function from Section 7.2.3. Since this token will be placed on the user's browser, it needs to be both *unique* and *secure*. We accomplish this by encrypting the `salt` with the `id`, thereby guaranteeing that the token is different for each user and is impervious to a dictionary attack (Section 7.2.2).

In order to establish the connection between the remember token and the user, we'll store the remember token in the database and at some point (Listing 9.19) use the `find_by_remember_token` method (created automatically by Active Record) to retrieve the user. This means that we must *save* the user after setting the

remember token, but recall that users have a virtual `password` attribute which is required to be present (Section 7.1.1). When calling the `remember_me!` method, there is no password, so the validations will fail if we simply use `save`. We bypass this in Listing 9.14 using the special `save_without_validation` method, which saves the user without running the validations.

Finally, note that we've updated the `encrypt_password` method from Chapter 6. Even though we skip the validation step using `save_without_validation`, the `before_save` callback still fires; with the original code from Listing 7.10,

```
def encrypt_password
  self.salt = make_salt
  self.encrypted_password = encrypt(password)
end
```

when `password` is `nil` the user's encrypted password would be reset to `encrypt(nil)`, thus preventing the user from ever signing in again. (Section 9.6 has an exercise to write an integration test that catches this problem.) To prevent this issue, we skip the encryption step when `password` is `nil`:⁹

```
def encrypt_password
  unless password.nil?
    self.salt = make_salt
    self.encrypted_password = encrypt(password)
  end
end
```

With the code in Listing 9.14, we've pushed the `remember_token` onto the stack. Since it needs to be stored in the database, it should be an attribute on the User model; we can add it with a migration:

```
$ script/generate migration add_remember_token_to_users remember_token:string
```

This generates a nearly complete migration; the only extra step (shown in Listing 9.15) is to add an index on the `remember_token` column. (Recall from Box 6.1 that an index is important to avoid a full-table scan when finding by a particular attribute.)

Listing 9.15. The migration for the `remember_token` attribute.

```
db/migrations/<timestampl>.add_remember_token_to_users.rb
```

```
class AddRememberTokenToUsers < ActiveRecord::Migration
  def self.up
    add_column :users, :remember_token, :string
    add_index   :users, :remember_token
  end
```

⁹There's a second way to do this, by passing an `:if` parameter to the `before_save` filter. The syntax is a `smell`, though, and I expect it to change in future versions of Active Record.

users	
id	integer
name	string
email	string
encrypted_password	string
salt	string
remember_token	string
created_at	datetime
updated_at	datetime

Figure 9.7: The User model with an added remember token.

```
def self.down
  remove_column :users, :remember_token
end
end
```

Then run the migration as usual.

```
$ rake db:migrate
```

The resulting data model is shown in [Figure 9.7](#).

Since we've changed the data model by adding a new column, we also need to prepare the test database:

```
$ rake db:test:prepare
```

At this point, all the tests in [Listing 9.13](#) should pass.

9.3.3 Cookies

Let's look again at the `sign_in` function to see where we are:

```
module SessionsHelper

  def sign_in(user)
    user.remember_me!
    cookies[:remember_token] = { :value => user.remember_token,
                                 :expires => 20.years.from_now.utc }
```

```
    self.current_user = user
  end
end
```

The whole point of the previous section was to get the first line to work by adding the `remember_me!` method to the User model. Now that we've done that, the cookie line will work as well, since after the call to `remember_me!` the user has valid a remember token.

As you can see, the `cookies` utility can be used like a hash; in this case, we define a hash-of-hashes—the `cookies` assignment makes a hash with key `:remember_token` and value equal to another hash, consisting of the user's remember token that expires twenty years from now (effectively lasting forever). (See Box 9.2 for more on the `20.years.from_now` syntax.) Of course, `cookies` isn't *really* a hash, since assigning to `cookies` actually *saves* a piece of text on the browser, but part of the beauty of Rails is that it lets you forget about that detail and concentrate on writing the application.

Box 9.2. Cookies expire `20.years.from_now`

You may recall from Section 4.4.2 that Ruby lets you add methods to *any* class, even built-in ones. In that section, we added a `palindrome?` method to the `String` class (and discovered as a result that "`deified`" is a palindrome), and also saw how Rails adds a `blank?` method to class `Object` (so that "`".blank?,` " `".blank?,` and `nil.blank?` are all `true`). The cookie code in Listing 9.12 shows yet another example of this practice, through one of Rails' *time helpers*, which are methods added to `Fixnum` (the base class for numbers):

```
$ script/console
>> 1.year.from_now
=> Sun, 13 Mar 2011 03:38:55 UTC +00:00
>> 10.weeks.ago
=> Sat, 02 Jan 2010 03:39:14 UTC +00:00
```

Rails adds other helpers, too:

```
>> 1.kilobyte
=> 1024
>> 5.megabytes
=> 5242880
```

These are useful for upload validations, making it easy to restrict, say, image uploads to `5.megabytes`.

Though it must be used with caution, the flexibility to add methods to built-in classes allows for extraordinarily natural additions to plain Ruby. Indeed, much of the elegance of Rails ultimately derives from the malleability of the underlying Ruby language.

Now that we know how to manipulate cookies, there's only one piece left: it's time to pop `current_user` off the stack.

9.3.4 Current user

Let's look at the `sign_in` function one last time:

```
module SessionsHelper

  def sign_in(user)
    user.remember_me!
    cookies[:remember_token] = { :value => user.remember_token,
                                 :expires => 20.years.from_now.utc }
    self.current_user = user
  end
end
```

Our focus now is the final line:

```
self.current_user = user
```

The purpose of this line is to create `current_user`, accessible in both controllers and views, which will allow constructions such as

```
<%= current_user.name %>
```

and

```
redirect_to current_user
```

In the context of the Sessions helper, the `self` in `self.current_user` is the Sessions *controller*, since the module is being included into the Application controller, which is the base class for all the other controllers (including Sessions). But there is no such variable as `current_user` right now; we need to define it.

To describe the behavior of the remaining signin machinery, we'll first fill in the test for signing a user in (Listing 9.16).

Listing 9.16. Filling in the test for signing the user in.

```
spec/controllers/sessions_controller_spec.rb
```

```
describe SessionsController do
  .
  .
  .
  describe "POST 'create'" do
    .
    .
    .
    describe "with valid email and password" do
```

```
before(:each) do
  @user = Factory(:user)
  @attr = { :email => @user.email, :password => @user.password }
  User.should_receive(:authenticate).
    with(@user.email, @user.password).
    and_return(@user)
end

it "should sign the user in" do
  post :create, :session => @attr
  controller.current_user.should == @user
  controller.should be_signed_in
end

it "should redirect to the user show page" do
  post :create, :session => @attr
  redirect_to user_path(@user)
end
end
end
end
```

The new test uses the `controller` variable (which is available inside Rails tests) to check that the `current_user` variable is set to the signed-in user, and that the user is signed in:

```
it "should sign the user in" do
  post :create, :session => @attr
  controller.current_user.should == @user
  controller.should be_signed_in
end
```

The second line may be a little confusing at this point, but you can guess based on the RSpec convention for boolean methods that

```
controller.should be_signed_in
```

is equivalent to

```
controller.signed_in?.should be_true
```

This is a hint that we will be defining a `signed_in?` method that returns `true` if a user is signed in and `false` otherwise. Moreover, the `signed_in?` method will be attached to the `controller`, not to a user, which is why we write `controller.signed_in?` instead of `current_user.signed_in?`. (If no user is signed in, how could we call `signed_in?` on it?)

To start writing the code for `current_user`, note that the line

```
self.current_user = user
```

is an *assignment*. Ruby has a special syntax for defining such an assignment function, shown in Listing 9.17.

Listing 9.17. Defining assignment to `current_user`.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper

  def sign_in(user)
    .
    .
    .
  end

  def current_user=(user)
    @current_user = user
  end
end
```

This might look confusing, but it simply defines a method `current_user=` expressly designed to handle assignment to `current_user`. Its one argument is the right-hand side of the assignment, in this case the user to be signed in. The one-line method body just sets an instance variable `@current_user`, effectively storing the user for later use.

In ordinary Ruby, we could define a second method, `current_user`, designed to return the value of `@current_user` (Listing 9.18).

Listing 9.18. A tempting but useless definition for `current_user`.

```
module SessionsHelper

  def sign_in(user)
    .
    .
    .
  end

  def current_user=(user)
    @current_user = user
  end

  def current_user
    @current_user      # Useless! Don't use this code.
  end
end
```

If we did this, we would effectively replicate the functionality of `attr_accessor`, first seen in Section 4.4.5 and used to make the virtual `password` attribute in Section 7.1.1.¹⁰ The problem is that it utterly fails to solve our problem: with the code in Listing 9.18, the user’s signin status would be forgotten: as soon as the user went to another page—poof!—the session would end and the user would be automatically signed out.

Luckily, thanks to our work with the remember token in Section 9.3.3, we can find the user in the database using the `cookies` hash and the `find_by_remember_token` (Listing 9.19).¹¹

Listing 9.19. Finding the current user by `remember_token`.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper
  .
  .
  .
  def current_user
    @current_user ||= user_from_remember_token
  end

  def user_from_remember_token
    remember_token = cookies[:remember_token]
    User.find_by_remember_token(remember_token) unless remember_token.nil?
  end
end
```

This code uses the common but initially obscure `||=` (“or equals”) assignment operator (Box 9.3). Its effect is to set the `@current_user` instance variable to the user corresponding to the remember token, but only if `@current_user` is undefined. In other words, the construction

```
@current_user ||= user_from_remember_token
```

calls the `user_from_remember_token` method the first time `current_user` is called, but on subsequent invocations returns `@current_user` without calling `user_from_remember_token`.¹²

Box 9.3. What the *\$@! is `||=` ?

The `||=` construction is very Rubyish—that is, it is highly characteristic of the Ruby language—and hence important to learn if you plan on doing much Ruby programming. Though at first it may seem mysterious, *or equals* is easy to understand by analogy.

We start by noting a common idiom for changing a currently defined variable. Many computer programs involve incrementing a variable, as in

```
x = x + 1
```

¹⁰In fact, the two are exactly equivalent; `attr_accessor` is merely a convenient way to create just such getter/setter methods automatically.

¹¹Recall from Section 6.1.4 that Active Record synthesizes find methods for all the user attributes, including `remember_token`.

¹²This optimization technique to avoid repeated function calls is known as `memoization`.

Most languages provide a syntactic shortcut for this operation; in Ruby (and in C, C++, Perl, Python, Java, etc.) it appears as follows:

```
x += 1
```

Analogous constructs exist for other operators as well:

```
$ script/console
>> x = 1
=> 1
>> x += 1
=> 2
>> x *= 3
=> 6
>> x -= 7
=> -1
```

In each case, the pattern is that `x = x O y` and `x O= y` are equivalent for any operator `O`.

Another common Ruby pattern is assigning to a variable if it's `nil` but otherwise leaving it alone. Recalling the *or* operator `||` seen in [Section 4.2.3](#), we can write this as follows:

```
>> @user
=> nil
>> @user = @user || "the user"
=> "the user"
>> @user = @user || "the user"
=> "the user"
```

Since `nil` is false in a boolean context, the first assignment is `nil || "the user"`, which evaluates to `"the user"`; similarly, the second assignment is `"the user" || "the user"`, which also evaluates to `"the user"`—since strings are `true` in a boolean context, the series of `||` expressions terminates after the first expression is evaluated. (This practice of evaluating `||` expressions from left to right and stopping on the first true value is known as *short-circuit evaluation*.)

Comparing the console sessions for the various operators, we see that `@user = @user || value` follows the `x = x O y` pattern with `||` in the place of `O`, which suggests the following equivalent construction:

```
>> @user ||= "the user"
=> "the user"
```

Voilà!

At this point, the signin test is almost passing; the only thing remaining is to define the required `signed_in?` boolean method. Happily, it's easy with the use of the “not” operator `!`: a user is signed in if `current_user` is not `nil` ([Listing 9.20](#)).

Listing 9.20. The `signed_in?` helper method.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper
  .
  .
  .
  def signed_in?
    !current_user.nil?
  end
end
```

Though it's already useful for the test, we'll put the `signed_in?` method to even better use in [Section 9.4.3](#) and again in [Chapter 10](#).

With that, all the tests should pass.

9.4 Signing out

As discussed in [Section 9.1](#), our authentication model is to keep users signed in until they sign out explicitly. In this section we'll add this necessary signout capability. Once we're done, we'll add some integration tests to put our authentication machinery through its paces.

9.4.1 Destroying sessions

So far, the Sessions controller actions have followed the RESTful convention of using `new` for a signin page and `create` to complete the signin. We'll continue this theme by using a `destroy` action to delete sessions, i.e., to sign out.

In order to test the signout action, we first need a way to sign in within a test. The easiest way to do this is to use the `controller` object we saw in [Section 9.3.4](#) and set its `current_user` to the given user. In order to use the resulting `test_sign_in` function in all our tests, we need to put it in the spec helper file, as shown in [Listing 9.21](#).

Listing 9.21. A `test_sign_in` function to simulate user signin inside tests.

`spec/spec_helper.rb`

```
Spec::Runner.configure do |config|
  .
  .
  .
  def test_sign_in(user)
    controller.current_user = user
  end
end
```

After running `test_sign_in`, the `current_user` will not be `nil`, so `signed_in?` will be `true`.

With this spec helper in hand, the test for signout is straightforward: sign in as a (factory) user, check that it is signed in, and then hit the `destroy` action and verify that the user gets signed out ([Listing 9.22](#)).

Listing 9.22. A test for destroying a session (user signout).

`spec/controllers/sessions_controller_spec.rb`

```
describe SessionsController do
  .
  .
  .
  describe "DELETE 'destroy'" do
    it "should sign a user out" do
      test_sign_in(Factory(:user))
      controller.should be_signed_in
      delete :destroy
      controller.should_not be_signed_in
      response.should redirect_to(root_path)
    end
  end
end
```

The only novel element here is the `delete` method, which issues an HTTP DELETE request (in analogy with the `get` and `post` methods seen in previous tests), as required by the REST conventions (Table 9.1).

As with user signin, which relied on the `sign_in` function, user signout just defers the hard work to a `sign_out` function (Listing 9.23).

Listing 9.23. Destroying a session (user signout).

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController
  .
  .
  .
  def destroy
    sign_out
    redirect_to root_path
  end
end
```

As with the other authentication elements, we'll put `sign_out` in the Sessions helper module (Listing 9.24).

Listing 9.24. The `sign_out` method in the Sessions helper module.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper
  def sign_in(user)
    user.remember_me!
```

```
cookies[:remember_token] = { :value => user.remember_token,
                             :expires => 20.years.from_now.utc }
self.current_user = user
end
.
.
.
def sign_out
  cookies.delete(:remember_token)
  self.current_user = nil
end
end
```

As you can see, the `sign_out` method effectively undoes the `sign_in` method by deleting the remember token and by setting the current user to `nil`.¹³

9.4.2 Signin upon signup

In principle, we are now done with authentication, but as currently constructed there are no links to the signin or signout actions. Moreover, newly registered users might be confused, as they are not signed in by default.

We'll fix the second problem first, starting with testing that a new user is automatically signed in (Listing 9.25).

Listing 9.25. Testing that newly signed-up users are also signed in.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views
  .
  .
  .
  describe "POST 'create'" do
    .
    .
    .
    describe "success" do
      .
      .
      .
      it "should sign the user in" do
        post :create, :user => @attr
        controller.should be_signed_in
      end
    .
    .
    .
  end
end
```

¹³You can learn about things like `cookies.delete` by reading the [cookies entry in the Rails API](#).

```
    end
  end
end
```

With the `sign_in` method from [Section 9.3](#), getting this test to pass by actually signing in the user is easy: just add `sign_in @user` right after saving the user to the database ([Listing 9.26](#)).

Listing 9.26. Signing in the user upon signup.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
  def create
    @user = User.new(params[:user])
    if @user.save
      sign_in @user
      flash[:success] = "Welcome to the Sample App!"
      redirect_to @user
    else
      @title = "Sign up"
      render 'new'
    end
  end
end
```

9.4.3 Changing the layout links

We come finally to a practical application of all our signin/out work: we'll change the layout links based on signin status. In particular, as seen in the [Figure 9.6](#) mockup, we'll arrange for the links change when users sign in or sign out, and we'll also add a profile link to the user show page for signed-in users.

We start with two integration tests: one to check that a "`Sign in`" link appears for non-signed-in users, and one to check that a "`Sign out`" link appears for signed-in users; both cases verify that the link goes to the proper URL. We'll put these tests in the layout links test we created in [Section 5.2.1](#); the result appears in [Listing 9.27](#).

Listing 9.27. Tests for the signin/signout links on the site layout.

`spec/integration/layout_links_spec.rb`

```
describe "Layout links" do
  .
  .
  .
  describe "when not signed in" do
    it "should have a signin link" do
      visit root_path
      response.should have_tag("a[href=?]", signin_path, "Sign in")
```

```

    end
  end

  describe "when signed in" do

    before(:each) do
      @user = Factory(:user)
      visit signin_path
      fill_in :email, :with => @user.email
      fill_in :password, :with => @user.password
      click_button
    end

    it "should have a signout link" do
      visit root_path
      response.should have_tag("a[href=?]", signout_path, "Sign out")
    end

    it "should have a profile link"
  end
end

```

Here the `before (:each)` block signs in by visiting the signin page and submitting a valid email/password pair.¹⁴ We do this instead of using the `test_sign_in` function from Listing 9.21 because `test_sign_in` doesn't work inside integration tests for some reason. (See Section 9.6 for an exercise to make an `integration_sign_in` function for use in integration tests.)

The application code uses an if-then branching structure inside of Embedded Ruby, using the `signed_in?` method defined in Listing 9.20:

```

<% if signed_in? %>
<li><%= link_to "Sign out", signout_path, :method => :delete %></li>
<% else %>
<li><%= link_to "Sign in", signin_path %></li>
<% end %>

```

Notice that the signout link passes a hash argument indicating that it should submit with an HTTP DELETE request.¹⁵ With this snippet added, the full header partial appears as in Listing 9.28.

Listing 9.28. Changing the layout links for signed-in users.

`app/views/layouts/_header.html.erb`

```

<div id="header" class="round">
  <%= link_to logo, root_path %>
  <ul class="navigation round">

```

¹⁴Note that we can use symbols in place of strings for the labels, e.g., `fill_in :email` instead of `fill_in "Email"`. We used the latter in Listing 8.22, but by now it shouldn't surprise you that Rails allows us to use symbols instead.

¹⁵Web browsers can't actually issue DELETE requests; Rails fakes it with JavaScript.

```

<li><%= link_to "Home", root_path %></li>
<li><%= link_to "Help", help_path %></li>
<% if signed_in? %>
<li><%= link_to "Sign out", signout_path, :method => :delete %></li>
<% else %>
<li><%= link_to "Sign in", signin_path %></li>
<% end %>
</ul>
</div>

```

In Listing 9.28 we've used the `logo` helper from the Chapter 5 exercises (Section 5.5); in case you didn't work that exercise, the answer appears in Listing 9.29.

Listing 9.29. A helper for the site logo.

`app/helpers/application_helper.rb`

```

module ApplicationHelper
  .
  .
  .
  def logo
    image_tag("logo.png", :alt => "Sample App", :class => "round")
  end
end

```

Finally, let's add a profile link. The test (Listing 9.30) and application code (Listing 9.31) are both straightforward. Notice that the profile link's URL is simply `current_user`,¹⁶ which is our first use of that helpful method. (It won't be our last.)

Listing 9.30. A test for a profile link.

`spec/integration/layout_links_spec.rb`

```

describe "Layout links" do
  .
  .
  .
  describe "when signed in" do
    .
    .
    .
    it "should have a profile link" do
      visit root_path
      response.should have_tag("a[href=?]", user_path(@user), "Profile")
    end
  end
end

```

¹⁶Recall from Section 7.3.3 that we can link directly to a user object and allow Rails to figure out the appropriate URL.

Listing 9.31. Adding a profile link.`app/views/layouts/_header.html.erb`

```
<div id="header" class="round">
  <%= link_to logo, root_path %>
  <ul class="navigation round">
    <li><%= link_to "Home", root_path %></li>
    <% if signed_in? %>
      <li><%= link_to "Profile", current_user %></li>
    <% end %>
    <li><%= link_to "Help", help_path %></li>
    <% if signed_in? %>
      <li><%= link_to "Sign out", signout_path, :method => :delete %></li>
    <% else %>
      <li><%= link_to "Sign in", signin_path %></li>
    <% end %>
  </ul>
</div>
```

With the code in this section, a signed-in user now sees both signout and profile links, as expected (Figure 9.8).

9.4.4 Signin/out integration tests

As a capstone to our hard work on authentication, we'll finish with integration tests for signin and signout (placed in the `users_spec.rb` file for convenience). The combination of RSpec and Webrat is expressive enough that Listing 9.32 should need little explanation; I especially like the use of `click_link "Sign out"`, which not only simulates a browser clicking the signout link, but also raises an error if no such link exists—thereby testing the URL, the named route, the link text, and the changing of the layout links, all in one line. If that's not an *integration* test, I don't know what is.

Listing 9.32. An integration test for signing in and out.`spec/integration/users_spec.rb`

```
require 'spec_helper'

describe "Users" do
  describe "signup" do
    .
    .
    .
  end

  describe "sign in/out" do
    describe "failure" do
      it "should not sign a user in" do
```

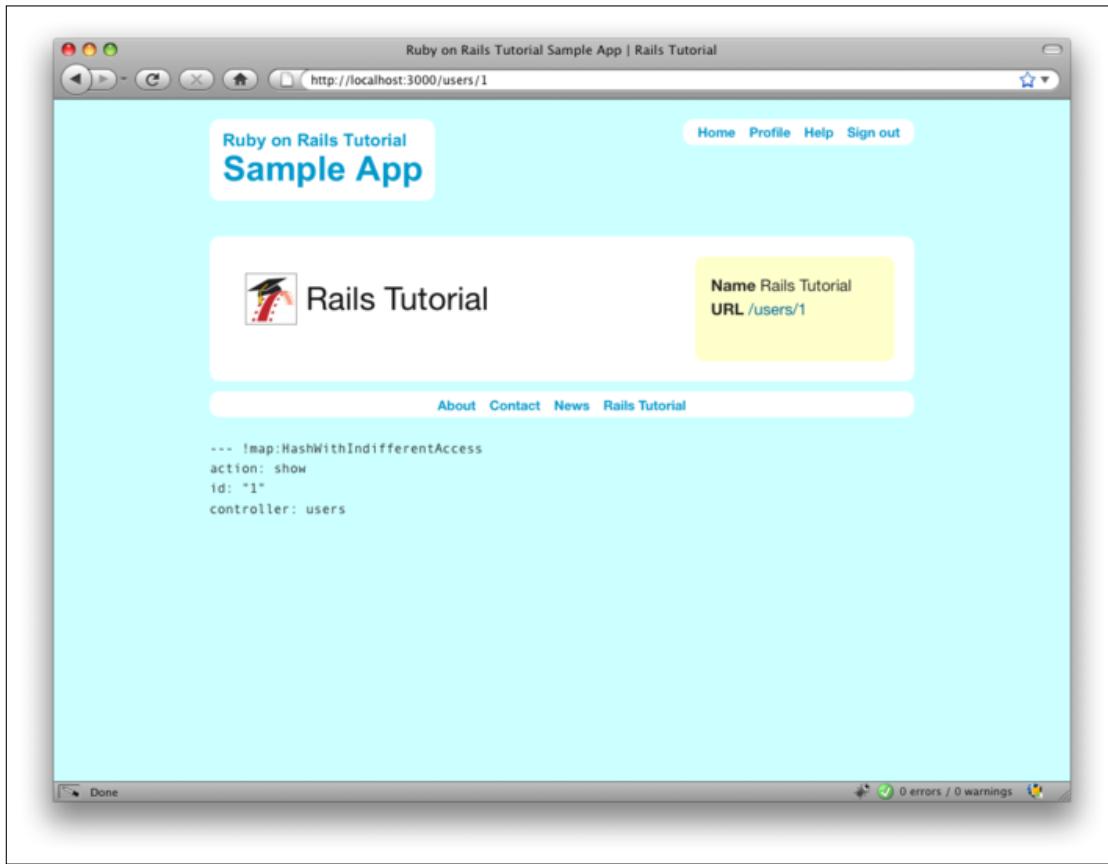


Figure 9.8: A signed-in user with signout and profile links. ([full size](#))

```
visit signin_path
fill_in :email,      :with => ""
fill_in :password,  :with => ""
click_button
response.should render_template('sessions/new')
response.should have_tag("div.flash.error", /invalid/i)
end
end

describe "success" do
  it "should sign a user in and out" do
    user = Factory(:user)
    visit signin_path
    fill_in :email,      :with => user.email
    fill_in :password,  :with => user.password
    click_button
    controller.should be_signed_in
    click_link "Sign out"
    controller.should_not be_signed_in
  end
end
end
end
```

9.5 Conclusion

We've covered a lot of ground in this chapter, transforming our promising but unformed application into a site capable of the full suite of registration and login behaviors. All that is needed to complete the authentication functionality is to restrict access to pages based on signin status and user identity. We'll accomplish this task en route to giving users the ability to edit their information and giving administrators the ability to remove users from the system.

Before moving on, merge your changes back into the master branch:

```
$ git add .
$ git commit -am "Done with sign in"
$ git checkout master
$ git merge sign-in-out
```

9.6 Exercises

The second and third exercises are more difficult than usual. Solving them will require some outside research (e.g., Rails API reading and Google searches), and they can be skipped without loss of continuity.

1. Several of the integration specs use the same code to sign a user in. Replace that code with the `integration_sign_in` function in Listing 9.33 and verify that the tests still pass.

2. Write an integration test for signing in, signing out, and then signing back in, which will catch the password re-encryption issue mentioned after Listing 9.14.
3. Use `session` instead of `cookies` so that users are automatically signed out when they close their browsers.¹⁷ Hint: Do a Google search on “Rails session”.
4. (advanced) Some sites use secure HTTP (HTTPS) for their signin pages. Search online to learn how to use HTTPS in Rails, and then secure the Sessions controller `new` and `create` actions. Extra challenge: Write tests for the HTTPS functionality.

Listing 9.33. A function to sign users in inside of integration tests.

`spec/spec_helper.rb`

```
.  
. .  
Spec::Runner.configure do |config|  
. .  
. .  
def test_sign_in(user)  
  controller.current_user = user  
end  
  
def integration_sign_in(user)  
  visit signin_path  
  fill_in :email, :with => user.email  
  fill_in :password, :with => user.password  
  click_button  
end  
end
```

¹⁷Somewhat confusingly, we've used `cookies` to implement sessions, and `session` is implemented with cookies!

Chapter 10

Updating, showing, and deleting users

In this chapter, we will complete the REST actions for the Users resource (Table 6.2) by adding `edit`, `update`, `index`, and `destroy` actions. We'll start by giving users the ability to update their profiles, which will also provide a natural opportunity to enforce a security model (made possible by the authentication code in Chapter 9). Then we'll make a listing of all users (also requiring authentication), which will motivate the introduction of sample data and pagination. Finally, we'll add the ability to destroy users, wiping them from the database. Since we can't allow just any user to have such dangerous powers, we'll take care to create a privileged class of administrative users (admins) along the way.

To get started, let's start work on an `updating-users` topic branch:

```
$ git checkout -b updating-users
```

10.1 Updating users

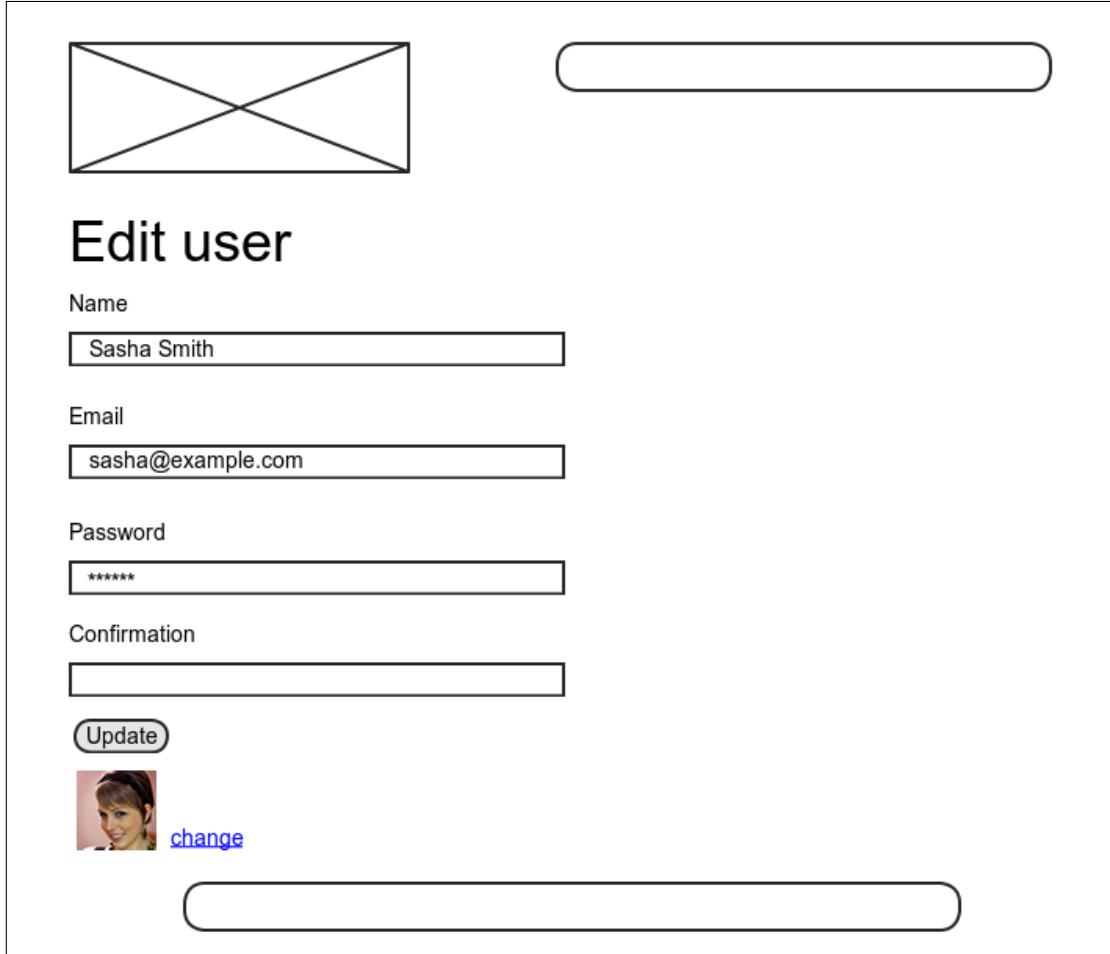
The pattern for editing user information closely parallels that for creating new users (Chapter 8). Instead of a `new` action rendering a view for new users, we have an `edit` action rendering a view to edit users; instead of `create` responding to a POST request, we have an `update` action responding to a PUT request (Box 3.1). The biggest difference is that, while anyone can sign up, only the current user should be able to update their information. This means that we need to enforce access control so that only authorized users can edit and update; the authentication machinery from Chapter 9 will allow us to use a *before filter* to ensure that this is the case.

10.1.1 Edit form

We start with tests for the edit form, whose mockup appears in Figure 10.1.¹ Two are analogous to tests we saw for the `new` user page (Listing 8.1), checking for the proper response and title; the third test makes sure that there is a link to edit the user's Gravatar image (Section 7.3.2). If you poke around the Gravatar site, you'll see that the page to add or edit images is (somewhat oddly) located at <http://gravatar.com/emails>, so we test the `edit` page for a link with that URL.² The result is shown in Listing 10.1.

¹Image from <http://www.flickr.com/photos/sashawolff/4598355045/>.

²The Gravatar site actually redirects this to <http://en.gravatar.com/emails>, which is for English language users, but I've omitted the `en` part to account for the use of other languages.



The image shows a wireframe-style mockup of a user edit page. At the top left is a placeholder image of a person's face with an 'X' drawn over it. To its right is a long, thin, rounded rectangular placeholder for a profile picture. Below these, the title "Edit user" is centered. The form consists of several input fields and labels:

- Name: A text input field containing "Sasha Smith".
- Email: A text input field containing "sasha@example.com".
- Password: A text input field containing "*****".
- Confirmation: An empty text input field.
- Update: A button labeled "Update" enclosed in an oval shape.
- User Picture: A small placeholder image of a woman's face with a "change" link next to it.

Below the "Confirmation" field is another long, thin, rounded rectangular placeholder for a profile picture.

Figure 10.1: A mockup of the user edit page. ([full size](#))

Listing 10.1. Tests for the user `edit` action.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views
  .

  .

  .

  describe "GET 'edit'" do
    before(:each) do
      @user = Factory(:user)
      test_sign_in(@user)
    end

    it "should be successful" do
      get :edit, :id => @user
      response.should be_success
    end

    it "should have the right title" do
      get :edit, :id => @user
      response.should have_tag("title", /edit user/i)
    end

    it "should have a link to change the Gravatar" do
      get :edit, :id => @user
      gravatar_url = "http://gravatar.com/emails"
      response.should have_tag("a[href=?]", gravatar_url, /change/i)
    end
  end
end
```

Here we've made sure to use `test_sign_in(@user)` to sign in as the user in anticipation of protecting the edit page from unauthorized access (Section 10.2). Otherwise, these tests would break as soon as we implemented our authentication code.

Note from [Table 6.2](#) that the proper URL for a user's edit page is `/users/1/edit` (assuming the user's id is 1). Recall that the id of the user is available in the `params[:id]` variable, which means that we can get the first test to pass using the code in [Listing 10.2](#). This uses `find` to find the relevant user in the database, and then sets the `@title` variable to the proper value.

Listing 10.2. An incomplete user `edit` action.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
```

```

.
def edit
  @user = User.find(params[:id])
  @title = "Edit user"
end
end

```

Getting the second (title) and third (Gravatar) tests to pass requires making the actual edit view, shown in Listing 10.3. Note how closely this resembles the new user view from Listing 8.2; the large overlap suggests factoring the repeated code into a partial, which is left as an exercise (Section 10.6).

Listing 10.3. The user edit view.

`app/views/users/edit.html.erb`

```

<h2>Edit user</h2>

<% form_for(@user) do |f| %>
  <%= f.error_messages %>
  <div class="field">
    <%= f.label :name %><br />
    <%= f.text_field :name %>
  </div>
  <div class="field">
    <%= f.label :email %><br />
    <%= f.text_field :email %>
  </div>
  <div class="field">
    <%= f.label :password %><br />
    <%= f.password_field :password %>
  </div>
  <div class="field">
    <%= f.label :password_confirmation, "Confirmation" %><br />
    <%= f.password_field :password_confirmation %>
  </div>
  <div class="actions">
    <%= f.submit "Update" %>
  </div>
<% end %>

<div>
  <%= gravatar_for @user %>
  <a href="http://gravatar.com/emails">change</a>
</div>

```

We'll also add a link to the site navigation for the user edit page (which we'll call "Settings"), as mocked up in Figure 10.2³ and shown in Listing 10.4.

Listing 10.4. Adding a Settings link.

`app/views/layouts/_header.html.erb`

³Image from <http://www.flickr.com/photos/sashawolff/4598355045/>.

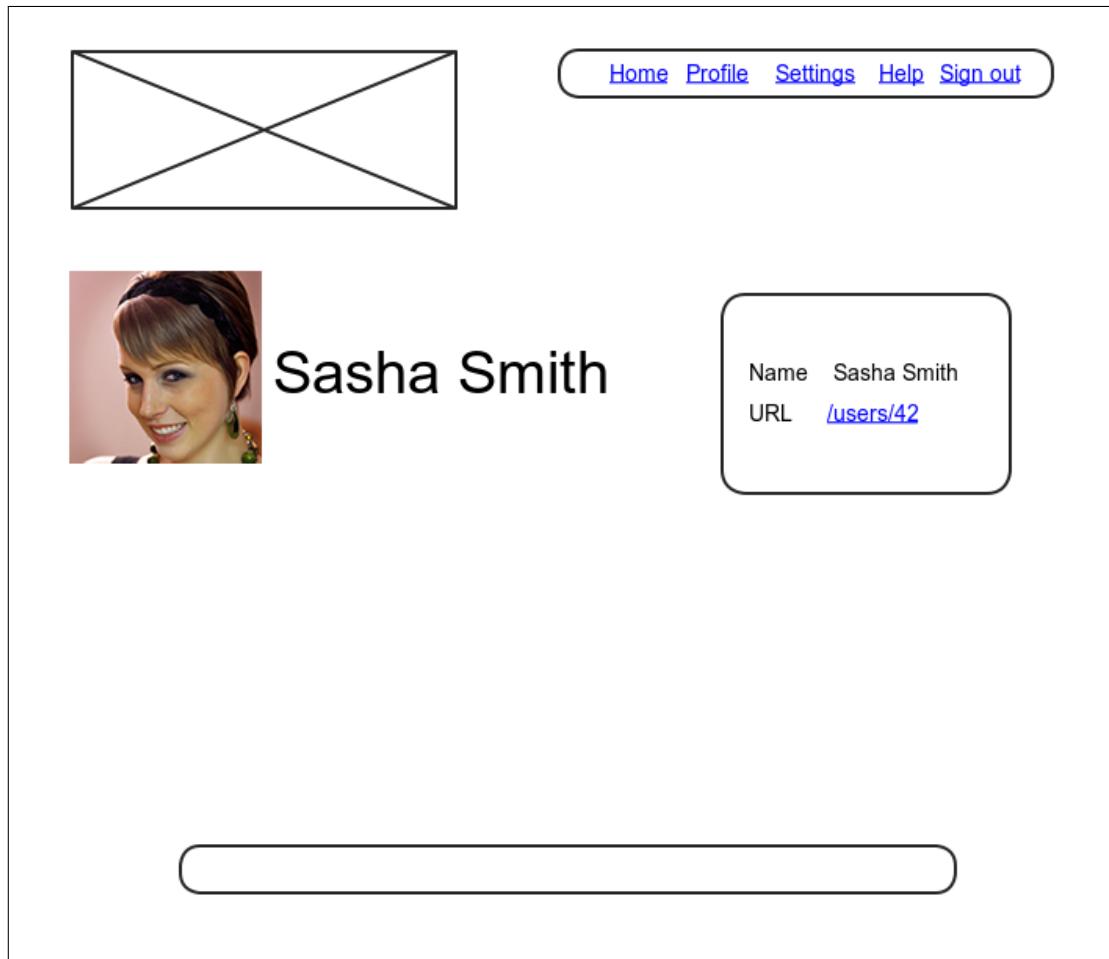


Figure 10.2: A mockup of the user profile page with a “Settings” link. ([full size](#))

```
<div id="header" class="round">
  <%= link_to logo, root_path %>
  <ul class="navigation round">
    <li><%= link_to "Home", root_path %></li>
    <% if signed_in? %>
      <li><%= link_to "Profile", current_user %></li>
      <li><%= link_to "Settings", edit_user_path(current_user) %></li>
    <% end %>
    .
    .
    .
  </ul>
</div>
```

Here we use the named route `edit_user_path` from Table 6.2, together with the handy `current_user` helper method defined in Listing 9.19.

With the `@user` instance variable from Listing 10.2, the tests from Listing 10.1 pass. As seen in Figure 10.3, the `new` page renders, though it doesn't yet work.

Looking at the HTML source for Figure 10.3, we see a form tag as expected (Listing 10.5).

Listing 10.5. HTML for the edit form defined in Listing 10.3 and shown in Figure 10.3.

```
<form action="/users/1" class="edit_user" id="edit_user_1" method="post">
  <input name="_method" type="hidden" value="put" />
  .
  .
  .
</form>
```

Note here the hidden input field

```
<input name="_method" type="hidden" value="put" />
```

Since web browsers can't natively send PUT requests (as required by the REST conventions from Table 6.2), Rails fakes it with a POST request and a hidden `input` field.⁴

There's another subtlety to address here: the code `form_for (@user)` in Listing 10.3 is *exactly* the same as the code in Listing 8.2—so how does Rails know to use a POST request for new users and a PUT for editing users? The answer is that it is possible to tell whether a user is new or already exists in the database via the `new_user?` boolean method:

```
$ script/console
>> User.new.new_record?
```

⁴Don't be worried about how this works; the details are of interest to developers of the Rails framework itself, but by design are not important for Rails application developers.

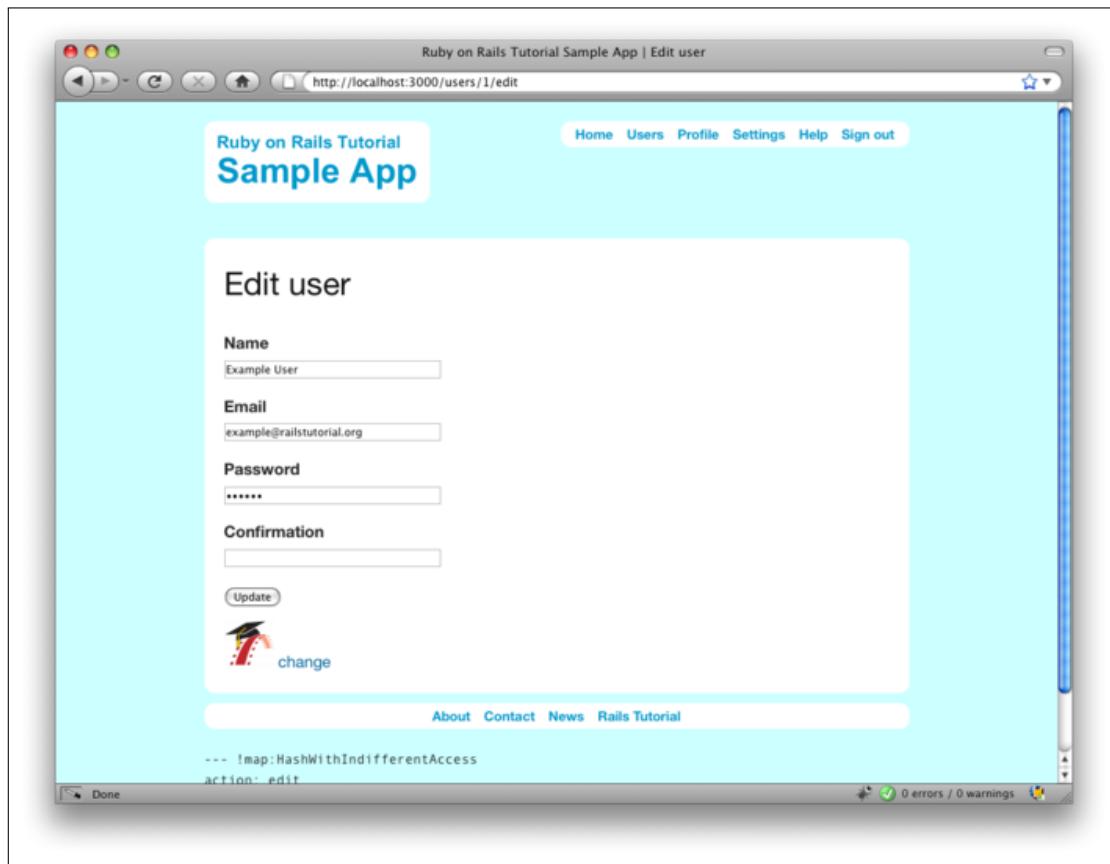


Figure 10.3: Editing user settings (</users/1/edit>). [\(full size\)](#)

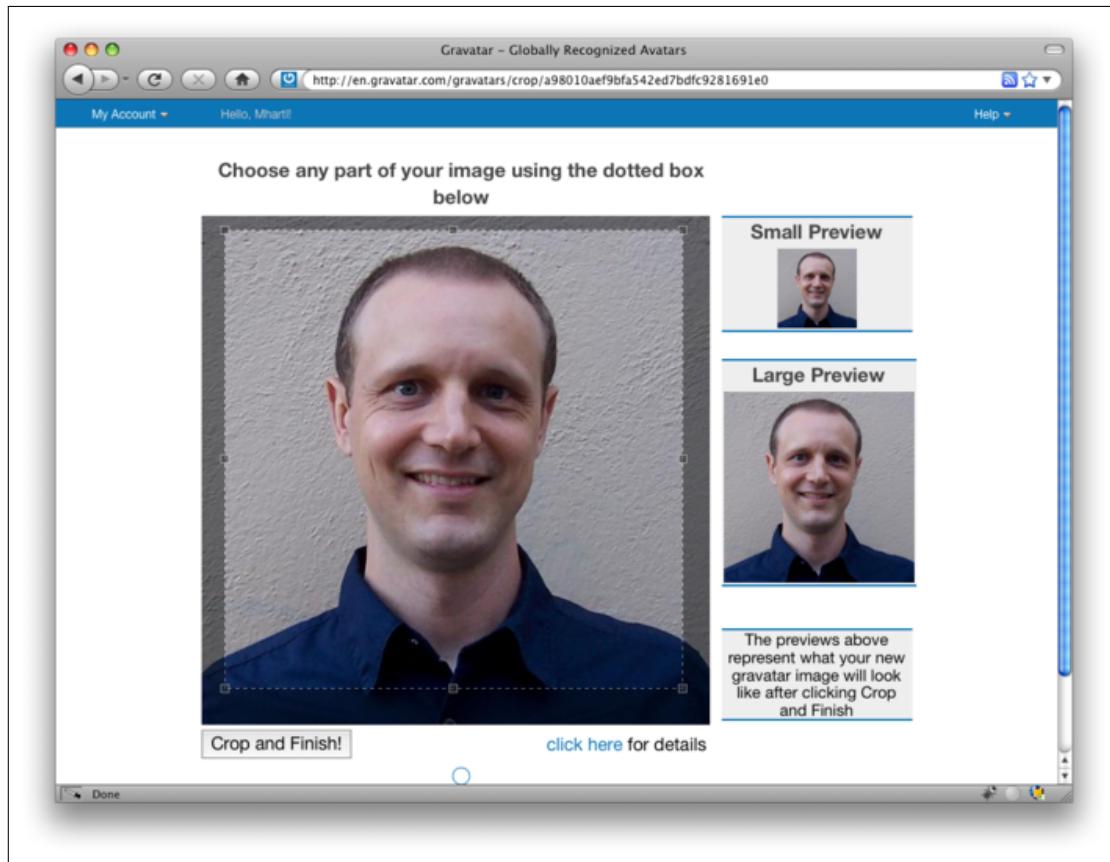


Figure 10.4: The [Gravatar](#) image-cropping interface, with a picture of [some dude.](#) ([full size](#))

```
=> true
>> User.first.new_record?
=> false
```

When constructing a form using `form_for (@user)`, Rails uses POST if `@user.new_record?` is `true` and PUT if it is `false`.

10.1.2 Enabling edits

Although the edit form doesn't yet work, we've outsourced image upload to Gravatar, so it works straightaway by clicking on the "change" link from Figure 10.3, as shown in Figure 10.4.

Let's get the rest of the user edit functionality working as well.

The tests for the `update` action are similar to those for `create`. In particular, we test both update failure and update success (Listing 10.6). (This is a lot of code; see if you can work through it by referring back to the tests in Chapter 8.)

Listing 10.6. Tests for the user `update` action.`spec/controllers/users_controller_spec.rb`

```
describe UsersController do
  integrate_views
  .
  .
  .
  describe "PUT 'update'" do
    before(:each) do
      @user = Factory(:user)
      test_sign_in(@user)
      User.should_receive(:find).with(@user).and_return(@user)
    end

    describe "failure" do
      before(:each) do
        @invalid_attr = { :email => "", :name => "" }
        @user.should_receive(:update_attributes).and_return(false)
      end

      it "should render the 'edit' page" do
        put :update, :id => @user, :user => {}
        response.should render_template('edit')
      end

      it "should have the right title" do
        put :update, :id => @user, :user => {}
        response.should have_tag("title", /edit user/i)
      end
    end

    describe "success" do
      before(:each) do
        @attr = { :name => "New Name", :email => "user@example.org",
                  :password => "barbaz", :password_confirmation => "barbaz" }
        @user.should_receive(:update_attributes).and_return(true)
      end

      it "should redirect to the user show page" do
        put :update, :id => @user, :user => @attr
        response.should redirect_to(user_path(@user))
      end

      it "should have a flash message" do
        put :update, :id => @user, :user => @attr
        flash[:success].should =~ '/updated/'
      end
    end
  end
end
```

```
    end
  end
end
```

Here we've set up the message expectations to require that `User.find` be called, as in

```
User.should_receive(:find).with(@user).and_return(@user)
```

which also arranges to return the proper user. We also require that `update_attributes` be called, as in

```
@user.should_receive(:update_attributes).and_return(false)
```

and

```
@user.should_receive(:update_attributes).and_return(true)
```

Compare this to the code from [Chapter 8](#), where we required calls to `save`, e.g.,

```
@user.should_receive(:save).and_return(false)
```

and

```
@user.should_receive(:save).and_return(true)
```

(The expectations are located in `before` blocks so that they execute before every test, thereby preventing any of the `put` calls from actually hitting the database. In a previous version of this book, the code in [Chapter 8](#) and [Chapter 9](#) didn't put the expectations in the before block; if you have been reading chapters immediately after their release, you might want to review those chapters before proceeding.)

The `update` action needed to get the tests in [Listing 10.6](#) to pass is similar to the final form of the `create` action ([Listing 9.26](#)), as seen in [Listing 10.7](#).

Listing 10.7. The user `update` action.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
```

```
def update
  @user = User.find(params[:id])
  if @user.update_attributes(params[:user])
    flash[:success] = "Profile updated."
    redirect_to @user
  else
    @title = "Edit user"
    render 'edit'
  end
end
end
```

With that, the user edit page should be working. As presently constructed, every edit requires the user to reconfirm the password (as implied by the empty confirmation text box in [Figure 10.3](#)), which makes updates more secure but is a minor annoyance. (I plan to show how to allow user updates without a password confirmation in the [Rails Tutorial screencasts](#).)

10.2 Protecting pages

Although the edit and update actions from [Section 10.1](#) are functionally complete, they suffer from a ridiculous security flaw: they allow anyone (even non-signed-in users) to access either action, and any signed-in user can update the information for any other user.⁵ In this section, we'll implement a security model that requires users to be signed in and prevents them from updating any information other than their own. Users who aren't signed in and who try to access protected pages will be forwarded to the signin page with a helpful message, as mocked up in [Figure 10.5](#).

10.2.1 Requiring signed-in users

Since the security restrictions for the `edit` and `update` actions are identical, we'll handle them in a single RSpec `describe` block. Starting with the sign-in requirement, our initial tests verify that non-signed-in users attempting to access either action are simply redirected to the signin page, as seen in [Listing 10.8](#).

Listing 10.8. The first tests for authentication.

`spec/controllers/users_controller_spec.rb`

```
describe UsersController do
  integrate_views
  .

  .

  .

  describe "authentication of edit/update pages" do
    before(:each) do
      @user = Factory(:user)
    end
  end
end
```

⁵To be fair, they would need the user's password, but if we ever made the password unnecessary (as planned for the screencasts) it would open up a *huge* security hole.

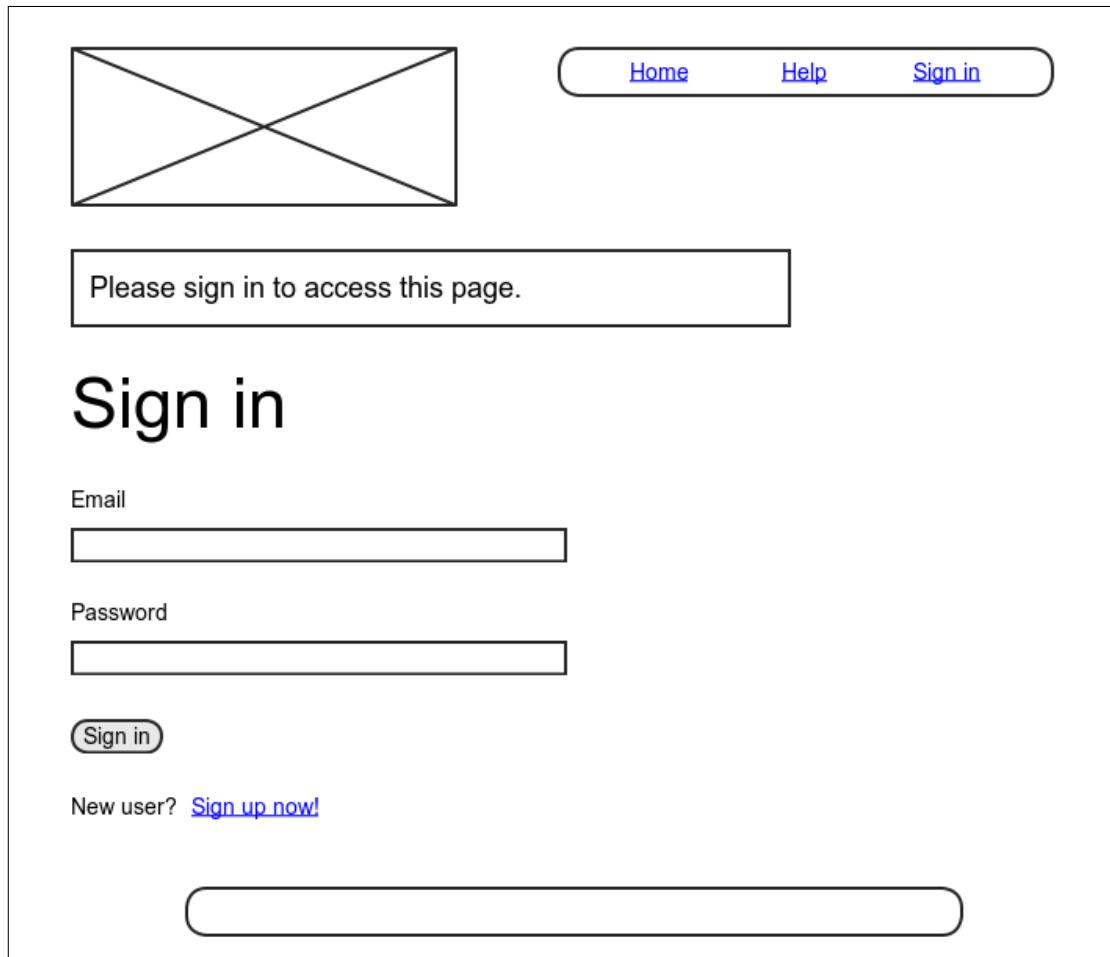


Figure 10.5: A mockup of the result of visiting a protected page ([full size](#))

```

describe "for non-signed-in users" do

  it "should deny access to 'edit'" do
    get :edit, :id => @user
    response.should redirect_to(signin_path)
  end

  it "should deny access to 'update'" do
    put :update, :id => @user, :user => {}
    response.should redirect_to(signin_path)
  end
end
end

```

The application code gets these tests to pass using a *before filter*, which arranges for a particular method to be called before the given actions. In this case, we define an `authenticate` method and invoke it using `before_filter :authenticate`, as shown in Listing 10.9.

Listing 10.9. Adding an `authenticate` before filter.

`app/controllers/users_controller.rb`

```

class UsersController < ApplicationController
  before_filter :authenticate, :only => [:edit, :update]
  .

  .

  .

  private

  def authenticate
    deny_access unless signed_in?
  end
end

```

By default, before filters apply to *every* action in a controller, so here we restrict the filter to act only on the `:edit` and `:update` actions by passing the `:only` options hash.

This code won't work yet, because `deny_access` hasn't been defined. Since access denial is part of authentication, we'll put it in the Sessions helper from Chapter 9. All `deny_access` does is put a message in `flash[:notice]` and then redirect to the signin page (Listing 10.10).

Listing 10.10. The `deny_access` method for user authentication.

`app/helpers/sessions_helper.rb`

```

module SessionsHelper
  .
  .
  .

```

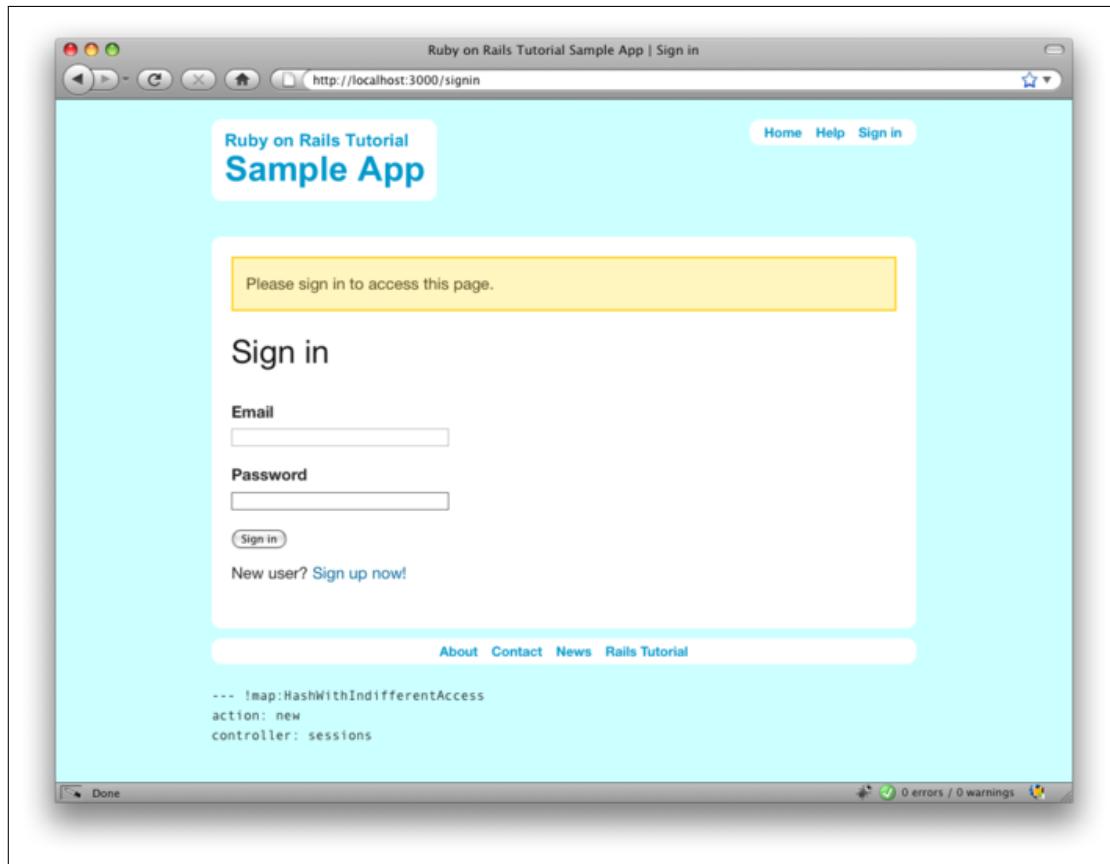


Figure 10.6: The signin form after trying to access a protected page. ([full size](#))

```
def deny_access
  flash[:notice] = "Please sign in to access this page."
  redirect_to signin_path
end
end
```

Together with `:success` and `:failure`, the `:notice` key completes our triumvirate of `flash` styles, all of which are supported natively by Blueprint CSS. By signing out and attempting to access the user edit page `/users/1/edit`, we can see the resulting yellow "`notice`" box, as seen in Figure 10.6.

10.2.2 Requiring the right user

Of course, requiring users to sign in isn't quite enough; users should only be allowed to edit their *own* information. We can test for this by first signing in as an incorrect user and then hitting the `edit` and `update` actions (Listing 10.11). Note that, since users should never even *try* to edit another user's profile, we redirect not to the signin page but to the root url.

Listing 10.11. Authentication tests for signed-in users.`spec/controllers/users_controller_spec.rb`

```
describe UsersController do
  integrate_views
  .
  .
  .
  describe "authentication of edit/update pages" do
    .
    .
    .
    describe "for signed-in users" do
      before(:each) do
        wrong_user = Factory(:user, :email => "user@example.net")
        test_sign_in(wrong_user)
      end

      it "should require matching users for 'edit'" do
        get :edit, :id => @user
        response.should redirect_to(root_path)
      end

      it "should require matching users for 'update'" do
        put :update, :id => @user, :user => {}
        response.should redirect_to(root_path)
      end
    end
  end
end
```

The application code is simple: we add a second before filter to call the `correct_user` method (which we have to write), as shown in Listing 10.12.

Listing 10.12. A `correct_user` before filter to protect the edit/update pages.`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  before_filter :authenticate, :only => [:index, :edit]
  before_filter :correct_user, :only => [:edit, :update]
  .
  .
  .
  private

  def authenticate
    deny_access unless signed_in?
  end
```

```
def correct_user
  @user = User.find(params[:id])
  redirect_to(root_path) unless current_user?(@user)
end
end
```

This uses the `current_user?` method, which (as with `deny_access`) we will define in the Sessions helper ([Listing 10.13](#)).

Listing 10.13. The `current_user?` method.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper
  .
  .
  .

  def current_user?(user)
    user == current_user
  end

  def deny_access
    flash[:notice] = "Please sign in to access this page."
    redirect_to signin_path
  end
end
```

Although this implementation now works, your test suite should fail. We'll understand why it fails (and get it to pass) in the next section.

10.2.3 An expectation bonus

As you can verify by filling in the edit form and submitting, the edit/update code currently works, but the test suite fails. This seems like a bug, but it's actually a feature: message expectations like

```
User.should_receive(:find)
```

require that the method be called *exactly once*. The test failure is thus a hint that `find` is being called too many times. Indeed, upon inspecting the current Users controller, we see that `User.find` is called both in the `edit/user` action *and* in the `current_user` before filter. To remove the superfluous call, we can simply rely on the before filter to find `@user` and remove the other lines, as shown in [Listing 10.14](#); the resulting `@user` instance variable will then be available in all the actions and views as usual.

Listing 10.14. Removing the superfluous `@user` variable assignments.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  .
  .
  .
  def edit
    @title = "Edit user"
  end

  def update
    if @user.update_attributes(params[:user])
      flash[:success] = "Profile updated."
      redirect_to @user
    else
      @title = "Edit user"
      render 'edit'
    end
  end

  private
  .
  .
  .
  def correct_user
    @user = User.find(params[:id])
    redirect_to(root_path) unless current_user?(@user)
  end
end
```

At this point, all the tests should pass, and we are nearly done with protecting pages.

10.2.4 Friendly forwarding

Our page protection is complete as written, but there is one minor blemish: when users try to access a protected page, they are currently redirected to their profile pages regardless of where they were trying to go. In other words, if a non-logged-in user tries to visit the edit page, after signing in the user will be redirected to /users/1 instead of /users/1/edit. It would be much friendlier to redirect them to their intended destination instead.

The sequence of attempted page visitation, signin, and redirect to destination page is a perfect job for an integration test, so let's make one for friendly forwarding:

```
$ script/generate integration_spec friendly_forwarding
```

The code then appears as in Listing 10.15.

Listing 10.15. An integration test for friendly forwarding.

```
spec/integration/friendly_forwardings_spec.rb
```

```

require 'spec_helper'

describe "FriendlyForwardings" do

  it "should forward to the requested page after signin" do
    user = Factory(:user)
    visit edit_user_path(user)
    # Webrat automatically follows the redirect to the signin page.
    fill_in :email,      :with => user.email
    fill_in :password,   :with => user.password
    click_button
    # Webrat follows the redirect again, this time to users/edit.
    response.should render_template('users/edit')
  end
end

```

(As indicated by the comments, Webrat *follows* redirects, so testing that the response **should redirect_to** some URL won't work. I learned this the hard way.)

Now for the implementation.⁶ In order to forward users to their intended destination, we need to store the location of the requested page somewhere, and then redirect there instead. The storage mechanism is the **session** facility provided by Rails, which you can think of as being like an instance of the **cookies** variable from Section 9.3.3 that automatically expires upon browser close.⁷ We also use the **request** object to get the **request_uri**, i.e., the URL of the requested page. The resulting application code appears in Listing 10.16.

Listing 10.16. Code to implement friendly forwarding.

`app/helpers/sessions_helper.rb`

```

module SessionsHelper
  .
  .
  .

  def deny_access
    store_location
    flash[:notice] = "Please sign in to access this page."
    redirect_to signin_path
  end

  def store_location
    session[:return_to] = request.request_uri
  end

  def redirect_back_or(default)
    redirect_to(session[:return_to] || default)
    clear_return_to
  end
end

```

⁶The code in this section is adapted from the [Clearance](#) gem by [thoughtbot](#).

⁷Indeed, as noted in Section 9.6, **session** is implemented in just this way.

```
def clear_return_to
  session[:return_to] = nil
end
end
```

Here we've added a line to the `deny_access` method, first storing the location of the request with `store_location` and then proceeding as before. The `store_location` method puts the requested URL in the `session` variable under the key `:return_to`.

We've also defined the `redirect_back_or` method to redirect to the requested URL if it exists, or some default URL otherwise. This method is needed in the Session `create` action to after successful signin (Listing 10.17).

Listing 10.17. The Sessions `create` action with friendly forwarding.

`app/controllers/sessions_controller.rb`

```
class SessionsController < ApplicationController
  .
  .
  .
  def create
    user = User.authenticate(params[:session][:email],
                            params[:session][:password])
    if user.nil?
      flash.now[:error] = "Invalid email/password combination."
      @title = "Sign in"
      render 'new'
    else
      sign_in user
      redirect_back_or user
    end
  end
  .
  .
  .
end
```

With that, the friendly forwarding integration test in Listing 10.15 should pass, and the basic user authentication and page protection implementation is complete.

10.3 Showing users

In this section, we'll add the penultimate user action, the `index` action, which is designed to display *all* the users, not just one. Along the way, we'll learn about populating the database with sample users and *paginating* the user output so that the index page can scale up to display a potentially large number of users. A mockup of the result—users, pagination links, and a “Users” navigation link—appears in Figure 10.7. In Section 10.4, we'll add an administrative interface to the user index so that (presumably troublesome) users can be destroyed.

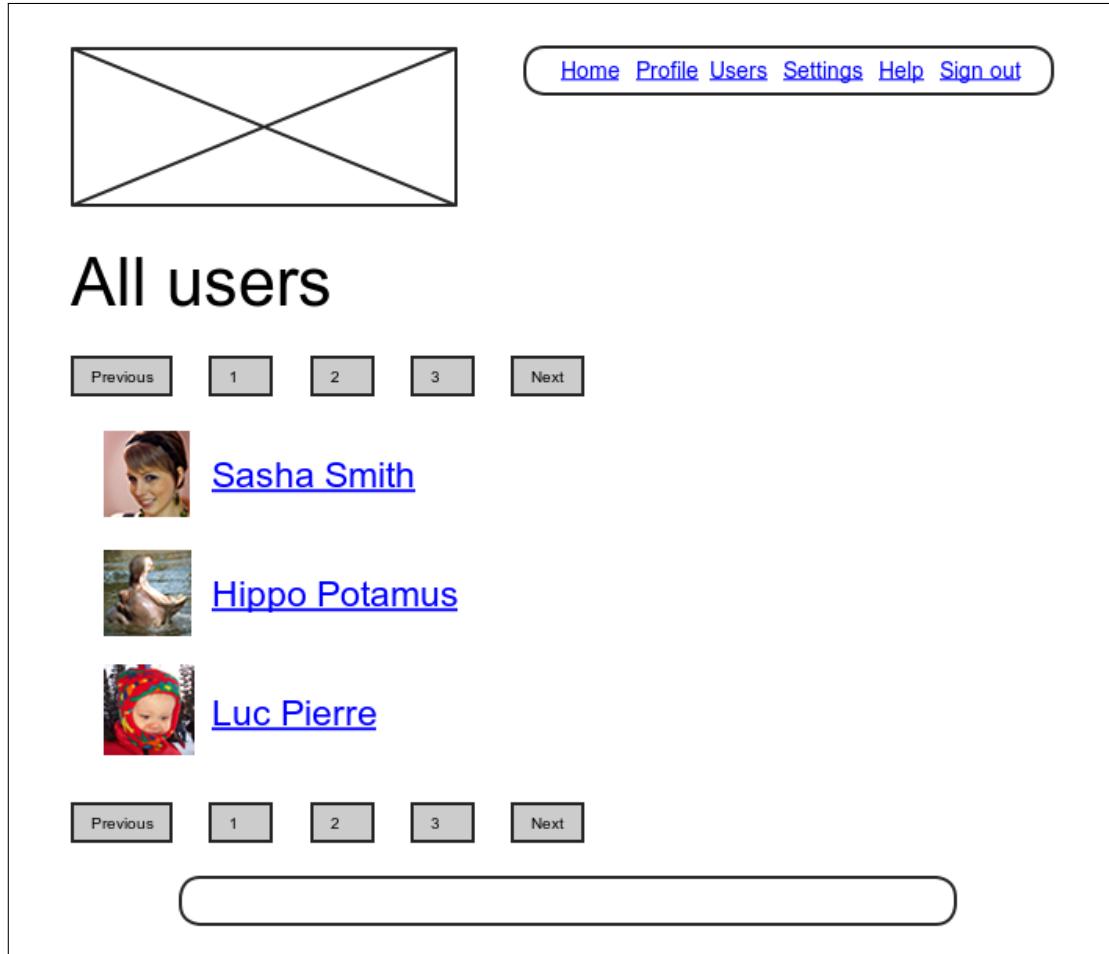


Figure 10.7: A mockup of the user index, with pagination and a “Users” nav link. ([full size](#))

10.3.1 User index

Although we'll keep individual user `show` pages visible to all site visitors, the user `index` will be restricted to signed-in users so that there's a limit to how much unregistered users can see by default. Our `index` tests check for this, and also verify that for signed-in users all the site's users are listed.

Listing 10.18. Tests for the user index page.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views

  describe "GET 'index'" do
    describe "for non-signed-in users" do
      it "should deny access" do
        get :index
        response.should redirect_to(signin_path)
        flash[:notice].should =~ /sign in/i
      end
    end

    describe "for signed-in users" do
      before(:each) do
        @user = test_sign_in.Factory(:user)
      end

      it "should be successful" do
        get :index
        response.should be_success
      end

      it "should have the right title" do
        get :index
        response.should have_tag("title", /all users/i)
      end

      it "should have an element for each user" do
        second_user = Factory(:user, :email => "another@example.com")
        third_user = Factory(:user, :email => "another@example.net")
        get :index
        [@user, second_user, third_user].each do |user|
          response.should have_tag("li", user.name)
        end
      end
    end
  end
.
```

```
.
end
```

As you can see, the method for checking the index page is to make three factory users (signing in as the first one) and then verify that the index page has a list element (`li`) tag for the name of each one.

To build the page, we first use `User.all` from Section 6.1.4 to make an `@users` instance variable in the `index` action of the Users controller (Listing 10.19).

Listing 10.19. The user `index` action.

`app/controllers/users_controller.rb`

```
class UsersController < ApplicationController
  before_filter :authenticate, :only => [:index, :edit, :update]
  .

  .

  def index
    @title = "All users"
    @users = User.all
  end
  .

  .
end
```

Note that we have added `:index` to the list of controllers protected by the `authenticate` before filter, thereby getting the first test from Listing 10.18 to pass.

To make the actual page, we need to make a view that iterates through the users and wraps each one in an `li` tag. We do this with the `each` method, displaying each user's Gravatar and name, while wrapping the whole thing in an unordered list (`ul`) tag (Listing 10.20). Note that each user's name is both linked to and escaped with `link_to h(user.name)`, `user`.⁸

Listing 10.20. The user index view.

`app/views/users/index.html.erb`

```
<h2>All users</h2>

<ul class="users">
  <% @users.each do |user| %>
    <li>
      <%= gravatar_for user, :size => 30 %>
      <%= link_to h(user.name), user %>
    </li>
  <% end %>
</ul>
```

⁸Note that Ruby allows the omission of *one* pair of parentheses, but not two, so instead of `link_to h user.name` we write `link_to h(user.name)`.

We'll then add a little CSS for style (Listing 10.21).

Listing 10.21. CSS for the user index.

`public/stylesheets/custom.css`

```
.  
. .  
  
ul.users {  
  margin-top: 1em;  
}  
  
.users li {  
  list-style: none;  
}
```

Finally, we'll add a "Users" link to the site's navigation header (Listing 10.22). This puts to use the `users_path` named route from Table 6.2.

Listing 10.22. A layout link to the user index.

`app/views/layouts/_header.html.erb`

```
<div id="header" class="round">  
  <%= link_to logo, root_path %>  
  <ul class="navigation round">  
    <li><%= link_to "Home", root_path %></li>  
    <% if signed_in? %>  
    <li><%= link_to "Users", users_path %></li>  
    <li><%= link_to "Profile", current_user %></li>  
    <li><%= link_to "Settings", edit_user_path(current_user) %></li>  
    <% end %>  
    .  
    .  
    .  
  </ul>  
</div>
```

With that, the user index is fully functional (with all tests passing), but it is a bit...lonely (Figure 10.8).

10.3.2 Sample users

In this section we'll give our lonely sample user some company. Of course, to create enough users to make a decent user index, we *could* use our web browser to visit the signup page and make the new users one by one, but far a better solution is to use Ruby (and Rake) to make the users for us.

First, we'll install the Faker gem, which will allow us to make sample users with semi-realistic names and email addresses:

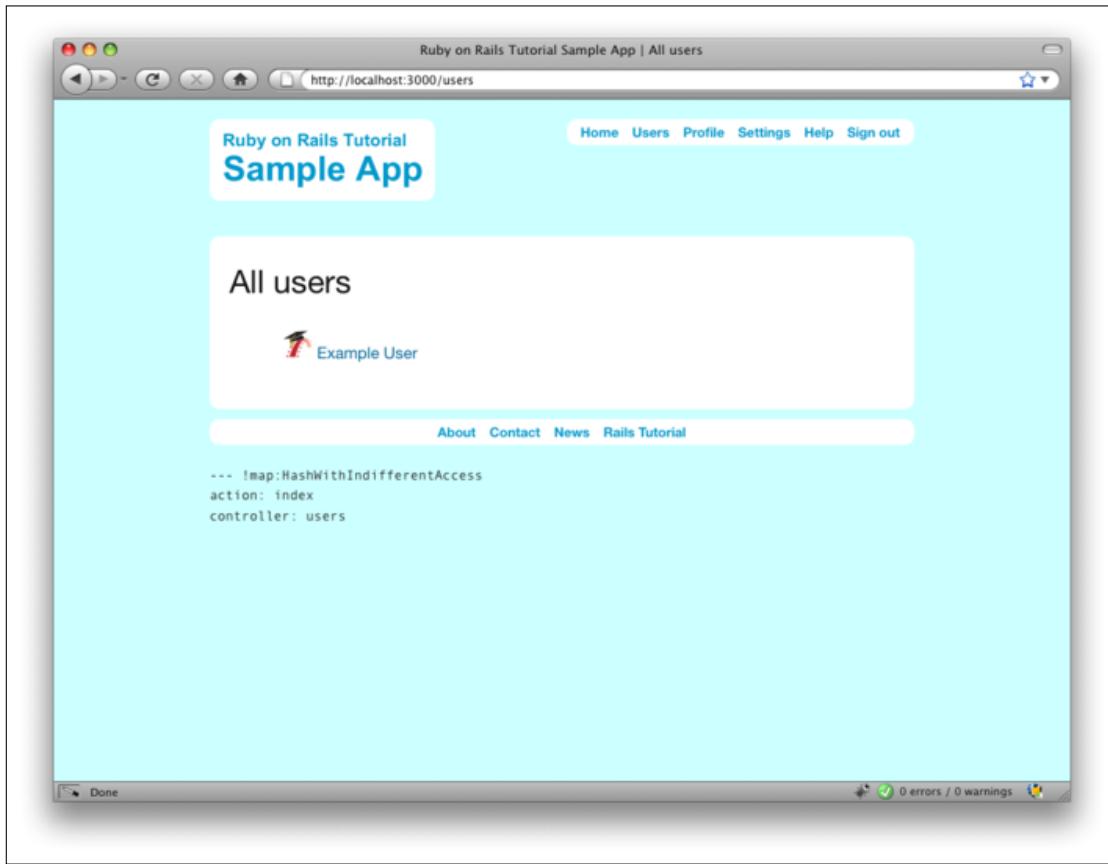


Figure 10.8: The user index page `/users` with only one user. ([full size](#))

```
$ [sudo] gem install faker -v 0.3.1
```

Next, we'll add a Rake task to create sample users. Rake tasks live in **lib/tasks**, and are defined using *namespaces* (in this case, **:db**), as seen in Listing 10.23.

Listing 10.23. A Rake task for populating the database with sample users.

lib/tasks/sample_data.rake

```
require 'faker'

namespace :db do
  desc "Fill database with sample data"
  task :populate => :environment do
    Rake::Task['db:reset'].invoke
    User.create!(:name => "Example User",
                 :email => "example@railstutorial.org",
                 :password => "foobar",
                 :password_confirmation => "foobar")
    99.times do |n|
      name = Faker::Name.name
      email = "example-#{n+1}@railstutorial.org"
      password = "password"
      User.create!(:name => name,
                  :email => email,
                  :password => password,
                  :password_confirmation => password)
    end
  end
end
```

This defines a task **db:populate** that resets the development database using **db:reset** (using slightly weird syntax you shouldn't worry about too much), creates an example user with name and email address replicating our previous one, and then makes 99 more. The line

```
task :populate => :environment do
```

ensures that the Rake task has access to the local Rails environment, including the User model (and hence **User.create!**).

With the **:db** namespace as in Listing 10.23, we can invoke the Rake task as follows:

```
$ rake db:populate
```

After running the Rake task, our application has 100 sample users, as seen in Figure 10.9. (I've taken the liberty of associating the first few sample addresses with photos so that they're not all the default Gravatar image.)

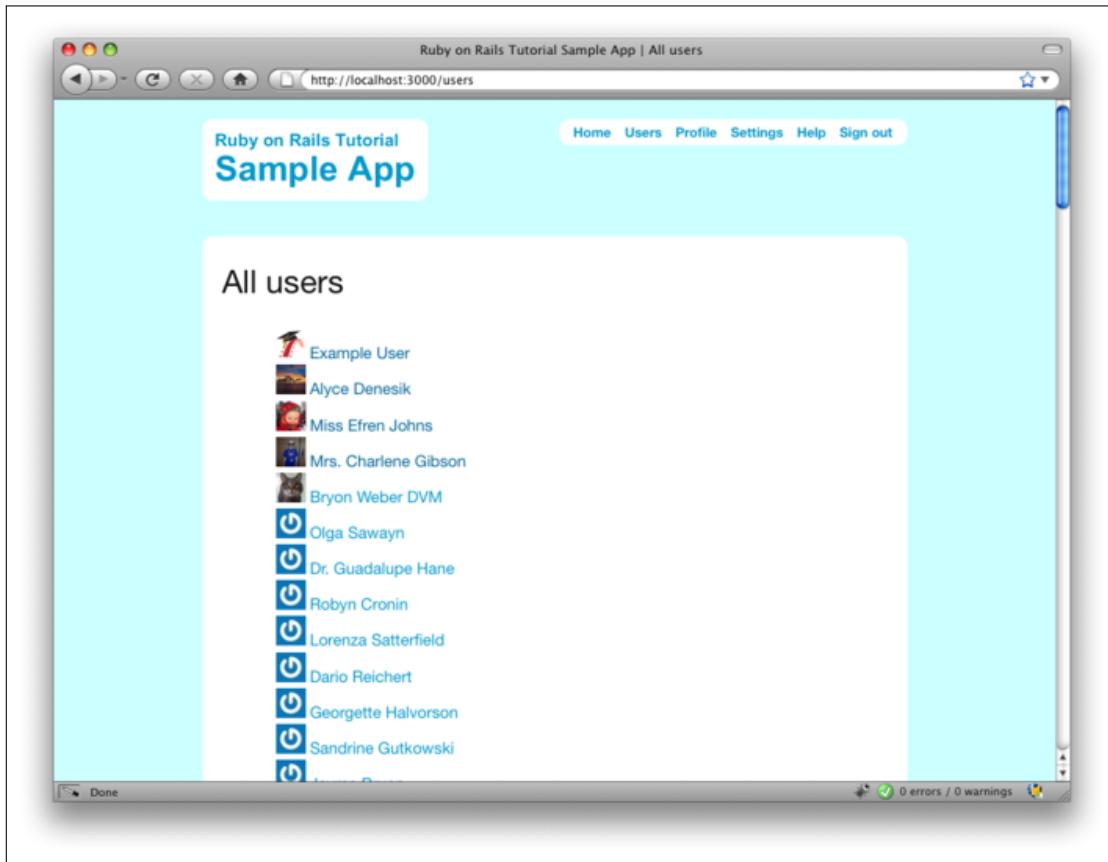


Figure 10.9: The user index page `/users` with 100 sample users. ([full size](#))

10.3.3 Pagination

Having solved the problem of too few sample users, we now encounter the opposite problem: having too many users on a page. Right now there are a hundred, which is already a reasonably large number, and on a real site it could be thousands. The solution is to *paginate* the users, so that (for example) only 30 show up on a page at any one time.

There are several pagination methods in Rails; we'll use one of the simplest and most robust, called `will_paginate`. It comes as a gem, which means we could install it using `gem install`; on the other hand, unlike previous gems such as RSpec and Factory Girl, `will_paginate` is needed in the application itself, not just the tests. There is a convenient way to indicate this requirement, and include the gem at the same time, using `config.gem` in the `environment.rb` file (Listing 10.24).⁹

Listing 10.24. Code to include the `will_paginate` gem.

```
config/environment.rb
```

```
.
.
.

Rails::Initializer.run do |config|
  .
  .
  .
  config.gem 'will_paginate', :version => '2.3.12'
end
```

With the configuration line in Listing 10.24, we can install the required gem dependencies as follows:

```
$ [sudo] rake gems:install
```

If you're also deploying the sample application to Heroku, you'll need to include `will_paginate` in the *gems manifest*, which involves creating a file called `.gems` in the Rails root.¹⁰ Navigate to the Rails root (if you're not already there) and make a new `.gems` file using your favorite text editor:

```
$ cd ~/rails_projects/sample_app
$ mate .gems
```

Then fill the `.gems` with contents of Listing 10.25.

Listing 10.25. The Heroku gems manifest.

```
.gems
```

```
will_paginate --version 2.3.12
faker
```

⁹In Rails 3, the `config.gem` method for including gems will be replaced by `Bundler`.

¹⁰The Heroku gems manifest will also be replaced (or, rather, supplemented) by Bundler in Rails 3.

(Here we've included the `faker` gem from [Section 10.3.2](#) since otherwise Heroku complains when running `heroku rake db:migrate`. The other gems in our application, such as Factory Girl and Webrat, are only needed for tests and so are not needed on the Heroku server.)

With `will_paginate` installed, we are now ready to paginate the results of finding users. We'll start by adding the special `will_paginate` method in the view ([Listing 10.26](#)); we'll see in a moment why the code appears both above and below the user list.

Listing 10.26. The user index with pagination.

`app/views/users/index.html.erb`

```
<h2>All users</h2>

<%= will_paginate %>

<ul class="users">
  <% @users.each do |user| %>
    <li>
      <%= gravatar_for user, :size => 30 %>
      <%= link_to h(user.name), user %>
    </li>
  <% end %>
</ul>

<%= will_paginate %>
```

The `will_paginate` method is a little magical; inside a `users` view, it automatically looks for an `@users` object, and then displays pagination links to access other pages. The view in [Listing 10.26](#) doesn't work yet, though, because currently `@users` contains the results of `User.all` ([Listing 10.19](#)), which is of class Array, whereas `will_paginate` expects an object of class `WillPaginate::Collection`. Happily, this is just the kind of object returned by the `paginate` method supplied by the `will_paginate` gem:

```
$ script/console
>> User.all.class
=> Array
>> User.paginate(:page => 1).class
=> WillPaginate::Collection
```

Note that `paginate` takes a hash argument with key `:page` and value equal to the page requested. `User.paginate` pulls the users out of the database one chunk at a time (30 by default), based on the `:page` parameter. So, for example, page 1 is users 1–30, page 2 is users 31–60, etc.

We can paginate the users in the sample application by using `paginate` in place of `all` in the `index` action ([Listing 10.27](#)). Here the `:page` parameter comes from `params[:page]`, which is generated automatically by `will_paginate`.

Listing 10.27. Paginating the users in the `index` action.

`app/controllers/users_controller.rb`

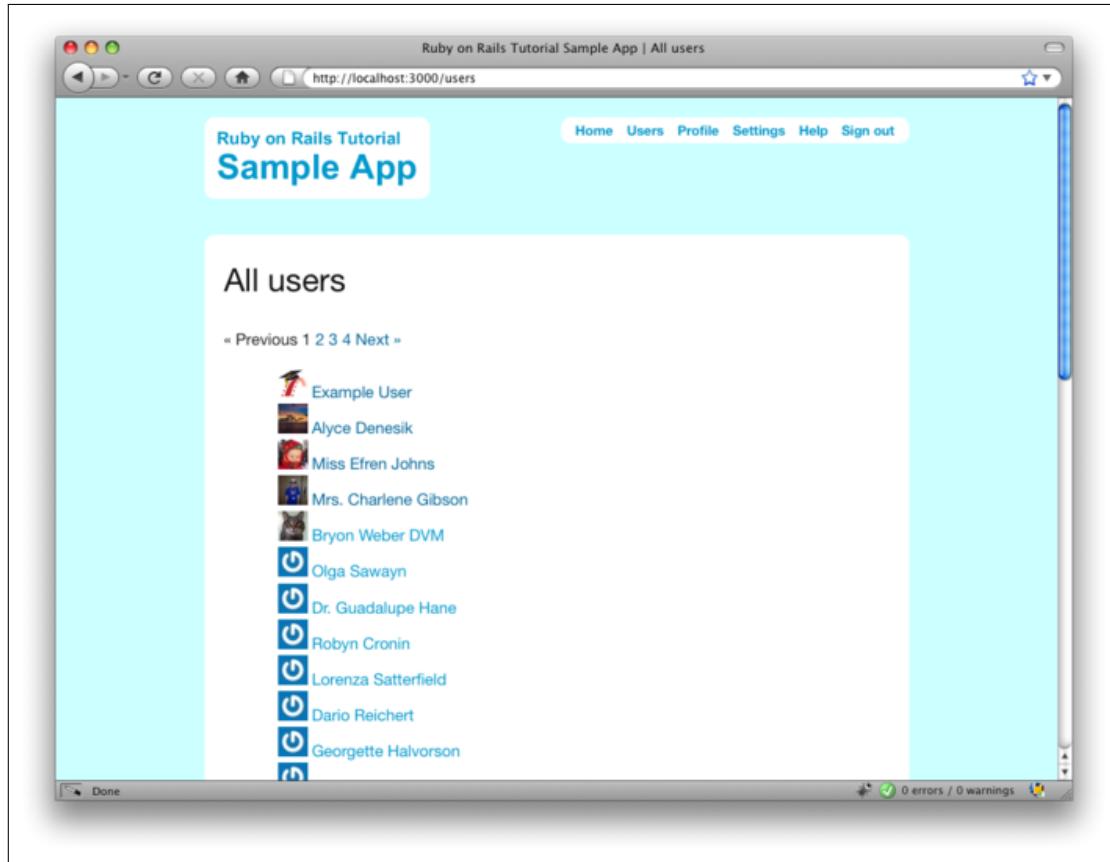


Figure 10.10: The user index page `/users` with pagination. ([full size](#))

```
class UsersController < ApplicationController
  before_filter :authenticate, :only => [:index, :edit, :update]
  .

  .

  def index
    @title = "All users"
    @users = User.paginate(:page => params[:page])
  end
  .

  .
end
```

The user index page should now be working, appearing as in Figure 10.10; because we included `will_paginate` both above and below the user list, the pagination links appear in both places.

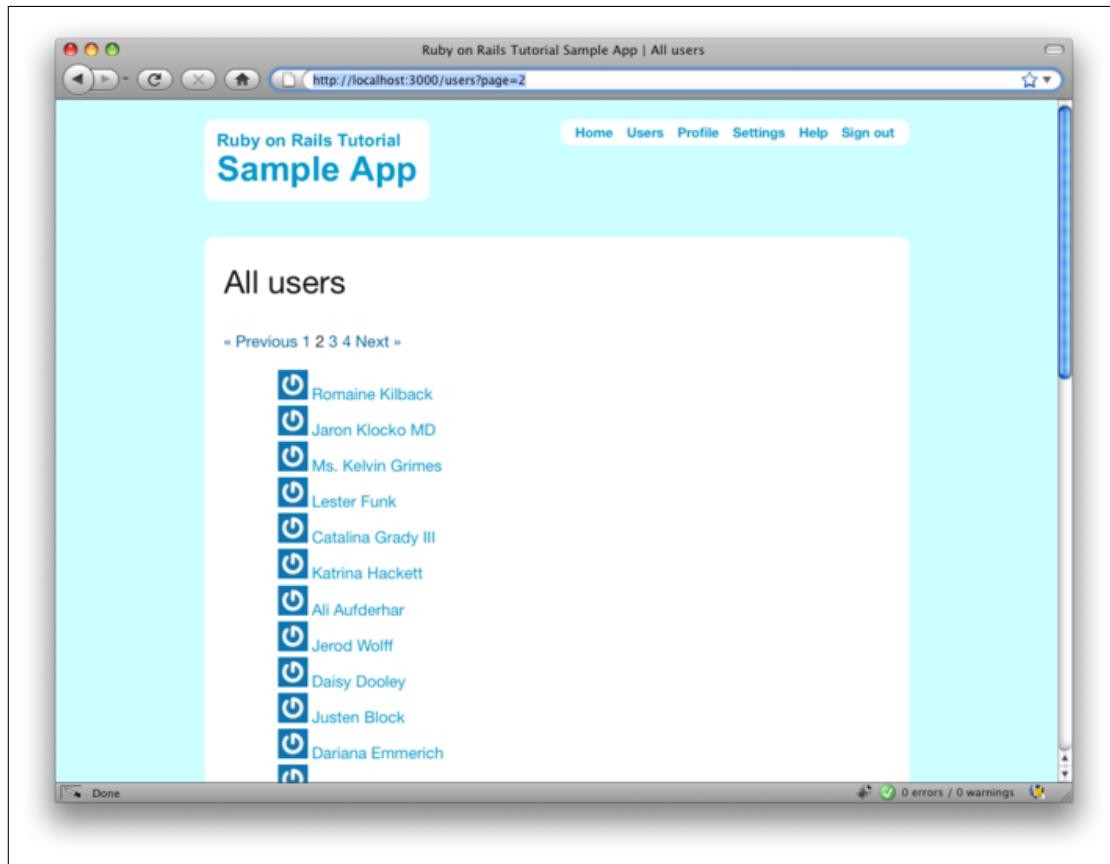


Figure 10.11: Page 2 of the user index (</users?page=2>). ([full size](#))

If you now click on either the [2](#) link or [Next](#) link, you'll get the second page of results, as shown in Figure 10.11.

Testing pagination

Testing pagination requires detailed knowledge of how `will_paginate` works, so we did the implementation first, but it's still a good idea to test it. To do this, we need to invoke pagination in a test, which means making more than 30 (factory) users.

As before, we'll use Factory Girl to simulate users, but immediately we have a problem: user email addresses must be unique, which would appear to require creating more than 30 users by hand—a terribly cumbersome job. Fortunately, Factory Girl anticipates this issue, and provides *sequences* to solve it, as shown in Listing 10.28.

Listing 10.28. Defining a Factory Girl sequence.

`spec/factories.rb`

```
Factory.define :user do |user|
  user.name          "Michael Hartl"
  user.email         "mhartl@example.com"
  user.password      "foobar"
  user.password_confirmation "foobar"
end

Factory.sequence :email do |n|
  "person-#{n}@example.com"
end
```

This arranges to return email addresses like `person-1@example.com`, `person-2@example.com`, etc., which we invoke using the `next` method:

```
Factory(:user, :email => Factory.next(:email))
```

Applying the idea of factory sequences, we can make 31 users (the original `@user` plus 30 more) inside a test, and then verify that the response has the HTML expected from `will_paginate` (which you should be able to determine using Firebug or by viewing the page source). The result appears in Listing 10.29.

Listing 10.29. A test for pagination.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe "UsersController" do
  integrate_views

  describe "GET 'index'" do
    .
    .
    .
    describe "for signed-in users" do
      .
      .
      .
      it "should paginate users" do
        30.times { Factory(:user, :email => Factory.next(:email)) }
        get :index
        response.should have_tag("div.pagination")
        response.should have_tag("span", "&laquo; Previous")
        response.should have_tag("span", "1")
        response.should have_tag("a[href=?]", "/users?page=2", "2")
        response.should have_tag("a[href=?]", "/users?page=2", "Next &raquo;")
      end
    end
  end
end
```

```
.
.
.
end
```

Note the compact notation `have_tag("div.pagination")`, which borrows the class convention from CSS (first seen in Listing 5.3) to check for a `div` tag with class `pagination`.

With that, our pagination code is well-tested, and there's only one minor detail left.

10.3.4 Partial refactoring

The paginated user index is now complete, but there's one improvement I can't resist including: Rails has some incredibly slick tools for making compact views, and in this section we'll refactor the index page to use them. Because our code is well-tested, we can refactor with confidence, assured that we are unlikely to break our site's functionality.

The first step in our refactoring is to replace the user `li` from Listing 10.26 with a `render` call (Listing 10.30).

Listing 10.30. The first refactoring attempt at the index view.

`app/views/users/index.html.erb`

```
<h2>All users</h2>

<%= will_paginate %>

<ul class="users">
  <% @users.each do |user| %>
    <%= render user %>
  <% end %>
</ul>

<%= will_paginate %>
```

Here we call `render` not on a string with the name of a partial, but rather on a `user` variable of class `User`:¹¹ in this context, Rails automatically looks for a partial called `_user.html.erb`, which we must create (Listing 10.31).

Listing 10.31. A partial to render a single user.

`app/views/users/_user.html.erb`

```
<li>
  <%= gravatar_for user, :size => 30 %>
  <%= link_to h(user.name), user %>
</li>
```

¹¹The name `user` is immaterial—we could have written `@users.each do |foobar|` and then used `render foobar`. The key is the `class` of the object—in this case, `User`.

This is a definite improvement, but we can do even better: we can call `render` directly on the `@users` variable ([Listing 10.32](#)).

Listing 10.32. The fully refactored user index.

`app/views/users/index.html.erb`

```
<h2>All users</h2>

<%= will_paginate %>

<ul class="users">
  <%= render @users %>
</ul>

<%= will_paginate %>
```

Here Rails infers that `@users` is an array of `User` objects; moreover, when called with a collection of users, Rails automatically iterates through them and renders each one with the `_user.html.erb` partial. The result is the impressively compact code in [Listing 10.32](#).

10.4 Destroying users

Now that the user index is complete, there's only one canonical REST action left: `destroy`. In this section, we'll add links to delete users, as mocked up in [Figure 10.12](#), and define the `destroy` action necessary to accomplish the deletion. But first, we'll create the class of administrative users authorized to do so.

10.4.1 Administrative users

We will identify privileged administrative users with a boolean `admin` attribute in the User model, which will lead to an `admin?` method to test for admin status. Following the model of the tests for “remember me” in [Listing 9.13](#), we can write tests for this attribute as in [Listing 10.33](#).

Listing 10.33. Tests for an `admin` attribute.

`spec/models/user_spec.rb`

```
.
.
.

describe "admin attribute" do

  before(:each) do
    @user = User.create! (@attr)
  end

  it "should respond to admin" do
    @user.should respond_to(:admin)
  end
end
```

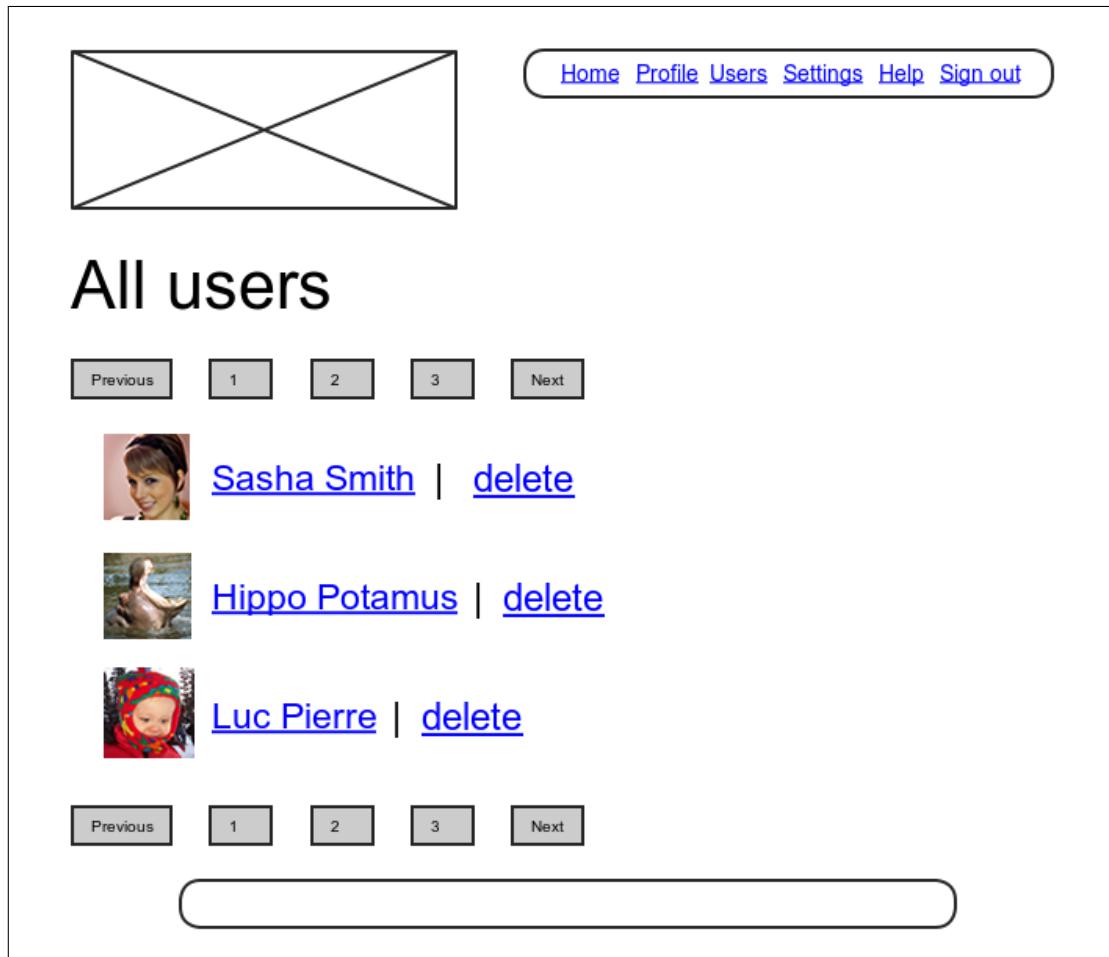


Figure 10.12: A mockup of the user index with delete links. ([full size](#))

```
it "should not be an admin by default" do
  @user.should_not be_admin
end

it "should be convertible to an admin" do
  @user.toggle!(:admin)
  @user.should be_admin
end
end
end
```

Here we've used the `toggle!` method to flip the `admin` attribute from `true` to `false`. Also note that the line

```
@user.should be_admin
```

implies (via the RSpec boolean convention) that the user should have an `admin?` boolean method.

We add the `admin` attribute with a migration as usual, indicating the `boolean` type on the command line:

```
$ script/generate migration add_admin_to_users admin:boolean
```

The migration simply adds the `admin` column to the `users` table (Listing 10.34), yielding the data model in Figure 10.13.

Listing 10.34. The migration to add a boolean `admin` attribute to users.

```
db/migrate/<timestamp>.add_admin_to_users.rb
```

```
class AddAdminToUsers < ActiveRecord::Migration
  def self.up
    add_column :users, :admin, :boolean
  end

  def self.down
    remove_column :users, :admin
  end
end
```

Finally, we migrate the development database and prepare the test database:

```
$ rake db:migrate
$ rake db:test:prepare
```

users	
id	integer
name	string
email	string
encrypted_password	string
salt	string
remember_token	string
admin	boolean
created_at	datetime
updated_at	datetime

Figure 10.13: The User model with an added `admin` boolean attribute.

As expected, Rails figures out the boolean nature of the `admin` attribute and automatically adds the question-mark method `admin?`:

```
$ script/console
>> user = User.first
>> user.admin?
=> false
>> user.toggle!(:admin)
=> true
>> user.admin?
=> true
```

As a final step, let's update our sample data populator to make the first user an admin (Listing 10.35).

Listing 10.35. The sample data populator code with an admin user.

`lib/tasks/sample_data.rake`

```
require 'faker'

namespace :db do
  desc "Fill database with sample data"
  task :populate => :environment do
    Rake::Task['db:reset'].invoke
    admin = User.create!(:name => "Example User",
                         :email => "example@railstutorial.org",
                         :password => "foobar",
                         :password_confirmation => "foobar")
    admin.toggle!(:admin)
  end
end
```

```
    .
  end
end
```

Finally, re-run the populator to reset the database and then rebuild it from scratch:

```
$ rake db:populate
```

Revisiting attr_accessible

You might have noticed that Listing 10.35 makes the user an admin with `toggle! (:admin)`, but why not just add `:admin => true` to the initialization hash? The answer is, it won't work, and this is by design: only `attr_accessible` attributes can be assigned through mass assignment, and the `admin` attribute isn't accessible. Listing 10.36 reproduces the most recent list of `attr_accessible` attributes—note that `:admin` is *not* on the list.

Listing 10.36. A review of the `attr_accessible` attributes in the User model.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessor :password
  attr_accessible :name, :email, :password, :password_confirmation
  .
  .
  .
end
```

Explicitly defining accessible attributes is crucial for good site security. If we omitted the `attr_accessible` list in the User model (or foolishly added `:admin` to the list), a malicious user could send a PUT request as follows:¹²

```
put /users/17?admin=1
```

This request would make user 17 an admin, which could be a potentially serious security breach, to say the least. Because of this danger, it is a good practice to define `attr_accessible` for every model.

10.4.2 The `destroy` action

The final step needed to complete the Users resource is to add delete links and a `destroy` action. We'll start by adding a delete link for each user on the user index page (Listing 10.37).

¹²Command-line tools such as `curl` (seen in Box 3.1) can issue PUT requests of this form.

Listing 10.37. User delete links (viewable only by admins).

`app/views/users/_user.html.erb`

```
<li>
  <%= gravatar_for user, :size => 30 %>
  <%= link_to h(user.name), user %>
  <% if current_user.admin? %>
    | <%= link_to "delete", user, :method => :delete, :confirm => "You sure?" %>
    <% end %>
</li>
```

Note the `:method => :delete` argument, which arranges for the link to issue the necessary DELETE request;¹³ we've also wrapped each link inside an `if` statement so that only admins can see them. The result for our admin user appears in [Figure 10.14](#).

Even though only admins can see the delete links, there's still a terrible security breach: any sufficiently sophisticated attacker could simply issue DELETE requests from the command line and delete any user on the site. To secure the site properly, we also need access control, so our tests should check not only that admins *can* delete users, but also that other users *can't*. The results appear in [Listing 10.38](#). Note that, in analogy with the `get`, `post`, and `put` methods, we use `delete` to issue DELETE requests inside of tests.

Listing 10.38. Tests for destroying users.

`spec/controllers/users_controller_spec.rb`

```
describe UsersController do
  integrate_views
  .

  .

  .

  describe "DELETE 'destroy'" do
    before(:each) do
      @user = Factory(:user)
    end

    describe "as a non-signed-in user" do
      it "should deny access" do
        delete :destroy, :id => @user
        response.should redirect_to(signin_path)
      end
    end

    describe "as a non-admin user" do
      it "should protect the page" do
        test_sign_in(@user)
```

¹³Web browsers can't send DELETE requests natively, so Rails fakes them with JavaScript. This means that the delete links won't work if the user has JavaScript disabled. For the purposes of this tutorial, we can just require our site admins to enable JavaScript, but if you must support non-JavaScript-enabled browsers you can fake a DELETE request using a form and a POST request, which works even without JavaScript; see the Railscast on [Destroy Without JavaScript](#) for details.

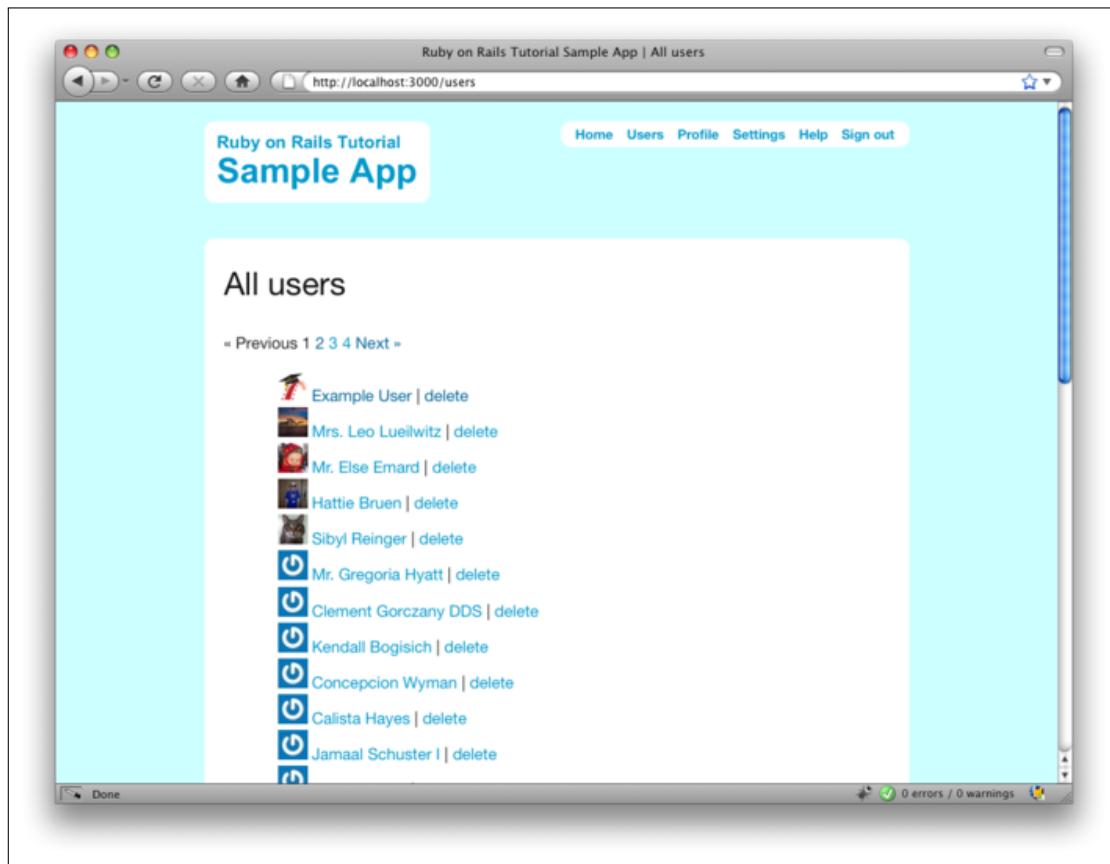


Figure 10.14: The user index `/users` with delete links. ([full size](#))

```

    delete :destroy, :id => @user
    response.should redirect_to(root_path)
end
end

describe "as an admin user" do

  before(:each) do
    admin = Factory(:user, :email => "admin@example.com", :admin => true)
    test_sign_in(admin)
    User.should_receive(:find).with(@user).and_return(@user)
    @user.should_receive(:destroy).and_return(@user)
  end

  it "should destroy the user" do
    delete :destroy, :id => @user
    response.should redirect_to(users_path)
  end
end
end
end

```

(You might notice that we've set an admin user using `:admin => true`; user factories are not bound by the rules of `attr_accessible` parameters.)

As you might suspect by now, the implementation uses a before filter, this time to restrict access to the `destroy` action to admins. The `destroy` action itself finds the user, destroys it, and then redirects to user index (Listing 10.39).

Listing 10.39. A before filter restricting the `destroy` action to admins.

`app/controllers/users_controller.rb`

```

class UsersController < ApplicationController
  before_filter :authenticate, :only => [:index, :edit, :update, :destroy]
  before_filter :correct_user, :only => [:edit, :update]
  before_filter :admin_user,   :only => :destroy
  .

  .

  .

  def destroy
    User.find(params[:id]).destroy
    flash[:success] = "User destroyed."
    redirect_to users_path
  end

  private
  .

  .

  .

  def admin_user
    redirect_to(root_path) unless current_user.admin?
  end
end

```

```
    end  
end
```

Note that the `destroy` action uses *method chaining* (seen briefly in [Section 4.2.3](#)) in the line

```
User.find(params[:id]).destroy
```

which saves a line of code.

At this point, all the tests should be passing, and the Users resource—with its controller, model, and views—is functionally complete.

10.5 Conclusion

We've come a long way since introducing the Users controller way back in [Section 5.3](#). Those users couldn't even sign up; now users can sign up, sign in, sign out, view their profiles, edit their settings, and see an index of all users—and some can even destroy other users.

The rest of this book builds on the foundation of the Users resource (and associated authentication system) to make a site with Twitter-like microposts ([Chapter 11](#)) and user following ([Chapter 12](#)). These chapters will introduce some of the most powerful features of Rails, including data modeling with `has_many` and `has-many :through`.

Before moving on, be sure to merge all the changes into the master branch:

```
$ git add .  
$ git commit -am "Done with user edit/update, index, and destroy actions"  
$ git checkout master  
$ git merge updating-users
```

10.6 Exercises

- Arrange for the Gravatar “change” link in [Listing 10.3](#) to open in a new window (or tab). *Hint:* Search the web; you should find one particularly robust method involving something called `blank`. This solution is not valid HTML (so if you installed the HTML Validator in [Chapter 1](#) it will indicate an error), but it works in every browser I've ever seen.
- Remove the duplicated form code by refactoring the `new.html.erb` and `edit.html.erb` views to use the partial in [Listing 10.40](#). Note that you will have to pass the form variable `f` explicitly as a local variable, as shown in [Listing 10.41](#).
- Signed-in users have no reason to access the `new` and `create` actions in the Users controller. Arrange for such users to be redirected to the root url if they do try to hit those pages.
- Add tests to check that the delete links in [Listing 10.37](#) appear for admins but not for normal users.
- Modify the `destroy` action to prevent admin users from destroying themselves. (Write a test first.)

Listing 10.40. A partial for the new and edit form fields.

`app/views/users/_fields.html.erb`

```
<%= f.error_messages %>


<%= f.label :name %><br />
  <%= f.text_field :name %>



<%= f.label :email %><br />
  <%= f.text_field :email %>



<%= f.label :password %><br />
  <%= f.password_field :password %>



<%= f.label :password_confirmation, "Confirmation" %><br />
  <%= f.password_field :password_confirmation %>


```

Listing 10.41. The new user view with partial.

`app/views/users/new.html.erb`

```
<h2>Sign up</h2>

<% form_for(@user) do |f| %>
  <%= render :partial => 'fields', :locals => { :f => f } %>
  <div class="actions">
    <%= f.submit "Sign up" %>
  </div>
<% end %>
```

Chapter 11

User microposts

Chapter 10 saw the completion of the REST actions for the Users resource, so the time has finally come to add a second resource: user *microposts*.¹ These are short messages associated with a particular user, first seen in larval form in Chapter 2. In this chapter, we will make a full-strength version of the sketch from Section 2.3 by constructing the Micropost data model, associating it with the User model using the `has-many` and `belongs_to` methods, and then making the forms and partials needed to manipulate and display the results. In Chapter 12, we will complete our tiny Twitter clone by adding the notion of *following* users in order to receive a *feed* of their microposts.

If you’re using Git for version control, I suggest making a topic branch as usual:

```
$ git checkout -b user-microposts
```

11.1 A Micropost model

We begin the Microposts resource by creating a Micropost model, which captures the essential characteristics of microposts. What follows builds on the work from Section 2.3; as with the model in that section, our new Micropost model will include data validations and an association with the User model. Unlike that model, the present Micropost model will be fully tested, and will also have a default *ordering* and automatic *destruction* if its parent user is destroyed.

11.1.1 The basic model

The Micropost model needs only two attributes: a `content` attribute to hold the micropost’s content,² and a `user_id` to associate a micropost with a particular user. As with the case of the User model (Listing 6.1), we generate it using `generate rspec_model`:

```
$ script/generate rspec_model Micropost content:string user_id:integer
```

¹Technically, we treated sessions as a resource in Chapter 9, but they are not saved to the database the way users and microposts are.

²The `content` attribute will be a `string`, but, as noted briefly in Section 2.1.2, for longer text fields you should use the `text` data type.

This produces a migration to create a `microposts` table in the database (Listing 11.1); compare it to the analogous migration for the `users` table from Listing 6.2.

Listing 11.1. The Micropost migration. (Note the index on `user_id`.)

`db/migrate/<timestamp>.create_microposts.rb`

```
class CreateMicroposts < ActiveRecord::Migration
  def self.up
    create_table :microposts do |t|
      t.string :content
      t.integer :user_id

      t.timestamps
    end
    add_index :microposts, :user_id
  end

  def self.down
    drop_table :microposts
  end
end
```

Note that, since we expect to retrieve all the microposts associated with a given user id, Listing 11.1 adds an index (Box 6.1) on the `user_id` column:

```
add_index :microposts, :user_id
```

Note also the `t.timestamps` line, which (as mentioned in Section 6.1.1) adds the magic `created_at` and `updated_at` columns. We'll put the `created_at` column to work in Section 11.1.3 and Section 11.2.1.

We can run the microposts migration as usual (taking care to prepare the test database since the data model has changed):

```
$ rake db:migrate
$ rake db:test:prepare
```

The result is a Micropost model with the structure shown in Figure 11.1.

Accessible attribute

Before fleshing out the Micropost model, it's important first to use `attr_accessible` to indicate the attributes editable through the web. As discussed in Section 6.1.2 and Section 10.4.1, failing to define accessible attributes means that anyone could change any aspect of a micropost object simply by using a command-line client to issue malicious requests. For example, a malicious user could change the `user_id` attributes on microposts, thereby associating microposts with the wrong users.

In the case of the Micropost model, there is only *one* attribute that needs to be editable through the web, namely, the `content` attribute (Listing 11.2).

microposts	
id	integer
content	string
user_id	integer
created_at	datetime
updated_at	datetime

Figure 11.1: The Micropost data model. ([full size](#))**Listing 11.2.** Making the `content` attribute (and *only* the `content` attribute) accessible.`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  attr_accessible :content
end
```

Since `user_id` isn't listed as an `attr_accessible` parameter, it can't be edited through the web, because a `user_id` parameter in a mass assignment such as

```
Micropost.new(:content => "foo bar", :user_id => 17)
```

will simply be ignored.

The `attr_accessible` declaration in Listing 11.2 is necessary for site security, but it introduces a problem in the default Micropost model spec (Listing 11.3).

Listing 11.3. The initial (lightly edited) Micropost spec.`spec/models/micropost_spec.rb`

```
require 'spec_helper'

describe Micropost do
  before(:each) do
    @attr = {
      :content => "value for content",
      :user_id => 1
    }
  end

  it "should create a new instance given valid attributes" do
```

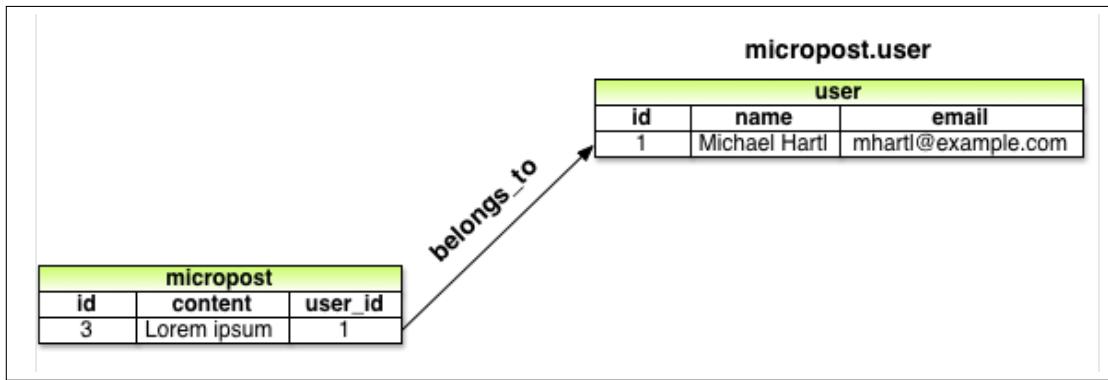


Figure 11.2: The `belongs_to` relationship between a micropost and its user. ([full size](#))

```

Micropost.create!(@attr)
end
end

```

(As in Listing 6.9, I've replaced `@valid_attributes` with the more succinct `@attr`.) This test currently passes, but there's something fishy about it. (See if you can figure out what before proceeding.)

The problem is that the `before(:each)` block in Listing 11.3 assigns the user id through mass assignment, which is exactly what `attr_accessible` is designed to prevent; in particular, as noted above, the `:user_id => 1` part of the initialization hash is simply ignored. The solution is to avoid using `Micropost.new` directly; instead, we will create the new micropost through its *association* with the User model, which sets the user id automatically. Accomplishing this is the task of the next section.

11.1.2 User/Micropost associations

The goal of this section is to establish an *association* between the Micropost model and the User model—a relationship seen briefly in Section 2.3.3 and shown schematically in Figure 11.2 and Figure 11.3. Along the way, we'll write tests for the Micropost model that, unlike Listing 11.3, are compatible with the use of `attr_accessible` in Listing 11.2.

We start with tests for the Micropost model association. First, we want to replicate the `Micropost.create!` test shown in Listing 11.3 without the invalid mass assignment. Second, we see from Figure 11.2 that a `micropost` object should have a `user` method. Finally, `micropost.user` should be the user corresponding to the micropost's `user_id`. We can express these requirements in RSpec with the code in Listing 11.4.

Listing 11.4. Tests for the micropost's user association.

`spec/models/micropost_spec.rb`

```

require 'spec_helper'

describe Micropost do

```

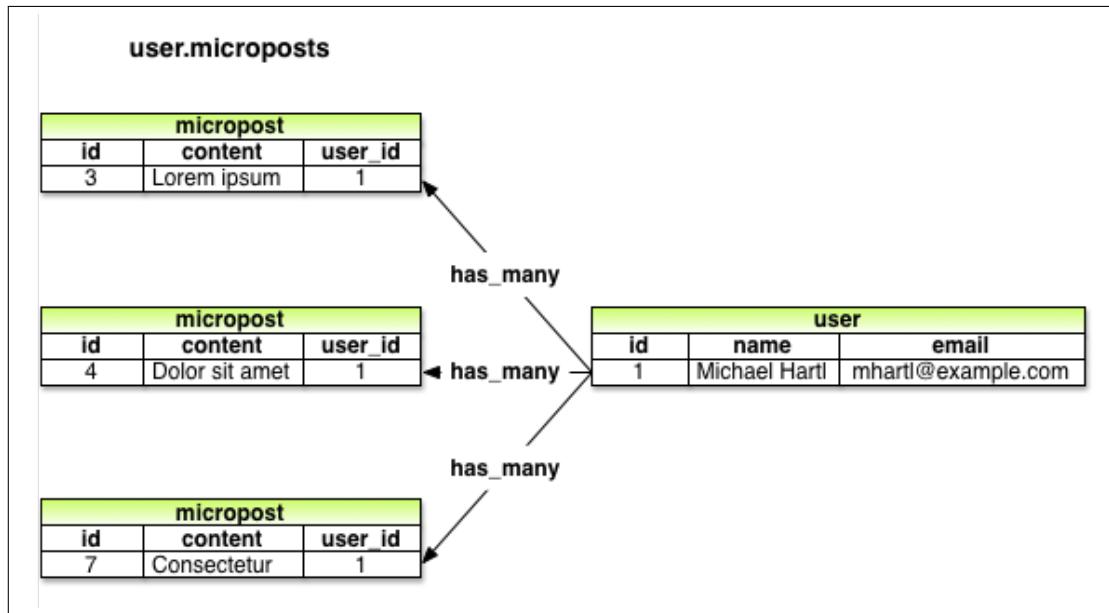


Figure 11.3: The `has_many` relationship between a user and its microposts. ([full size](#))

```

before(:each) do
  @user = Factory(:user)
  @attr = { :content => "value for content" }
end

it "should create a new instance given valid attributes" do
  @user.microposts.create!(@attr)
end

describe "user associations" do

  before(:each) do
    @micropost = @user.microposts.create(@attr)
  end

  it "should have a user attribute" do
    @micropost.should respond_to(:user)
  end

  it "should have the right associated user" do
    @micropost.user_id.should == @user.id
    @micropost.user.should == @user
  end
end
end

```

Note that, rather than using `Micropost.create` or `Micropost.create!` to create a micropost, Listing 11.4 uses

```
@user.microposts.create(@attr)
```

and

```
@user.microposts.create!(@attr)
```

This pattern is the [canonical](#) way to create a micropost through its association with users. (We use a factory user because these tests are for the Micropost model, not the User model.) When created in this way, the micropost object *automatically* has its `user_id` set to the right value, which fixes the issue noted in Section 11.1.1. In particular, the code

```
before(:each) do
  @attr = {
    :content => "value for content",
    :user_id => 1
  }
end

it "should create a new instance given valid attributes" do
  Micropost.create!(@attr)
end
```

from Listing 11.3 is defective because `:user_id => 1` does nothing when `user_id` is not one of the Micropost model's accessible attributes. By going through the user association, on the other hand, the code

```
it "should create a new instance given valid attributes" do
  @user.microposts.create!(@attr)
end
```

from Listing 11.4 has the right `user_id` by construction.

These special `create` methods won't work yet; they require the proper `has_many` association in the User model. We'll defer the more detailed tests for this association to Section 11.1.3; for now, we'll simply test for the presence of a `microposts` attribute (Listing 11.5).

Listing 11.5. A test for the user's `microposts` attribute.

`spec/models/user_spec.rb`

```
require 'spec_helper'
```

```
describe User do
  .
  .
  .
  describe "micropost associations" do
    before(:each) do
      @user = User.create(@attr)
    end

    it "should have a microposts attribute" do
      @user.should respond_to(:microposts)
    end
  end
end
```

We can get the tests in both Listing 11.4 and Listing 11.5 to pass using the `belongs_to/has_many` association illustrated in Figure 11.2 and Figure 11.3, as shown in Listing 11.6 and Listing 11.7.

Listing 11.6. A micropost `belongs_to` a user.

`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  attr_accessible :content

  belongs_to :user
end
```

Listing 11.7. A user `has_many` microposts.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  attr_accessor :password
  attr_accessible :name, :email, :password, :password_confirmation

  has_many :microposts
  .
  .
  .
end
```

Using this `belongs_to/has_many` association, Rails constructs the methods shown in Table 11.1. You should compare the entries in Table 11.1 with the code in Listing 11.4 and Listing 11.5 to satisfy yourself that you understand the basic nature of the associations. (There is one method in Table 11.1 we haven't used so far, the `build` method; it will be put to good use in Section 11.1.4 and especially in Section 11.3.2.)

Method	Purpose
<code>micropost.user</code>	Return the User object associated with the micropost.
<code>user.microposts</code>	Return an array of the user's microposts.
<code>user.microposts.create(arg)</code>	Create a micropost (<code>user.id = user.id</code>).
<code>user.microposts.create!(arg)</code>	Create a micropost (exception on failure).
<code>user.microposts.build(arg)</code>	Return a new Micropost object (<code>user.id = user.id</code>).

Table 11.1: A summary of user/micropost association methods.

11.1.3 Micropost refinements

The test in Listing 11.5 of the `has_many` association doesn't test for much—it merely verifies the *existence* of a `microposts` attribute. In this section, we'll add *ordering* and *dependency* to microposts, while also testing that the `user.microposts` method actually returns an array of microposts.

We will need to construct some microposts in the User model spec, which means that we should make a micropost factory at this point. To do this, we need a way to make an association in Factory Girl. Happily, this is easy—we just use the Factory Girl method `micropost.association`, as seen in Listing 11.8.³

Listing 11.8. The complete factory file, including a new factory for microposts.

`spec/factories.rb`

```
# By using the symbol ':user', we get Factory Girl to simulate the User model.
Factory.define :user do |user|
  user.name          "Michael Hartl"
  user.email         "mhartl@example.com"
  user.password      "foobar"
  user.password_confirmation "foobar"
end

Factory.sequence :email do |n|
  "person-#{n}@example.com"
end

Factory.define :micropost do |micropost|
  micropost.content "Foo bar"
  micropost.association :user
end
```

Default scope

We can put the micropost factory to work in a test for the ordering of microposts. By default, using `user.microposts` to pull a user's microposts from the database makes no guarantees about the order of the posts, but (following the convention of blogs and Twitter) we want the microposts to come out in reverse order of when they were created, i.e., most recent first. To test this ordering, we first create a couple of microposts as follows:

³For more on Factory Girl associations, including the many options available, see the [Factory Girl documentation](#).

```
@mp1 = Factory(:micropost, :user => @user, :created_at => 1.day.ago)
@mp2 = Factory(:micropost, :user => @user, :created_at => 1.hour.ago)
```

Here we indicate that the second post was created more recently, `1.hour.ago`, with the first post created `1.day.ago`. Note how convenient the use of Factory Girl is: not only can we assign the user using mass assignment (since factories bypass `attr_accessible`), we can also set `created_at` manually, which Active Record won't allow us to do.⁴

Most database adapters (including the one for SQLite) return the microposts in order of their ids, so we can arrange for an initial test that almost certainly fails using the code in Listing 11.9.

Listing 11.9. Testing the order of a user's microposts.

`spec/models/user_spec.rb`

```
require 'spec_helper'

describe User do
  .
  .
  .
  describe "micropost associations" do
    before(:each) do
      @user = User.create(@attr)
      @mp1 = Factory(:micropost, :user => @user, :created_at => 1.day.ago)
      @mp2 = Factory(:micropost, :user => @user, :created_at => 1.hour.ago)
    end

    it "should have a microposts attribute" do
      @user.should respond_to(:microposts)
    end

    it "should have the right microposts in the right order" do
      @user.microposts.should == [@mp2, @mp1]
    end
  end
end
```

The key line here is

```
@user.microposts.should == [@mp2, @mp1]
```

indicating that the posts should be ordered newest first. This should fail because by default the posts will be ordered by id, i.e., `[@mp1, @mp2]`. This test also verifies the basic correctness of the `has_many` association itself, by checking (as indicated in Table 11.1) that `user.microposts` is an array of microposts.

⁴Recall that `created_at` and `updated_at` are “magic” columns, so any explicit initialization values are overwritten by the magic.

To get the ordering test to pass, we use a Rails facility called `default_scope` with an `:order` parameter, as shown in Listing 11.10. (This is our first example of the notion of *scope*. We will learn about scope in a more general context in Chapter 12.)

Listing 11.10. Ordering the microposts with `default_scope`.

`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  .
  .
  .
  default_scope :order => 'created_at DESC'
end
```

The order here is '`created_at DESC`', where `DESC` is SQL for “descending”, i.e., in descending order from newest to oldest.

Dependent: destroy

Apart from proper ordering, there is a second refinement we’d like to add to microposts. Recall from Section 10.4 that site administrators have the power to *destroy* users. It stands to reason that if a user is destroyed, the user’s microposts should be destroyed as well. We can test for this by first destroying a micropost’s user and then verifying that the associated microposts are no longer in the database (Listing 11.11).

Listing 11.11. Testing that microposts are destroyed when users are.

`spec/models/user_spec.rb`

```
describe User do
  .
  .
  .
  describe "micropost associations" do
    before(:each) do
      @user = User.create(@attr)
      @mp1 = Factory(:micropost, :user => @user, :created_at => 1.day.ago)
      @mp2 = Factory(:micropost, :user => @user, :created_at => 1.hour.ago)
    end
    .
    .
    .
    it "should destroy associated microposts" do
      @user.destroy
      [@mp1, @mp2].each do |micropost|
        Micropost.find_by_id(micropost.id).should be_nil
      end
    end
  end
end
```

```
  .
  .
end
```

Here we have used `Micropost.find_by_id`, which returns `nil` if the record is not found, whereas `Micropost.find` raises an exception on failure, which is a bit harder to test for. (In case you're curious,

```
lambda do
  Micropost.find(micropost.id)
end.should raise_error(ActiveRecord::RecordNotFound)
```

does the trick in this case.)

The application code to get Listing 11.11 to pass is less than one line; in fact, it's just an option to the `has_many` association method, as shown in Listing 11.12.

Listing 11.12. Ensuring that a user's microposts are destroyed along with the user.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
  .
  has_many :microposts, :dependent => :destroy
  .
  .
  .
end
```

With that, the final form of the user/micropost association is in place.

11.1.4 Micropost validations

Before leaving the Micropost model, we'll tie off a couple loose ends by adding validations (following the example from Section 2.3.2). Both the `user_id` and `content` attributes are required, and `content` is further constrained to be shorter than 140 characters, which we test for using the code in Listing 11.13.

Listing 11.13. Tests for the Micropost model validations.

`spec/models/micropost_spec.rb`

```
require 'spec_helper'

describe Micropost do
  before(:each) do
    @user = Factory(:user)
    @attr = { :content => "value for content" }
```

```

end
.
.
.
describe "validations" do

  it "should require a user id" do
    Micropost.new(@attr).should_not be_valid
  end

  it "should require nonblank content" do
    @user.microposts.build(:content => " ") .should_not be_valid
  end

  it "should reject long content" do
    @user.microposts.build(:content => "a" * 141) .should_not be_valid
  end
end
end

```

These generally follow the examples from the User model validation tests from [Section 6.2](#). (The analogous tests were broken into multiple lines in that section, but you should be comfortable enough reading RSpec code by now to digest the more compact formulation above.)

As in [Section 6.2](#), the code in [Listing 11.13](#) uses string multiplication to test the micropost length validation:

```

$ script/console
>> "a" * 10
=> "aaaaaaaaaa"
>> "a" * 141
=> "aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa
aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa"

```

In contrast, instead of using the default **new** constructor as in

```
User.new(...)
```

the code in [Listing 11.13](#) uses the **build** method:

```
@user.microposts.build
```

Recall from [Table 11.1](#) that this is essentially equivalent to **Micropost.new**, except that it automatically sets the micropost's **user_id** to **@user.id**.

The validations themselves are straightforward analogues of the User model validations, as seen in [Listing 11.14](#).

Listing 11.14. The Micropost model validations.

`app/models/micropost.rb`

```
class Micropost < ActiveRecord::Base
  attr_accessible :content

  belongs_to :user

  validates_presence_of :content, :user_id
  validates_length_of   :content, :maximum => 140

  default_scope :order => 'created_at DESC'
end
```

This completes the data modeling for users and microposts. It's time now to build the web interface.

11.2 Showing microposts

Although we don't yet have a way to create microposts through the web—that comes in [Section 11.3.2](#)—that won't stop us from displaying them (and testing that display). Following Twitter's lead, we'll plan to display a user's microposts not on a separate microposts `index` page, but rather directly on the user `show` page itself, as mocked up in [Figure 11.4](#). We'll start with fairly simple ERb templates for adding a micropost display to the user profile, and then we'll add microposts to the sample data populator from [Section 10.3.2](#) so that we have something to display.

As with the discussion of the signin machinery in [Section 9.3.2](#), [Section 11.2.1](#) will often push several elements onto the `stack` at a time, and then pop them off one by one. If you start getting bogged down, be patient; there's some nice payoff in [Section 11.2.2](#).

11.2.1 Augmenting the user show page

We begin with a test for displaying the user's microposts. We work in the Users controller spec, since it is the Users controller that contains the user `show` action. Our strategy is to create a couple of factory microposts associated with the user, and then verify that the show page has a `span` tag with CSS class "`content`" containing each post's content. The resulting RSpec example appears in [Listing 11.15](#).

Listing 11.15. A test for showing microposts on the user `show` page.

`spec/controllers/users_controller_spec.rb`

```
require 'spec_helper'

describe UsersController do
  integrate_views
  .
  .
  .
  describe "GET 'show'" do
```

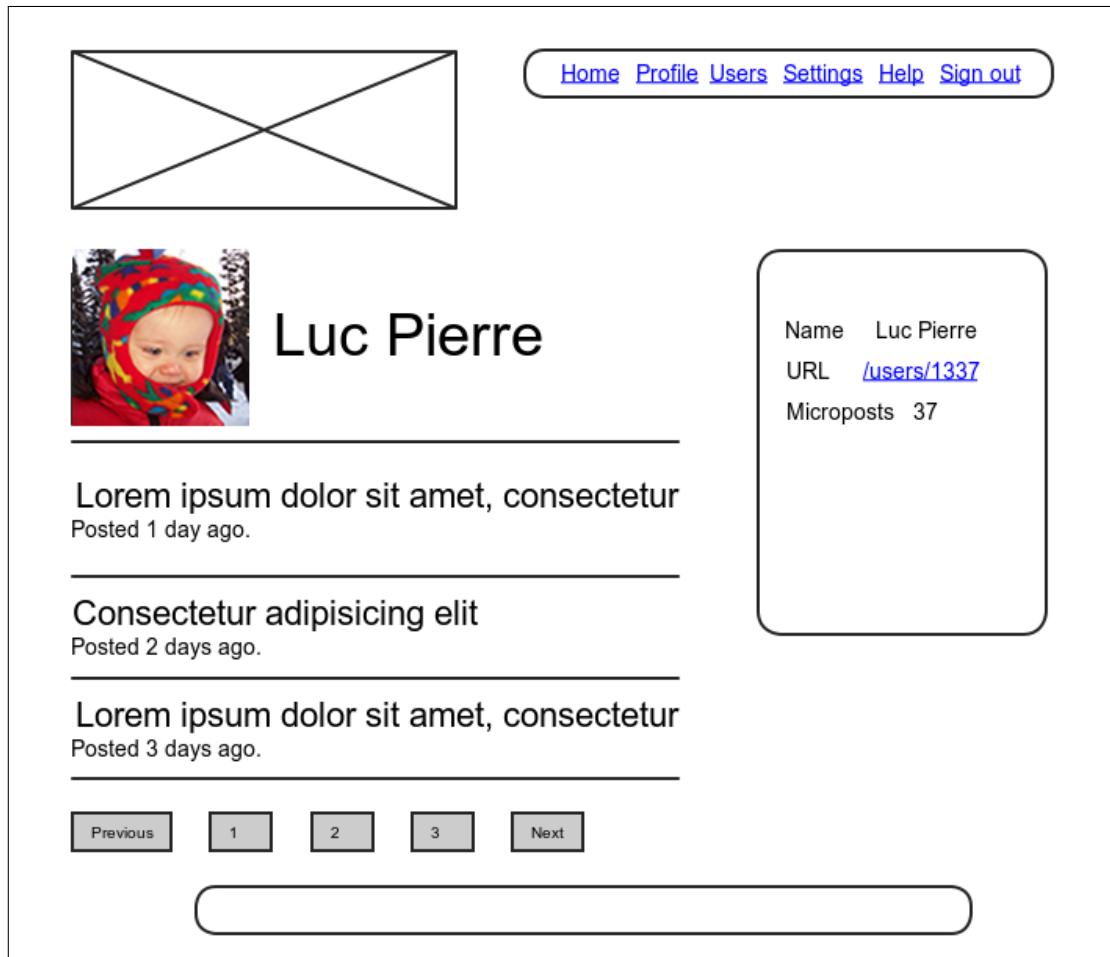


Figure 11.4: A mockup of a profile page with microposts. ([full size](#))

```

before(:each) do
  @user = Factory(:user)
  # Arrange for User.find(params[:id]) to find the right user.
  User.stub!(:find, @user.id).and_return(@user)
end

.
.

it "should show the user's microposts" do
  mp1 = Factory(:micropost, :user => @user, :content => "Foo bar")
  mp2 = Factory(:micropost, :user => @user, :content => "Baz quux")
  get :show, :id => @user
  response.should have_tag("span.content", mp1.content)
  response.should have_tag("span.content", mp2.content)
end
end

.
.

end

```

Although these tests won't pass until Listing 11.17, we'll get started on the application code by inserting a table of microposts into the user profile page, as shown in Listing 11.16.⁵

Listing 11.16. Adding microposts to the user `show` page.

`app/views/users/show.html.erb`

```


|                                                                                                                                                           |                                                                                                                                                                             |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| . . . <% unless @user.microposts.empty? %> <table class="microposts"> &lt;%= render @microposts %&gt; </table> <%= will_paginate @microposts %> <% end %> | <strong>Name</strong> <%= h @user.name %><br /> <strong>URL</strong> <%= link_to user_path(@user), @user %><br /> <strong>Microposts</strong> <%= @user.microposts.count %> |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|


```

⁵In the sense of [semantic markup](#), it would probably be better to use an [ordered list](#), but in that case the vertical alignment of text and images is much more difficult than with tables. See the exercise in Section 11.5 if you insist on struggling with the semantic version. :-)

We'll deal with the microposts **table** momentarily, but there are several other things to note first. One new idea is the use of `empty?` in the line

```
@user.microposts.empty?
```

This applies the `empty?` method, seen before in the context of strings (e.g., [Section 4.2.3](#)), to an array:

```
$ script/console
>> [1, 2].empty?
=> false
>> [].empty?
=> true
```

By using the conditional `unless` clause,

```
<% unless @user.microposts.empty? %>
```

we make sure that an empty table won't be displayed when the user has no microposts.

You'll also note from [Listing 11.16](#) that we've preemptively added pagination for microposts through

```
<%= will_paginate @microposts %>
```

If you compare this with the analogous line on the user index page, [Listing 10.26](#), you'll see that before we had just

```
<%= will_paginate %>
```

This worked because, in the context of the Users controller, `will_paginate` assumes the existence of an instance variable called `@users` (which, as we saw in [Section 10.3.3](#), should be of class `WillPaginate::Collection`). In the present case, since we are still in the Users controller but want to paginate `microposts` instead, we pass an explicit `@microposts` variable to `will_paginate`. Of course, this means that we will have to define such a variable in the user `show` action ([Listing 11.18](#) below).

Finally, note that we have taken this opportunity to add a count of the current number of microposts to the profile sidebar:

```
<td class="sidebar round">
  <strong>Name</strong> <%= h @user.name %><br />
  <strong>URL</strong> <%= link_to user_path(@user), @user %><br />
  <strong>Microposts</strong> <%= @user.microposts.count %>
</td>
```

Here `@user.microposts.count` is the analogue of the `User.count` method, except that it counts the microposts belonging to a given user through the user/micropost association.⁶

Now for the microposts `table` itself:

```
<table class="microposts">
  <%= render @microposts %>
</table>
```

This code is responsible for generating the table of microposts, but you can see that it just defers the heavy lifting to a micropost partial. We saw in [Section 10.3.4](#) that the code

```
<%= render @users %>
```

automatically renders each of the users in the `@users` variable using the `_user.html.erb` partial. Similarly, the code

```
<%= render @microposts %>
```

does exactly the same thing for microposts. This means that we must define a `micropost.html.erb` partial, as shown in [Listing 11.17](#).

Listing 11.17. A partial for showing a single micropost.

`app/views/microposts/_micropost.html.erb`

```
<tr>
  <td class="micropost">
    <span class="content"><%= h micropost.content %></span>
    <span class="timestamp">
      Posted <%= time_ago_in_words(micropost.created_at) %> ago.
    </span>
  </td>
</tr>
```

This uses the awesome `time_ago_in_words` helper method, whose effect we will see in [Section 11.2.2](#).

Thus far, despite defining all the relevant ERb templates, the test in [Listing 11.15](#) should have been failing for want of an `@microposts` variable. We can get it to pass with [Listing 11.18](#).

Listing 11.18. Adding an `@microposts` instance variable to the user `show` action.

`app/controllers/users_controller.rb`

⁶In case you're wondering, the association `count` method is smart, and performs the count directly in the database. In particular, it does *not* pull all the microposts out of the database and then call `length` on the resulting array, as this would become terribly inefficient as the number of microposts grew. Instead, it asks the database to count the microposts with the given `user_id`. By the way, in the unlikely event that finding the count is still a bottleneck in your application, you can make it even faster with a `counter cache`.

```
class UsersController < ApplicationController
  .
  .
  .
  def show
    @user = User.find(params[:id])
    @microposts = @user.microposts.paginate(:page => params[:page])
    @title = CGI.escapeHTML(@user.name)
  end
end
```

Notice here how clever `paginate` is—it even works with the microposts association, converting the array into a `WillPaginate::Collection` object on the fly.

Upon adding the CSS from Listing 11.19 to our `custom.css` stylesheet,⁷ we can get a look at our new user profile page in Figure 11.5. It's rather...disappointing. Of course, this is because there are not currently any microposts. It's time to change that.

Listing 11.19. The CSS for microposts (includes all the CSS for this chapter).

`public/stylesheets/custom.css`

```
.
.
.

h2.micropost {
  margin-bottom: 0;
}

table.microposts tr {
  height: 70px;
}

table.microposts tr td.gravatar {
  border-top: 1px solid #ccc;
  vertical-align: top;
  width: 50px;
}

table.microposts tr td.micropost {
  border-top: 1px solid #ccc;
  vertical-align: top;
  padding-top: 10px;
}

table.microposts tr td.micropost span.timestamp {
  display: block;
  font-size: 85%;
  color: #666;
}
```

⁷For convenience, Listing 11.19 actually has *all* the CSS needed for this chapter.

```
div.user_info img {
  padding-right: 0.1em;
}

div.user_info a {
  text-decoration: none;
}

div.user_info span.user_name {
  position: absolute;
}

div.user_info span.microposts {
  font-size: 80%;
}

form.new_micropost {
  margin-bottom: 2em;
}

form.new_micropost textarea {
  height: 4em;
  margin-bottom: 0;
}
```

11.2.2 Sample microposts

With all the work making templates for user microposts in [Section 11.2.1](#), the ending was rather anticlimactic. We can rectify this sad situation by adding microposts to the sample populator from [Section 10.3.2](#). Adding sample microposts for *all* the users actually takes a rather long time, so first we'll select just the first six users⁸ using the `:limit` option to the `User.all` method.⁹

```
User.all(:limit => 6)
```

We then make 50 microposts for each user (plenty to overflow the pagination limit of 30), generating sample content for each micropost using the Faker gem's handy `Lorem.sentence` method. (`Faker::Lorem.sentence` returns *lorem ipsum* text; as noted in [Chapter 6](#), *lorem ipsum* has a [fascinating back story](#).) The result is the new sample data populator shown in [Listing 11.20](#).

Listing 11.20. Adding microposts to the sample data.

```
lib/tasks/sample_data.rake
```

⁸i.e., the five users with custom Gravatars, and one with the default Gravatar

⁹Tail your `log/development.log` file if you're curious about the SQL this method generates.

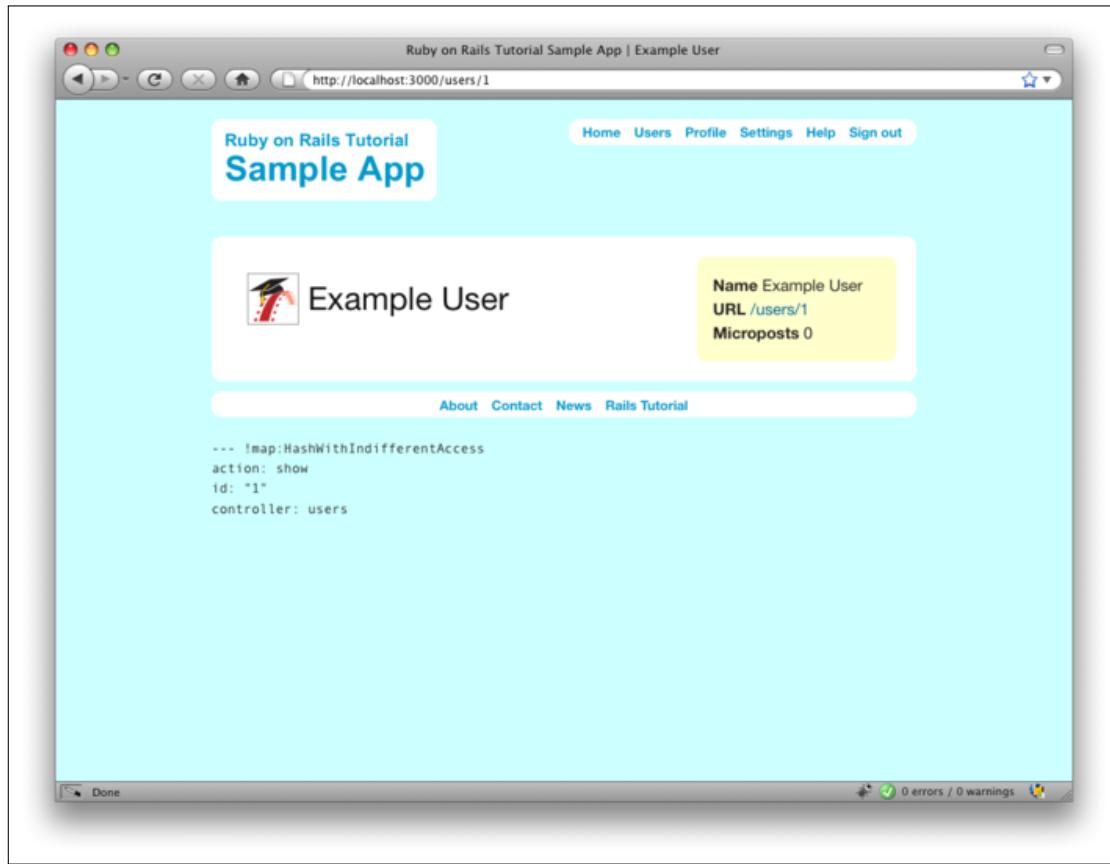


Figure 11.5: The user profile page with code for microposts—but no microposts. ([full size](#))

```
require 'faker'

namespace :db do
  desc "Fill database with sample data"
  task :populate => :environment do
    .
    .
    .
    User.all(:limit => 6).each do |user|
      50.times do
        user.microposts.create! (:content => Faker::Lorem.sentence(5))
      end
    end
  end
end
```

Of course, to generate the new sample data we have to run the `db:populate` Rake task:

```
$ rake db:populate
```

With that, we are in a position to enjoy the fruits of our [Section 11.2.1](#) labors by displaying information for each micropost.¹⁰ [Figure 11.6](#) shows the user profile page for the first (signed-in) user, while [Figure 11.7](#) shows the profile for a second user. (Note in particular the disappearance of the delete links in the latter case.) Finally, [Figure 11.8](#) shows the *second* page of microposts for the first user, along with the pagination links at the bottom of the display. In all three cases, observe that each micropost display indicates the time since it was created (e.g., “Posted 1 minute ago.”); this is the work of the `time_ago_in_words` method from [Listing 11.17](#). If you wait a couple minutes and reload the pages, you’ll see how the text gets automatically updated based on the new time.

11.3 Manipulating microposts

Having finished both the data modeling and display templates for microposts, we now turn our attention to the interface for creating them through the web. The result will be our third example of using an HTML form to create a resource—in this case, a Microposts resource.¹¹ In this section we’ll also see the first hint of a *status feed*—a notion brought to full fruition in [Chapter 12](#). Finally, as with users, we’ll make it possible to destroy microposts through the web (thus fulfilling the promise implicit in the delete links from [Listing 11.17](#)).

There is one break with past convention worth noting: the interface to the Microposts resource will run principally through the Users and Pages controllers, rather than relying a controller of its own. This means that the routes for the Microposts resource are unusually simple, as seen in [Listing 11.21](#). The code in [Listing 11.21](#) leads in turn to the RESTful routes show in [Table 11.2](#), which is a small subset of the full set of routes seen in [Table 2.3](#). Of course, this simplicity is a sign of being *more* advanced, not less—we’ve come a long way since our reliance on scaffolding in [Chapter 2](#), and we no longer need most of its complexity.

¹⁰By design, the Faker gem’s *lorem ipsum* text is randomized, so the contents of your sample microposts will differ.

¹¹The other two resources are Users in [Section 8.1](#) and Sessions in [Section 9.1](#).

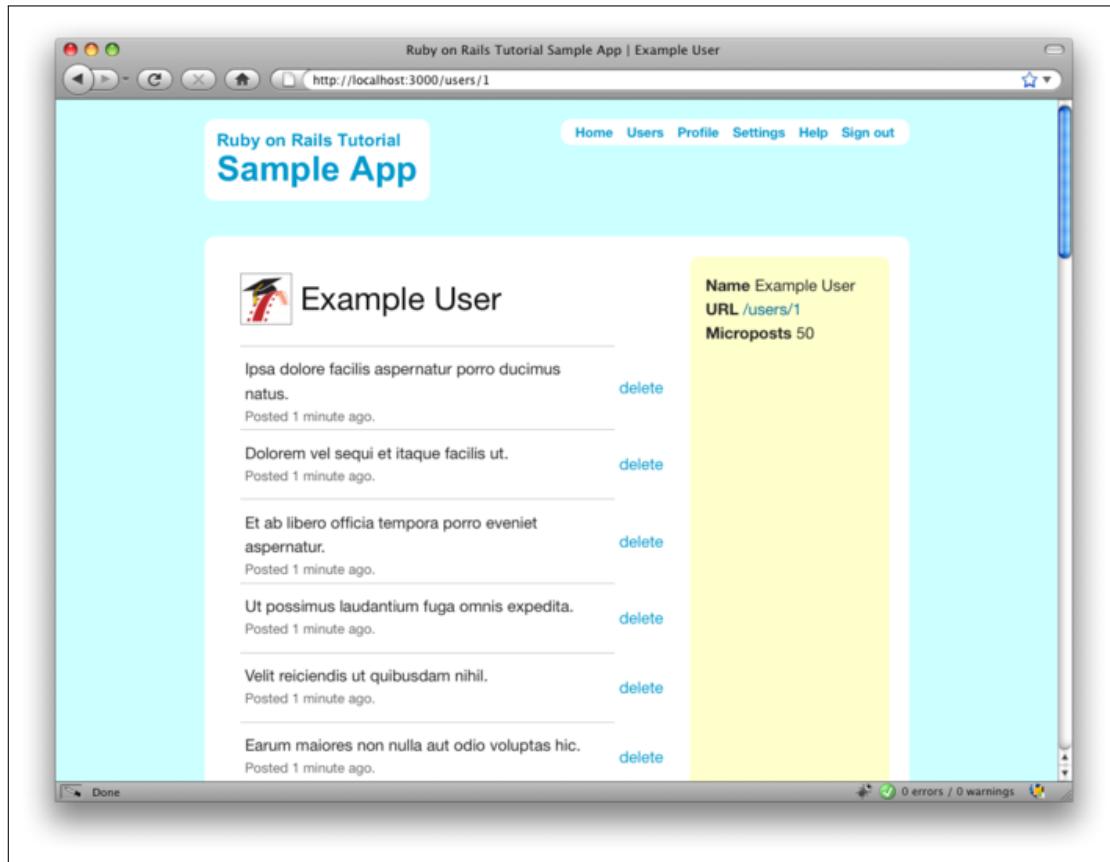


Figure 11.6: The user profile (</users/1>) with microposts. ([full size](#))

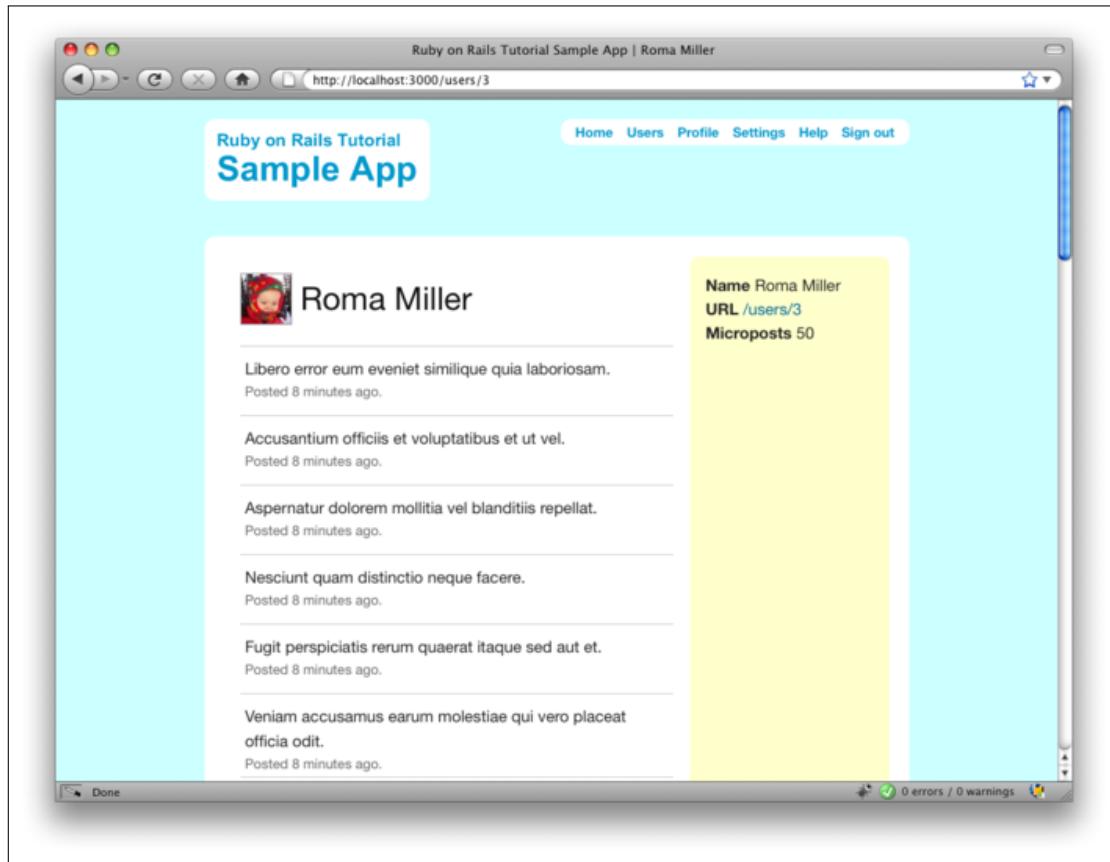


Figure 11.7: The profile of a different user, also with microposts ([/users/3](#)). [\(full size\)](#)

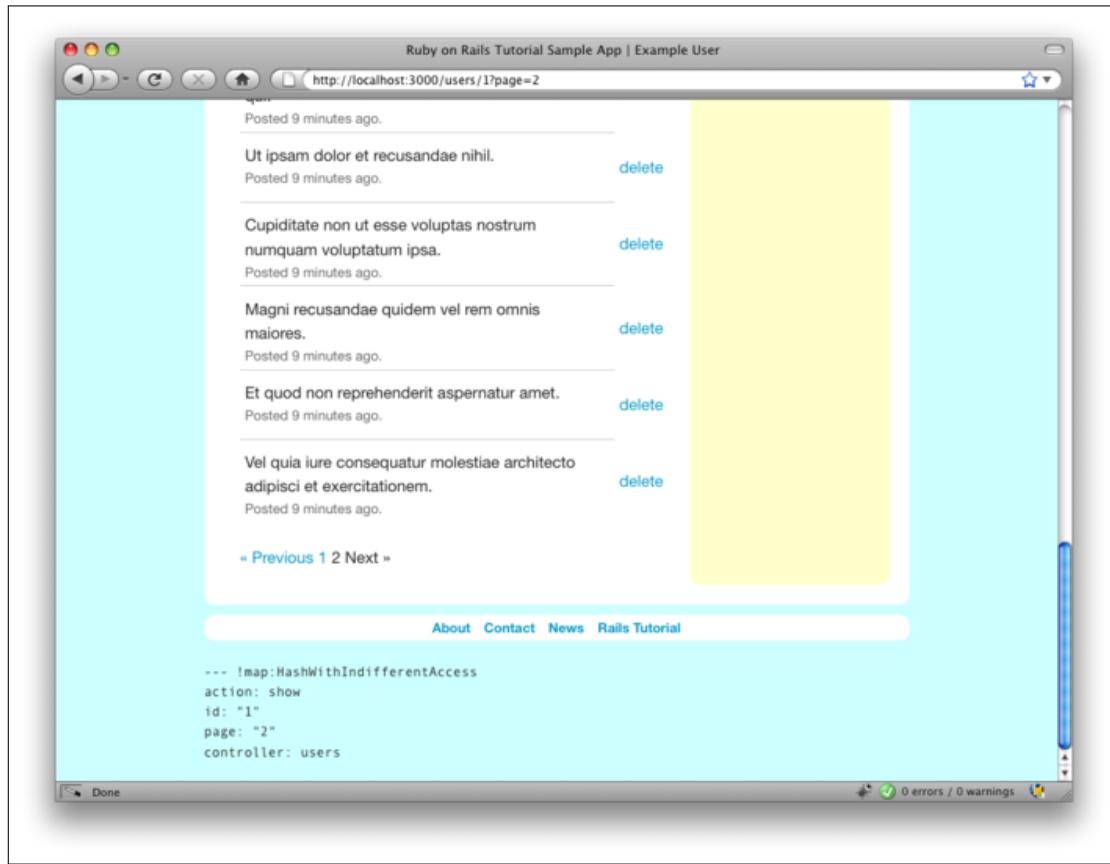


Figure 11.8: A second of profile microposts, with pagination links (</users/1?page=2>). [\(full size\)](#)

HTTP request	URL	Action	Purpose
POST	/microposts	create	create a new micropost
DELETE	/microposts/1	destroy	delete micropost with id 1

Table 11.2: RESTful routes provided by the Microposts resource in Listing 11.21.

Listing 11.21. Routes for the Microposts resource.`config/routes.rb`

```
ActionController::Routing::Routes.draw do |map|
  map.resources :users
  map.resources :sessions, :only => [:new, :create, :destroy]
  map.resources :microposts, :only => [:create, :destroy]
  .
  .
  .
end
```

11.3.1 Access control

We begin our development of the Microposts resource with some access control in the Microposts controller. The idea is simple: both the **create** and **destroy** actions should require users to be signed in. The RSpec code to test for this appears in Listing 11.22. (We'll test for and add a third protection—ensuring that only a micropost's user can destroy it—in Section 11.3.4.)

Listing 11.22. Access control tests for the Microposts controller.`spec/controllers/microposts_controller_spec.rb`

```
require 'spec_helper'

describe MicropostsController do
  integrate_views

  describe "access control" do
    it "should deny access to 'create'" do
      post :create
      response.should redirect_to(signin_path)
    end

    it "should deny access to 'destroy'" do
      delete :destroy
      response.should redirect_to(signin_path)
    end
  end
end
```

Writing the application code needed to get the tests in Listing 11.22 to pass requires a little refactoring first. Recall from Section 10.2.1 that we enforced the signin requirement using a before filter that called the `authenticate` method (Listing 10.9). At the time, we only needed `authenticate` in the Users controller, but now we find that we need it in the Microposts controller as well, so we'll move `authenticate` into the Sessions helper, as shown in Listing 11.23.¹²

Listing 11.23. Moving the `authenticate` method into the Sessions helper.

`app/helpers/sessions_helper.rb`

```
module SessionsHelper
  .
  .
  .

  def authenticate
    deny_access unless signed_in?
  end

  def deny_access
    store_location
    flash[:notice] = "Please sign in to access this page."
    redirect_to(signin_path)
  end
  .
  .
  .
end
```

(To avoid code repetition, you should also remove `authenticate` from the Users controller at this time.)

With the code in Listing 11.23, the `authenticate` method is now available in the Microposts controller, which means that we can restrict access to the `create` and `destroy` actions with the before filter shown in Listing 11.24.

Listing 11.24. Adding authentication to the Microposts controller actions.

`app/controllers/microposts_controller.rb`

```
class MicropostsController < ApplicationController
  before_filter :authenticate

  def create
  end

  def destroy
  end
end
```

¹²We noted in Section 9.14 that helper methods are available only in `views` by default, but we arranged for the Sessions helper methods to be available in the controllers as well by adding `include SessionsHelper` to the Application controller (Listing 9.11).

Note that we haven't restricted the actions the before filter applies to, since presently it applies to them both. If we were to add, say, an `index` action accessible even to non-signed-in users, we would need to specify the protected actions explicitly:

```
class MicropostsController < ApplicationController
  before_filter :authenticate, :only => [:create, :destroy]

  def index
  end

  def create
  end

  def destroy
  end
end
```

11.3.2 Creating microposts

In [Chapter 8](#), we implemented user signup by making an HTML form that issued an HTTP POST request to the `create` action in the Users controller. The implementation of micropost creation is similar; the main difference is that, rather than using a separate page at `/microposts/new`, we will (following Twitter's convention) put the form on the Home page itself (i.e., the root path `/`), as mocked up in [Figure 11.9](#).

When we last left the Home page, it appeared as in [Figure 5.7](#)—that is, it had a big, fat “Sign up now!” button in the middle. Since a micropost creation form only makes sense in the context of a particular signed-in user, one goal of this section will be to serve different versions of the Home page depending on a visitor's signin status. We'll implement this in [Listing 11.27](#) below, but for now the only implication is that the tests for the Microposts controller `create` action should sign a (factory) user in before attempting to make a post.

With that caveat in mind, the micropost creation tests parallel those for user creation from [Listing 8.6](#) and [Listing 8.13](#); the result appears in [Listing 11.25](#).

Listing 11.25. Tests for the Microposts controller `create` action.

```
spec/controllers/microposts_controller_spec.rb
```

```
require 'spec_helper'

describe MicropostsController do
  .
  .
  .
  describe "POST 'create'" do
    before(:each) do
      @user = test_sign_in(Factory(:user))
      @attr = { :content => "Lorem ipsum" }
      @micropost = Factory(:micropost, @attr.merge(:user => @user))

      @user.microposts.stub! (:build).and_return(@micropost)
    end
  end
end
```

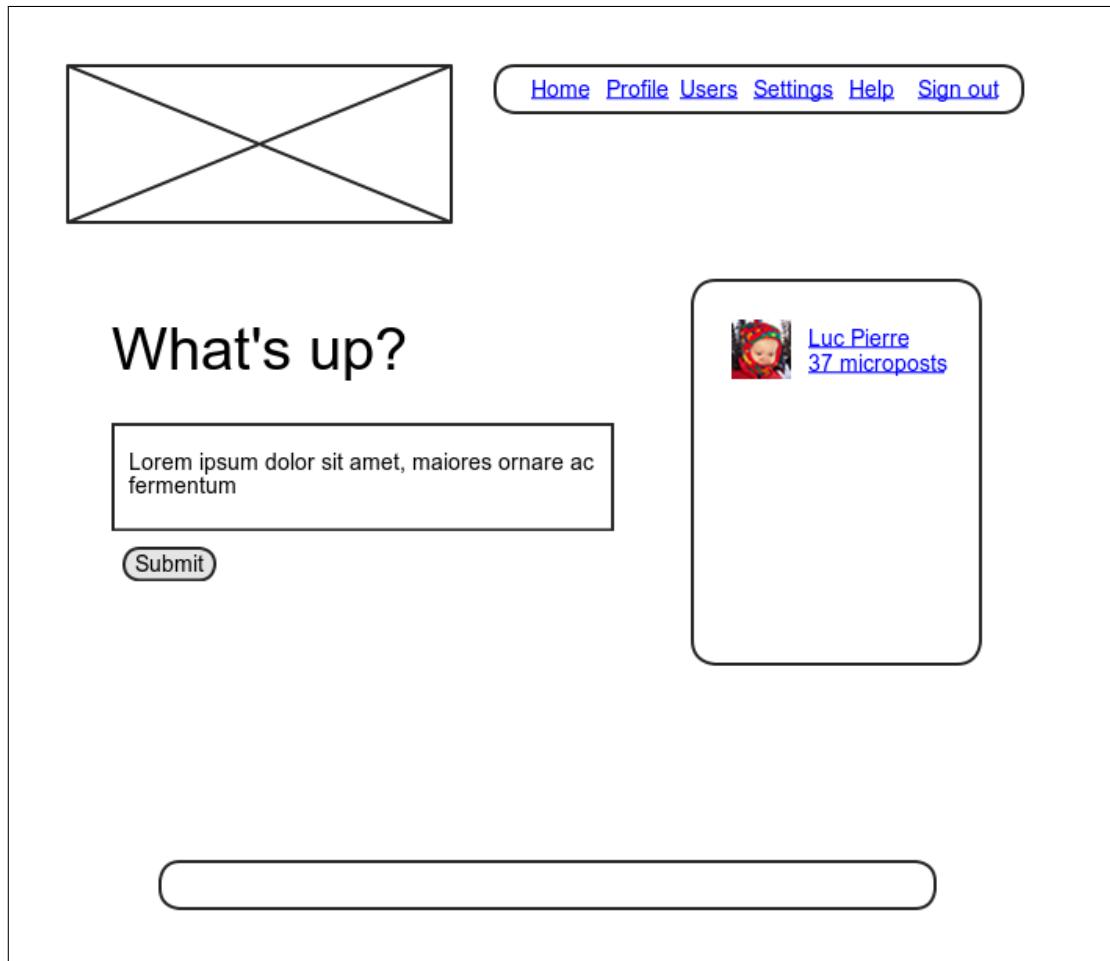


Figure 11.9: A mockup of the Home page with a form for creating microposts. ([full size](#))

```
end

describe "failure" do

  before(:each) do
    @micropost.should_receive(:save).and_return(false)
  end

  it "should render the home page" do
    post :create, :micropost => @attr
    response.should render_template('pages/home')
  end
end

describe "success" do

  before(:each) do
    @micropost.should_receive(:save).and_return(true)
  end

  it "should redirect to the home page" do
    post :create, :micropost => @attr
    response.should redirect_to(root_path)
  end

  it "should have a flash message" do
    post :create, :micropost => @attr
    flash[:success].should =~ /micropost created/i
  end
end
end
end
```

Here we've stubbed out `@user.microposts.build` instead of `User.new` as in, e.g., Listing 8.6, but the basic idea is the same: the test intercepts calls to `build` and returns instead the factory micropost created by

```
@micropost = Factory(:micropost, @attr.merge(:user => @user))
```

The micropost creation failure and success tests then use the message expectations

```
@micropost.should_receive(:save)
```

to ensure that the code attempts to save the micropost, in each case either rendering the Home page (on failure) or redirecting (on success).¹³

¹³Note that '`pages/home`' and `root_path` are actually the same page in the present case, but this needn't be so in general, so we use different notation to keep them conceptually distinct.

The `create` action for microposts is similar to its user analogue ([Listing 8.14](#)); the principal difference lies in using the user/micropost association to `build` the new micropost, as seen in [Listing 11.26](#).

Listing 11.26. The Microposts controller `create` action.

`app/controllers/microposts_controller.rb`

```
class MicropostsController < ApplicationController
  .
  .
  .

  def create
    @micropost = current_user.microposts.build(params[:micropost])
    if @micropost.save
      flash[:success] = "Micropost created!"
      redirect_to root_path
    else
      render 'pages/home'
    end
  end
end
```

At this point, the tests in [Listing 11.26](#) should all be passing, but of course we still don't have a form to create microposts. We can rectify this with [Listing 11.27](#), which serves up different HTML based on whether the site visitor is signed in or not.

Listing 11.27. Adding microposts creation to the Home page ([/](#)).

`app/views/pages/home.html.erb`

```
<% if signed_in? %>


|                                                                            |                                 |
|----------------------------------------------------------------------------|---------------------------------|
| <h2 class="micropost">What's up?</h2> <%= render 'pages/micropost_form' %> | <%= render 'pages/user_info' %> |
|----------------------------------------------------------------------------|---------------------------------|


<% else %>


# Sample App



This is the home page for the
    Ruby on Rails Tutorial


```

```
sample application.  
</p>  
  
<%= link_to "Sign up now!", signup_path, :class => "signup_button round" %>  
<% end %>
```

Having so much code in each branch of the `if-else` conditional is a bit messy, and cleaning it up using partials is left as an exercise (Section 11.5). Filling in the necessary partials from Listing 11.27 isn't an exercise, though; we fill in the micropost form partial in Listing 11.28 and the new Home page sidebar in Listing 11.29.

Listing 11.28. The form partial for creating microposts.

`app/views/pages/_micropost_form.html.erb`

```
<% form_for @micropost do |f| %>  
  <%= f.error_messages %>  
  <div class="field">  
    <%= f.text_area :content %>  
  </div>  
  <div class="actions">  
    <%= f.submit "Submit" %>  
  </div>  
<% end %>
```

Listing 11.29. The partial for the user info sidebar.

`app/views/pages/_user_info.html.erb`

```
<div class="user_info">  
  <a href="<%= user_path(current_user) %>">  
    <%= gravatar_for current_user, :size => 30 %>  
    <span class="user_name">  
      <%= h current_user.name %>  
    </span>  
    <span class="microposts">  
      <%= pluralize(current_user.microposts.count, "micropost") %>  
    </span>  
  </a>  
</div>
```

Note that, as in the profile sidebar (Listing 11.16), the user info in Listing 11.29 displays the total number of microposts for the user. There's a slight difference in the display, though; in the profile sidebar, **Microposts** is a label, and showing **Microposts 1** makes perfect sense. In the present case, though, saying “1 microposts” is ungrammatical, so we arrange to display “1 micropost” (but “2 microposts”) using the convenient `pluralize` helper method.

The form defined in Listing 11.28 is an exact analogue of the signup form in Listing 8.2, which means that it needs an `@micropost` instance variable. This is supplied in Listing 11.30—but only when the user is signed in.

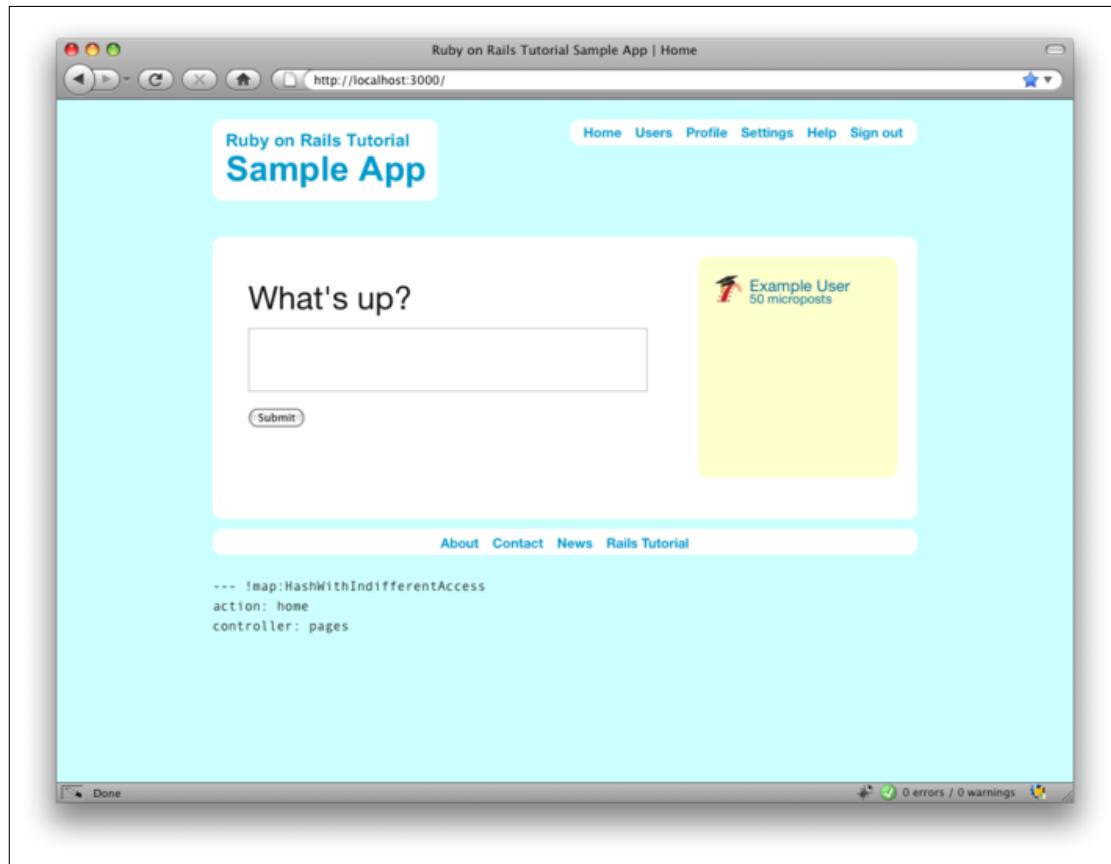


Figure 11.10: The Home page ([/](#)) with a new micropost form. ([full size](#))

Listing 11.30. Adding a micropost instance variable to the `home` action.
`app/controllers/pages_controller.rb`

```
class PagesController < ApplicationController
  def home
    @title = "Home"
    @micropost = Micropost.new if signed_in?
  end
end
```

Now the HTML should render properly, showing the form as in Figure 11.10, and a form with a submission error as in Figure 11.11. You are invited at this point to create a new post for yourself and verify that everything is working—but you should probably wait until after Section 11.3.3.

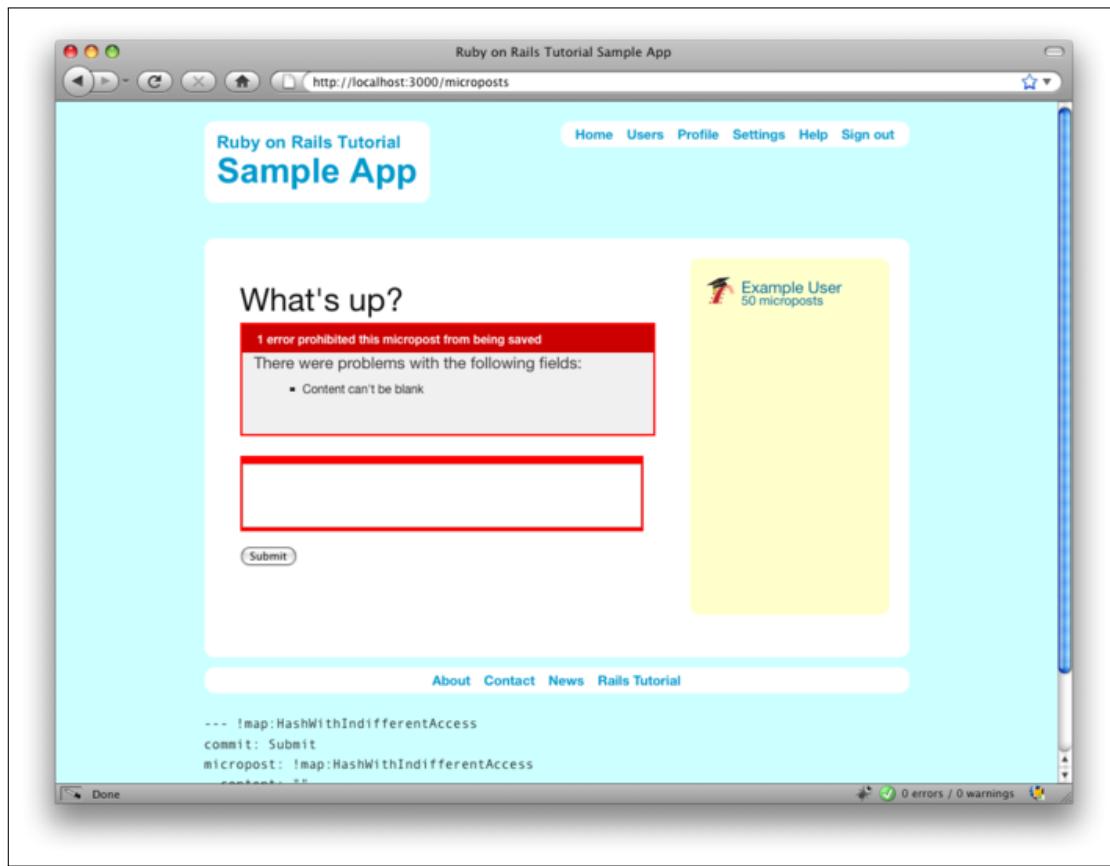


Figure 11.11: The home page with form errors. (full size)

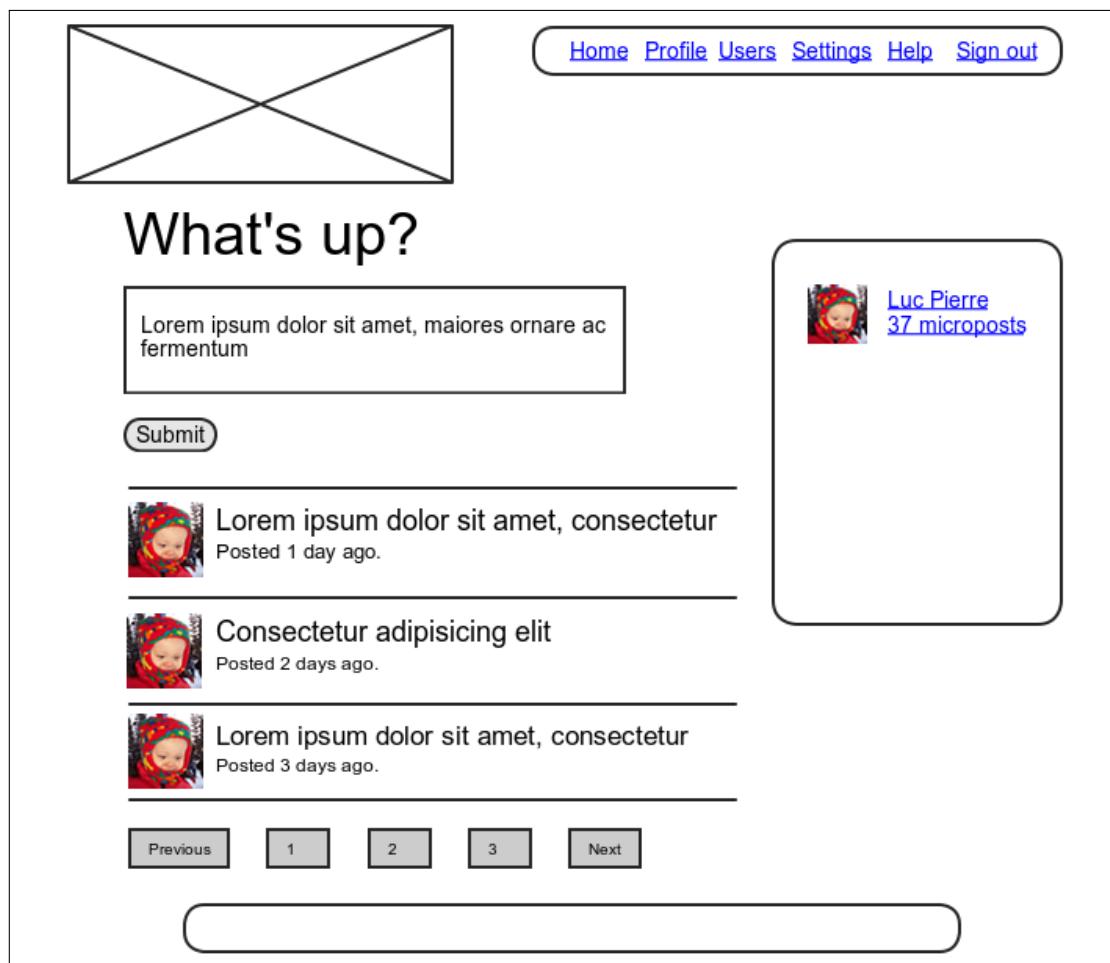


Figure 11.12: A mockup of the Home page with a proto-feed. ([full size](#))

11.3.3 A proto-feed

The comment at the end of [Section 11.3.2](#) alluded to a problem: the current Home page doesn't display any microposts. If you like, you can verify that the form shown in [Figure 11.10](#) is working by submitting a valid entry and then navigating to the [profile page](#) to see the post, but that's rather cumbersome. It would be far better to have a *feed* of microposts that includes the user's own posts, as mocked up in [Figure 11.12](#). (In [Chapter 12](#), we'll generalize this feed to include the microposts of users being *followed* by the current user.)

Since each user should have a feed, we are led naturally to a **feed** method in the User model. Eventually, we will test that the feed returns the microposts of the users being followed, but for now we'll just test that the **feed** method *includes* the current user's microposts but *excludes* the posts of a different user. We can express these requirements in code with [Listing 11.31](#).

Listing 11.31. Tests for the (proto-)status feed.

`specs/models/user_spec.rb`

```
require 'spec_helper'

describe User do
  .
  .
  .
  describe "micropost associations" do
    before(:each) do
      @user = User.create(@attr)
      @mp1 = Factory(:micropost, :user => @user, :created_at => 1.day.ago)
      @mp2 = Factory(:micropost, :user => @user, :created_at => 1.hour.ago)
    end
  .
  .
  .
  describe "status feed" do
    it "should have a feed" do
      @user.should respond_to(:feed)
    end

    it "should include the user's microposts" do
      @user.feed.include?(@mp1).should be_true
      @user.feed.include?(@mp2).should be_true
    end

    it "should not include a different user's microposts" do
      mp3 = Factory(:micropost,
                    :user => Factory(:user, :email => Factory.next(:email)))
      @user.feed.include?(mp3).should be_false
    end
  end
end
```

These tests introduce the array `include?` method, which simply checks if an array includes the given element:¹⁴

```
$ script/console
>> a = [1, "foo", :bar]
>> a.include?("foo")
=> true
>> a.include?(:bar)
=> true
>> a.include?("baz")
=> false
```

¹⁴Learning about methods such as `include?` is one reason why, as noted in Section 1.1.1, I recommend reading a pure Ruby book after finishing this one.

We can arrange for an appropriate micropost `feed` by selecting all the microposts with `user_id` equal to the current user's id, which we accomplish using a `:conditions` option to `Micropost.all`, as shown in Listing 11.32.¹⁵

Listing 11.32. A preliminary implementation for the micropost status feed.

`app/models/user.rb`

```
class User < ActiveRecord::Base
  .
  .
  .
  def feed
    # This is preliminary. See Chapter 12 for the full implementation.
    Micropost.all(:conditions => ["user_id = ?", id])
  end
  .
  .
  .
end
```

The question mark in

```
:conditions => ["user_id = ?", id]
```

ensures that `id` is properly *escaped* before being included in the underlying SQL query, thereby avoiding a serious security breach called [SQL injection](#). (The `id` attribute here is just an integer, so there is no danger in this case, but *always* escaping variables injected into SQL statements is a good habit to cultivate.)

Alert readers might note at this point that the code in Listing 11.32 is essentially equivalent to writing

```
def feed
  microposts
end
```

We've used the code in Listing 11.32 instead because it generalizes much more naturally to the full status feed needed in Chapter 12.

To use the feed in the sample application, we add an `@feed_items` instance variable for the current user's (paginated) feed, as in Listing 11.33, and then add a feed partial (Listing 11.34) to the Home page (Listing 11.36).

Listing 11.33. Adding a feed instance variable to the `home` action.

`app/controllers/pages_controller.rb`

¹⁵See the Rails Guide on the [Active Record Query Interface](#) for more on `:conditions` and the like.

```
class PagesController < ApplicationController
  def home
    @title = "Home"
    if signed_in?
      @micropost = Micropost.new
      @feed_items = current_user.feed.paginate(:page => params[:page])
    end
  end
.
```

Listing 11.34. The status feed partial.
`app/views/pages/_feed.html.erb`

```
<% unless @feed_items.empty? %>





<%= will_paginate @feed_items %>
<% end %>
```

The status feed partial defers the feed item rendering to a feed item partial using the code

```
<%= render :partial => 'feed_item', :collection => @feed_items %>
```

Here we pass a `:collection` parameter with the feed items, which causes `render` to use the given partial (`'feed_item'` in this case) to render each item in the collection. (We have omitted the `:partial` parameter in previous renderings, writing, e.g., `render 'pages/micropost'`, but with a `:collection` parameter that syntax doesn't work.) The feed item partial itself appears in Listing 11.35.

Listing 11.35. A partial for a single feed item.
`app/views/pages/_feed_item.html.erb`

```
<tr>
  <td class="gravatar">
    <%= link_to gravatar_for(feed_item.user), feed_item.user %>
  </td>
  <td class="micropost">
    <span class="user">
      <%= link_to h(feed_item.user.name), feed_item.user %>
    </span>
    <span class="content"><%= h feed_item.content %></span>
    <span class="timestamp">
```

```

  Posted <%= time_ago_in_words(feed_item.created_at) %> ago.
</span>
</td>
<% if current_user?(feed_item.user) %>
<td>
  <%= link_to "delete", feed_item, :method => :delete,
              :confirm => "You sure?" %>
</td>
<% end %>
</tr>

```

We can then add the feed to the Home page by rendering the feed partial as usual ([Listing 11.36](#)). The result is a display of the feed on the Home page, as required ([Figure 11.13](#)).

Listing 11.36. Adding a status feed to the Home page.

`app/views/pages/home.html.erb`

```

<% if signed_in? %>


|                                                                                                       |
|-------------------------------------------------------------------------------------------------------|
| <h2 class="micropost">What's up?</h2> <%= render 'pages/micropost_form' %> <%= render 'pages/feed' %> |
|-------------------------------------------------------------------------------------------------------|


```

At this point, creating a new micropost works as expected, as seen in [Figure 11.14](#). (We'll write an integration test to this effect in [Section 11.3.5](#).) There is one subtlety, though: on *failed* micropost submission, the Home page expects an `@feed_items` instance variable, so failed submissions currently break (as you should be able to verify by running your test suite). The easiest solution is to suppress the feed entirely by assigning it an empty array, as shown in [Listing 11.37](#).¹⁶

Listing 11.37. Adding an (empty) `@feed_items` instance variable to the `create` action.

`app/controllers/microposts_controller.rb`

¹⁶Unfortunately, returning a paginated feed doesn't work in this case. Implement it and click on a pagination link to see why. (The screencasts will cover this issue in more depth.)

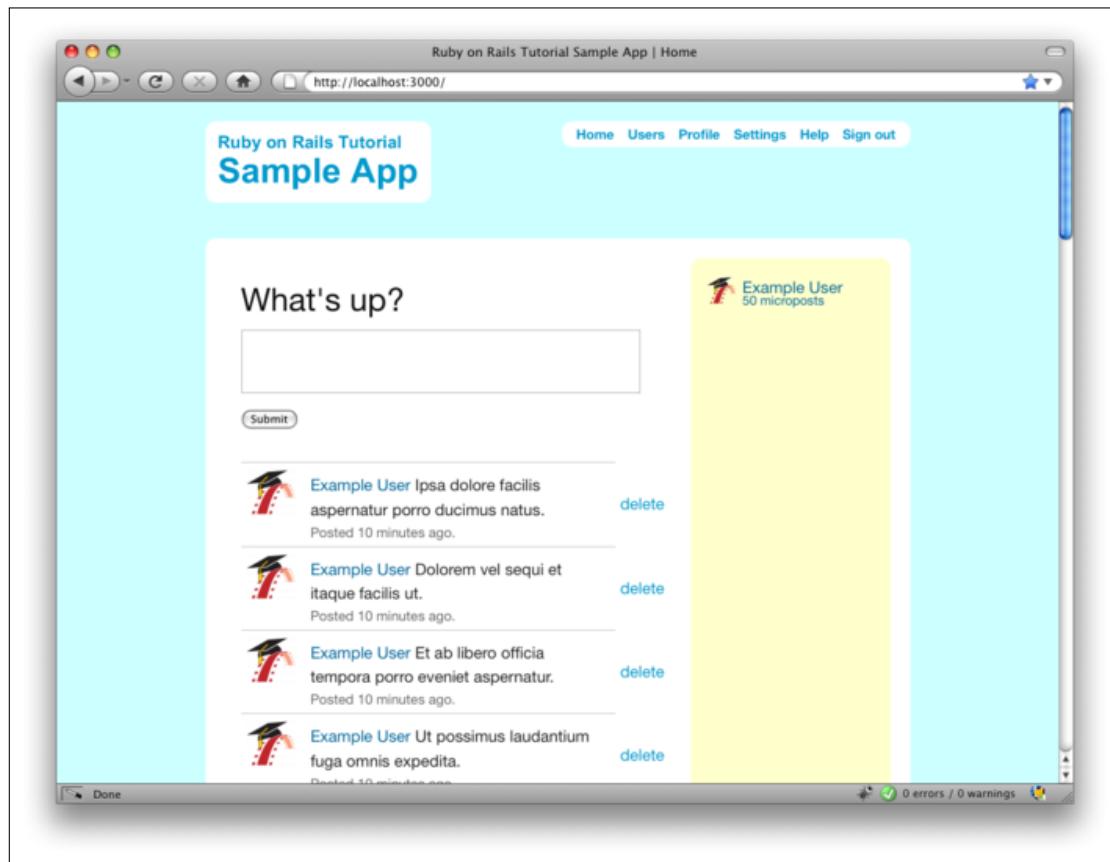


Figure 11.13: The Home page (/) with a proto-feed. ([full size](#))

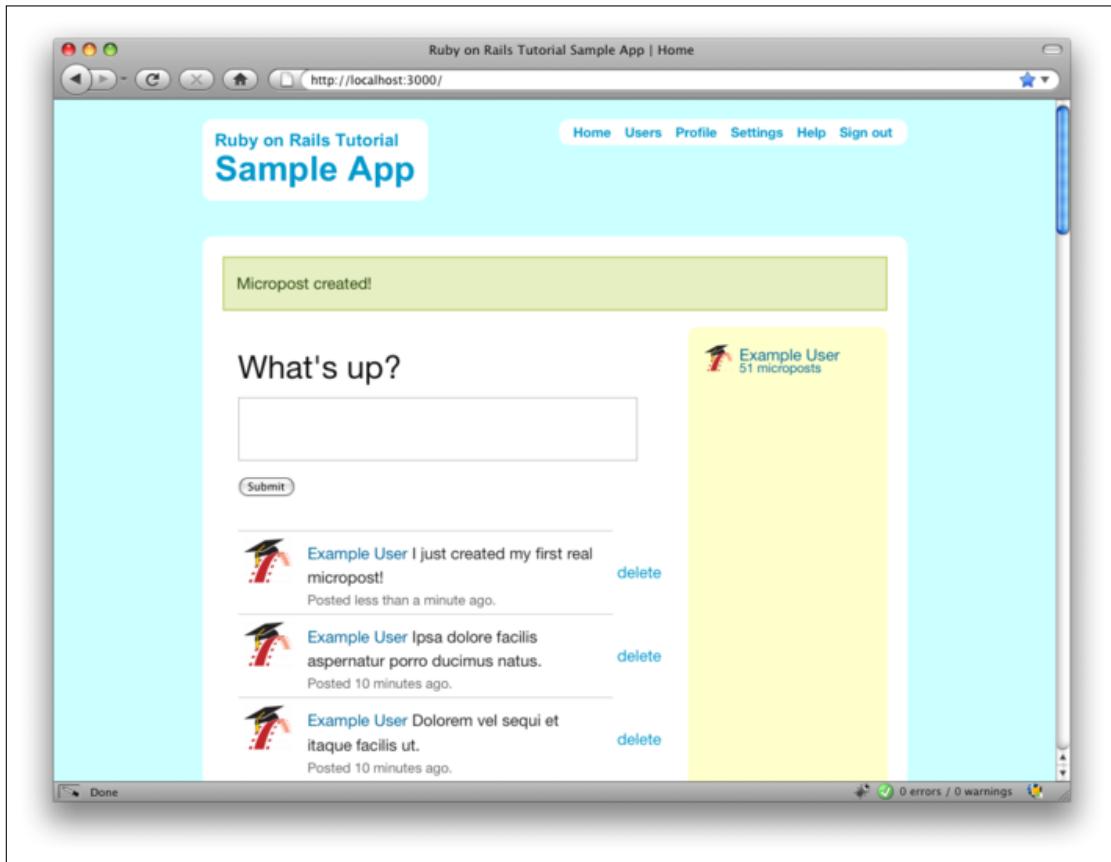


Figure 11.14: The Home page after creating a new micropost. ([full size](#))

```
class MicropostsController < ApplicationController
  .
  .
  .
  def create
    @micropost = current_user.microposts.build(params[:micropost])
    if @micropost.save
      flash[:success] = "Micropost created!"
      redirect_to root_path
    else
      @feed_items = []
      render 'pages/home'
    end
  end
  .
  .
  .
end
```

11.3.4 Destroying microposts

The last piece of functionality to add to the Microposts resource is the ability to destroy posts. As with user deletion (Section 10.4.2), we accomplish this with “delete” links, as mocked up in Figure 11.15. Unlike that case, which restricted user destruction to admin users, the delete links will work only for microposts created by the current user.

Our first step is to add a delete link to the micropost partial, following the example from Listing 10.37. The result appears in Listing 11.38.

Listing 11.38. A partial for showing a single micropost.

`app/views/microposts/_micropost.html.erb`

```
<tr>
  <td class="micropost">
    <span class="content"><%= h micropost.content %></span>
    <span class="timestamp">
      Posted <%= time_ago_in_words(micropost.created_at) %> ago.
    </span>
  </td>
  <% if current_user?(micropost.user) %>
  <td>
    <%= link_to "delete", micropost, :method => :delete,
               :confirm => "You sure?" %>
  </td>
  <% end %>
</tr>
```

The tests for the `destroy` action are straightforward generalizations of the similar tests for destroying users (Listing 10.38), as seen in Listing 11.39.

What's up?

Lore ipsum dolor sit amet, maiores ornare ac fermentum

[Submit](#)



Lore ipsum dolor sit amet, consectetur
Posted 1 day ago.

[delete](#)



Consectetur adipisicing elit
Posted 2 days ago.

[delete](#)



Lore ipsum dolor sit amet, consectetur
Posted 3 days ago.

[delete](#)

[Previous](#) [1](#) [2](#) [3](#) [Next](#)

Figure 11.15: A mockup of the proto-feed with micropost delete links. ([full size](#))

Listing 11.39. Tests for the Microposts controller `destroy` action.

`spec/controllers/microposts_controller_spec.rb`

```
describe MicropostsController do
  .
  .
  .
  describe "DELETE 'destroy'" do
    describe "for an unauthorized user" do
      before(:each) do
        @user = Factory(:user)
        wrong_user = Factory(:user, :email => Factory.next(:email))
        test_sign_in(wrong_user)
        @micropost = Factory(:micropost, :user => @user)
      end

      it "should deny access" do
        @micropost.should_not_receive(:destroy)
        delete :destroy, :id => @micropost
        response.should redirect_to(root_path)
      end
    end
  end

  describe "for an authorized user" do
    before(:each) do
      @user = test_sign_in.Factory(:user)
      @micropost = Factory(:micropost, :user => @user)
      Micropost.should_receive(:find).with(@micropost).and_return(@micropost)
    end

    it "should destroy the micropost" do
      @micropost.should_receive(:destroy).and_return(@micropost)
      delete :destroy, :id => @micropost
    end
  end
end
end
```

The application code is also analogous to the user case in Listing 10.39; the main difference is that, rather than using an `admin_user` before filter, in the case of microposts we have an `authorized_user` before filter to check that the current user is the micropost's user. The code appears in Listing 11.40, and the result of destroying the second-most-recent post appears in Figure 11.16.

Listing 11.40. The Microposts controller `destroy` action.

`app/controllers/microposts_controller.rb`

```

class MicropostsController < ApplicationController
  before_filter :authenticate, :only => [:create, :destroy]
  before_filter :authorized_user, :only => :destroy
  .
  .
  .

  def destroy
    @micropost.destroy
    redirect_back_or root_path
  end

  private

  def authorized_user
    @micropost = Micropost.find(params[:id])
    redirect_to root_path unless current_user?(@micropost.user)
  end
end

```

11.3.5 Testing the new home page

Before leaving micropost creation and destruction, we'll write some RSpec integration specs to test that our forms are working properly. As in the case of users (Section 8.4), we start by generating a microposts integration spec:

```
$ script/generate integration_spec microposts
```

Tests for failed and successful micropost creation and for micropost destruction appear in Listing 11.41.

Listing 11.41. An integration test for the microposts on the **home** page.

specs/integration/microposts_spec.rb

```

require 'spec_helper'

describe "Microposts" do

  before(:each) do
    @user = Factory(:user)
    visit signin_path
    fill_in :email,      :with => @user.email
    fill_in :password,  :with => @user.password
    click_button
  end

  describe "creation" do

```

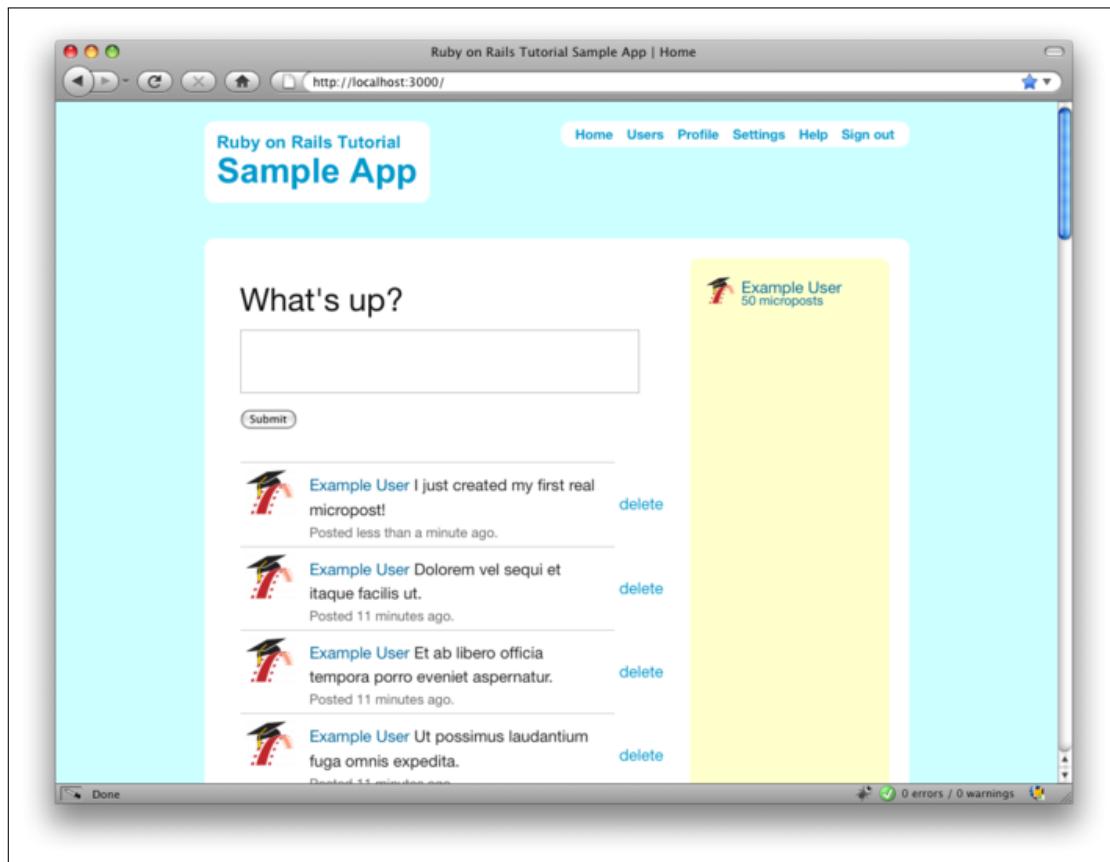


Figure 11.16: The user home page after deleting the second-most-recent micropost. ([full size](#))

```

describe "failure" do

  it "should not make a new micropost" do
    lambda do
      visit root_path
      click_button
      response.should render_template('pages/home')
      response.should have_tag("div#errorExplanation")
    end.should_not change(Micropost, :count)
  end
end

describe "success" do

  it "should make a new micropost" do
    content = "Lorem ipsum dolor sit amet"
    lambda do
      visit root_path
      fill_in :micropost_content, :with => content
      click_button
      response.should have_tag("span.content", content)
    end.should change(Micropost, :count).by(1)
  end
end
end

describe "destruction" do

  it "should destroy a micropost" do
    # Create a micropost.
    visit root_path
    fill_in :micropost_content, :with => "lorem ipsum"
    click_button
    # Destroy it.
    lambda { click_link "delete" }.should change(Micropost, :count).by(-1)
  end
end
end

```

In the final test, we use the curly brace syntax for Ruby blocks (Section 4.3.2) and the Webrat `click_link` method to verify that clicking on the delete link removes one micropost:

```
lambda { click_link "delete" }.should change(Micropost, :count).by(-1)
```

11.4 Conclusion

With the addition of the Microposts resource, we are nearly finished with our sample application. All that remains is to add a social layer by letting users follow each other. We'll learn how to model such user

relationships, and see the implications for the status feed, in [Chapter 12](#).

Before proceeding, be sure to commit and merge your changes if you’re using Git for version control:

```
$ git add .
$ git commit -am "Added user microposts"
$ git checkout master
$ git merge user-microposts
```

You can also push the app up to Heroku at this point. Because the data model has changed through the addition of the `microposts` table, you will also need to migrate the production database:

```
$ git push heroku
$ heroku rake db:migrate
```

11.5 Exercises

We’ve covered enough material now that there is a combinatorial explosion of possible extensions to the application. Here are just a few of the many possibilities:

1. **(challenging)** Add a JavaScript display to the Home page to count down from 140 characters.
2. Add tests for the sidebar micropost counts (including proper pluralization).
3. **(mainly for designers)** Modify the microposts listing to use an ordered list instead of a table. (*Note:* this is how Twitter displays its status updates.) Then add the appropriate CSS to make the resulting feed not look like crap.
4. Add tests for micropost pagination.
5. Refactor the Home page to use separate partials for the two branches of the `if-else` statement.
6. Write a test to make sure delete links do not appear for microposts not created by the current user.
7. Define a `display` helper using the `sanitize` and `auto_link` methods, as shown in Listing 11.42, and use it to replace the `h` method in the micropost partial ([Listing 11.17](#)). Examine the results for a micropost whose content includes a URL (such as “I’m reading an awesome Rails book at [http://www.railstutorial.org/!](http://www.railstutorial.org/)”).
8. Add a `nested route` so that `/users/1/microposts` shows all the microposts for user 1. (You will also have to add a Microposts controller `index` action and corresponding view.)
9. Very long words currently break our layout, as shown in [Figure 11.17](#). Fix this problem using the `wrap` helper defined in [Listing 11.43](#).

Listing 11.42. Adding a `display` helper method for micropost content.

`app/helpers/microposts_helper.rb`

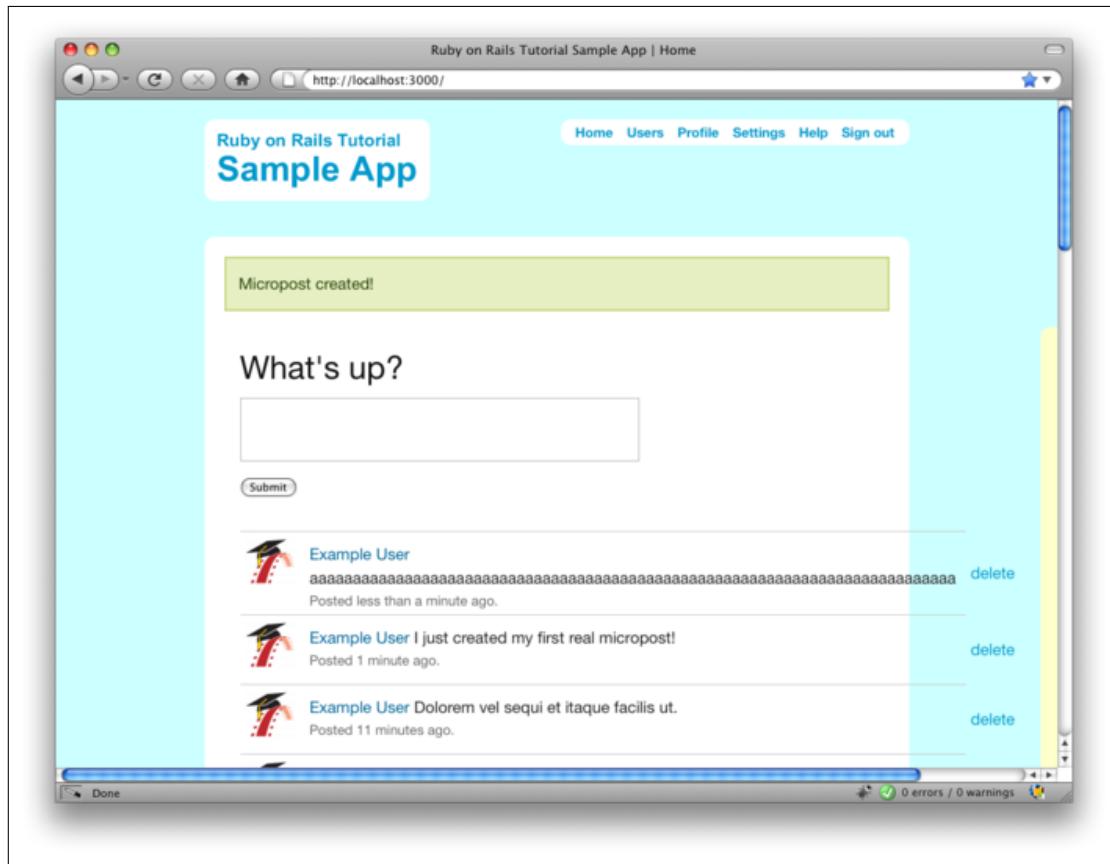


Figure 11.17: The (broken) site layout with a particularly long word. ([full size](#))

```
module MicropostsHelper

  def display(content)
    auto_link(sanitize(content))
  end
end
```

Listing 11.43. A helper to wrap long words.
`app/helpers/microposts_helper.rb`

```
module MicropostsHelper

  def wrap(content)
    content.split.map{ |s| wrap_long_string(s) }.join(' ')
```

```
end

private

def wrap_long_string(text, max_width = 30)
  zero_width_space = "\u200b"
  regex = /.{1,#{max_width}}/
  (text.length < max_width) ? text :
    text.scan(regex).join(zero_width_space)
end
end
```


Chapter 12

Following users

This chapter has been started.